

The NEORV32 RISC-V Processor - Datasheet

Version v1.12.5-r58-gafcde96d



Documentation

The online documentation of the project (a.k.a. the **data sheet**) is available on GitHub-pages: <https://stnolting.github.io/neorv32/>

The online documentation of the **software framework** is also available on GitHub-pages: <https://stnolting.github.io/neorv32/sw/files.html>

Table of Contents

1. Overview	7
1.1. Project Key Features.....	8
1.2. Project Folder Structure	10
1.3. VHDL File Hierarchy	11
1.3.1. File-List Files	12
1.4. VHDL Coding Style	14
1.5. Performance	15
2. NEORV32 Processor (SoC).....	16
2.1. Processor Top Entity - Signals	18
2.2. Processor Top Entity - Generics	22
2.3. Processor Clocking	28
2.3.1. Peripheral Clocks	28
2.4. Processor Reset	29
2.5. Processor Interrupts	30
2.5.1. RISC-V Standard Interrupts	30
2.5.2. NEORV32-Specific Fast Interrupt Requests.....	30
2.6. Address Space	32
2.6.1. Bus System	33
2.6.2. Bus Gateway.....	35
2.6.3. IO Switch.....	35
2.6.4. Atomic Memory Operations Controller.....	36
2.6.5. Memory Coherence	37
2.7. Boot Configuration	40
2.7.1. Booting via Bootloader	40
2.7.2. Boot from Custom Address	40
2.7.3. Boot IMEM Image	41
2.8. Processor-Internal Modules.....	42
2.8.1. Instruction Memory (IMEM)	43
2.8.2. Data Memory (DMEM)	45
2.8.3. Bootloader ROM (BOOTROM).....	47
2.8.4. Processor-Internal Instruction Cache (iCache)	48
2.8.5. Processor-Internal Data Cache (dCache).....	50
2.8.6. Direct Memory Access Controller (DMA)	52
2.8.7. Processor-External Bus Interface (XBUS)	56
2.8.8. Stream Link Interface (SLINK).....	61
2.8.9. General Purpose Input and Output Port (GPIO)	65
2.8.10. Watchdog Timer (WDT).....	68

2.8.11. Core-Local Interruptor (CLINT).....	70
2.8.12. Primary Universal Asynchronous Receiver and Transmitter (UART0).....	72
2.8.13. Secondary Universal Asynchronous Receiver and Transmitter (UART1).....	76
2.8.14. Serial Peripheral Interface Controller (SPI).....	77
2.8.15. Serial Data Interface Controller (SDI)	81
2.8.16. Two-Wire Serial Interface Controller (TWI).....	84
2.8.17. Two-Wire Serial Device Controller (TWD)	88
2.8.18. One-Wire Serial Interface Controller (ONEWIRE)	92
2.8.19. Pulse-Width Modulation Controller (PWM).....	98
2.8.20. True Random-Number Generator (TRNG)	102
2.8.21. Custom Functions Subsystem (CFS)	104
2.8.22. Smart LED Interface (NEOLED).....	106
2.8.23. General Purpose Timer (GPTMR)	110
2.8.24. Execution Trace Buffer (TRACER).....	113
2.8.25. System Configuration Information Memory (SYSINFO)	118
3. NEORV32 Central Processing Unit (CPU)	122
3.1. CPU Top Entity - Signals	123
3.2. CPU Top Entity - Generics	124
3.3. RISC-V Compatibility	124
3.4. Architecture	126
3.4.1. CPU Register File	126
3.4.2. CPU Arithmetic Logic Unit	127
3.4.3. CPU Bus Unit	128
3.4.4. CPU Control Unit	128
3.4.5. Execution Trace Port	129
3.4.6. CPU Tuning Options	130
3.4.7. Sleep Mode	132
3.4.8. Full Virtualization.....	133
3.5. Bus Interface	134
3.5.1. Bus Interface Protocol	135
3.5.2. Locked Bus Accesses and Bursts	136
3.6. Instruction Sets and Extensions	139
3.6.1. Latency Definitions	141
3.6.2. A ISA Extension	141
3.6.3. B ISA Extension	142
3.6.4. C ISA Extension	142
3.6.5. E ISA Extension	142
3.6.6. I ISA Extension	142
3.6.7. M ISA Extension	143

3.6.8. U ISA Extension	144
3.6.9. X ISA Extension	144
3.6.10. Zaamo ISA Extension	144
3.6.11. Zalrsc ISA Extension	144
3.6.12. Zcb ISA Extension	145
3.6.13. Zifencei ISA Extension	145
3.6.14. Zfinx ISA Extension	145
3.6.15. Zibi ISA Extension	146
3.6.16. Zicntr ISA Extension	146
3.6.17. Zicond ISA Extension	147
3.6.18. Zicsr ISA Extension	147
3.6.19. Zihpm ISA Extension	147
3.6.20. Zimop ISA Extension	148
3.6.21. Zba ISA Extension	148
3.6.22. Zbb ISA Extension	148
3.6.23. Zbs ISA Extension	149
3.6.24. Zbkb ISA Extension	149
3.6.25. Zbkc ISA Extension	150
3.6.26. Zbkx ISA Extension	150
3.6.27. Zkn ISA Extension	150
3.6.28. Zknd ISA Extension	150
3.6.29. Zkne ISA Extension	151
3.6.30. Zknh ISA Extension	151
3.6.31. Zks ISA Extension	151
3.6.32. Zksed ISA Extension	152
3.6.33. Zksh ISA Extension	152
3.6.34. Zkt ISA Extension	152
3.6.35. Zmmul - ISA Extension	153
3.6.36. Zxcfu ISA Extension	153
3.6.37. Smpmp ISA Extension	154
3.6.38. Sdext ISA Extension	154
3.6.39. Sdtrig ISA Extension	155
3.7. Custom Functions Unit (CFU)	156
3.7.1. CFU Instruction Formats	156
3.7.2. Using Custom Instructions in Software	157
3.7.3. Custom Instructions Hardware	158
3.8. Control and Status Registers (CSRs)	160
3.8.1. Floating-Point CSRs	163
3.8.2. Machine Trap Setup CSRs	165

3.8.3. Machine Trap Handling CSRs	169
3.8.4. Machine Configuration CSRs	173
3.8.5. Machine Physical Memory Protection CSRs	174
3.8.6. (Machine) Counter and Timer CSRs	176
3.8.7. Hardware Performance Monitors (HPM) CSRs	178
3.8.8. Machine Counter Setup CSRs	181
3.8.9. Machine Information CSRs	182
3.8.10. NEORV32-Specific CSRs	184
3.9. Traps, Exceptions and Interrupts	187
3.9.1. Memory Access Exceptions	187
3.9.2. Nested Interrupts	188
3.9.3. Custom Fast Interrupt Request Lines	188
3.9.4. NEORV32 Trap Listing	188
3.10. Dual-Core Configuration	192
3.10.1. Dual-Core Boot	193
4. Software Framework	195
4.1. Compiler Toolchain	195
4.2. Core Libraries	196
4.3. System View Description File (SVD)	198
4.4. Application Makefile	199
4.4.1. Makefile Targets	200
4.4.2. Default Compiler Flags	201
4.5. Linker Script	203
4.5.1. RAM Layout	204
4.5.2. ROM Layout	205
4.6. C Standard Library	206
4.7. Start-Up Code (crt0)	207
4.8. Executable Image Formats	208
4.9. Bootloader	210
4.9.1. Bootloader Console	210
4.9.2. Auto Boot Sequence	212
4.9.3. Uploading an Executable	213
4.9.4. Programming an SPI (/TWI) Flash	213
4.9.5. Booting from SD Card	214
4.9.6. Customizing the Internal Bootloader	214
4.9.7. Bootloader Error Codes	216
4.10. NEORV32 Runtime Environment	218
4.10.1. RTE Operation	218
4.10.2. Using the RTE	218

4.10.3. Default RTE Trap Handlers	220
4.10.4. Application Context Handling.....	221
5. On-Chip Debugger (OCD)	222
5.1. openOCD	225
5.2. Semihosting	226
5.3. Debug Transport Module (DTM).....	228
5.4. Debug Module (DM).....	230
5.4.1. DM Registers	231
5.4.2. DM CPU Access	236
5.5. Debug Authentication	239
5.5.1. Default Authentication Mechanism	239
5.6. CPU Debug Mode	241
5.6.1. CPU Debug Mode CSRs.....	242
5.7. Trigger Module	245
5.7.1. Trigger Module CSRs	245
6. Legal	248
About	248
License	248
Proprietary Notice	250
Disclaimer	250
Limitation of Liability for External Links	250
Acknowledgments	250

Chapter 1. Overview

The NEORV32 RISC-V Processor is an open-source RISC-V compatible processor system that is intended as **ready-to-go** auxiliary processor within a larger SoC designs or as stand-alone custom / customizable microcontroller.

The system is highly configurable and provides optional common peripherals like embedded memories, timers, serial interfaces, general purpose IO ports and an external bus interface to connect custom IP like memories, NoCs and other peripherals. On-line and in-system debugging is supported by an OpenOCD/gdb compatible on-chip debugger accessible via JTAG.

Special focus is paid on **execution safety** to provide defined and predictable behavior at any time. Therefore, the CPU ensures that all memory access are acknowledged and no invalid/malformed instructions are executed. Whenever an unexpected situation occurs, the application code is informed via hardware exceptions.

The software framework of the processor comes with application makefiles, software libraries for all CPU and processor features, a bootloader, a runtime environment and several example programs - including a port of the CoreMark MCU benchmark and the official RISC-V architecture test suite. RISC-V GCC is used as default toolchain.

Check out the processor's [online User Guide](#) that provides hands-on tutorials to get you started.

Structure

2. [NEORV32 Processor \(SoC\)](#)
3. [NEORV32 Central Processing Unit \(CPU\)](#)
4. [Software Framework](#)
5. [On-Chip Debugger \(OCD\)](#)
6. [Legal](#)

1.1. Project Key Features

Project

- all-in-one package: **CPU + SoC + Software Framework & Tooling**
- completely described in behavioral, platform-independent VHDL - no vendor- or technology-specific primitives, attributes, macros, libraries, etc. are used at all
- all-Verilog "version" available (auto-generated by GHDL)
- extensive configuration options for adapting the processor to the requirements of the application
- highly extensible hardware - on CPU, SoC and system level
- aims to be as small as possible while being as RISC-V-compliant as possible - with a reasonable area-vs-performance trade-off
- FPGA friendly (e.g. all internal memories can be mapped to block RAM - including the register file)
- optimized for high clock frequencies to ease timing closure and integration
- from zero to "*hello world!*" - completely open source and documented
- easy to use even for FPGA/RISC-V starters – intended to *work out of the box*

NEORV32 CPU (the core)

- 32-bit RISC-V CPU
- fully compatible to the RISC-V ISA specs. - checked by the [official RISCOF architecture tests](#)
- base ISA + privileged ISA + several optional standard and custom ISA extensions
- option to add user-defined RISC-V instructions as custom ISA extension
- rich set of customization options (ISA extensions, design goal: performance / area / energy, tuning options, ...)
- [Full Virtualization](#) capabilities to increase execution safety
- official RISC-V open source architecture ID

NEORV32 Processor (the SoC)

- highly-configurable full-scale microcontroller-like processor system
- based on the NEORV32 CPU
- optional standard serial interfaces (UART, TWI, SPI (host and device), 1-Wire)
- optional timers and counters (watchdog, system timer)
- optional general purpose IO and PWM; a native NeoPixel(c)-compatible smart LED interface
- optional embedded memories and caches for data, instructions and bootloader
- optional external memory interface for custom connectivity

- optional DMA controller for CPU-independent data transfers
- on-chip debugger compatible with OpenOCD and GDB including hardware trigger module and optional authentication

Software framework

- written in C and based on GCC
- internal bootloader with serial user interface (via UART)
- core libraries and HAL for high-level usage of the provided functions and peripherals
- processor-specific runtime environment and several example programs
- Doxygen-based documentation of the software framework; a deployed version is available at <https://stnolting.github.io/nerv32/sw/files.html>
- Ada support at <https://github.com/GNAT-Academic-Program/nerv32-hal>

OS Support

- FreeRTOS port: <https://github.com/stnolting/nerv32-freertos>
- Upstream Zephyr support: <https://docs.zephyrproject.org/latest/boards/others/nerv32/doc/index.html>
- MicroPython port: <https://github.com/stnolting/nerv32-micropython>

Extensibility and Customization

The NEORV32 processor is designed to ease customization and extensibility and provides several options for adding application-specific custom hardware modules and accelerators. The three most common options for adding custom on-chip modules are listed below.

- [Processor-External Bus Interface \(XBUS\)](#) to attach processor-external IP modules (memories and peripherals)
- [Custom Functions Subsystem \(CFS\)](#) for tightly-coupled processor-internal co-processors
- [Custom Functions Unit \(CFU\)](#) for custom RISC-V instructions



A more detailed comparison of the extension/customization options can be found in section [Adding Custom Hardware Modules](#) of the user guide.

1.2. Project Folder Structure

The root directory of the repository is considered the NEORV32 base or home folder (i.e. [neorv32/](#)).

Folder Structure

neorv32	- Project home folder
docs	- Documentation
datasheet	- AsciiDoc sources for the NEORV32 data sheet
figures	- Figures and logos
references	- Data sheets and RISC-V specs
userguide	- AsciiDoc sources for the NEORV32 user guide
rtl	- HDL sources
core	- Core sources of the CPU & SoC
system_integration	- System wrappers and bridges for advanced SoC
connectivity	
test_setups	- Minimal test setup "SoCs" used in the User Guide
verilog	- Scripts and examples for converting NEORV32 to Verilog
sim	- Simulation files
sw	- Software framework
bootloader	- Sources of the processor-internal bootloader
common	- Linker script, crt0.S start-up code and central
makefile	
example	- Example programs for the core and the SoC modules
eclipse	- Pre-configured Eclipse IDE project
...	- Several example programs
image_gen	- Helper program to generate executables & memory images
lib	- NEORV32 core library
ocd_firmware	- Firmware for the on-chip debugger "park loop"
openocd	- OpenOCD configuration files
svd	- Processor system view description file (CMSIS-SVD)

1.3. VHDL File Hierarchy

All required VHDL hardware source files are located in the project's `rtl/core` folder.



VHDL Library
All core VHDL files from the list below have to be assigned to a **new library** named `neorv32`.



Compilation Order

See section [File-List Files](#) for more information.



Replacing Modules for Customization or Optimization

Any module of the core can be replaced for customization purpose. For example, the default IMEM and DMEM modules or the CPU's register file can be replaced by technology-specific primitives to optimize energy, speed and area utilization.

RTL File List (in alphabetical order)

```
rtl/core
|
├── neorv32_application_image.vhd - IMEM application initialization image (package)
├── neorv32_boot_rom.vhd           - Bootloader ROM
├── neorv32_bootloader_image.vhd   - Bootloader ROM memory image (package)
├── neorv32_bus.vhd               - SoC bus infrastructure modules
├── neorv32_cache.vhd             - Generic cache module
├── neorv32_cfs.vhd               - Custom functions subsystem
├── neorv32_clint.vhd             - Core local interruptor
├── neorv32_cpu.vhd               - NEORV32 CPU TOP ENTITY
├── neorv32_cpu_alu.vhd           - Arithmetic/logic unit
├── neorv32_cpu_control.vhd       - CPU control, exception system and CSRs
├── neorv32_cpu_counters.vhd      - Hardware counters (Zicntr & Zihpm ext.)
├── neorv32_cpu_cp_bitmanip.vhd   - Bit-manipulation co-processor (B ext.)
├── neorv32_cpu_cp_cfu.vhd        - Custom instructions co-processor (Zxfcuh ext.)
├── neorv32_cpu_cp_cond.vhd       - Integer conditional co-processor (Zicond ext.)
├── neorv32_cpu_cp_crypto.vhd     - Scalar cryptography co-processor (Zk*/Zbk* ext.)
├── neorv32_cpu_cp_fpu.vhd        - Floating-point co-processor (Zfinx ext.)
├── neorv32_cpu_cp_muldiv.vhd     - Mul/Div co-processor (M ext.)
├── neorv32_cpu_cp_shifter.vhd    - Bit-shift co-processor (base ISA)
├── neorv32_cpu_decompressor.vhd   - Compressed instructions decoder (C ext.)
├── neorv32_cpu_frontend.vhd       - Instruction fetch and issue
├── neorv32_cpu_hwtrig.vhd         - Hardware trigger module (Sdtrig ext.)
├── neorv32_cpu_lsu.vhd            - Load/store unit
├── neorv32_cpu_pmp.vhd            - Physical memory protection unit (Smpmp ext.)
├── neorv32_cpu_regfile.vhd        - Data register file
├── neorv32_cpu_trace.vhd          - Trace generator
└── neorv32_debug_auth.vhd         - On-chip debugger: authentication module
```

└── neorv32_debug_dm.vhd	- On-chip debugger: debug module
└── neorv32_debug_dtm.vhd	- On-chip debugger: debug transfer module
└── neorv32_dma.vhd	- Direct memory access controller
└── neorv32_dmem.vhd	- Processor-internal data memory
└── neorv32_gpio.vhd	- General purpose input/output port unit
└── neorv32_gptmr.vhd	- General purpose 32-bit timer
└── neorv32_imem.vhd	- Processor-internal instruction memory
└── neorv32_neoled.vhd	- NeoPixel (TM) compatible smart LED interface
└── neorv32_onewire.vhd	- One-Wire serial interface controller
└── neorv32_package.vhd	- Main VHDL package file
└── neorv32_prim.vhd	- Generic RTL primitives
└── neorv32_pwm.vhd	- Pulse-width modulation controller
└── neorv32_sdi.vhd	- Serial data interface controller (SPI device)
└── neorv32_slink.vhd	- Stream link interface
└── neorv32_spi.vhd	- Serial peripheral interface controller (SPI host)
└── neorv32_sys.vhd	- System infrastructure modules
└── neorv32_sysinfo.vhd	- System configuration information memory
└── neorv32_top.vhd	- NEORV32 PROCESSOR/SOC TOP ENTITY
└── neorv32_tracer.vhd	- Instruction trace buffer
└── neorv32_trng.vhd	- True random number generator
└── neorv32_twd.vhd	- Two wire serial device controller
└── neorv32_twi.vhd	- Two wire serial interface controller
└── neorv32_uart.vhd	- Universal async. receiver/transmitter
└── neorv32_wdt.vhd	- Watchdog timer
└── neorv32_xbus.vhd	- External bus interface gateway

1.3.1. File-List Files

Most of the RTL sources use **entity instantiation**. Hence, the RTL compile order might be relevant (depending on the synthesis/simulation tool). Therefore, two file-list files are provided in the `rtl` folder that list all required HDL files for the CPU core and for the entire processor and also represent their recommended compile order. These file-list files can be consumed by EDA tools to simplify project setup.

- `file_list_cpu.f` - HDL files and compile order for the CPU core; top module: `neorv32_cpu`
- `file_list_soc.f` - HDL files and compile order for the entire processor/SoC; top module: `neorv32_top`

A simple bash script `generate_file_lists.sh` is provided for regenerating the file-lists (using GHDL's `elaborate` command). This script can also be invoked using the default application makefile (see [Makefile Targets](#)).

By default, the file-list files include a **placeholder** in the path of each included hardware source file. These placeholders need to be replaced by the actual path before being used. Example:

- default: `NEORV32_RTL_PATH_PLACEHOLDER/core/neorv32_package.vhd`
- adjusted: `path/to/neorv32/rtl/core/neorv32_package.vhd`

Listing 1. Example: Processing the File-List Files in a Makefile

```
NEORV32_HOME = path/to/neorv32 ①  
NEORV32_SOC_FILE = $(shell cat $(NEORV32_HOME)/rtl/file_list_soc.f) ②  
NEORV32_SOC_SRCS = $(subst NEORV32_RTL_PATH_PLACEHOLDER, $(NEORV32_HOME)/rtl,  
$(NEORV32_SOC_FILE)) ③
```

- ① Path to the NEORV32 home folder (i.e. the root folder of the GitHub repository).
- ② Load the content of the `file_list_soc.f` file-list into a new variable `NEORV32_SOC_FILE`.
- ③ Substitute the file-list file's path placeholder “`NEORV32_RTL_PATH_PLACEHOLDER`” by the actual path.

Listing 2. Example: Processing the File-List Files in a TCL Script

```
set file_list_file [read [open "$neorv32_home/rtl/file_list_soc.f" r]]  
set file_list [string map [list "NEORV32_RTL_PATH_PLACEHOLDER" "$neorv32_home/rtl"]  
$file_list_file]  
puts "NEORV32 source files:"  
puts $file_list
```

1.4. VHDL Coding Style

- The entire processor, including the CPU core, is written in platform-/technology-independent VHDL. The code makes minimal use of VHDL 2008 features in order to remain compatible with older EDA tools.
- A specific VHDL library `neorv32` is used for all sources.
- All registers / flip-flops provide an *asynchronous* reset (see [Processor Reset](#)).
- The entire setup uses a single clock domain. External "clock" signals are synchronized into this clock domain using 2-stage shift registers.
- A single package/library file (`neorv32_package.vhd`) is used to provide global definitions and auxiliary functions. The user-defined configuration is done entirely via the generics of the top entity.
- Internally, all generics are checked to ensure correct configuration. Asserts are used to inform the user about the actual processor configuration and possible invalid settings.
- The code uses entity instantiation for all internal modules. However, if multiple submodules are specified within the same source file, component instantiation is used for them. * When instantiating the top-level processor module (`neorv32_top.vhd`) in a user-defined design, either entity instantiation or component instantiation can be used, as the NEORV32 package file/library file already provides a corresponding component declaration.



Verilog "Version"

Scripts and instructions for an automatic to-Verilog conversion can be found in [rtl/verilog](#). See [UG: NEORV32 in Verilog](#) for more information.

1.5. Performance

Area Utilization

The NEORV32 processor is optimized for minimal size. However, the actual size (silicon area or FPGA resources) depends on the specific configuration. For example, an RTOS-capable setup based on a `rv32imc_Zicsr_Zicntr` CPU configuration including peripheral and memories requires about 2300 LUTs and 1000 FFs and can run at up to 130 MHz (implementation results for an Altera Cyclone IV E `EP4CE22F17C6` FPGA).

NEORV32 Setups



The processor has been successfully ported to AMD, Altera, Lattice, Microchip, Gowin, Cologne Chip and NanoXplore FPGAs. Some pre-configured example setup are available online: <https://github.com/stnolting/neorv32-setups>

Processing Speed

The computational performance of the NEORV32 is evaluated using [Core Mark CPU benchmark](#). The according sources can be found in the `sw/example/coremark` folder.

Table 1. CoreMark results (-O3 optimization)

CPU	CoreMark Score	CoreMark s/MHz	Average
			CPI
<i>small</i> (<code>rv32i_Zicsr_Zifencei</code>)	33.89	0.3389	4.04
<i>medium</i> (<code>rv32imc_Zicsr_Zifencei</code>)	62.50	0.6250	5.34
<i>performance</i> (<code>rv32imc_Zicsr_Zifencei</code> + perf. options)	95.23	0.9523	3.54

Performance can be further improved by enabling additional ISA extensions. Note that the compiler also needs to support these extensions (make sure to use a `multilib` GCC configuration for optimal code results).

The NEORV32 CPU is based on a multi-cycle architecture. Each instruction is executed in a sequence of several consecutive micro operations. The average CPI (cycles per instruction) depends on the instruction mix of a specific applications and also on the available CPU extensions, tuning options and memory latency.

Chapter 2. NEORV32 Processor (SoC)

The NEORV32 Processor is build around the [NEORV32 Central Processing Unit \(CPU\)](#). Together with common peripheral interfaces and embedded memories it provides a RISC-V-based full-scale microcontroller-like SoC platform.

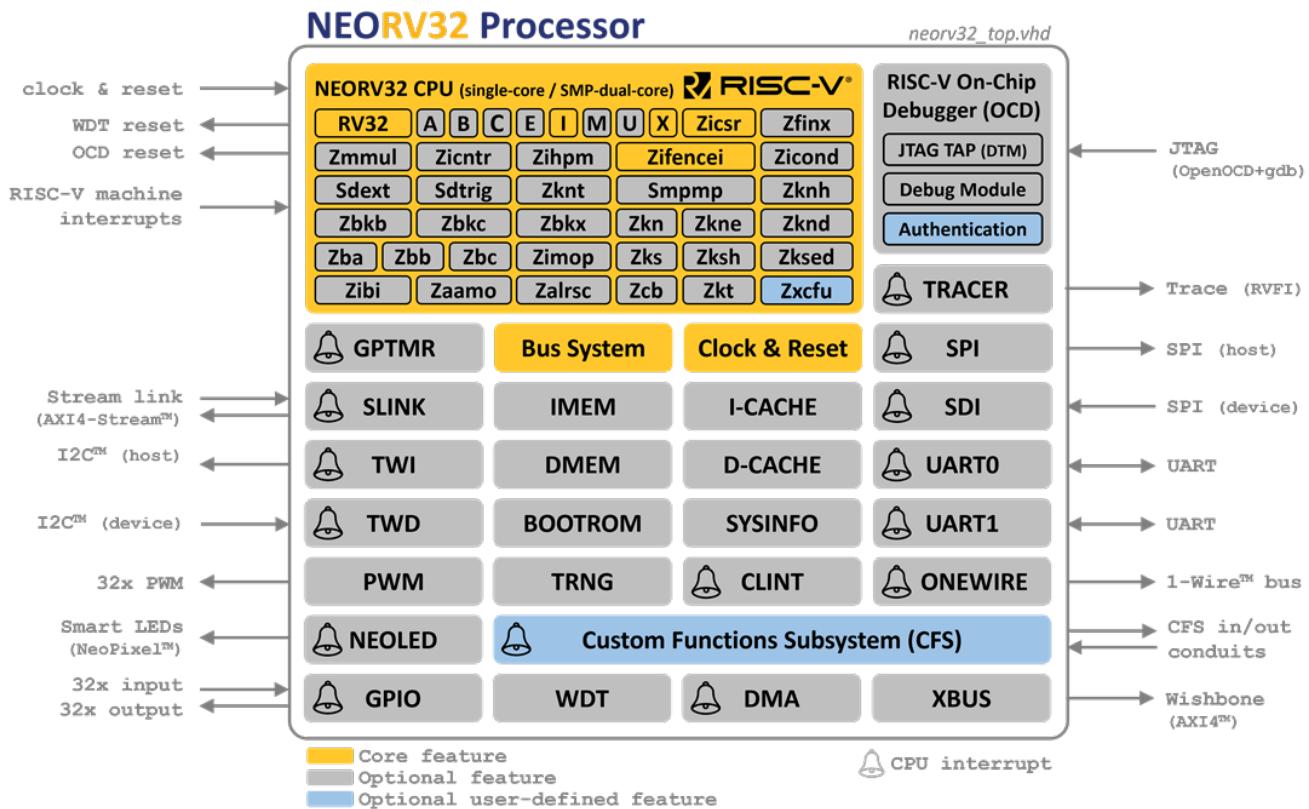


Figure 1. The NEORV32 Processor (Block Diagram)

Section Structure

- Processor Top Entity - Signals and Processor Top Entity - Generics
- Processor Clocking and Processor Reset
- Processor Interrupts
- Address Space and Boot Configuration
- Processor-Internal Modules

Key Features

- optional* SMP Dual-Core Configuration
- optional* processor-internal data and instruction memories (**DMEM/IMEM**)
- optional* caches (**I-CACHE, D-CACHE**)
- optional* internal bootloader (**BOOTROM**) with UART console & SPI/TWI flash and SD card boot options

- *optional* RISC-V-compatible core local interruptor (**CLINT**)
- *optional* two independent universal asynchronous receivers and transmitters (**UART0**, **UART1**) with optional hardware flow control (RTS/CTS)
- *optional* serial peripheral interface host controller (**SPI**) with 8 dedicated CS lines
- *optional* 8-bit serial data device interface (**SDI**)
- *optional* two-wire serial interface controller (**TWI**), compatible to the I²C standard
- *optional* two-wire serial device controller (**TWD**), compatible to the I²C standard
- *optional* general purpose parallel IO port (**GPIO**), 32 inputs (interrupt capable), 32 outputs
- *optional* 32-bit external bus interface, Wishbone-compatible (**XBUS**), AXI4-compatible bridge available
- *optional* watchdog timer (**WDT**)
- *optional* PWM controller with up to 32 individual channels (**PWM**)
- *optional* ring-oscillator-based true random number generator (**TRNG**)
- *optional* custom functions subsystem for custom co-processor extensions (**CFS**)
- *optional* NeoPixel™/WS2812-compatible smart LED interface (**NEOLED**)
- *optional* general purpose 32-bit timer (**GPTMR**) with up to 16 individual timer slices
- *optional* 1-wire serial interface controller (**ONEWIRE**), compatible to the 1-wire standard
- *optional* autonomous direct memory access controller (**DMA**)
- *optional* stream link interface (**SLINK**), AXI4-Stream-compatible
- *optional* on-chip debugger with JTAG TAP (**OCD**), optional authentication and hardware breakpoint
- *optional* execution trace buffer to debug program flow (**TRACER**) via branch tracing
- *optional* RVFI-compatible **Execution Trace Port** for advanced debugging, profiling and verification
- *optional* system configuration information memory to determine hardware configuration via software (**SYSINFO**)

2.1. Processor Top Entity - Signals

The following table shows all interface signals of the processor top entity (`rtl/core/neorv32_top.vhd`). All signals are of type `std_ulogic` or `std_logic_vector`, respectively (except for the `trace_*` ports).

Default Values of Inputs



All *optional* input signals provide default values in case they are not explicitly assigned during instantiation. The weak driver strengths of VHDL ('L' and 'H') are used to model a pull-down or pull-up resistor.

Variable-Sized Ports



Some peripherals allow to configure the number of channels to-be-implemented by a generic (for example the number of PWM channels). The according input/output signals have a fixed sized regardless of the actually configured amount of channels. If less than the maximum number of channels is configured, only the LSB-aligned channels are used: in case of an *input port* the remaining bits/channels are left unconnected; in case of an *output port* the remaining bits/channels are hardwired to zero.

Tri-State Interfaces



Some interfaces (like the TWI, the TWD and the 1-Wire bus) require explicit tri-state drivers in the final top module.

Input/Output Registers



By default all output signals are driven by register and all input signals are synchronized into the processor's clock domain also using registers. However, for ASIC implementations it is recommended to add another register state to all inputs and output so the synthesis tool can insert an explicit IO (boundary) scan chain.

Table 2. NEORV32 Processor Signal List

Name	Width	Direction	Default	Description
Global Control (Processor Clocking and Processor Reset)				
<code>clk_i</code>	1	in	none	global clock line, all registers triggering on rising edge
<code>rstn_i</code>	1	in	none	global reset, asynchronous, low-active
<code>rstn_ocd_o</code>	1	out	none	On-Chip Debugger (OCD) reset output, synchronous, low-active
<code>rstn_wdt_o</code>	1	out	none	Watchdog Timer (WDT) reset output, synchronous, low-active
Execution Trace Port				

Name	Width	Direction	Default	Description
trace_cpu0_o	*	out	-	CPU core 0 trace port (type <code>trace_port_t</code>).
trace_cpu1_o	*	out	-	CPU core 1 trace port (type <code>trace_port_t</code>).
JTAG Access Port for On-Chip Debugger (OCD)				
jtag_tck_i	1	in	'L'	serial clock
jtag_tdi_i	1	in	'L'	serial data input
jtag_tdo_o	1	out	-	serial data output
jtag_tms_i	1	in	'L'	mode select
Processor-External Bus Interface (XBUS)				
xbus_adr_o	32	out	-	destination address
xbus_dat_o	32	out	-	read data
xbus_cti_o	3	out	-	cycle type
xbus_tag_o	3	out	-	access tag
xbus_we_o	1	out	-	write enable ('0' = read transfer)
xbus_sel_o	4	out	-	byte enable
xbus_stb_o	1	out	-	strobe
xbus_cyc_o	1	out	-	valid cycle
xbus_dat_i	32	in	'L'	write data
xbus_ack_i	1	in	'L'	transfer acknowledge
xbus_err_i	1	in	'L'	transfer error
Stream Link Interface (SLINK)				
slink_rx_dat_i	32	in	'L'	RX data
slink_rx_src_i	4	in	'L'	RX source routing information
slink_rx_val_i	1	in	'L'	RX data valid
slink_rx_lst_i	1	in	'L'	RX last element of stream
slink_rx_rdy_o	1	out	-	RX ready to receive
slink_tx_dat_o	32	out	-	TX data
slink_tx_dst_o	4	out	-	TX destination routing information
slink_tx_val_o	1	out	-	TX data valid
slink_tx_lst_o	1	out	-	TX last element of stream
slink_tx_rdy_i	1	in	'L'	TX allowed to send
General Purpose Input and Output Port (GPIO)				
gpio_o	32	out	-	general purpose parallel output

Name	Width	Direction	Default	Description
gpio_i	32	in	'L'	general purpose parallel input (interrupt-capable)
Primary Universal Asynchronous Receiver and Transmitter (UART0)				
uart0_txd_o	1	out	-	serial transmitter
uart0_rxd_i	1	in	'L'	serial receiver
uart0_rtsn_o	1	out	-	RX ready to receive new char
uart0_ctsn_i	1	in	'L'	TX allowed to start sending, low-active
Secondary Universal Asynchronous Receiver and Transmitter (UART1)				
uart1_txd_o	1	out	-	serial transmitter
uart1_rxd_i	1	in	'L'	serial receiver
uart1_rtsn_o	1	out	-	RX ready to receive new char
uart1_ctsn_i	1	in	'L'	TX allowed to start sending, low-active
Serial Peripheral Interface Controller (SPI)				
spi_clk_o	1	out	-	controller clock line
spi_dat_o	1	out	-	serial data output
spi_dat_i	1	in	'L'	serial data input
spi_csn_o	8	out	-	select (low-active)
Serial Data Interface Controller (SDI)				
sdi_clk_i	1	in	'L'	controller clock line
sdi_dat_o	1	out	-	serial data output
sdi_dat_i	1	in	'L'	serial data input
sdi_csn_i	1	in	'H'	chip select, low-active
Two-Wire Serial Interface Controller (TWI)				
twi_sda_i	1	in	'H'	serial data line sense input
twi_sda_o	1	out	-	serial data line output (pull low only)
twi_scl_i	1	in	'H'	serial clock line sense input
twi_scl_o	1	out	-	serial clock line output (pull low only)
Two-Wire Serial Device Controller (TWD)				
twd_sda_i	1	in	'H'	serial data line sense input
twd_sda_o	1	out	-	serial data line output (pull low only)
twd_scl_i	1	in	'H'	serial clock line sense input
twd_scl_o	1	out	-	serial clock line output (pull low only)

Name	Width	Direction	Default	Description
One-Wire Serial Interface Controller (ONEWIRE)				
onewire_i	1	in	'H'	1-wire bus sense input
onewire_o	1	out	-	1-wire bus output (pull low only)
Pulse-Width Modulation Controller (PWM)				
pwm_o	16	out	-	pulse-width modulated channels
Custom Functions Subsystem (CFS)				
cfs_in_i	256	in	'L'	custom CFS input signal conduit
cfs_out_o	256	out	-	custom CFS output signal conduit
Smart LED Interface (NEOLED)				
neoled_o	1	out	-	asynchronous serial data output
Core-Local Interruptor (CLINT)				
mtime_time_o	64	out	-	CLINT.MTICK system time output
RISC-V Machine-Mode Processor Interrupts				
mtime_irq_i	1	in	'L'	machine timer interrupt (RISC-V), high-level-active; for chip-internal usage only
msw_irq_i	1	in	'L'	machine software interrupt (RISC-V), high-level-active; for chip-internal usage only
mext_irq_i	1	in	'L'	machine external interrupt (RISC-V), high-level-active; for chip-internal usage only

2.2. Processor Top Entity - Generics

This section lists all configuration generics of the NEORV32 processor top entity ([rtl/neorv32_top.vhd](#)). These generics allow to configure the system according to your needs. The generics are used to control implementation of certain CPU extensions and peripheral modules and even allow to optimize the system for certain design goals like minimal area or maximum performance.



Default Values

All configuration generics provide default values in case they are not explicitly assigned during instantiation. By default, all configuration options are **disabled**.



Software Discovery of Configuration

Software can determine the actual CPU configuration via the `misa` and `mxisa` CSRs. The Soc/Processor and can be determined via the `SYSINFO` memory-mapped registers.



Excluded Modules and Extensions

If optional modules (like CPU extensions or peripheral devices) are not enabled the according hardware will not be synthesized at all. Hence, the disabled modules do not increase area and power requirements and do not impact timing.



Table Abbreviations

The generic type “`suv(x:y)`” is an abbreviation for “`std_ulogic_vector(x downto y)`”.

Table 3. NEORV32 Processor Generic List

Name	Type	Default	Description
General			
<code>CLOCK_FREQUENCY</code>	natural	0	The clock frequency of the processor's <code>clk_i</code> input port in Hertz (Hz). See Processor Clocking .
<code>TRACE_PORT_EN</code>	boolean	false	Enable external CPU execution Execution Trace Port .
<code>DUAL_CORE_EN</code>	boolean	false	Enable the SMP Dual-Core Configuration .
Boot Configuration			
<code>BOOT_MODE_SELECT</code>	natural	0	Boot mode select; see Boot Configuration .
<code>BOOT_ADDR_CUSTOM</code>	<code>suv(31:0)</code>	x"00000000"	Custom CPU boot address (available if " <code>BOOT_MODE_SELECT</code> = 1").
On-Chip Debugger (OCD)			
<code>OCD_EN</code>	boolean	false	Implement the on-chip debugger and the CPU debug mode (Sdext ISA Extension).

Name	Type	Default	Description
OCD_NUM_HW_TRIGGER_S	natural	0	Number of implemented HW triggers (Trigger Module / Sdtrig ISA Extension) for hardware break-/watchpoints (0..16).
OCD_AUTHENTICATION	boolean	false	Implement Debug Authentication module.
OCD_JEDEC_ID	suv(10:0)	"000000000000"	JEDEC ID; continuation codes plus vendor ID (passed to the JTAG Debug Transport Module (DTM)).
CPU Instruction Sets and Extensions			
RISCV_ISA_C	boolean	false	Enable C ISA Extension (compressed instructions).
RISCV_ISA_E	boolean	false	Enable E ISA Extension (reduced register file size).
RISCV_ISA_M	boolean	false	Enable M ISA Extension (hardware-based integer multiplication and division).
RISCV_ISA_U	boolean	false	Enable U ISA Extension (less-privileged user mode).
RISCV_ISA_Zaamo	boolean	false	Enable Zaamo ISA Extension (atomic read-modify-write operations).
RISCV_ISA_Zalrsc	boolean	false	Enable Zalrsc ISA Extension (atomic reservation-set operations).
RISCV_ISA_Zba	boolean	false	Enable Zba ISA Extension (shifted-add bit-manipulation instructions).
RISCV_ISA_Zbb	boolean	false	Enable Zbb ISA Extension (basic bit-manipulation instructions).
RISCV_ISA_Zbkb	boolean	false	Enable Zbkb ISA Extension (scalar cryptography bit manipulation instructions).
RISCV_ISA_Zbkc	boolean	false	Enable Zbkc ISA Extension (scalar cryptography carry-less multiplication instructions).
RISCV_ISA_Zbkx	boolean	false	Enable Zbkx ISA Extension (scalar cryptography crossbar permutations).
RISCV_ISA_Zbs	boolean	false	Enable Zbs ISA Extension (single-bit bit-manipulation instructions).
RISCV_ISA_Zcb	boolean	false	Enable Zcb ISA Extension (additional code size reduction instruction; builds upon C).
RISCV_ISA_Zfinx	boolean	false	Enable Zfinx ISA Extension (single-precision floating-point unit).
RISCV_ISA_Zibi	boolean	false	Enable Zibi ISA Extension (CPU base counters).
RISCV_ISA_Zicntr	boolean	false	Enable Zicntr ISA Extension (CPU base counters).

Name	Type	Default	Description
RISCV_ISA_Zicond	boolean	false	Enable Zicond ISA Extension (integer conditional instructions).
RISCV_ISA_Zihpm	boolean	false	Enable Zihpm ISA Extension (hardware performance monitors).
RISCV_ISA_Zimop	boolean	false	Enable Zimop ISA Extension (may-be-operations).
RISCV_ISA_Zknd	boolean	false	Enable Zknd ISA Extension (scalar cryptography NIST AES decryption instructions).
RISCV_ISA_Zkne	boolean	false	Enable Zkne ISA Extension (scalar cryptography NIST AES encryption instructions).
RISCV_ISA_Zknh	boolean	false	Enable Zknh ISA Extension (scalar cryptography NIST hash instructions).
RISCV_ISA_Zksed	boolean	false	Enable Zksed ISA Extension (scalar cryptography ShangMi block cyphers).
RISCV_ISA_Zksh	boolean	false	Enable Zksh ISA Extension (scalar cryptography ShangMi hash functions).
RISCV_ISA_Zmmul	boolean	false	Enable Zmmul - ISA Extension (hardware-based integer multiplication).
RISCV_ISA_Zxfcu	boolean	false	Enable NEORV32-specific Zxfcu ISA Extension (custom RISC-V instructions).
CPU Tuning Options			
CPU_CONSTT_BR_EN	boolean	false	Implement constant-time branches (same execution times for taken and not-taken branches).
CPU_FAST_MUL_EN	boolean	false	Implement fast but large full-parallel multipliers (trying to infer DSP blocks); see section CPU Arithmetic Logic Unit .
CPU_FAST_SHIFT_EN	boolean	false	Implement fast but large full-parallel barrel shifters; see section CPU Arithmetic Logic Unit .
CPU_RF_HW_RST_EN	boolean	false	Implement full hardware reset for register file (use individual FFs instead of BRAM); see section CPU Register File .
Physical Memory Protection (Smpmp ISA Extension)			
PMP_NUM_REGIONS	natural	0	Number of implemented PMP regions (0..16).
PMP_MIN_GRANULARITY	natural	4	Minimal region granularity in bytes. Has to be a power of two, min 4.
PMP_TOR_MODE_EN	boolean	false	Implement support for top-of-region (TOR) mode.
PMP_NAP_MODE_EN	boolean	false	Implement support for naturally-aligned power-of-two (NAPOT & NA4) modes.

Name	Type	Default	Description
Hardware Performance Monitors (Zihpm ISA Extension)			
HPM_NUM_CNTS	natural	0	Number of implemented hardware performance monitor counters (0..13).
HPM_CNT_WIDTH	natural	40	Total LSB-aligned size of each HPM counter. Min 0, max 64.
Internal Instruction Memory (IMEM)			
IMEM_EN	boolean	false	Implement the processor-internal instruction memory.
IMEM_SIZE	natural	16*1024	Size in bytes of the processor internal instruction memory (use a power of 2).
IMEM_OUTREG_EN	boolean	false	Add IMEM output register stage (improves mapping/timing at the expense of latency).
Internal Data Memory (DMEM)			
DMEM_EN	boolean	false	Implement the processor-internal data memory.
DMEM_SIZE	natural	8*1024	Size in bytes of the processor-internal data memory (use a power of 2).
DMEM_OUTREG_EN	boolean	false	Add DMEM output register stage (improves mapping/timing at the expense of latency).
CPU Caches (instruction-cache & data-cache)			
ICACHE_EN	boolean	false	Implement the instruction cache (I\$).
ICACHE_NUM_BLOCKS	natural	4	Number of blocks ("lines") Has to be a power of two.
DCACHE_EN	boolean	false	Implement the data cache (D\$)
DCACHE_NUM_BLOCKS	natural	4	Number of blocks ("lines"). Has to be a power of two.
CACHE_BLOCK_SIZE	natural	64	Size in bytes of each block (I\$ and D\$). Has to be a power of two, min 8.
CACHE_BURSTS_EN	boolean	true	Enable burst transfers for cache updates.
Processor-External Bus Interface (XBUS) (Wishbone / AXI4-Compatible Bridging)			
XBUS_EN	boolean	false	Implement the external bus interface.
XBUS_TIMEOUT	natural	2048	Number of clock cycles after which an unacknowledged external bus access will auto-terminate (0 = disabled).
XBUS_REGSTAGE_EN	boolean	false	Implement XBUS register stages to ease timing closure.
Peripheral/IO Modules			

Name	Type	Default	Description
IO_DISABLE_SYSINFO	boolean	false	Disable System Configuration Information Memory (SYSINFO) module; not recommended - for advanced users only!
IO_GPIO_NUM	natural	0	Number of general purpose input/output pairs of the General Purpose Input and Output Port (GPIO) , max 32.
IO_CLINT_EN	boolean	false	Implement the Core-Local Interruptor (CLINT) .
IO_UART0_EN	boolean	false	Implement the Primary Universal Asynchronous Receiver and Transmitter (UART0) .
IO_UART0_RX_FIFO	natural	1	UART0 RX FIFO depth, has to be a power of two, minimum value is 1, max 32768.
IO_UART0_TX_FIFO	natural	1	UART0 TX FIFO depth, has to be a power of two, minimum value is 1, max 32768.
IO_UART1_EN	boolean	false	Implement the Secondary Universal Asynchronous Receiver and Transmitter (UART1) .
IO_UART1_RX_FIFO	natural	1	UART1 RX FIFO depth, has to be a power of two, minimum value is 1, max 32768.
IO_UART1_TX_FIFO	natural	1	UART1 TX FIFO depth, has to be a power of two, minimum value is 1, max 32768.
IO_SPI_EN	boolean	false	Implement the Serial Peripheral Interface Controller (SPI) .
IO_SPI_FIFO	natural	1	Depth of the Serial Peripheral Interface Controller (SPI) FIFO. Has to be a power of two, min 1, max 32768.
IO_SDI_EN	boolean	false	Implement the Serial Data Interface Controller (SDI) .
IO_SDI_FIFO	natural	1	Depth of the Serial Data Interface Controller (SDI) FIFO. Has to be a power of two, min 1, max 32768.
IO_TWI_EN	boolean	false	Implement the Two-Wire Serial Interface Controller (TWI) .
IO_TWI_FIFO	natural	1	Depth of the Two-Wire Serial Interface Controller (TWI) FIFO. Has to be a power of two, min 1, max 32768.
IO_TWD_EN	boolean	false	Implement the Two-Wire Serial Device Controller (TWD) .
IO_TWD_RX_FIFO	natural	1	Depth of the Two-Wire Serial Device Controller (TWD) RX FIFO. Has to be a power of two, min 1, max 32768.

Name	Type	Default	Description
IO_TWD_TX_FIFO	natural	1	Depth of the Two-Wire Serial Device Controller (TWD) TX FIFO. Has to be a power of two, min 1, max 32768.
IO_PWM_NUM	natural	0	Number of channels of the Pulse-Width Modulation Controller (PWM) to implement (0..32).
IO_WDT_EN	boolean	false	Implement the Watchdog Timer (WDT).
IO_TRNG_EN	boolean	false	Implement the True Random-Number Generator (TRNG).
IO_TRNG_FIFO	natural	1	Depth of the TRNG data FIFO. Has to be a power of two, min 1, max 32768.
IO_CFS_EN	boolean	false	Implement the Custom Functions Subsystem (CFS).
IO_NEOLED_EN	boolean	false	Implement the Smart LED Interface (NEOLED).
IO_NEOLED_TX_FIFO	natural	1	TX FIFO depth of the Smart LED Interface (NEOLED). Has to be a power of two, min 1, max 32768.
IO_GPTMR_NUM	natural	0	Number of individual General Purpose Timer (GPTMR) timer slices (0..16).
IO_ONEWIRE_EN	boolean	false	Implement the One-Wire Serial Interface Controller (ONEWIRE).
IO_ONEWIRE_FIFO	natural	1	Depth of the One-Wire Serial Interface Controller (ONEWIRE) FIFO. Has to be a power of two, min 1, max 32768.
IO_DMA_EN	boolean	false	Implement the Direct Memory Access Controller (DMA).
IO_DMA_DSC_FIFO	natural	4	Depth of the DMA transfer descriptor FIFO. Has to be a power of two, min 4, max 512.
IO_SLINK_EN	boolean	false	Implement the Stream Link Interface (SLINK) (AXI4-Stream-Compatible).
IO_SLINK_RX_FIFO	natural	1	SLINK RX FIFO depth, has to be a power of two, minimum value is 1, max 32768.
IO_SLINK_TX_FIFO	natural	1	SLINK TX FIFO depth, has to be a power of two, minimum value is 1, max 32768.
IO_TRACER_EN	boolean	false	Implement the Execution Trace Buffer (TRACER).
IO_TRACER_BUFFER	natural	1	Depth of the Execution Trace Buffer (TRACER). Has to be a power of two, min 1, max 32768.
IO_TRACER_SIMLOG_EN	boolean	false	Write full trace log to file (simulation-only).

2.3. Processor Clocking

The processor is implemented as fully-synchronous logic design using a single clock domain that is driven entirely by the top's `clk_i` signal. This clock signal is used by all internal registers and memories. All of them trigger on the **rising edge** of this clock signal. External "clocks" like the OCD's JTAG clock or the SDI's serial clock are synchronized into the processor's clock domain before being used as "general logic signal" (and not as a dedicated clock).

2.3.1. Peripheral Clocks

Many processor modules like the UARTs or the timers provide a programmable time base for operations. In order to simplify the hardware, the processor implements a global "clock generator" (`neorv32_sys.vhd`) that provides single-cycle *clock enables* for certain frequencies which are derived from the main clock. These clock enable signals are synchronous to the system's main clock. The processor modules can use these enables for sub-main-clock operations while still providing a single clock domain only.

In total, 8 sub-main-clock signals are available. All processor modules, which feature a time-based configuration, provide a programmable three-bit prescaler select in their control register to select one of the 8 available clocks. The mapping of the prescaler select bits to the according clock source is shown in the table below. Here, f represents the processor main clock from the top entity's `clk_i` signal.

Prescaler bits:	0b000	0b001	0b010	0b011	0b100	0b101	0b110	0b111
Resulting clock:	$f/2$	$f/4$	$f/8$	$f/64$	$f/128$	$f/1024$	$f/2048$	$f/4096$

2.4. Processor Reset

The NEORV32 processor includes a central reset sequencer (`neorv32_sys.vhd`) that handles all reset requests and controls the internal reset nets. The processor-wide reset (aka "system reset") can be triggered by any of the following sources:

- the asynchronous low-active `rstn_i` top entity input signal (External source)
- the [On-Chip Debugger \(OCD\)](#) (internal source)
- the [Watchdog Timer \(WDT\)](#) (internal source)

Processor Reset Signal



Make sure to connect the processor's reset signal `rstn_i` to a valid reset source (a button, the "locked" signal of a PLL, a dedicated reset controller, etc.).



Reset Cause

The actual reset cause can be determined via the [Watchdog Timer \(WDT\)](#).

If any of these sources triggers a reset, the internal system-wide reset will be active for at least 4 clock cycles ensuring a valid reset of the entire processor. This system reset is asserted *asynchronously* if triggered by the external `rstn_i` signal and is asserted *synchronously* if triggered by an internal reset source. However, the system reset is always de-asserted *synchronously* at the next rising clock edge.

Internally, **all registers** that are not meant for mapping to blockRAM (like the register file) do provide a dedicated and **low-active asynchronous** hardware reset. This asynchronous reset ensures that the entire processor logic is reset to a defined state even if the main clock is not operational yet.

2.5. Processor Interrupts

The NEORV32 Processor provides several interrupt request signals (IRQs) for custom platform use.



Trigger Type

All interrupt request lines are **level-triggered and high-active**. Once set, the signal should remain high until the interrupt request is explicitly acknowledged (e.g. writing to a memory-mapped register).

2.5.1. RISC-V Standard Interrupts

The processor supports the standard RISC-V machine-level interrupts for "machine timer interrupt", "machine software interrupt" and "machine external interrupt". Their usage is defined by the RISC-V privileged architecture specifications. However, these interrupt could also be repurposed for custom applications. See CPU section [Traps, Exceptions and Interrupts](#) for more information.

Table 4. Processor Top RISC-V Machine-Level Interrupts

Top signal	Description
<code>mtime_irq_i</code>	Machine timer interrupt from <i>processor-external</i> CLINT (MTI). This IRQ is only available if the processor-internal Core-Local Interruptor (CLINT) unit is not implemented.
<code>msw_irq_i</code>	Machine software interrupt from <i>processor-external</i> CLINT (MSI). This IRQ is only available if the processor-internal Core-Local Interruptor (CLINT) unit is not implemented.
<code>mext_irq_i</code>	Machine external interrupt from <i>processor-external</i> PLIC (MEI). This interrupt is always available as top-entity port.

2.5.2. NEORV32-Specific Fast Interrupt Requests

As part of the NEORV32-specific CPU extensions, the processor core features 16 fast interrupt request signals ([FIRQ0](#) to [FIRQ15](#)) providing dedicated bits in the `mip` and `mie` CSRs and custom `mcause` trap codes. The FIRQ signals are reserved for *processor-internal* modules only (for example for the communication interfaces to signal "available incoming data" or "ready to send new data").

The mapping of the 16 FIRQ channels to the according processor-internal modules is shown in the following table (the channel number also corresponds to the according FIRQ priority: 0 = highest, 15 = lowest):

Table 5. NEORV32 Fast Interrupt Request (FIRQ) Mapping

Channel / priority	Source	Description
0	TWD	TWD FIFO level interrupt
1	CFS	Custom functions subsystem (CFS) interrupt (user-defined)

Channel / priority	Source	Description
2	UART0	UART0 FIFO level interrupt
3	UART1	UART1 FIFO level interrupt
4	-	<i>reserved</i>
5	TRACER	Tracing stop-address match interrupt
6	SPI	SPI FIFO level interrupt
7	TWI	TWI FIFO level interrupt
8	GPIO	GPIO input pin(s) interrupt
9	NEOLED	NEOLED TX FIFO level interrupt
10	DMA	DMA transfer done interrupt
11	SDI	SDI FIFO level interrupt
12	GPTMR	General purpose timer interrupt
13	ONEWIRE	1-wire idle interrupt
14	SLINK	SLINK FIFO level interrupt
15	TRNG	TRNG FIFO level interrupt

2.6. Address Space

As a 32-bit architecture the NEORV32 can access a 4GB physical address space. By default, this address space is split into three main regions. All accesses to "unmapped" addresses (a.k.a. "the void") are redirected to the [Processor-External Bus Interface \(XBUS\)](#). For example, if the internal IMEM is disabled, the accesses to the *entire* address space between `0x00000000` and `0x7FFFFFFF` are converted into XBUS requests. If the XBUS interface is not enabled any access to the void will raise a bus error exception.

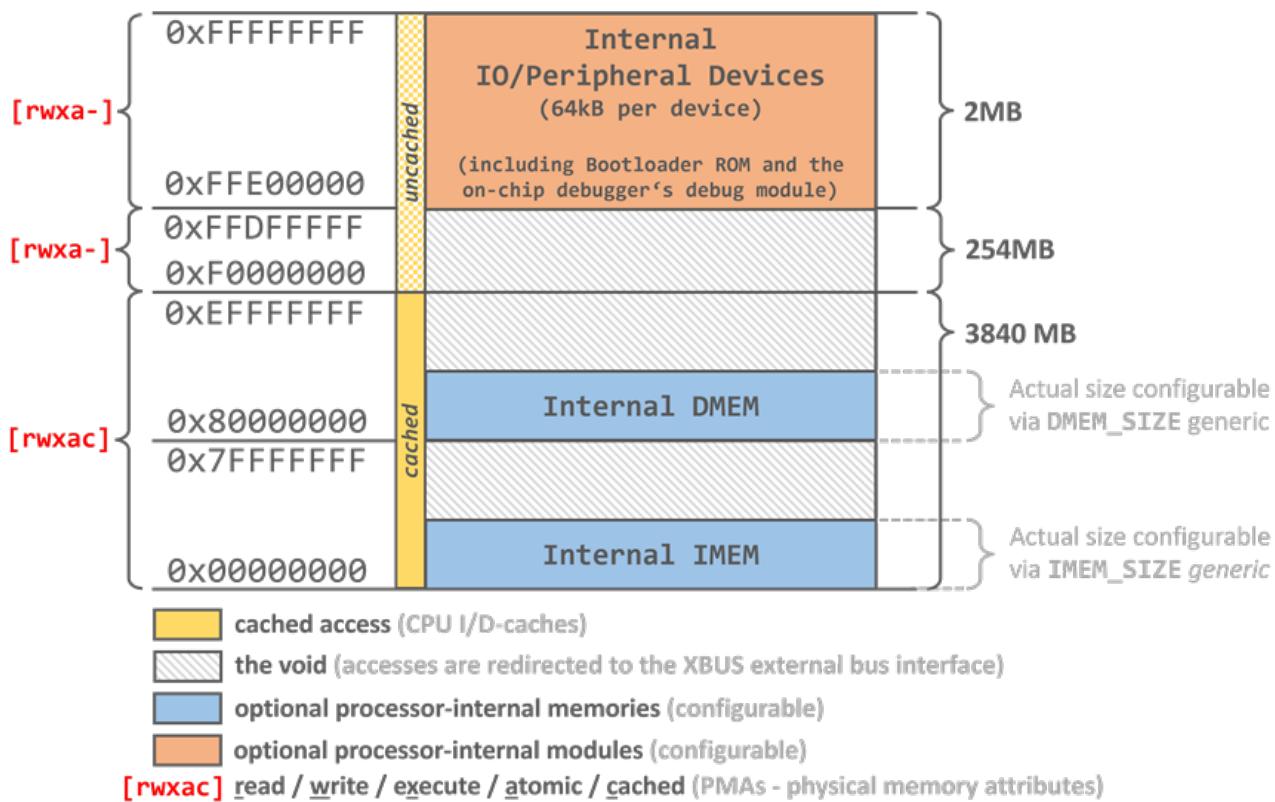


Figure 2. NEORV32 Processor Address Space (Default Configuration)

Each region provides specific *physical memory attributes* ("PMAs") that define the access capabilities (`rwxac`; `r` = read access, `w` = write access, `x` - execute access, `a` = atomic access, `c` = cached CPU access).



Custom PMAs

Custom physical memory attributes enforced by the CPU's *physical memory protection* ([Smpmp ISA Extension](#)) can be used to further constrain the physical memory attributes.

Table 6. Main Address Regions

#	Region	PMAs	Description
1	Internal IMEM address space	<code>rwxac</code>	For instructions / code and constants; mapped to the internal Instruction Memory (IMEM) if implemented.

#	Region	PMAs	Description
2	Internal DMEM address space	<code>rwxac</code>	For application runtime data (heap, stack, etc.); mapped to the internal Data Memory (DMEM) if implemented.
3	IO/peripheral address space	<code>rwxar</code>	Processor-internal peripherals / IO devices including the Bootloader ROM (BOOTROM) .
-	The "void"	<code>rwxac[c]</code>	Unmapped address space. All accesses to this region(s) are redirected to the Processor-External Bus Interface (XBUS) if implemented.

2.6.1. Bus System

The CPU provides individual interfaces for instruction fetch and data access. It can access all of the 32-bit address space from each of the interface. Both of them can be equipped with optional caches (**Processor-Internal Data Cache (dCache)** and **Processor-Internal Instruction Cache (iCache)**).

The two CPU interfaces are multiplexed by a simple bus switch into a *single processor-internal bus*. Optionally, this bus is further multiplexed by another instance of the bus switch so the **Direct Memory Access Controller (DMA)** controller can also access the entire address space. Accesses via the resulting SoC bus are split by the **Bus Gateway** that redirects accesses to the according main address regions (see table above). Accesses to the processor-internal IO/peripheral devices are further redirected via a dedicated **IO Switch**.

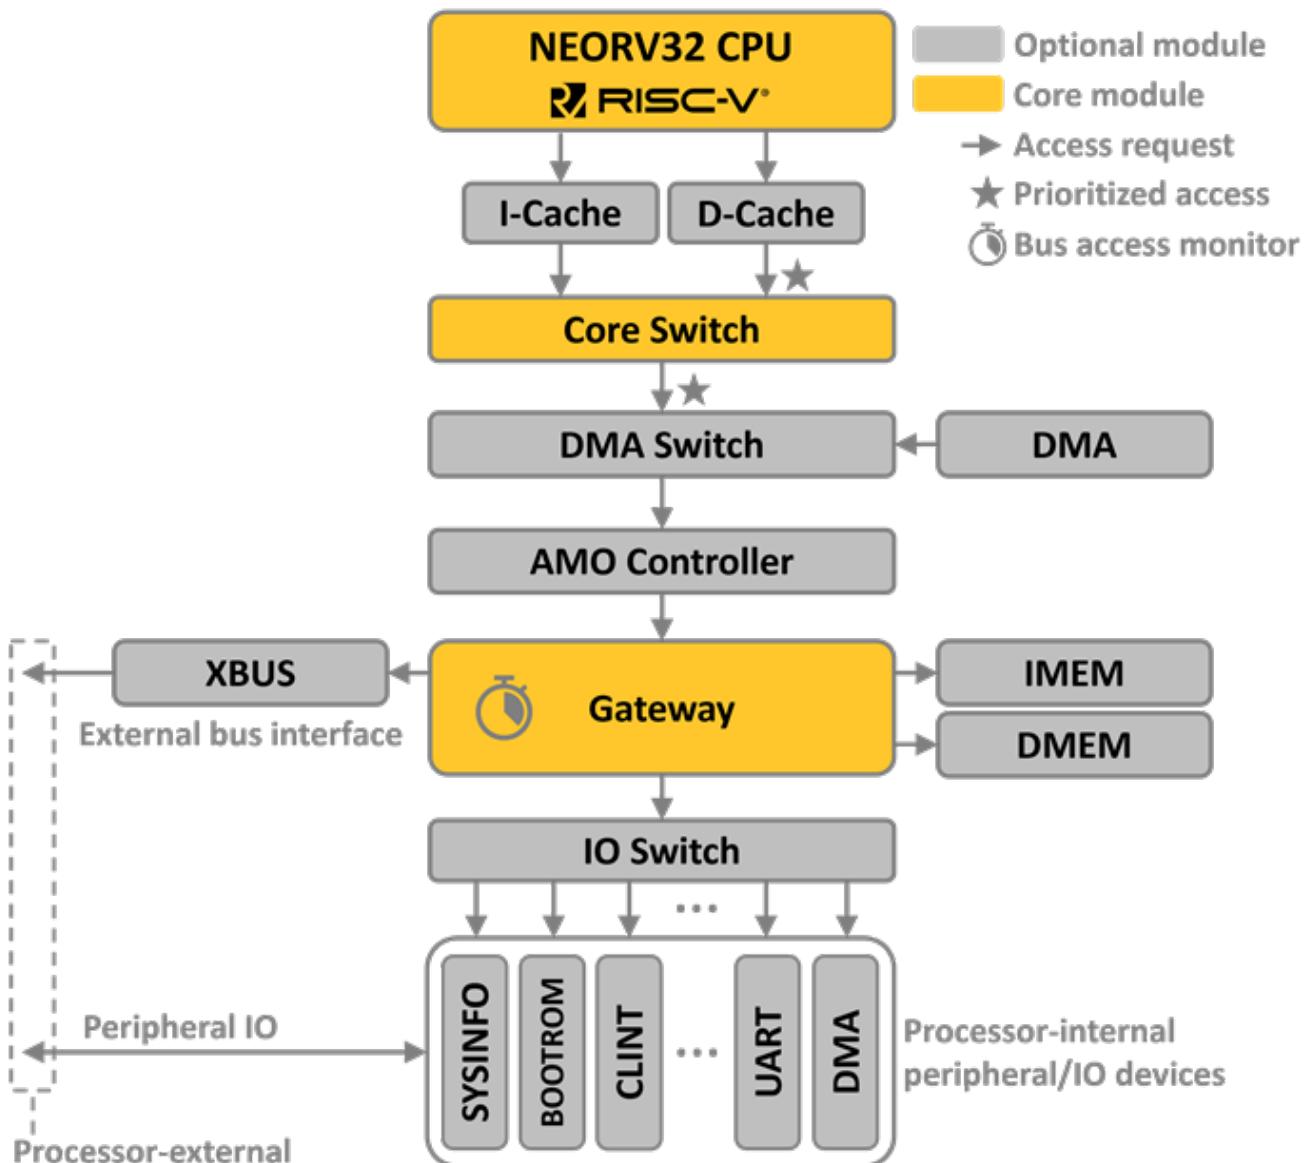


Figure 3. Processor-Internal Bus Architecture

Bus System Infrastructure

The components of the processor's bus system infrastructure are located in [rtl/core/neorv32_bus.vhd](#).

Bus Interface

See sections [CPU Architecture](#) and [Bus Interface](#) for more information regarding the CPU bus accesses.

SMP Dual-Core Configuration

The dual-core configuration adds a second CPU core complex in parallel to the first one. See section [Dual-Core Configuration](#) for more information.

2.6.2. Bus Gateway

The central bus gateway serves two purposes: it **redirects** accesses to the according modules (e.g. memory accesses vs. memory-mapped IO accesses) and also **monitors** all bus transactions. The redirection of access request is based on a customizable memory map implemented via VHDL constants in the main package file ([rtl/core/neorv32_package.vhd](#)):

Listing 3. Main Address Regions Configuration in the VHDL Package File

```
-- Main Address Regions ---
constant mem_imem_base_c : std_ulegic_vector(31 downto 0) := x"00000000"; -- IMEM size
via generic
constant mem_dmem_base_c : std_ulegic_vector(31 downto 0) := x"80000000"; -- DMEM size
via generic
constant mem_io_base_c   : std_ulegic_vector(31 downto 0) := x"ffe00000";
constant mem_io_size_c   : natural := 32*64*1024; -- = 32 * iodev_size_c
```

Bus Monitor and Timeout

Besides the redirecting of bus requests the gateway also implements a **bus monitor** (aka "the bus keeper") that tracks all bus transactions to ensure *safe* and *deterministic* operations. Whenever a memory-mapped device is accessed the bus monitor starts an internal countdown. The accessed module has to respond ("ACK") to the bus request within a bound time window. For **processor-internal** accesses this time windows is defined by a constant in the main NEORV32 package file ([neorv32_package.vhd](#)). For **processor-external** accesses via the **Processor-External Bus Interface (XBUS)** this time window is defined by the **XBUS_TIMEOUT** top configuration generic.

Listing 4. Internal Bus Timeout Configuration (package constant)

```
constant int_bus_tmo_c : natural := 16;
```

Listing 5. External Bus Timeout Configuration (top generic)

```
XBUS_TIMEOUT : natural := 2048;
```

The according time window defines the *maximum* number of cycles after which a non-responding bus request will time out raising a bus access fault exception. For example this can happen when accessing "address space holes" (aka *the void*) - addresses that are not mapped to any physical module. The specific bus access exception type corresponds to the according access type, i.e. instruction fetch bus fault, load bus fault or store bus fault.

2.6.3. IO Switch

The IO switch further decodes the address when accessing the processor-internal IO/peripheral devices and forwards the access request to the according module. Note that a total address space size of 256 bytes is assigned to each IO module in order to simplify address decoding. The IO-

specific address map is also defined in the main VHDL package file ([rtl/core/neorv32_package.vhd](#)).

Listing 6. Exemplary Cut-Out from the IO Address Map

```
-- IO Address Map --
constant iodev_size_c    : natural := 256; -- size of a single IO device (bytes)
constant base_io_cfs_c   : std_ulexic_vector(31 downto 0) := x"fffffeb00";
constant base_io_slink_c : std_ulexic_vector(31 downto 0) := x"fffffec00";
constant base_io_dma_c   : std_ulexic_vector(31 downto 0) := x"fffffed00";
```

2.6.4. Atomic Memory Operations Controller

The atomic memory operations controller is split into two individual modules. Each module implements a specific sub-extensions of the A ISA extension:

Hardware Module	ISA Extensions	Description
<code>neorv32_bus_amo_rmw</code>	Zaamo ISA Extension	Atomic read-modify-write operations
<code>neorv32_bus_amo_rvs</code>	Zalrsc ISA Extension	Atomic reservation-set operations



Direct Access

Atomic operations **always bypass** the CPU's **data cache** using direct/uncached accesses. Care must be taken to maintain data **Memory Coherence**.



Physical Memory Attributes

Atomic memory operations can be executed for *any* address. This also includes cached memory, memory-mapped IO devices and processor-external address spaces.

Atomic Read-Modify-Write Controller

This module converts a single atomic memory operations request into a set of bus transactions to execute an un-interruptable read-modify-write (RMW) operation. For each request, the controller executes an atomic set of three operations:

Table 7. Simplified AMO Controller Operation

Step	Pseudo Code	Description
1	<code>tmp1 ← MEM[address];</code>	Perform a read operation accessing the addressed memory cell and store the loaded data into an internal buffer (<code>tmp1</code>).
2	<code>tmp2 ← tmp1 OP cpu_wdata</code>	The buffered data from the first step is processed using the write data provided by the CPU. The result is stored to another internal buffer (<code>tmp2</code>).

Step	Pseudo Code	Description
3	<code>MEM[address] ← tmp2; cpu_rdata ← tmp1;</code>	The data from the second buffer (<code>tmp2</code>) is written to the addressed memory cell. In parallel, the data from the first buffer (<code>tmp1</code> = original content of the addresses memory cell) is sent back to the requesting CPU.

The controller performs two bus transactions: a read operations and a write operation. Only the acknowledge/error handshake of the last transaction is sent back to the CPU. As the RMW controller is the memory-nearest instance (see [Bus System](#)) the previously described set of operations cannot be interrupted. Hence, they execute in an atomic way.

Atomic Reservation-Set Controller

A "reservation" defines an address or address range that provides a guarding mechanism to support atomic accesses. A new reservation is registered by the `LR` instruction. The address provided by this instruction defines the memory location that is now monitored for atomic accesses. The according `SC` instruction evaluates the state of this reservation. If the reservation is still valid the write access triggered by the `SC` instruction is finally executed and the instruction return a "success" state (`rd` = 0). If the reservation has been invalidated the `SC` instruction will not write to memory and will return a "failed" state (`rd` = 1).

The reservation-set controller implements the *strong semantics*. A reservation is created when executing the `LR` instruction. Any other data memory access (for example by the DMA) will invalidate any active reservation (instruction fetches do not affect the reservations at all). Furthermore, executing a `fence[.i]` instruction will also invalidate any reservation.

The controller implements a **single, global reservation set only**, which is used by both cores in the SMP [Dual-Core Configuration](#). Since the reservation is global, no access address is stored or compared. Furthermore, no CPU IDs are used at all. The `LR` operation of one CPU generates a reservation (for the entire address space, since it is global) that is also used by the other CPU. However, only one CPU will receive a success state for the according `SC` operation, which invalidates the reservation. This guarantees mutually exclusive access.



External Reservations

Note that the reservation-set controller cannot track memory accesses outside of the NEORV32 processor that are executed by processor-external agents.

2.6.5. Memory Coherence

Depending on the configuration, the NEORV32 processor provides several *layers* of memory consisting of caches, buffers and storage.

1. The CPU pipeline and its instruction prefetch buffer
2. The [Processor-Internal Data Cache \(dCache\)](#) and [Processor-Internal Instruction Cache \(iCache\)](#)
3. Internal and external memories

All caches and buffers operate transparently for the software. Hence, special attention must therefore be paid to maintain memory coherence. Note that coherence and cache *synchronization* is **not** performed automatically by the hardware itself as there is no snooping implemented.

NEORV32 uses two instructions for manual memory synchronization which are always available regardless of the actual CPU/ISA configuration:

- **fence.i** (*Zifencei ISA Extension*): flush the CPU's instruction prefetch buffer and clear the CPU's **instruction cache**.
- **fence** (*I ISA Extension / E ISA Extension*): clear and reload the CPU's **data cache**. Flushing is not required as the data cache uses the **write-through** strategy. Hence, write operations are always synchronized with main memory.

Weak Coherence Model

The NEORV32-specific implementation of the **fence[.i]** ordering instructions only provides a rather **weak** coherence model. A core's **fence** just orders all memory accesses towards main memory. Hence, they *can* become visible by other agents (the secondary CPU core, the DMA, processor-external modules) if these agents also synchronize (e.g. reload) their cache(s).



Coherence Example

The following C example shows how to declare and use an atomic variable using the `__Atomic` specifier:

Listing 7. Atomic Variables - C Source Code

```
_Atomic int counter = 0;

counter = 3;
counter++;
```

Listing 8. Atomic Variables - According RISC-V Assembly Code

```
li a3,3
li a4,1

fence rw,w
sw    a3,0(a2)
fence rw,rw

amoadd.w.aqrl zero,a4,(a2)
```

The initial assignment **counter = 0** is translated into a store-word instruction (**sw**) that is automatically encapsulated within two **fence** instructions. This guarantees memory coherence as each **fence** will synchronize the CPU's data cache with upstream/main memory.

The counter increment (`counter++`) is implemented as RISC-V atomic memory operation (`amoadd`). However, the compiler does not encapsulate this in within FENCE instructions. Hence, the altered atomic variable is **not** updated in the CPU's data cache (but in upstream/main memory).

The above example clearly shows that special attention must be paid to memory coherence when using atomic memory operations.

2.7. Boot Configuration

The NEORV32 processor provides some pre-defined boot configurations to adjust system start-up to the requirements of the application. The actual boot configuration is defined by the `BOOT_MODE_SELECT` generic (see [Processor Top Entity - Generics](#)).

Table 8. NEORV32 Boot Configurations

<code>BOOT_MODE_SELECT</code>	Name	Boot address	Description
0 (default)	Bootloader	Base of internal BOOTROM	Implement the processor-internal Bootloader ROM (BOOTROM) as pre-initialized ROM and boot from there.
1	Custom Address	<code>BOOT_ADDR_CUSTOM</code> generic	Start booting at user-defined address (<code>BOOT_ADDR_CUSTOM</code> top generic).
2	IMEM Image	Base of internal IMEM	Implement the processor-internal Instruction Memory (IMEM) as pre-initialized ROM and boot from there.



Dual-Core Boot

For the SMPA dual-core CPU configuration boot procedure see section [Dual-Core Boot](#).

2.7.1. Booting via Bootloader

This is the most common and thus, the default boot configuration. When selected, the processor-internal [Bootloader ROM \(BOOTROM\)](#) is enabled. This ROM contains the executable image (`rtl/core/neorv32_bootloader_image.vhd`) of the default NEORV32 [Bootloader](#) that will be executed right after reset. The bootloader provides an interactive user console for executable upload as well as an automatic boot-configuration targeting external (SPI) memories.

If the processor-internal [Instruction Memory \(IMEM\)](#) is enabled it will be implemented as *blank* RAM.

2.7.2. Boot from Custom Address

This is the most flexible boot configuration as it allows the user to specify a custom boot address via the `BOOT_ADDR_CUSTOM` generic. Note that this address has to be aligned to 4-byte boundary. The processor will start executing from the defined address right after reset. For example, this boot configuration can be used to execute a *custom bootloader* from a memory that is attached via the [Processor-External Bus Interface \(XBUS\)](#).

The [Bootloader ROM \(BOOTROM\)](#) is not enabled / implemented at all. If the processor-internal [Instruction Memory \(IMEM\)](#) is enabled it will be implemented as *blank* RAM.

2.7.3. Boot IMEM Image

This configuration will implement the **Instruction Memory (IMEM)** as *pre-initialized read-only memory* (ROM). The ROM is initialized during synthesis with the according application image file ([rtl/core/neorv32_application_image.vhd](#)). After reset, the CPU will directly start executing this image. Since the IMEM is implemented as ROM, the executable cannot be altered at runtime at all.

The **Bootloader ROM (BOOTROM)** is not enabled / implemented at all.



Internal IMEM is Required

This boot configuration requires the IMEM to be enabled (`IMEM_EN = true`).



Simulation Setup

This boot configuration is handy for simulations as the application software is executed right away without the need for an explicit initialization / executable upload.

2.8. Processor-Internal Modules

Full-Word Write Accesses Only



All peripheral/IO devices should only be accessed in full-word mode (i.e. 32-bit). Byte or half-word (8/16-bit) write accesses might cause undefined behavior.

IO Module Address Space



Each peripheral/IO module occupies an address space of 64kB bytes. Most devices do not fully utilize this address space and will *mirror* the available memory-mapped registers across the entire 64kB address space. However, accessing memory-mapped registers other than the specified ones should be avoided.

Unimplemented Modules / Address Holes



When accessing an IO device that has not been implemented (disabled via the according generic) or when accessing an address that is actually unused, a load/store access fault exception is raised.

Writing to Read-Only Registers



Unless otherwise specified, writing to registers that are listed as read-only does not trigger an exception as the write access is simply ignored by the corresponding hardware module.

IO Access Latency



In order to shorten the critical path of the IO system, the IO switch provides register stages for the request and response buses. Hence, accesses to the processor-internal IO region require two additional clock cycles to complete.

Module Interrupts



Several peripheral/IO devices provide some kind of interrupt. These interrupts are mapped to the CPU's [Custom Fast Interrupt Request Lines](#). See section [Processor Interrupts](#) for more information.



CMSIS System Description View (SVD)

A CMSIS-compatible **System View Description (SVD)** file including all peripherals is available in [sw/svd](#).

2.8.1. Instruction Memory (IMEM)

Hardware source files:	neorv32_imem.vhd neorv32_application_image.vhd	generic processor-internal instruction memory (RAM or ROM) memory image (VHDL package)
Software driver files:	none	
Top entity ports:	none	
Configuration generics:	<code>IMEM_EN</code> <code>IMEM_SIZE</code> <code>IMEM_OUTREG_EN</code> <code>BOOT_MODE_SELECT</code>	implement processor-internal IMEM when <code>true</code> IMEM size in bytes (use a power of 2) add IMEM output register stage implement IMEM as ROM when <code>BOOT_MODE_SELECT</code> = 2; see Boot Configuration
CPU interrupts:	none	

Key Features

- Burst-capable tightly-coupled on-chip instruction memory
- Configurable RAM size
- Accessible at the byte level
- Can be mapped to blockRAM primitives
- Optional output register for improved timing
- Can be implemented as pre-initialized ROM containing application firmware

Overview

Implementation of the processor-internal instruction memory is enabled by the processor's `IMEM_EN` generic. The total memory size in bytes is defined via the `IMEM_SIZE` generic. Note that this size should be a power of two to optimize physical implementation. If enabled, the IMEM is mapped to base address `0x00000000` (see section [Address Space](#)).

By default the IMEM is implemented as true RAM so the content can be modified during run time. This is required when using the [Bootloader](#) (or the [On-Chip Debugger \(OCD\)](#)) so it can update the content of the IMEM at any time.

Alternatively, the IMEM can be implemented as **pre-initialized read-only memory (ROM)**, so the processor can directly boot from it after reset. This option is configured via the `BOOT_MODE_SELECT` generic. See section [Boot Configuration](#) for more information. The initialization image is embedded into the bitstream during synthesis. The software framework provides an option to generate and override the default VHDL initialization file `rtl/core/neorv32_application_image.vhd`, which is

automatically inserted into the IMEM (see [Makefile Targets](#)). If the IMEM is implemented as RAM (default), the memory block will not be initialized at all.



Platform-Specific Memory Primitives

If required, the default IMEM can be replaced by a platform-/technology-specific primitive to optimize area utilization, timing and power consumption.



Memory Size

If the configured memory size (via the `IMEM_SIZE` generic) is not a power of two the actual memory size will be auto-adjusted to the next power of two (e.g. configuring a memory size of 60kB will result in a physical memory size of 64kB).



Output Register Stage

An optional output register stage can be enabled via `IMEM_OUTREG_EN`. For FPGA targets this might improve mapping/timing results. Note that this option will increase the read latency by one clock cycle. Write accesses are not affected by this at all.



Read-Only Access

If the IMEM is implemented as ROM any write attempt to it will raise a *store access fault* exception.

2.8.2. Data Memory (DMEM)

Hardware source files:	neorv32_dmem.vhd	generic processor-internal data memory (RAM)
Software driver files:	none	
Top entity ports:	none	
Configuration generics:	DMEM_EN DMEM_SIZE DMEM_OUTREG_EN	implement processor-internal DMEM when true DMEM size in bytes (use a power of 2) add DMEM output register stage
CPU interrupts:	none	

Key Features

- Burst-capable tightly-coupled on-chip data memory
- Configurable RAM size
- Accessible at the byte level
- Can be mapped to blockRAM primitives
- Optional output register for improved timing

Overview

Implementation of the processor-internal data memory is enabled by the processor's **DMEM_EN** generic. The total memory size in bytes is defined via the **DMEM_SIZE** generic. Note that this size should be a power of two to optimize physical implementation. If the DMEM is implemented, it is mapped to base address **0x80000000** by default (see section [Address Space](#)). The DMEM is always implemented as true RAM.

Platform-Specific Memory Primitives



If required, the default DMEM can be replaced by a platform-/technology-specific primitive to optimize area utilization, timing and power consumption.

Memory Size



If the configured memory size (via the **DMEM_SIZE** generic) is not a power of two the actual memory size will be auto-adjusted to the next power of two (e.g. configuring a memory size of 60kB will result in a physical memory size of 64kB).

Output Register Stage



An optional output register stage can be enabled via **DMEM_OUTREG_EN**. For FPGA targets this might improve mapping/timing results. Note that this option will increase the read latency by one clock cycle. Write accesses are not affected by this

at all.



Execute from RAM

The CPU is capable of executing code also from arbitrary data memory.

2.8.3. Bootloader ROM (BOOTROM)

Hardware source files:	neorv32_boot_rom.vhd neorv32_bootloader_image.vhd	default platform-agnostic bootloader ROM memory image (VHDL package)
Software driver files:	none	<i>implicitly used</i>
Top entity ports:	none	
Configuration generics:	<code>BOOT_MODE_SELECT</code>	implement BOOTROM when <code>BOOT_MODE_SELECT</code> = 0; see Boot Configuration
CPU interrupts:	none	

Key Features

- Burst-capable tightly-coupled on-chip ROM
- Accessible at the byte level
- Can be mapped to blockRAM primitives
- Pre-initialized with default NEORV32 bootloader

Overview

The boot ROM contains the pre-compiled executable image of the default NEORV32 [Bootloader](#). When the [Boot Configuration](#) is set to *bootloader* mode (0) via the `BOOT_MODE_SELECT` generic, the boot ROM is automatically enabled and the CPU boot address is adjusted to the base address of the boot ROM. Note that the entire boot ROM is read-only.

Bootloader Image



The bootloader ROM is initialized during synthesis with the default bootloader image ([rtl/core/neorv32_bootloader_image.vhd](#)). The physical size of the ROM is automatically adjusted to the next power of two of the image size. However, note that the BOOTROM is constrained to a maximum size of 64kB.

2.8.4. Processor-Internal Instruction Cache (iCache)

Hardware source files:	neorv32_cache.vhd	Generic cache module
Software driver files:	none	
Top entity ports:	none	
Configuration generics:	ICACHE_EN ICACHE_NUM_BLOCKS CACHE_BLOCK_SIZE CACHE_BURSTS_EN	implement processor-internal, CPU-exclusive instruction cache (I\$) when true number of cache blocks ("cache lines"); has to be a power of two size of a cache block in bytes; has to be a power of two (global configuration for I\$ and D\$), min 8 enable burst transfers for cache update
CPU interrupts:	none	

Key Features

- Direct-mapped read-only cache
- Configurable number of lines
- Configurable line size
- Allows bypassing for *uncached* accesses

Overview

The processor features an optional CPU instruction cache. The cache is connected directly to the CPU's instruction fetch interface and provides full-transparent accesses. The cache is direct-mapped and read-only. The cache uses [Locked Bus Accesses and Bursts](#) to download cache blocks/lines from main memory. In the [Dual-Core Configuration](#) each CPU core is equipped with a private instruction cache.

The instruction cache is implemented if **ICACHE_EN** is *true*. The total cache memory size in bytes is defined by **ICACHE_NUM_BLOCKS** x **CACHE_BLOCK_SIZE**. **ICACHE_NUM_BLOCKS** defines the number of cache blocks (or "cache lines") and **CACHE_BLOCK_SIZE** defines the block size in bytes; note that this configuration is global for all caches.

Burst Transfers



Cache update operations (e.g. to resolve a cache miss) can use [burst transfers](#) to increase performance. Burst operations are enabled (for all caches) by the **CACHE_BURSTS_EN** top generic. Note that when bursts are enabled all cache block transfers are **always executed as burst transfers**. Hence, all devices, memories and endpoints that can be accessed by the cache must also be able to process

bursts (including the [Processor-External Bus Interface \(XBUS\)](#)).

Uncached Accesses

The cache provides direct/uncached accesses to memory (bypassing the cache) in order to access memory-mapped IO (like the processor-internal IO/peripheral modules). All accesses that target the address range from `0xF0000000` to `0xFFFFFFF` will bypass the cache. Hence, access will not be cached. See section [Address Space](#) for more information. Furthermore, the atomic memory operations of the [Zaamo ISA Extension](#) will always **bypass** the cache.



Manual Cache Clear/Reload and Memory Coherence

By executing the `fence.i` instruction the instruction cache is cleared and reloaded. See section [Memory Coherence](#) for more information.



Cache Block Update Bus Error Handling

If the cache encounters a bus error while downloading a new block from main memory, only the according data word is marked as faulty (instead of invalidating the entire cache block). A bus access fault exception is only raised if the CPU accesses a cached word that has actually been marked as faulty.



Retrieve Cache Configuration from Software

Software can retrieve the cache configuration/layout from the [SYSINFO - Cache Configuration](#) register.



2.8.5. Processor-Internal Data Cache (dCache)

Hardware source files:	neorv32_cache.vhd	Generic cache module
Software driver files:	none	
Top entity ports:	none	
Configuration generics:	DCACHE_EN DCACHE_NUM_BLOCKS CACHE_BLOCK_SIZE CACHE_BURSTS_EN	implement processor-internal, CPU-exclusive data cache (D\$) when true number of cache blocks ("cache lines"); has to be a power of two size of a cache block in bytes; has to be a power of two (global configuration for I\$ and D\$), min 8 enable burst transfers for cache update
CPU interrupts:	none	

Key Features

- Direct-mapped write-through cache
- Configurable number of lines
- Configurable line size
- Allows bypassing for *uncached* accesses

Overview

The processor features an optional CPU data cache. The cache is connected directly to the CPU's data access interface and provides full-transparent accesses. The cache is direct-mapped and uses "write-through" as write strategy. The cache uses [Locked Bus Accesses and Bursts](#) to download cache blocks/lines from main memory. In the [Dual-Core Configuration](#) each CPU core is equipped with a private data cache.

The data cache is implemented if **DCACHE_EN** is *true*. The total cache memory size in bytes is defined by **DCACHE_NUM_BLOCKS** \times **CACHE_BLOCK_SIZE**. **DCACHE_NUM_BLOCKS** defines the number of cache blocks (or "cache lines") and **CACHE_BLOCK_SIZE** defines the block size in bytes; note that this configuration is global for all caches.

Burst Transfers

Cache update operations (e.g. to resolve a cache miss) can use [burst transfers](#) to increase performance. Burst operations are enabled (for all caches) by the **CACHE_BURSTS_EN** top generic. Note that when bursts are enabled all cache block transfers are **always executed as burst transfers**. Hence, all devices, memories and endpoints that can be accessed by the cache must also be able to process bursts (including the [Processor-External Bus Interface \(XBUS\)](#)).



Uncached Accesses

The cache provides direct/uncached accesses to memory (bypassing the cache) in order to access memory-mapped IO (like the processor-internal IO/peripheral modules). All accesses that target the address range from `0xF0000000` to `0xFFFFFFFF` will bypass the cache. Hence, access will not be cached. See section [Address Space](#) for more information. Furthermore, the atomic memory operations of the [Zaamo ISA Extension](#) will always **bypass** the cache.



Manual Cache Flush/Clear/Reload and Memory Coherence

By executing the `fence` instruction the data cache is flushed, cleared and reloaded. See section [Memory Coherence](#) for more information.



Cache Block Update Bus Error Handling

If the cache encounters a bus error while downloading a new block from main memory, only the according data word is marked as faulty (instead of invalidating the entire cache block). A bus access fault exception is only raised if the CPU accesses a cached word that has actually been marked as faulty.



Retrieve Cache Configuration from Software

Software can retrieve the cache configuration/layout from the [SYSINFO - Cache Configuration](#) register.



2.8.6. Direct Memory Access Controller (DMA)

Hardware source files:	neorv32_dma.vhd	
Software driver files:	neorv32_dma.c neorv32_dma.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	none	
Configuration generics:	<code>IO_DMA_EN</code> <code>IO_DMA_DSC_FIFO</code>	implement DMA when <code>true</code> descriptor FIFO depth, has to be a power of 2, min 4, max 512
CPU interrupts:	fast IRQ channel 10	DMA transfer(s) done (see Processor Interrupts)

Key Features

- CPU-independent data movement
- Byte-wide or word-wide data transfers
- Up to 16MB (bytes) or 64MB (words) per transfer
- Optional Endianness conversion
- Optional descriptor FIFO
- Chaining of pre-programmed transfers
- Transfer-done interrupt

Overview

The NEORV32 DMA features a lightweight direct memory access controller that allows to move and modify data independently of the CPU. Only a single read/write channel is implemented. So only one programmed transfer can be in progress at a time. However, a configurable descriptor FIFO is provided which allows to program several transfers so the DMA can execute them one after the other.

The DMA is connected to the central processor-internal bus system (see section [Address Space](#)) and can access the entire/same address space as the CPU core. It uses *interleaving mode* accessing the central processor bus only if the CPU does not currently request a bus access. The DMA controller can handle different data quantities (e.g. read bytes and write them back as zero-extended words) and can also change the Endianness of data while transferring. It supports reading/writing data from/to fixed or auto-incrementing addresses.

DMA Bus Access



Transactions performed by the DMA are executed as bus transactions with elevated **machine-mode** privilege level. Note that any physical memory protection rules ([Smpmp ISA Extension](#)) are not applied to DMA transfers. Furthermore, the DMA uses single-transfers only (.e. no burst transfers).

*DMA Demo Program*

A DMA example program can be found in [sw/example/demo_dma](#).

Theory of Operation

The DMA provides just two memory-mapped interface registers: A status and control register [CTRL](#) and another one for writing the transfer descriptor(s) to the internal descriptor FIFO.

The DMA is enabled by setting the [DMA_CTRL_EN](#) bit of the control register. Clearing this flag will abort any outstanding transfer and will also reset/clear the descriptor FIFO. A programmed DMA transfer is initiated by setting the control register's [DMA_CTRL_START](#) bit. Setting this bit while the descriptor FIFO is empty has no effect. The current status of the FIFO can be checked via the [DMA_CTRL_D*](#) flags.

The DMA uses an atomic read-modify-write transfer process. Data is read from the current source address, modified/aligned internally and then written back to the current destination address. If the DMA controller encounters a bus error during operation, it will set the [DMA_CTRL_ERROR](#) flag and will terminate the current transfer. A new transfer can only start if the [DMA_CTRL_ERROR](#) flag is cleared manually.

When the [DMA_CTRL_DONE](#) flag is set the DMA has completed all programmed transfers, i.e. all descriptors from the FIFO were executed. This flag also triggers the DMA controller's interrupt request signal. The application software has to clear [DMA_CTRL_DONE](#) in order to acknowledge the interrupt and to start further transfers.

DMA Descriptor

All DMA transfers are executed based on *descriptors*. A descriptor contains the data source and destination base addresses as well as the number of elements to transfer and the data type and handling configuration. A complete descriptor is encoded as 3 consecutive 32-bit words:

Table 9. DMA Descriptor

Index	Size	Description
0	32-bit	Source data base address
1	32-bit	Destination data base address
2	32-bit	Transfer configuration word (see next table)

Descriptor FIFO Size



The descriptor FIFO has a minimal size of 4 entries. This can be extended by the [IO_DMA_DSC_FIFO](#) generic.

Incomplete Descriptors



The DMA controller consumes 3 entries from the FIFO for each transfer. If the FIFO does not provide a complete DMA descriptor, the controller will wait until a

complete descriptor is available.

The source and destination data addresses can target any memory location in the entire 32-bit address space including memory-mapped peripherals. The number of elements to transfer as well as incrementing or constant byte- or word-level transfers are configured via the transfer configuration word (3rd descriptor word):

Table 10. DMA Descriptor - Transfer Configuration Word

Bit(s)	Name	Description
23:0	DMA_CONF_NUM	Number of elements to transfer; must be greater than zero
26:24	-	<i>reserved</i> , set to zero
27	DMA_CONF_BSWAP	Set to swap byte order ("Endianness" conversion)
29:28	DMA_CONF_SRC	Source data configuration (see list below)
31:30	DMA_CONF_DST	Destination data configuration (see list below)

Source and destination data accesses are configured by a 2-bit selector individually for the source and the destination data:

- **00**: Constant byte - transfer data as byte (8-bit); do not alter address during transfer
- **01**: Constant word - transfer data as word (32-bit); do not alter address during transfer
- **10**: Incrementing byte - transfer data as byte (8-bit); increment the source address by 1
- **11**: Incrementing word - transfer data as word (32-bit); increment the source address by 4

Optionally, the controller can automatically swap the logical byte order ("Endianness") of the transferred data when the **DMA_CONF_BSWAP** bit is set.

Register Map

Table 11. DMA Register Map ([struct NEORV32_DMA](#))

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffffe0000	CTRL	0 DMA_CTRL_EN	r/w	DMA module enable; reset module when cleared
		1 DMA_CTRL_START	-/w	Start programmed DMA transfer(s)
		15:2 reserved	r/-	reserved, read as zero
		19:16 DMA_CTRL_DFI0_MSB : DMA_CTRL_DFI0_LSB	r/-	Descriptor FIFO depth, log2(IO_DMA_DSC_FIFO)
		26:20 reserved	r/-	reserved, read as zero
		27 DMA_CTRL_ACK	-/w	Write 1 to clear DMA interrupt (also clears DMA_CTRL_ERROR and DMA_CTRL_DONE)
		27 DMA_CTRL_DEMPTY	r/-	Descriptor FIFO is empty
		28 DMA_CTRL_DFULL	r/-	Descriptor FIFO is full
		29 DMA_CTRL_ERROR	r/-	Bus access error during transfer or incomplete descriptor data
		30 DMA_CTRL_DONE	r/1	All transfers executed
0xffffe0004	DESC	31 DMA_CTRL_BUSY	r/-	DMA transfer(s) in progress
		31:0	-/w	Descriptor FIFO write access

2.8.7. Processor-External Bus Interface (XBUS)

Hardware source files:	neorv32_xbus.vhd	External bus gateway
Software driver files:	none	
Top entity ports:	<code>xbus_adr_o</code> <code>xbus_dat_i</code> <code>xbus_dat_o</code> <code>xbus_cti_o</code> <code>xbus_tag_o</code> <code>xbus_we_o</code> <code>xbus_sel_o</code> <code>xbus_stb_o</code> <code>xbus_cyc_o</code> <code>xbus_ack_i</code> <code>xbus_err_i</code>	address output (32-bit) data input (32-bit) data output (32-bit) cycle type (3-bit) access tag (3-bit) write enable (1-bit) byte enable (4-bit) bus strobe (1-bit) valid cycle (1-bit) acknowledge (1-bit) bus error (1-bit)
Configuration generics:	<code>XBUS_EN</code> <code>XBUS_TIMEOUT</code> <code>XBUS_REGSTAGE_EN</code> <code>(CACHE_BLOCK_SIZE)</code> <code>(CACHE_BURSTS_EN)</code>	enable external bus interface when <code>true</code> number of clock cycles after which an unacknowledged external bus access will auto-terminate (0 = disabled) implement XBUS register stages burst size enable burst transfers for cache update
CPU interrupts:	none	

Key Features

- Gateway for processor-external modules
- Wishbone-compatible bus protocol
- Optional burst support
- Optional AXI4-compatible bridging

Overview

The external bus interface provides a **Wishbone**-compatible on-chip bus interface. This bus interface can be used to attach processor-external modules like memories, custom hardware accelerators or additional peripheral devices.

Burst Transfers



If any cache (**i-cache** or **d-cache**) is implemented and bursts are globally enabled (by the **CACHE_BURSTS_EN** top generic) all cache block transfers are **always executed as burst transfers** with a burst size equal to the cache block size (**CACHE_BLOCK_SIZE** top generic). Burst transfers should **not** be enabled if any external module mapped to *cached Address Space* does not support bursts. Note that burst transactions need to set **ACK** or **ERR** for each burst element. Note that the cycle type identifier signal (**xbus_cti_o**) does not support the end-of-burst identifier.

Address Mapping



The external interface is **not** mapped to a specific address range. Instead, all CPU memory accesses that do not target a specific (and actually implemented) processor-internal address region (hence, accessing the "void"; see section **Address Space**) are **redirected** to the external bus interface.



Wishbone Specs

The official Wishbone specification can be found online: <https://wishbone-interconnect.readthedocs.io/en/latest/index.html>



AXI4-Compatible Interface Bridge

A bridge that converts the processor's XBUS interface into an AXI4-compatible host interface is available in [rtl/system_integration/xbus2axi4_bridge.vhd](#). This bridge is also used for the ENORV32 Vivado IP block: https://stnolting.github.io/neorv32/ug/#_packaging_the_processor_as_vivado_ip_block

XBUS Bus Protocol

The external bus interface implements a subset of the **pipelined Wishbone** protocol. Basically, three types of bus transfer are implemented which are illustrated in the following figures. The actual transfer type is indicated via the *cycle type identifier* (**xbus_cti_o**) signal.

1. **Single-access** (**xbus_cti_o = 000**) transfers perform a single read or write operation.
2. **Atomic-access** (**xbus_cti_o = 001**) transfers perform a read followed by a write operation. The bus is locked during the entire transfer (keeping **cyc** high) to maintain exclusive bus access. This transfer type is used by the CPU to perform atomic read-modify-write operations.
3. **Burst read** (**xbus_cti_o = 010**) transfers perform several consecutive read accesses. This transfer type is used by cache block operations.

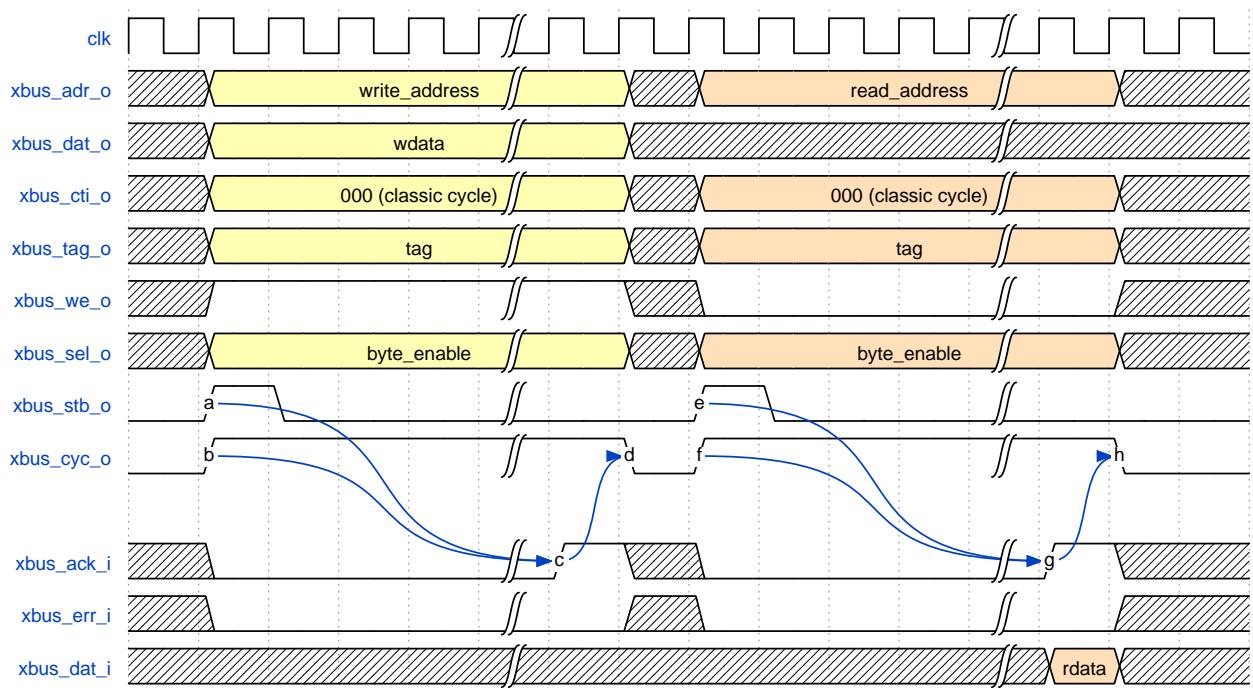


Figure 4. XBUS Single Access Transfers: Write (left) and Read (right)

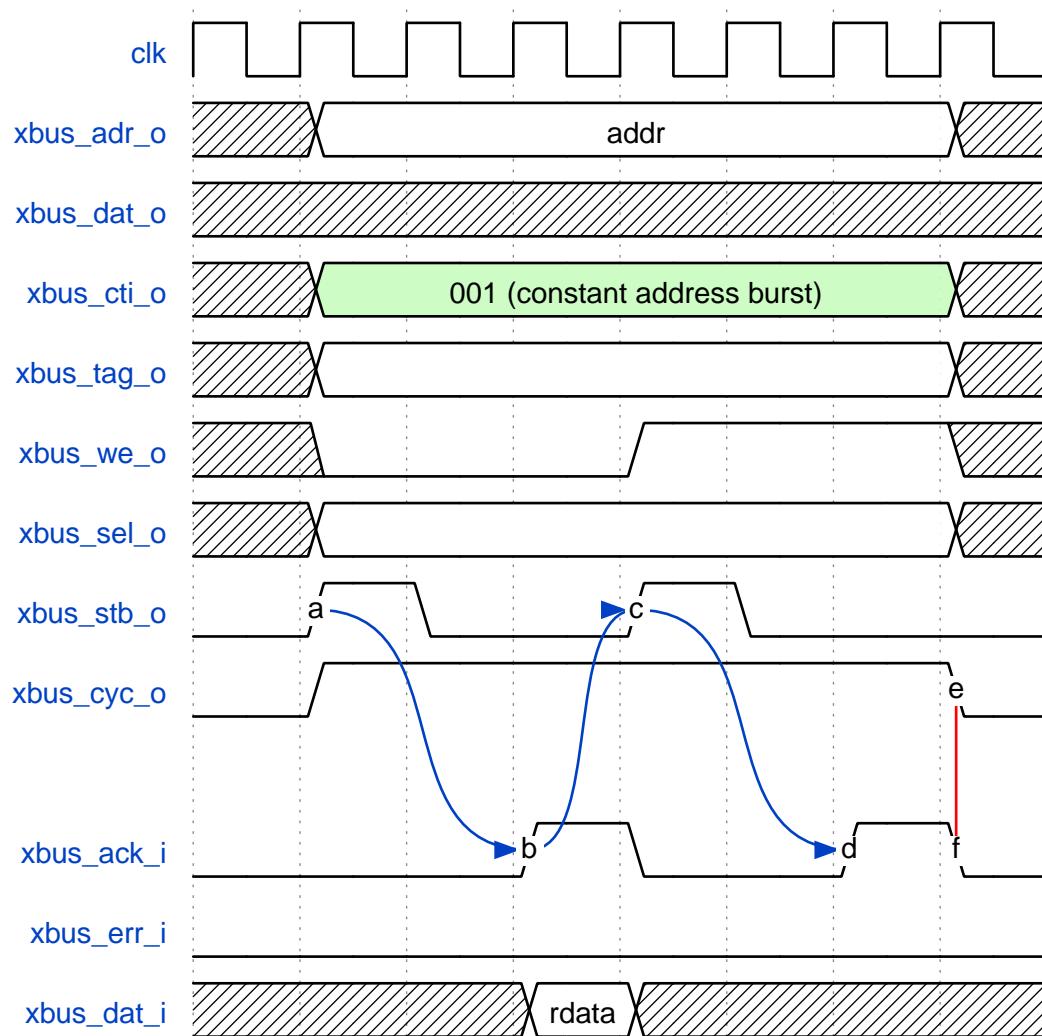


Figure 5. XBUS Atomic Access Transfer

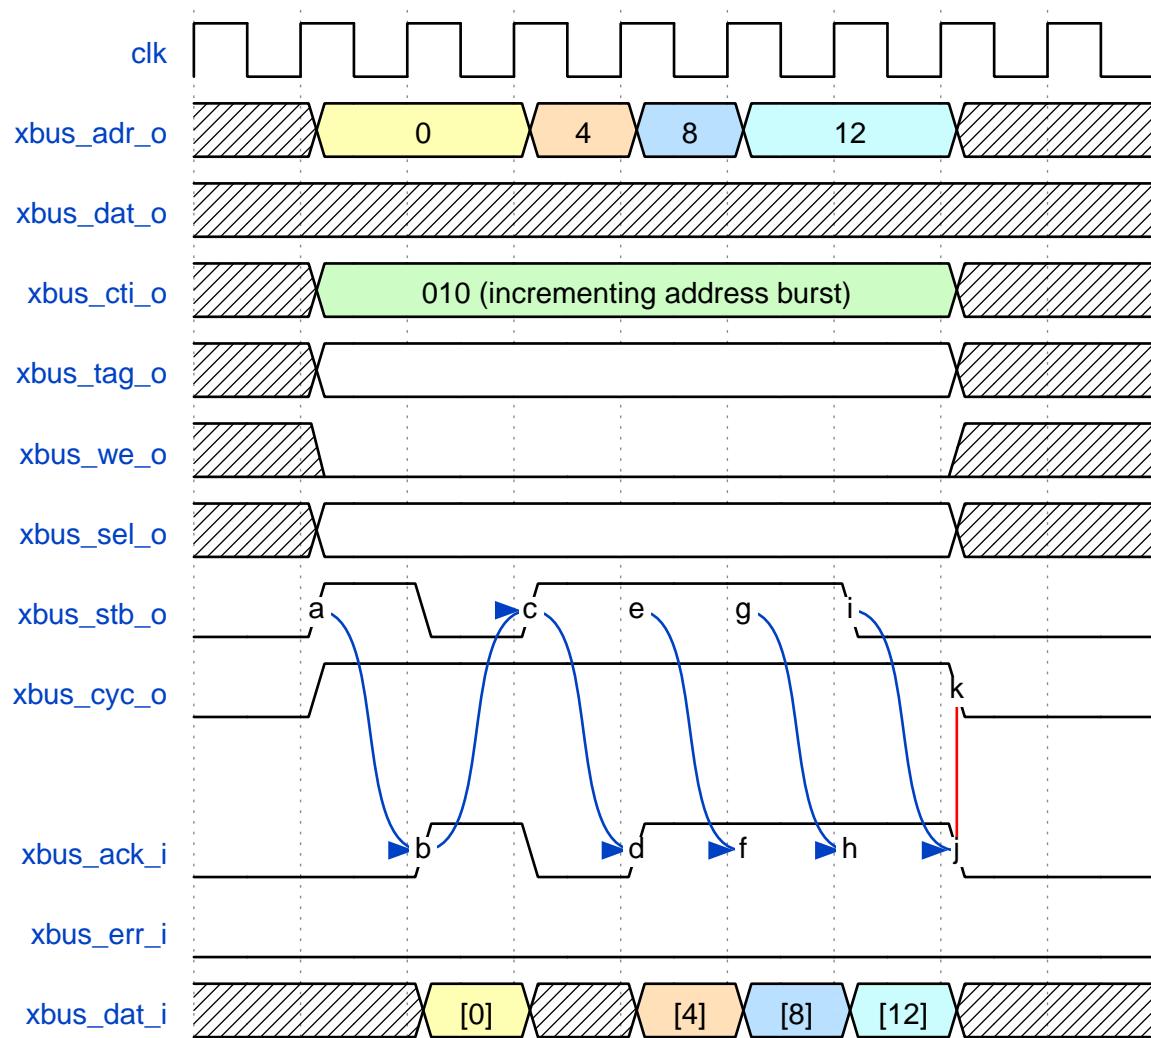


Figure 6. XBUS Burst Read Transfer (4-Words)

Bus Latency

An accessed XBUS device does not have to respond immediately to a bus request by sending an **ACK**. Instead, there is a **time window** where the device has to acknowledge the transfer. This time window is configured by the **XBUS_TIMEOUT** generic. Note that the value provided by this generic is internally extended to the next power of two.

All XBUS transactions have to be acknowledged within this time window. Otherwise the transfer is terminated and a bus fault exception is raised. See section [Bus Monitor and Timeout](#) for more information.

Furthermore, an accessed XBUS device can signal an error condition at any time by setting the **ERR** signal high for one cycle. This will also terminate the current bus transaction raising a CPU bus fault exception.

Register Stage



An optional register stage can be added to the XBUS gateway to break up the critical path easing timing closure. When **XBUS_REGSTAGE_EN** is *true* all outgoing and incoming XBUS signals are registered increasing access latency by two cycles. Furthermore, all outgoing signals (like the address) will be kept stable if there is no

bus access being initiated.

Access Tag

The XBUS tag signal `xbus_tag_o` provides additional information about the current access cycle. The encoding is compatible to the AXI4 `xPROT` signal.

- `xbus_tag_o(2)` **I**: access is an **instruction** fetch when set; access is a data access when cleared
- `xbus_tag_o(1)` **NS**: this bit is hardwired to `0` indicating a **secure** access
- `xbus_tag_o(0)` **P**: access is performed from **privileged** mode (machine-mode) when set

2.8.8. Stream Link Interface (SLINK)

Hardware source files:	neorv32_slink.vhd	
Software driver files:	neorv32_slink.c neorv32_slink.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>slink_rx_dat_i</code> <code>slink_rx_src_i</code> <code>slink_rx_val_i</code> <code>slink_rx_lst_i</code> <code>slink_rx_rdy_o</code> <code>slink_tx_dat_o</code> <code>slink_tx_dst_o</code> <code>slink_tx_val_o</code> <code>slink_tx_lst_o</code> <code>slink_tx_rdy_i</code>	RX link data (32-bit) RX routing information (4-bit, optional) RX link data valid (1-bit) RX link last element of stream (1-bit, optional) RX link ready to receive (1-bit) TX link data (32-bit) TX routing information (4-bit, optional) TX link data valid (1-bit) TX link last element of stream (1-bit, optional) TX link allowed to send (1-bit)
Configuration generics:	<code>IO_SLINK_EN</code> <code>IO_SLINK_RX_FIFO</code> <code>IO_SLINK_TX_FIFO</code>	implement SLINK when <i>true</i> RX FIFO depth, has to be a power of two, min 1 TX FIFO depth, has to be a power of two, min 1
CPU interrupts:	fast IRQ channel 14	SLINK IRQ (see Processor Interrupts)

Key Features

- Compatible to the AXI-4 stream standard
- Independent RX and TX channels
- Optional per-channel FIFOs
- Supports "last" and "source/destination" stream signals
- Interrupt based on FIFO status

Overview

The stream link interface provides independent RX and TX channels for sending and receiving stream data. Each channel features a configurable internal data FIFO (`IO_SLINK_RX_FIFO` and `IO_SLINK_TX_FIFO`). The SLINK interface provides higher bandwidth and less latency than the

external bus interface making it ideally suited for coupling custom stream processors or streaming peripherals.



Example Program

An example program for the SLINK module is available in [sw/example/demo_slink](#).

Interface & Protocol

Each link/channel consists of the following signals:

- `dat` provides the actual data word
- `val` indicates the current transmission cycle is valid
- `lst` indicates the current transmission cycle is the last element of a stream (optional)
- `rdy` indicates the receiver is ready to receive
- `src` and `dst` provide source/destination routing information (optional)

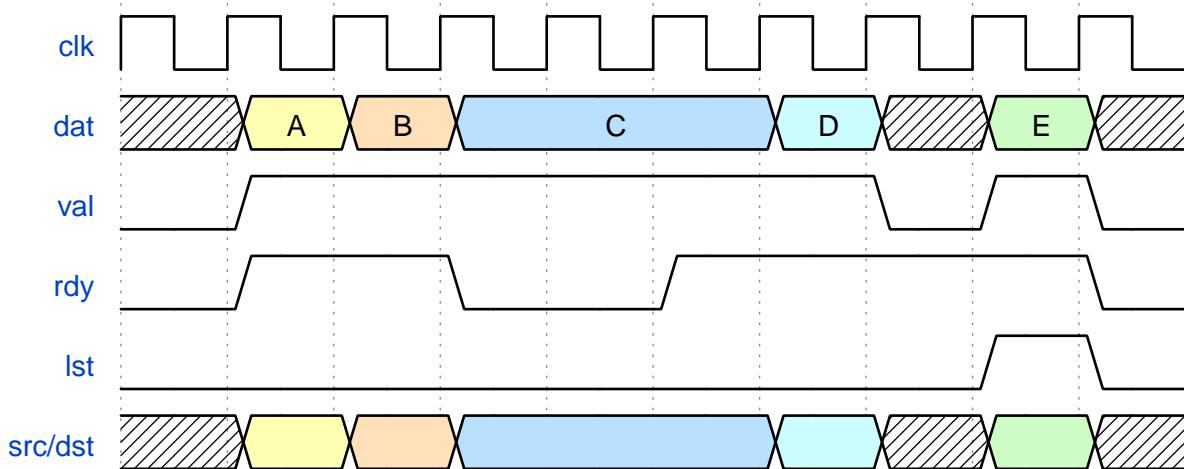


Figure 7. SLINK example transmissions



AXI4-Stream Compatibility

The interface names (except for `src` and `dst`) and the underlying protocol is compatible to the AXI4-Stream protocol standard. A processor top entity with AXI4-Stream-compatible interfaces can be found in [rtl/system_inegration](#). More information can be found in the user guide: https://stnolting.github.io/neorv32/ug/#_packaging_the_processor_as_vivado_ip_block

Theory of Operation

The SLINK provides four interface register. The control register (`CTRL`) is used to configure the module and to check its status. The `ROUTE` register can be used to configure a 4-bit routing ID. Two individual data registers (`DATA` and `DATA_LAST`) are used to send and receive the actual data.

The `DATA` register provides direct access to the RX/TX FIFO buffers. Read accesses return data from the RX FIFO. The end-of-stream delimiter as well as the RX routing information are also buffered in

the RX FIFO and can be obtained from control register's `SLINK_CTRL_RX_LAST` flag and the `ROUTE` register, respectively, after reading the actual RX data.

Writing to the `DATA` register will immediately write to the TX link FIFO. Writing to the `DATA_LAST` register will also push data to the TX link FIFO and will also set the end-of-stream delimiter for that data word. TX routing information can be configured via the `ROUTE` register before writing data to `DATA` / `DATA_LAST`.

The configured FIFO sizes can be retrieved by software via the control register's `SLINK_CTRL_RX_FIFO_*` and `SLINK_CTRL_TX_FIFO_*` bits.

The SLINK is globally enabled by setting the control register's enable bit `SLINK_CTRL_EN`. Clearing this bit will reset all internal logic and will also clear both FIFOs. Writing to the TX channel's FIFO while it is *full* will have no effect. Reading from the RX channel's FIFO while it is *empty* will also have no effect and will return the last received data word. The current status of the RX and TX FIFOs can be determined via the control register's `SLINK_CTRL_RX_EMPTY`, `SLINK_CTRL_RX_FULL`, `SLINK_CTRL_TX_EMPTY` and `SLINK_CTRL_TX_FULL` flags.

Interrupt

The SLINK module provides a single interrupt request that can be used to signal certain RX/TX data FIFO conditions. The interrupt conditions are based on the RX/TX FIFO status flags `SLINK_CTRL_RX_*` / `SLINK_CTRL_TX_*` and are configured via the according `SLINK_CTRL_IRQ_RX_*` / `SLINK_CTRL_IRQ_TX_*` bits. The SLINK interrupt will fire when the module is enabled (`SLINK_CTRL_EN`) and **any** of the enabled interrupt conditions is met. Hence, all enabled interrupt conditions are logically OR-ed. The interrupt remains active until all interrupt-causing conditions are resolved.

Register Map

Table 12. SLINK register map (`struct NEORV32_SLINK`)

Address	Name [C]	Bit(s)	R/W	Function
0xffec0000	CTRL	0 SLINK_CTRL_EN	r/w	SLINK global enable
		7:1 reserved	r/-	<i>reserved</i> , read as zero
		8 SLINK_CTRL_RX_EMPTY	r/-	RX FIFO empty
		9 SLINK_CTRL_RX_FULL	r/-	RX FIFO full
		10 SLINK_CTRL_TX_EMPTY	r/-	TX FIFO empty
		11 SLINK_CTRL_TX_FULL	r/-	TX FIFO full
		12 SLINK_CTRL_RX_LAST	r/-	Last word read from DATA is end-of-stream
		15:13 reserved	r/-	<i>reserved</i> , read as zero
		16 SLINK_CTRL_IRQ_RX_NEMPTY	r/w	Fire interrupt if RX FIFO not empty
		17 SLINK_CTRL_IRQ_RX_FULL	r/w	Fire interrupt if RX FIFO full
		18 SLINK_CTRL_IRQ_TX_EMPTY	r/w	Fire interrupt if TX FIFO empty
		19 SLINK_CTRL_IRQ_TX_NFULL	r/w	Fire interrupt if TX FIFO not full
		23:20 reserved	r/-	<i>reserved</i> , read as zero
		27:24 SLINK_CTRL_RX_FIFO_MSB : SLINK_CTRL_RX_FIFO_LSB	r/-	$\log_2(\text{RX FIFO size})$
		31:28 SLINK_CTRL_TX_FIFO_MSB : SLINK_CTRL_TX_FIFO_LSB	r/-	$\log_2(\text{TX FIFO size})$
0xffec0004	ROUTE	3:0	-/w	TX destination routing information (<code>slink_tx_dst_o</code>)
		3:0	r/-	RX source routing information (<code>slink_rx_src_i</code>)
		31:4	r/-	<i>reserved</i> , read as zero
0xffec0008	DATA	31:0	r/w	Write data to TX FIFO; read data from RX FIFO
0xffec000c	DATA_LAST	31:0	r/w	Write data to TX FIFO (and also set end-of-stream delimiter); read data from RX FIFO

2.8.9. General Purpose Input and Output Port (GPIO)

Hardware source files:	neorv32_gpio.vhd	
Software driver files:	neorv32_gpio.c neorv32_gpio.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>gpio_o</code> <code>gpio_i</code>	32-bit parallel output port 32-bit parallel input port
Configuration generics:	<code>IO_GPIO_NUM</code>	number of input/output pairs to implement (0..32)
CPU interrupts:	fast IRQ channel 8	input-pin change interrupt (see Processor Interrupts)

Key Features

- Up to 32 inputs and 32 outputs
- Each input pin is interrupt-capable
- Configurable interrupt triggers (positive/negative level or edge)

Overview

The general purpose IO unit provides simple uni-directional input and output port. These ports can be used chip-externally (for example to drive status LEDs, connect buttons, etc.) or chip-internally to provide control signals for other IP modules. The input port features programmable pin-individual level or edge interrupts capabilities.

Data written to the `PORT_OUT` will appear on the processor's `gpio_o` port. Vice versa, the `PORT_IN` register represents the current state of the processor's `gpio_i`.

The actual number of input/output pairs is defined by the `IO_GPIO_NUM` generic. When set to zero, the GPIO module is excluded from synthesis and the output port `gpio_o` is tied to all-zero. If `IO_GPIO_NUM` is less than the maximum value of 32, only the LSB-aligned bits in `gpio_o` and `gpio_i` are actually connected while the remaining bits are tied to zero or are left unconnected, respectively. This also applies to all memory-mapped interface registers of the GPIO module (i.e. the according most-significant bits are hardwired to zero).

Input Pin Interrupts

Each input pin (`gpio_i`) provides an individual programmable interrupt trigger. The actual interrupt trigger type can be configured individually for each input pin using the `IRQ_TYPE` and `IRQ_POLARITY` registers. `IRQ_TYPE` defines the actual trigger type (level-triggered or edge-triggered), while `IRQ_POLARITY` defines the trigger's polarity (low-level/falling-edge or high-level/rising-edge). The position of each bit in these registers corresponds the according `gpio_i` input pin.

Each pin interrupt channel can be enabled or disabled individually using the `IRQ_ENABLE` register. Each bit in this register corresponds to the according input pin. If the programmed trigger of a

disabled input (`IRQ_ENABLE(i) = 0`) fires, the interrupt request is entirely ignored.

Table 13. GPIO Trigger Configuration for Pin *i*

<code>IRQ_ENABLE(i)</code>	<code>IRQ_TYPE(i)</code>	<code>IRQ_POLARITY(i)</code>	Resulting trigger of <code>gpio_i(i)</code>
1	0	0	low-level (<code>GPIO_TRIG_LEVEL_LOW</code>)
1	0	1	high-level (<code>GPIO_TRIG_LEVEL_HIGH</code>)
1	1	0	falling-edge (<code>GPIO_TRIG_EDGE_FALLING</code>)
1	1	1	rising-edge (<code>GPIO_TRIG_EDGE_RISING</code>)
0	-	-	interrupt disabled

If the configured trigger of an enabled input pin (`IRQ_ENABLE(i) = 1`) fires, the according interrupt request is buffered internally in the `IRQ_PENDING` register. When this register contains a non-zero value (i.e. any bit becomes set) an interrupt request is sent to the CPU via FIRQ channel 8 (see [Processor Interrupts](#)).

The CPU can determine the interrupt-triggering pins by reading the `IRQ_PENDING` register. Each set bit in this register indicates that the according input pin's interrupt trigger has fired. Then, the CPU can clear those pending interrupt pin by setting all set bits to zero.



GPIO Interrupts Demo Program

A demo program for the GPIO input interrupts can be found in `sw/example/demo_gpio`.

Register Map

Table 14. GPIO unit register map (`struct NEORV32_GPIO`)

Address	Name [C]	Bit(s)	R/W	Function
<code>0xffffc0000</code>	<code>PORT_IN</code>	31:0	r/-	Parallel input port; <code>PORT_IN(i)</code> corresponds to <code>gpio_i(i)</code>
<code>0xffffc0004</code>	<code>PORT_OUT</code>	31:0	r/w	Parallel output port; <code>PORT_OUT(i)</code> corresponds to <code>gpio_o(i)</code>
<code>0xffffc0008</code>	-	31:0	r/-	<i>reserved</i> , read as zero
<code>0xffffc000c</code>	-	31:0	r/-	<i>reserved</i> , read as zero
<code>0xffffc0010</code>	<code>IRQ_TYPE</code>	31:0	r/w	Trigger type select (0 = level trigger, 1 = edge trigger); <code>IRQ_TYPE(i)</code> corresponds to <code>gpio_i(i)</code>
<code>0xffffc0014</code>	<code>IRQ_POLARITY</code>	31:0	r/w	Trigger polarity select (0 = low-level/falling-edge, 1 = high-level/rising-edge); <code>IRQ_POLARITY(i)</code> corresponds to <code>gpio_i(i)</code>
<code>0xffffc0018</code>	<code>IRQ_ENABLE</code>	31:0	r/w	Per-pin interrupt enable; <code>IRQ_ENABLE(i)</code> corresponds to <code>gpio_i(i)</code>

Address	Name [C]	Bit(s)	R/W	Function
0xffffc001c	IRQ_PENDING	31:0	r/c	Per-pin interrupt pending, can be cleared by writing zero to the according bit(s); IRQ_PENDING(i) corresponds to gpio_i(i)

2.8.10. Watchdog Timer (WDT)

Hardware source files:	neorv32_wdt.vhd	
Software driver files:	neorv32_wdt.c neorv32_wdt.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	rstn_wdt_o	synchronous watchdog reset output, low-active
Configuration generics:	IO_WDT_EN	implement watchdog when <code>true</code>
CPU interrupts:	none	

Key Features

- 32-bit timeout counter with 3-bit clock divider
- Password-protected; lockable configuration
- Provides information about the cause of the last reset

Overview

The watchdog (WDT) provides a last resort for safety-critical applications. When a pre-programmed timeout value is reached a system-wide hardware reset is generated. The internal counter has to be reset explicitly by the application program every now and then to prevent a timeout.

Theory of Operation

The watchdog is enabled by setting the control register's `WDT_CTRL_EN` bit. When this bit is cleared, the internal timeout counter is reset to zero and no system reset can be triggered by this module.

The internal 24-bit timeout counter is clocked at 1/4096 of the processor's main clock. Whenever this counter reaches the programmed timeout value (24-bit `WDT_CTRL_TIMEOUT` value) a hardware reset is triggered. The timeout counter is reset by writing the reset **PASSWORD** to the **RESET** register ("feeding the watchdog"). The password is hardwired to hexadecimal `0x709D1AB3`.

$$T_{\text{timeout}} = (1 / f_{\text{main}}[\text{Hz}]) * 4096 * \text{WDT_CTRL_TIMEOUT}$$



Operation in Sleep-Mode and During Debugging

Once enabled, the watchdog keeps operating even if the CPU is in [Sleep Mode](#) or if the processor is being debugged via the [On-Chip Debugger \(OCD\)](#).

Configuration Lock

The watchdog control register can be *locked* to protect the current configuration from being modified. The lock is activated by setting the `WDT_CTRL_LOCK` bit. The lock bit can only be cleared by a *hardware* reset

In the locked state any write access to the control register will trigger an immediate hardware reset (read accesses are still possible and have no side effects). Additionally, writing an incorrect password to the **RESET** register will also trigger an immediate hardware reset.

Cause of last Hardware Reset

The cause of the last system hardware reset can be determined via the two **WDT_CTRL_RCAUSE** bits:

- **WDT_RCAUSE_EXT** (0b00): Reset caused by external reset signal pin
- **WDT_RCAUSE_OCD** (0b01): Reset caused by on-chip debugger
- **WDT_RCAUSE_TMO** (0b10): Reset caused by watchdog timeout
- **WDT_RCAUSE_ACC** (0b11): Reset caused by illegal watchdog access

External Reset Output

The WDT provides a dedicated output ([Processor Top Entity - Signals: rstn_wdt_o](#)) to reset processor-external modules when the watchdog times out. This signal is low-active and synchronous to the processor clock. It is available if the watchdog is implemented; otherwise it is hardwired to 1. Note that the signal also becomes active (low) when the processor's main reset signal is active (even if the watchdog is deactivated or disabled for synthesis).

Register Map

Table 15. WDT register map (struct NEORV32_WDT)

Address	Name [C]	Bit(s), Name [C]	R/W	Reset value	Writable if locked	Function
0xffffb0000	CTRL	0 WDT_CTRL_EN	r/w	0	no	Watchdog enable
		1 WDT_CTRL_LOCK	r/w	0	no	Lock configuration when set, clears only on system reset
		3:2 WDT_CTRL_RCAUSE_HI : WDT_CTRL_RCAUSE_LO	r/-	0	-	Cause of last system reset
		7:4 -	r/-	-	-	<i>reserved</i> , reads as zero
		31:8 WDT_CTRL_TIMEOUT_MSB : WDT_CTRL_TIMEOUT_LSB	r/w	0	no	Timeout value (24-bit)
0xffffb0004	RESET	31:0	-/w	-	yes	Write PASSWORD to reset WDT timeout counter

2.8.11. Core-Local Interruptor (CLINT)

Hardware source files:	neorv32_clint.vhd	
Software driver files:	neorv32_clint.c neorv32_clint.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>mtime_irq_i</code> <code>msw_irq_i</code> <code>mtime_time_o</code>	RISC-V machine timer IRQ if CLINT is not implemented RISC-V software IRQ if CLINT is not implemented Current system time (from CLINT's MTIMER)
Configuration generics:	<code>IO_CLINT_EN</code>	implement core local interruptor when <code>true</code>
CPU interrupts:	<code>MTI</code> <code>MSI</code>	machine timer interrupt (see Processor Interrupts) machine software interrupt (see Processor Interrupts)

Key Features

- Compatible to the RISC-V CLINT specification
- Global 64-bit system time
- Timer and software interrupts for each CPU core

Overview

The core local interruptor provides machine-level timer and software interrupts for a set of CPU cores (also called *harts*). *It is compatible to the original SiFive® CLINT specifications and it is also backwards-compatible to the upcoming RISC-V _Advanced Core Local Interruptor (ACLINT) specifications.* In terms of the ACLINT spec the NEORV32 CLINT implements three *devices* that are placed next to each other in the address space: an MTIMER and an MSWI device.

The CLINT can support up to 4095 harts. However, the NEORV32 CLINT is configured for a single hart only (yet). Hence, only the according registers are implemented while the remaining ones are hardwired to zero.

MTIMER Device

The MTIMER device provides a global 64-bit machine timer (`NEORV32_CLINT→MTIME`) that increments with every main processor clock tick. Upon reset the timer is reset to all zero. Each hart provides an individual 64-bit timer-compare register (`NEORV32_CLINT→MTIMECMP[0]` for hart 0). Whenever `MTIMECMP >= MTIME` the according machine timer interrupt is pending.

MSIW Device

The MSIV provides software interrupts for each hart via hart-individual memory-mapped registers (`NEORV32_CLINT→MSWI[0]` for hart 0). Setting bit 0 of this register will bring the machine software interrupt into pending state.

External Machine Timer and Software Interrupts



If the NEORV32 CLINT module is disabled (`IO_CLINT_EN = false`) the core's machine timer interrupt and machine software interrupt become available as processor-external signals (`mtime_irq_i` and `msw_irq_i`, respectively).

Register Map

Table 16. CLINT register map (`struct NEORV32_CLINT`)

Address	Name [C]	Bits	R/W	Function
<code>0xffff40000</code>	<code>MSWI[0]</code>	0	r/w	trigger machine software interrupt for hart 0 when set
		31:1	r/-	hardwired to zero
<code>0xffff40004</code>	<code>MSWI[1]</code>	0	r/w	trigger machine software interrupt for hart 1 when set
		31:1	r/-	hardwired to zero
<code>0xffff44000</code>	<code>MTIMECMP[0]</code>	63:0	r/w	64-bit time compare for hart 0
<code>0xffff44008</code>	<code>MTIMECMP[1]</code>	63:0	r/w	64-bit time compare for hart 1
<code>0xffff4bff8</code>	<code>MTIME</code>	63:0	r/w	64-bit global machine timer

2.8.12. Primary Universal Asynchronous Receiver and Transmitter (UART0)

Hardware source files:	neorv32_uart.vhd	
Software driver files:	neorv32_uart.c neorv32_uart.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>uart0_txd_o</code> <code>uart0_rxd_i</code> <code>uart0_rtsn_o</code> <code>uart0_ctsn_i</code>	serial transmitter output serial receiver input flow control: RX ready to receive, low-active flow control: RX ready to receive, low-active
Configuration generics:	<code>IO_UART0_EN</code> <code>UART0_RX_FIFO</code> <code>UART0_TX_FIFO</code>	implement UART0 when <code>true</code> RX FIFO depth (power of 2, min 1) TX FIFO depth (power of 2, min 1)
CPU interrupts:	fast IRQ channel 2	Programmable FIFO status interrupt (see Processor Interrupts)

Key Features

- Independent RX and TX lines
- Fixed format: 8 data bits, 1 stop bit, no parity bit
- Programmable baud rate
- Optional RX and TX FIFO buffers
- Optional support for hardware flow-control
- Interrupt based on FIFO buffer status

Overview

The NEORV32 UART provides a standard universal asynchronous serial interface with independent transmitter and receiver channels, each equipped with a configurable FIFO. The transmission frame is fixed to **8-N-1**: 8 data bits, no parity bit, 1 stop bit. The actual transmission baud rate is programmable via software.

Standard Console



All default example programs and software libraries of the NEORV32 software framework (including the bootloader and the runtime environment) use the primary UART (UART0) as default user console interface. Furthermore, UART0 is used to implement the standard consoles `STDIN`, `STDOUT` and `STDERR`.

Theory of Operation

The module is enabled by setting the `UART_CTRL_EN` bit in the control register `CTRL`. A new TX transmission is started by writing to the `DATA` register. RX data is available via the `DATA` register. The baud rate is configured via a 10-bit `UART_CTRL_BAUDx` baud divisor (`baud_div`) and a 3-bit `UART_CTRL_PRSCx` clock prescaler select (`clock_prescaler`).

Table 17. UART0 Clock Configuration

<code>UART_CTRL_PRSC[2:0]</code>	0b000	0b001	0b010	0b011	0b100	0b101	0b110	0b111
Resulting <code>clock_prescaler</code>	2	4	8	64	128	1024	2048	4096

$$\text{Baud rate} = (f_{\text{main}}[\text{Hz}] / \text{clock_prescaler}) / (\text{baud_div} + 1)$$

RX and TX FIFOs

The UART provides individual data FIFOs for RX and TX to allow data transmission without CPU intervention. The sizes of these FIFOs can be configured via the according configuration generics (`UART0_RX_FIFO` and `UART0_TX_FIFO`). Write to `DATA` and reads from `DATA` are automatically buffered by the according FIFO. Both FIFOs are automatically cleared when disabling the module via the `UART_CTRL_EN` flag. The control register's `UART_CTRL_RX_*` and `UART_CTRL_TX_*` flags provide information about the RX and TX FIFO fill level.

RX/TX FIFO Size



Software can retrieve the configured sizes of the RX and TX FIFO via the according `UART_DATA_RX_FIFO` and `UART_DATA_TX_FIFO` bits from the `DATA` register.

UART Interrupt

The UART module provides a single interrupt line that can be used to signal certain RX/TX data FIFO conditions. The interrupt conditions are based on the RX/TX FIFO status flags (`UART_CTRL_RX_*` / `UART_CTRL_TX_*`) and are configured via the according `UART_CTRL_IRQ_RX_*` / `UART_CTRL_IRQ_TX_*` bits. The UART interrupt will fire when the module is enabled (`SLINK_CTRL_EN`) and **any** of the enabled interrupt conditions is met. Hence, all enabled interrupt conditions are logically OR-ed. The interrupt remains active until all interrupt-causing conditions are resolved.

RTS/CTS Hardware Flow Control

The NEORV32 UART supports optional hardware flow control using the standard CTS `uart0_ctsn_i` ("clear to send") and RTS `uart0_rtsn_o` ("ready to send" / "ready to receive (RTR)") signals. Both signals are low-active. Hardware flow control is enabled by setting the `UART_CTRL_HWFC_EN` bit in the modules control register `CTRL`.

When hardware flow control is enabled:

- The UART's transmitter will not start a new transmission until the `uart0_ctsn_i` signal goes low. During this time the UART busy flag `UART_CTRL_TX_BUSY` remains set.
- The UART will set `uart0_rtsn_o` signal low if the RX FIFO is not already full. `uart0_rtsn_o` is

always low if hardware flow-control is disabled. Disabling the UART (setting `UART_CTRL_EN` low) while having hardware flow-control enabled, will set `uart0_rtsn_o` high to indicate that the UART is not capable of receiving data.

Simulation Mode

The simulation mode of the UART allows to redirect TX data to the simulator console instead of sending it via the physical `uart0_txd_o` signal. Simulation mode is enabled by setting the `UART_CTRL_SIM_MODE` bit. When enabled, all data written to the `DATA` register is sent to the simulator and printed as ASCII characters to the simulator console. Note that this feature is only available within a simulation.

Register Map

Table 18. UART0 register map (struct `NEORV32_UART0`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffff50000	CTRL	0 <code>UART_CTRL_EN</code>	r/w	UART enable
		1 <code>UART_CTRL_SIM_MODE</code>	r/w	enable simulation mode
		2 <code>UART_CTRL_HWFC_EN</code>	r/w	enable RTS/CTS hardware flow-control
		5:3 <code>UART_CTRL_PRSC_MSB : UART_CTRL_PRSC_LSB</code>	r/w	baud rate clock prescaler select
		15:6 <code>UART_CTRL_BAUD_MSB : UART_CTRL_BAUD_LSB</code>	r/w	12-bit baud value configuration value
		16 <code>UART_CTRL_RX_NEMPTY</code>	r/-	RX FIFO not empty (data available)
		17 <code>UART_CTRL_RX_FULL</code>	r/-	RX FIFO full
		18 <code>UART_CTRL_TX_EMPTY</code>	r/-	TX FIFO empty
		19 <code>UART_CTRL_TX_NFULL</code>	r/-	TX FIFO not full
		20 <code>UART_CTRL_IRQ_RX_NEMPTY</code>	r/w	fire RX-IRQ if RX FIFO not empty
		21 <code>UART_CTRL_IRQ_RX_FULL</code>	r/w	fire RX-IRQ if RX FIFO full
		22 <code>UART_CTRL_IRQ_TX_EMPTY</code>	r/w	fire TX-IRQ if TX FIFO empty
		23 <code>UART_CTRL_IRQ_TX_NHALF</code>	r/w	fire TX-IRQ if TX not full
		29:24 -	r/-	<i>reserved</i> , read as zero
		30 <code>UART_CTRL_RX_OVER</code>	r/-	RX FIFO overflow; cleared by disabling the module
		31 <code>UART_CTRL_TX_BUSY</code>	r/-	TX busy or TX FIFO not empty

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffff50004	DATA	7:0 UART_DATA_RTX_MSB : UART_DATA_RTX_LSB	r/w	receive/transmit data
		11:8 UART_DATA_RX_FIFO_MSB : UART_DATA_RX_FIFO_LSB	r/-	log2(RX FIFO size)
		15:12 UART_DATA_TX_FIFO_MSB : UART_DATA_TX_FIFO_LSB	r/-	log2(TX FIFO size)
		31:16	r/-	<i>reserved</i> , read as zero

2.8.13. Secondary Universal Asynchronous Receiver and Transmitter (UART1)

Hardware source files:	neorv32_uart.vhd	
Software driver files:	neorv32_uart.c neorv32_uart.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>uart1_txd_o</code> <code>uart1_rxd_i</code> <code>uart1_rtsn_o</code> <code>uart1_ctsn_i</code>	serial transmitter output serial receiver input flow control: RX ready to receive, low-active flow control: RX ready to receive, low-active
Configuration generics:	<code>IO_UART1_EN</code> <code>UART1_RX_FIFO</code> <code>UART1_TX_FIFO</code>	implement UART1 when <code>true</code> RX FIFO depth (power of 2, min 1) TX FIFO depth (power of 2, min 1)
CPU interrupts:	fast IRQ channel 3	Programmable FIFO status interrupt (see Processor Interrupts)

Overview

The secondary UART (UART1) is functionally identical to the primary UART ([Primary Universal Asynchronous Receiver and Transmitter \(UART0\)](#)). UART1 uses different addresses for the control register (`CTRL`) and the data register (`DATA`) and uses a different CPU fast interrupt (FIRQ) channel.

Register Map

Table 19. UART1 register map (`struct NEORV32_UART1`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
<code>0xffff60000</code>	<code>CTRL</code>	Same as UART0
<code>0xffff60004</code>	<code>DATA</code>	Same as UART0

2.8.14. Serial Peripheral Interface Controller (SPI)

Hardware source files:	neorv32_spi.vhd	
Software driver files:	neorv32_spi.c neorv32_spi.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>spi_clk_o</code> <code>spi_dat_o</code> <code>spi_dat_i</code> <code>spi_csn_o</code>	1-bit serial clock output 1-bit serial data output 1-bit serial data input 8-bit dedicated chip select output (low-active)
Configuration generics:	<code>IO_SPI_EN</code> <code>IO_SPI_FIFO</code>	implement SPI controller when <code>true</code> FIFO depth, has to be a power of two, min 1
CPU interrupts:	fast IRQ channel 6	configurable SPI interrupt (see Processor Interrupts)

Key Features

- SPI-compatible host controller
- Programmable clock phase and polarity
- Fine-grained programmable clock
- 8 dedicated chip-select lines
- Optional data/command FIFO (ring-buffer)
- Interrupt if programmed transfer sequences have completed

Overview

The NEORV32 SPI module provides a **host-mode** serial peripheral interface. The module operates on byte-wide data, supports all 4 standard SPI clock modes, provides a precise SPI clock generator and implements 8 dedicated chip select signals via the top entity's `spi_csn_o` signal. An optional receive/transmit ring-buffer/FIFO can be configured via the `IO_SPI_FIFO` generic to support programming of complete SPI transmissions without CPU interaction.

Host-Mode Only



The NEORV32 SPI module only supports *host mode*. Transmission are initiated only by the SPI module itself. If you are looking for a *device-mode* serial peripheral interface (transactions initiated by an external host) check out the [Serial Data Interface Controller \(SDI\)](#).

Theory of Operation

The SPI module provides a single control register [CTRL](#) to configure the module and to check its status and a single data register [DATA](#) for receiving/transmitting data and for issuing chip-select commands.

The SPI module is enabled by setting the [SPI_CTRL_EN](#) bit in the [CTRL](#) control register. No transfer can be initiated and no interrupt request will be triggered if this bit is cleared. Clearing this bit resets the entire module, clears the RX/TX FIFO and terminates any transfer being in process.

The actual SPI transfer (receiving one byte while also sending one byte) as well as control of the chip-select lines is handled by the module's [DATA](#) register. Note that this register will access the TX FIFO of the ring-buffer when writing to it and will access the RX FIFO of the ring-buffer when reading from it.

The most significant bit of the [DATA](#) register ([SPI_DATA_CMD](#)) is used to select the purpose of the data being written. When the [SPI_DATA_CMD](#) is cleared, the lowest 8-bit represent the actual SPI TX data that will be transmitted by the SPI engine. After completion, the according receive data is pushed to the RX FIFO.

If [SPI_DATA_CMD](#) is set, the lowest 4-bit control the chip-select lines. In this case, bits [2:0](#) select one of the eight chip-select lines. The selected line will become enabled when bit [3](#) is set. If bit [3](#) is cleared, all chip-select lines will be disabled at once. Note that the bits of the [spi_csn_o](#) port are low-active. Only one chip-select can be active at a time.

Since all SPI operations are controlled via the FIFOs, entire SPI sequences (chip-enable, data transmission(s), chip-disable) can be *programmed*. Thus, SPI operations can be executed without any CPU interaction at all and can also be coordinated by the system's DMA controller.

SPI Clock Configuration

The SPI module supports all standard SPI clock modes (0, 1, 2, 3), which are configured via the two control register bits [SPI_CTRL_CPHA](#) and [SPI_CTRL_CPOL](#). The [SPI_CTRL_CPHA](#) bit defines the *clock phase* and the [SPI_CTRL_CPOL](#) bit defines the *clock polarity*.

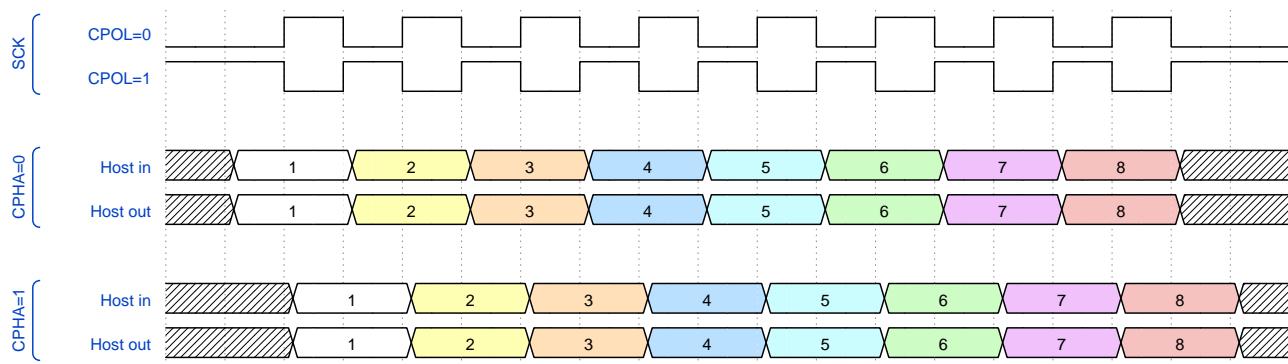


Figure 8. SPI Clock Modes

The SPI clock frequency ([spi_clk_o](#)) is programmed by the 3-bit [SPI_CTRL_PRSCx](#) clock prescaler for a coarse clock selection and a 4-bit clock divider [SPI_CTRL_CDIVx](#) for a fine clock configuration. The following clock prescalers ([SPI_CTRL_PRSCx](#)) are available:

Table 20. SPI prescaler configuration

SPI_CTRL_PRSC[2:0]	0b000	0b001	0b010	0b011	0b100	0b101	0b110	0b111
Resulting <code>clock_prescaler</code>	2	4	8	64	128	1024	2048	4096

Based on the programmed clock configuration, the actual SPI clock frequency f_{SPI} is derived from the processor's main clock f_{main} according to the following equation:

$$f_{\text{SPI}} = f_{\text{main}}[\text{Hz}] / (2 * \text{clock_prescaler} * (1 + \text{SPI_CTRL_CDIVx}))$$

Hence, the maximum SPI clock is $f_{\text{main}} / 4$ and the lowest SPI clock is $f_{\text{main}} / 131072$. The SPI clock is always symmetric having a duty cycle of exactly 50%.

SPI Interrupt

The SPI module provides a single interrupt that gets triggered when the programmed SPI sequence has completed (i.e. the TX FIFO is empty and the SPI engine is idle).

Register Map

Table 21. SPI register map (`struct NEORV32_SPI`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffff80000	CTRL	0 SPI_CTRL_EN	r/w	SPI module enable
		1 SPI_CTRL_CPHA	r/w	clock phase
		2 SPI_CTRL_CPOL	r/w	clock polarity
		5:3 SPI_CTRL_PRSC2 : SPI_CTRL_PRSC0	r/w	3-bit clock prescaler select
		9:6 SPI_CTRL_CDIV3 : SPI_CTRL_CDIV0	r/w	4-bit clock divider for fine-tuning
		15:10 -	r/-	<i>reserved</i> , read as zero
		16 SPI_CTRL_RX_AVAIL	r/-	RX FIFO data available (RX FIFO not empty)
		17 SPI_CTRL_TX_EMPTY	r/-	TX FIFO empty
		18 SPI_CTRL_TX_FULL	r/-	TX FIFO full
		23:19 -	r/-	<i>reserved</i> , read as zero
		27:24 SPI_CTRL_FIFO_MSB : SPI_CTRL_FIFO_LSB	r/-	FIFO depth; log2(<code>IO_SPI_FIFO</code>)
		29:28 -	r/-	<i>reserved</i> , read as zero
		30 SPI_CS_ACTIVE	r/-	Set if any chip-select line is active
		31 SPI_CTRL_BUSY	r/-	SPI module busy when set (serial engine operation in progress and TX FIFO not empty yet)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffff80004	DATA	7:0 SPI_DATA_MSB : SPI_DATA_LSB	r/w	receive/transmit data (FIFO), only for data mode (<code>SPI_DATA_CMD = 0</code>)
		3:0	-/w	chip-select-enable (bit 3) and chip-select (bit 2:0), only for command mode (<code>SPI_DATA_CMD = 1</code>)
		30:8 -	r/-	<i>reserved</i> , read as zero
		31 SPI_DATA_CMD	-/w	0 = data, 1 = chip-select-command

2.8.15. Serial Data Interface Controller (SDI)

Hardware source files:	neorv32_sdi.vhd	
Software driver files:	neorv32_sdi.c neorv32_sdi.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>sdi_clk_i</code> <code>sdi_dat_o</code> <code>sdi_dat_i</code> <code>sdi_csn_i</code>	1-bit serial clock input 1-bit serial data output 1-bit serial data input 1-bit chip-select input (low-active)
Configuration generics:	<code>IO_SDI_EN</code> <code>IO_SDI_FIFO</code>	implement SDI controller when <code>true</code> data FIFO size, has to a power of two, min 1
CPU interrupts:	fast IRQ channel 11	configurable SDI interrupt (see Processor Interrupts)

Key Features

- SPI-compatible device-side controller
- Programmable SPI clock polarity
- Optional data RX/TX FIFO
- Interrupt based on FIFO status

Overview

The serial data interface module provides a **device-side** SPI interface to receive communications from an **external SPI host**, which is responsible of performing the actual transmission (i.e. the host generates the SPI clock and control the chip-select line). An optional receive/transmit FIFO can be configured via the `IO_SDI_FIFO` generic to support transmissions without CPU interaction.

Device-Mode Only



The NEORV32 SDI module only supports *device mode*. If you are looking for a *host-mode* serial peripheral interface (transactions performed by the NEORV32) check out the [Serial Peripheral Interface Controller \(SPI\)](#).

The SDI module provides a single control register `CTRL` to configure the module and to check it's status and a single data register `DATA` for receiving/transmitting data. Any access to the `DATA` register actually accesses the internal RX/TX FIFO.

Theory of Operation

The SDI module is enabled by setting the `SDI_CTRL_EN` bit in the `CTRL` control register. Clearing this bit resets the entire module and will also clear the entire RX/TX FIFO.

The SDI operates on byte-level. Data bytes written to the `DATA` register will be pushed to the internal TX FIFO. This TX data will be sent back to the external host during an SPI transfer. If no data is available in the TX FIFO all-zero is sent. Received bytes are pushed to the RX FIFO and can be retrieved by reading the RX FIFO via the `DATA` register. Data is always transferred MSB-first. The current state of these FIFOs is available via the control register's `SDI_CTRL_RX_*` and `SDI_CTRL_TX_*` status flags. The RX/TX FIFOs can be cleared at any time by the `SDR_CTRL_CLR_*` control register bits.

Data is only transferred (and pushed to the RX FIFO / popped from the TX FIFO) when the chip-select input `sdi_csn_i` is active (low-active). If the external SPI host aborts the transmission by setting the chip-select signal high again *before* 8 data bits have been transferred, no data is written to the RX FIFO and no data is retrieved from the TX FIFO.

SDI Clocking

The SDI module supports both SPI clock polarity modes ("CPOL"), but only "CPHA=0"-clock-phase is *officially* supported yet. However, experiments have shown that the SDI module can also deal with both clock phase modes.

All SDI operations are clocked by the external `sdi_clk_i` signal. This signal is synchronized to the processor's clock domain to simplify timing behavior. This clock synchronization requires the external SDI clock (`sdi_clk_i`) to not **not exceed 1/4 of the processor's main clock**.

SDI Interrupt

The SDI module provides a single interrupt that is triggered by a set of programmable interrupt conditions that are based on the status of the RX & TX FIFOs. The different interrupt sources are enabled by setting the according `SDI_CTRL_IRQ_*` bits. All enabled interrupt conditions are logically OR-ed - so any enabled interrupt source will trigger the module's interrupt signal. Once the SDI interrupt has fired it will remain active until the actual cause of the interrupt is resolved.

Register Map

Table 22. SDI register map (`struct NEORV32_SDI`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffff70000	CTRL	0 SDI_CTRL_EN	r/w	SDI module enable
		1 SDR_CTRL_CLR_RX	r/-	clear RX FIFO, flag auto-clears
		2 SDR_CTRL_CLR_TX	r/-	clear TX FIFO, flag auto-clears
		3 -	r/-	<i>reserved</i> , read as zero
		7:4 SDI_CTRL_FIFO_MSB : SDI_CTRL_FIFO_LSB	r/-	FIFO depth; $\log_2(LO_SDI_FIFO)$
		15:8 -	r/-	<i>reserved</i> , read as zero
		16 SDI_CTRL_IRQ_RX_NEMPTY	r/w	fire interrupt if RX FIFO is not empty
		17 SDI_CTRL_IRQ_RX_FULL	r/w	fire interrupt if RX FIFO is full
		18 SDI_CTRL_IRQ_TX_EMPTY	r/w	fire interrupt if TX FIFO is empty
		23:19 -	r/-	<i>reserved</i> , read as zero
		24 SDI_CTRL_RX_EMPTY	r/-	RX FIFO empty
		25 SDI_CTRL_RX_FULL	r/-	RX FIFO full
		26 SDI_CTRL_TX_EMPTY	r/-	TX FIFO empty
		27 SDI_CTRL_TX_FULL	r/-	TX FIFO full
		30:28 -	r/-	<i>reserved</i> , read as zero
		31 SDI_CTRL_CS_ACTIVE	r/-	Chip-select is active when set
0xffff70004	DATA	7:0	r/w	receive/transmit data (FIFO)
		31:8 -	r/-	<i>reserved</i> , read as zero

2.8.16. Two-Wire Serial Interface Controller (TWI)

Hardware source files:	neorv32_twi.vhd	
Software driver files:	neorv32_twi.c neorv32_twi.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>twi_sda_i</code> <code>twi_sda_o</code> <code>twi_scl_i</code> <code>twi_scl_o</code>	1-bit serial data line sense input 1-bit serial data line output (pull low only) 1-bit serial clock line sense input 1-bit serial clock line output (pull low only)
Configuration generics:	<code>IO_TWI_EN</code> <code>IO_TWI_FIFO</code>	implement TWI controller when <code>true</code> FIFO depth, has to be a power of two, min 1
CPU interrupts:	fast IRQ channel 7	FIFO empty and module idle interrupt (see Processor Interrupts)

Key Features

- Philips I²C-compatible host controller
- Programmable clock speed
- Support for host-ACK generated by the controller
- Optional support for clock stretching
- Optional data/command FIFO
- Interrupt based on FIFO status

Overview

The NEORV32 TWI implements an I²C-compatible host controller to communicate with arbitrary I²C-devices. Note that multi-controller mode and bus arbitration are not supported.

Host-Mode Only



The NEORV32 TWI controller only supports **host mode**. Transmission are initiated by the processor's TWI controller and not by an external I²C module. If you are looking for a *device-mode* module (transactions initiated by an external host) check out the [Two-Wire Serial Device Controller \(TWD\)](#).

The TWI controller provides two memory-mapped registers that are used for configuration & status check (**CTRL**) and for accessing transmission data (**DATA**). The **DATA** register is transparently buffered by separate RX and TX FIFOs. The size of those FIFOs can be configured by the `IO_TWI_RX_FIFO` and `IO_TWI_TX_FIFO` generics. Software can determine the FIFO size via the control register's

`TWI_CTRL_FIFO_*` bits. The current status of the RX and TX FIFO can be polled by software via the `TWD_CTRL_RX_*` and `TWI_CTRL_TX_*` flags.

The module is globally enabled by setting the control register's `TWI_CTRL_EN` bit. Clearing this bit will disable and reset the entire module also clearing the internal RX and TX FIFOs.

Current Bus State



The current state of the I²C bus lines (SCL and SDA) can be checked by software via the `TWI_CTRL_SENSE_*` control register bits. Note that the TWI module needs to be enabled in order to sample the bus state.

TWI Clock Speed

The TWI clock frequency is programmed by two bit-fields in the device's control register `CTRL`: a 3-bit clock prescaler (`TWI_CTRL_PRSCx`) is used for a coarse clock configuration and a 4-bit clock divider (`TWI_CTRL_CDIVx`) is used for a fine clock configuration.

Table 23. TWI prescaler configuration

<code>TWI_CTRL_PRSC[2:0]</code>	0b000	0b001	0b010	0b011	0b100	0b101	0b110	0b111
Resulting <code>clock_prescaler</code>	2	4	8	64	128	1024	2048	4096

Based on the clock configuration, the actual TWI clock frequency f_{SCL} is derived from the processor's main clock f_{main} according to the following equation:

$$f_{SCL} = f_{main}[Hz] / (4 * \text{clock_prescaler} * (1 + \text{TWI_CTRL_CDIV}))$$

Hence, the maximum TWI clock is $f_{main} / 8$ and the lowest TWI clock is $f_{main} / 262144$. The generated TWI clock is always symmetric having a duty cycle of exactly 50% (if the clock is not halted by a device during clock stretching).

Clock Stretching



An accessed peripheral can slow down/halt the controller's bus clock by using clock stretching (= actively keeping the SCL line low). The controller will halt operation in this case. Clock stretching is enabled by setting the `TWI_CTRL_CLKSTR` bit in the module's control register `CTRL`.

TWI Transfers

All TWI operations are controlled by the `DCMD` register, which is buffered by a FIFO. This command/data FIFO is internally split into a TX FIFO and a RX FIFO. Writing to `DCMD` will write to the TX FIFO while reading from `DCMD` will read from the RX FIFO. Thus, complete TWI sequences can be programmed and executed without further CPU intervention. The actual operation is selected by a 2-bit value that is written to the register's `TWI_DCMD_CMD_*` bit-field:

- **00**: NOP (no-operation); all further bit-fields in `DCMD` are ignored; no bus operation is executed

- **01**: Generate a (repeated) START condition; all further bit-fields in **DCMD** are ignored
- **10**: Generate a STOP condition; all further bit-fields in **DCMD** are ignored
- **11**: Trigger a data transmission; see below

For any transmission the data to be send has to be written to the register's **TWI_DCMD_MSB** : **TWI_DCMD_LSB** bit-field. If **TWI_DCMD_ACK** is set the controller will generate an ACK by it's own. If **TWI_DCMD_ACK** is cleared the controller will sample ACK/NACK from the accessed device. Vice versa, the device's data response can be read from **TWI_DCMD_MSB** : **TWI_DCMD_LSB** bit-field. The ACK/NACK generated by the device can be read from the **TWI_DCMD_ACK** bit.

The control register's busy flag **TWI_CTRL_BUSY** is set as long as the TX FIFO contains data (i.e. programmed TWI operations that have not been executed yet) or if the TWI bus engine is still processing an operation. An active transmission can be terminated at any time by disabling the TWI module. This will also clear the data/command FIFO.

Tristate Drivers

The TWI module requires two tristate drivers (actually: open-drain drivers - signals can only be actively driven low) for the SDA and SCL lines, which have to be implemented by the user in the setup's top module / IO ring. A generic VHDL example is shown below (here, **sda_io** and **scl_io** are the actual I²C bus lines, which are of type **std_logic**).

Listing 9. TWI VHDL Tristate Driver Example

```
sda_io    <= '0' when (twi_sda_o = '0') else 'Z'; -- drive
scl_io    <= '0' when (twi_scl_o = '0') else 'Z'; -- drive
twi_sda_i <= std_ulogic(sda_io); -- sense
twi_scl_i <= std_ulogic(scl_io); -- sense
```

TWI Interrupt

The TWI module provides a single interrupt to signal "idle condition" to the CPU. The interrupt becomes active when the TWI module is enabled (**TWI_CTRL_EN** = 1) and the TX FIFO is empty and the TWI bus engine is idle.

Register Map

Table 24. TWI register map (struct NEORV32_TWI)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffff90000	CTRL	0 TWI_CTRL_EN	r/w	TWI enable, reset if cleared
		3:1 TWI_CTRL_PRSC2 : TWI_CTRL_PRSC0	r/w	3-bit clock prescaler select
		7:4 TWI_CTRL_CDIV3 : TWI_CTRL_CDIV0	r/w	4-bit clock divider
		8 TWI_CTRL_CLKSTR	r/w	Enable (allow) clock stretching
		14:9 -	r/-	<i>reserved</i> , read as zero
		18:15 TWI_CTRL_FIFO_MSB : TWI_CTRL_FIFO_LSB	r/-	FIFO depth; log2(IO_TWI_FIFO)
		26:12 -	r/-	<i>reserved</i> , read as zero
		27 TWI_CTRL_SENSE_SCL	r/-	current state of the SCL bus line
		28 TWI_CTRL_SENSE_SDA	r/-	current state of the SDA bus line
		29 TWI_CTRL_TX_FULL	r/-	set if the TWI bus is claimed by any controller
		30 TWI_CTRL_RX_AVAIL	r/-	RX FIFO data available
		31 TWI_CTRL_BUSY	r/-	TWI bus engine busy or TX FIFO not empty
0xffff90004	DCMD	7:0 TWI_DCMD_MSB : TWI_DCMD_LSB	r/w	write: TX data byte; read: RX data byte
		8 TWI_DCMD_ACK	r/w	write: ACK issued by controller; read: 1 = device NACK, 0 = device ACK
		10:9 TWI_DCMD_CMD_HI : TWI_DCMD_CMD_LO	-/w	TWI operation (00 = NOP, 01 = START conditions, 10 = STOP condition, 11 = data transmission)
		31:11 -	r/-	<i>reserved</i> , read as zero

2.8.17. Two-Wire Serial Device Controller (TWD)

Hardware source files:	neorv32_twd.vhd	
Software driver files:	neorv32_twd.c neorv32_twd.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>twd_sda_i</code> <code>twd_sda_o</code> <code>twd_scl_i</code> <code>twd_scl_o</code>	1-bit serial data line sense input 1-bit serial data line output (pull low only) 1-bit serial clock line sense input 1-bit serial clock line output (pull low only)
Configuration generics:	<code>IO_TWD_EN</code> <code>IO_TWD_RX_FIFO</code> <code>IO_TWD_TX_FIFO</code>	implement TWD controller when <code>true</code> RX FIFO depth, has to be a power of two, min 1 TX FIFO depth, has to be a power of two, min 1
CPU interrupts:	fast IRQ channel 0	FIFO status interrupt (see Processor Interrupts)

Key Features

- Philips I²C-compatible device-side controller
- Programmable 7-bit device address
- Optional RX/TX data FIFOs
- Programmable interrupt conditions

Overview

The NEORV32 TWD implements an I²C-compatible **device-side** interface. Processor-external I²C hosts can communicate with this module by issuing I²C transactions. The TWD is entirely passive and only reacts on those external transmissions.

Device-Mode Only



The NEORV32 TWD controller only supports **device mode**. Transmission are initiated by processor-external modules and not by an external TWD. If you are looking for a *host-mode* module (transactions initiated by the processor) check out the [Two-Wire Serial Interface Controller \(TWI\)](#).

Theory of Operation

The TWD module provides two memory-mapped registers that are used for configuration & status check (**CTRL**) and for accessing transmission data (**DATA**). The **DATA** register is transparently buffered

by separate RX and TX FIFOs. The size of those FIFOs can be configured by the `IO_TWD_RX_FIFO` and `IO_TWD_TX_FIFO` generics. Software can determine the FIFO size via the control register's `TWD_CTRL_FIFO_*` bits. The current status of the RX and TX FIFO can be polled by software via the `TWD_CTRL_RX_*` and `TWD_CTRL_TX_*` flags.

The module is globally enabled by setting the control register's `TWD_CTRL_EN` bit. Clearing this bit will disable and reset the entire module also clearing the internal RX and TX FIFOs. Each FIFO can also be cleared individually at any time by setting `TWD_CTRL_CLR_RX` or `TWD_CTRL_CLR_TX` control register bits, respectively.

The external two wire bus is sampled and synchronized into the processor's clock domain with a sampling frequency of 1/8 of the processor's main clock. In order to increase the resistance to glitches the sampling frequency can be lowered to 1/64 of the processor clock by setting the control register's `TWD_CTRL_FSEL` bit.

Current Bus State



The current state of the I2C bus lines (SCL and SDA) can be checked by software via the `TWD_CTRL_SENSE_*` control register bits. Note that the TWD module needs to be enabled in order to sample the bus state.

The TWD only responds to a single configurable 7bit device address that is programmed by the `TWD_CTRL_DEV_ADDR` bits. Specific general-call or broadcast addresses are not supported.

Depending on the transaction type, data is either read from the RX FIFO and transferred to the host (**read operation**) or data is received from the host and written to the TX FIFO (**write operation**).

Read Operation: host reads data from TWD

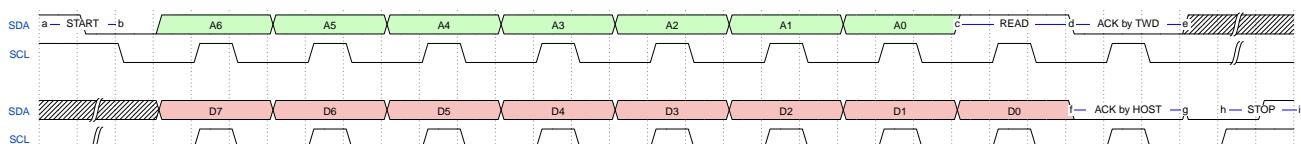


Figure 9. TWD read operation timing (not to scale, split across two lines)

For a read operation the host first generates START or REPEATED-START condition. Then, the host transmits the 7 bit device address (green signals `A6` to `A0`) plus the read-command bit. If the transferred address matches the one programmed to `TWD_CTRL_DEV_ADDR` control register bits the TWD module will response with an ACK by pulling the SDA bus line actively low during the 9th SCL clock pulse. If there is no address match the TWD will not interfere with the bus and will wait for the next START or REPEATED-START condition.

For the actual data transmission the host keeps the SDA line at high state while sending the clock pulses. The TWD will read a byte from the internal TX FIFO and will transmit it MSB-first to the host. During the 9th clock pulse the host has to acknowledged the transfer (ACK) by pulling SDA low. If no ACK is received by the TWD no data is taken from the TX FIFO and the same byte is transmitted in the next data phase. If the TX FIFO becomes empty while the host keeps reading data, all-one bytes are sent back to the host. The transaction is terminated by a STOP condition.

Write Operation: host writes data to TWD

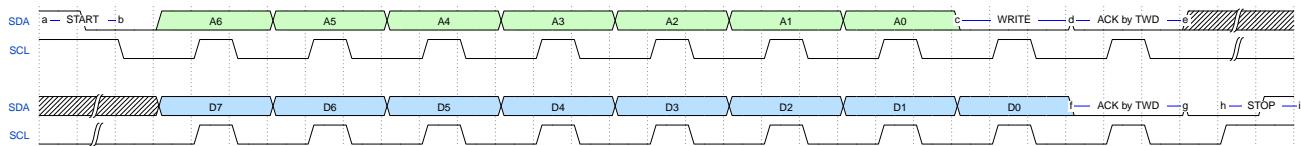


Figure 10. TWD write operation timing (not to scale, split across two lines)

For a write operation the host first generates START or REPEATED-START condition. Then, the host transmits the 7 bit device address (green signals **A6** to **A0**) plus the write-command bit. If the transferred address matches the one programmed to **TWD_CTRL_DEV_ADDR** control register bits the TWD module will response with an ACK by pulling the SDA bus line actively low during the 9th SCL clock pulse. If there is no address match the TWD will not interfere with the bus and will wait for the next START or REPEATED-START condition.

For the actual data transmission the host sends the write-data MSB first together with the clock pulses. During the 9th clock pulse the TWD will respond with an ACK by pulling SDA low. Note that the TWD will respond with a NACK if the RX FIFO is full. The transaction is terminated by a STOP condition.

Tristate Drivers

The TWD module requires two tristate drivers (actually: open-drain drivers - signals can only be actively driven low) for the SDA and SCL lines, which have to be implemented by the user in the setup's top module / IO ring. A generic VHDL example is shown below (here, **sda_io** and **scl_io** are the actual TWD bus lines, which are of type **std_logic**).

Listing 10. TWD Tristate Driver Example (VHDL)

```

sda_io      <= '0' when (twd_sda_o = '0') else 'Z'; -- drive
scl_io      <= '0' when (twd_scl_o = '0') else 'Z'; -- drive
twd_sda_i  <= std_ulogic(sda_io); -- sense
twd_scl_i  <= std_ulogic(scl_io); -- sense

```

TWD Interrupt

The TWD module provides a single interrupt to signal certain FIFO conditions to the CPU. The control register's **TWD_CTRL_IRQ_*** bits are used to enable individual interrupt conditions. Note that all enabled conditions are logically OR-ed.

- **TWD_CTRL_IRQ_RX_AVAIL**: trigger interrupt if at least one data byte is available in the RX FIFO
- **TWD_CTRL_IRQ_RX_FULL**: trigger interrupt if the RX FIFO is completely full
- **TWD_CTRL_IRQ_TX_EMPTY**: trigger interrupt if the TX FIFO is empty

The interrupt remains active until all enabled interrupt-causing conditions are resolved. The interrupt can only trigger if the module is actually enabled (**TWD_CTRL_EN** is set).

Register Map

Table 25. TWD register map (`struct NEORV32_TWD`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
<code>0xffea0000</code>	<code>CTRL</code>	<code>0 TWD_CTRL_EN</code>	r/w	TWD enable, reset if cleared
		<code>1 TWD_CTRL_CLR_RX</code>	-/w	Clear RX FIFO, flag auto-clears
		<code>2 TWD_CTRL_CLR_TX</code>	-/w	Clear TX FIFO, flag auto-clears
		<code>3 TWD_CTRL_FSEL</code>	r/w	Bus sample clock / filter select
		<code>10:4 TWD_CTRL_DEV_ADDR6 : TWD_CTRL_DEV_ADDR0</code>	r/w	Device address (7-bit)
		<code>11 TWD_CTRL_IRQ_RX_AVAIL</code>	r/w	IRQ if RX FIFO data available
		<code>12 TWD_CTRL_IRQ_RX_FULL</code>	r/w	IRQ if RX FIFO full
		<code>13 TWD_CTRL_IRQ_TX_EMPTY</code>	r/w	IRQ if TX FIFO empty
		<code>15:14 -</code>	r/w	<i>reserved</i> , read as zero
		<code>19:16 TWD_CTRL_RX_FIFO_MSB : TWD_CTRL_RX_FIFO_LSB</code>	r/-	FIFO depth; $\log_2(\text{IO_TWD_RX_FIFO})$
		<code>23:20 TWD_CTRL_TX_FIFO_MSB : TWD_CTRL_TX_FIFO_LSB</code>	r/-	FIFO depth; $\log_2(\text{IO_TWD_TX_FIFO})$
		<code>24 -</code>	r/-	<i>reserved</i> , read as zero
		<code>25 TWD_CTRL_RX_AVAIL</code>	r/-	RX FIFO data available
		<code>26 TWD_CTRL_RX_FULL</code>	r/-	RX FIFO full
		<code>27 TWD_CTRL_TX_EMPTY</code>	r/-	TX FIFO empty
		<code>28 TWD_CTRL_TX_FULL</code>	r/-	TX FIFO full
		<code>29 TWD_CTRL_SENSE_SCL</code>	r/-	current state of the SCL bus line
		<code>30 TWD_CTRL_SENSE_SDA</code>	r/-	current state of the SDA bus line
		<code>31 TWD_CTRL_BUSY</code>	r/-	bus engine is busy (transaction in progress)
<code>0xffea0004</code>	<code>DATA</code>	<code>7:0 TWD_DATA_MSB : TWD_DATA_LSB</code>	r/w	RX/TX data FIFO access
		<code>31:8 -</code>	r/-	<i>reserved</i> , read as zero

2.8.18. One-Wire Serial Interface Controller (ONEWIRE)

Hardware source files:	neorv32_onewire.vhd	
Software driver files:	neorv32_onewire.c neorv32_onewire.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	onewire_i onewire_o	1-bit 1-wire bus sense input 1-bit 1-wire bus output (pull low only)
Configuration generics:	IO_ONEWIRE_EN IO_ONEWIRE_FIFO	implement ONEWIRE interface controller when true RTX FIFO depth, has to be zero or a power of two, min 1
CPU interrupts:	fast IRQ channel 13	operation done interrupt (see Processor Interrupts)

Key Features

- Dallas® 1-Wire™-compatible host controller
- Hardware-based protocol handling (RESET, presence, data, etc.)
- Fine-grained programmable bus timings
- Optional data/command FIFO
- Interrupt based on FIFO status

Overview

The NEORV32 ONEWIRE module implements a single-wire interface controller that is compatible to the Dallas/Maxim 1-Wire protocol, which is an asynchronous half-duplex bus requiring only a single signal wire (plus ground) for communication.

The bus is based on a single open-drain signal. The controller as well as all devices on the bus can only pull-down the bus (similar to TWI/I2C). The default high-level is provided by a single pull-up resistor connected to the positive power supply close to the bus controller. Recommended values are between 1kΩ and 10kΩ depending on the bus characteristics (wire length, number of devices, etc.).

Tri-State Drivers

The ONEWIRE module requires a tristate driver (actually, just an open-drain driver) for the 1-wire bus line, which has to be implemented in the top module / IO ring of the design. A generic VHDL example is given below (**onewire_io** is the actual 1-wire bus signal, which is of type **std_logic**; **onewire_o** and **onewire_i** are the processor's ONEWIRE port signals).

Listing 11. ONEWIRE VHDL Tristate Driver Example

```
onewire_io <= '0' when (onewire_o = '0') else 'Z'; -- drive (low)
```

```
onewire_i <= std_ulogic(onewire_io); -- sense
```

Theory of Operation

The ONEWIRE controller provides two interface registers: **CTRL** and **DCMD**. The control register (**CTRL**) is used to configure the module and to monitor the current state. The **DCMD** register, which can optionally be buffered by a configurable FIFO (**IO_ONEWIRE_FIFO** generic), is used to read/write data from/to the bus and to trigger bus operations.

The module is enabled by setting the **ONEWIRE_CTRL_EN** bit in the control register. If this bit is cleared, the module is automatically reset, any bus operation is aborted, the bus is brought to high-level (due to the external pull-up resistor) and the internal FIFO is cleared. The basic timing configuration is programmed via a coarse clock prescaler (**ONEWIRE_CTRL_PRSCx** bits) and a fine clock divider (**ONEWIRE_CTRL_CLKDIVx** bits).

The controller can execute four basic bus operations, which are triggered by writing the according command bits in the **DCMD** register (**ONEWIRE_DCMD_DATA_*** bits) while also writing the actual data bits (**ONEWIRE_DCMD_CMD_*** bits).

1. **0b00 (ONEWIRE_CMD_NOP)** - no operation (dummy)
2. **0b01 (ONEWIRE_CMD_BIT)** - transfer a single-bit (read-while-write)
3. **0b10 (ONEWIRE_CMD_BYTE)** - transfer a full-byte (read-while-write)
4. **0b11 (ONEWIRE_CMD_RESET)** - generate reset pulse and check for device presence

Every command (except NOP) will result in a bus operation when dispatched from the data/command FIFO. Each command (except NOP) will also sample a bus response (a read bit, a read byte or a presence pulse) to a shadowed receive FIFO that is accessed when reading the **DCMD** register.

When the single-bit operation (**ONEWIRE_CMD_BIT**) is executed, the data previously written to **DCMD[0]** will be send to the bus and the response is sampled to **DCMD[7]**. Accordingly, a full-byte transmission (**ONEWIRE_CMD_BYTE**) will send the byte written to **DCMD[7:0]** to the bus and will sample the response to **DCMD[7:0]** (LSB-first). Finally, the reset command (**ONEWIRE_CMD_RESET**) will generate a bus reset and will also sample the "presence pulse" from the device(s) to the **DCMD[ONEWIRE_DCMD_PRESENCE]**.

Read from Bus



In order to read a single bit from the bus **DCMD[0]** has to be set to **1** before triggering the bit transmission operation to allow the accessed device to pull-down the bus. Accordingly, **DCMD[7:0]** has to be set to **0xFF** before triggering the byte transmission operation when the controller shall read a byte from the bus.

As soon as the current bus operation has completed (and there are no further operations pending in the FIFO) the **ONEWIRE_CTRL_BUSY** bit in the control registers clears.

Bus Timing

The control register provides a 2-bit clock prescaler select (`ONEWIRE_CTRL_PRSC`) and a 8-bit clock divider (`ONEWIRE_CTRL_CLKDIV`) for timing configuration. Both are used to define the elementary base time T_{base} . All bus operations are timed using multiples of this elementary base time.

Table 26. ONEWIRE Clock Prescaler Configurations

<code>ONEWIRE_CTRL_PRSC[2:0]</code>	<code>0b00</code>	<code>0b01</code>	<code>0b10</code>	<code>0b11</code>
Resulting <code>clock_prescaler</code>	2	4	8	64

Together with the clock divider value (`ONEWIRE_CTRL_PRSCx` bits = `clock_divider`) the base time is defined by the following formula:

$$T_{base} = (1 / f_{main}[Hz]) * \text{clock_prescaler} * (\text{clock_divider} + 1)$$

Example:

- $f_{main} = 100\text{MHz}$
- clock prescaler select = `0b01` → `clock_prescaler` = 4
- clock divider `clock_divider` = 249

$$T_{base} = (1 / 100000000\text{Hz}) * 4 * (249 + 1) = 10000\text{ns} = 10\mu\text{s}$$

The base time is used to coordinate all bus interactions. Hence, all delays, time slots and points in time are quantized as multiples of the base time T_{base} . The following images show the two basic operations of the ONEWIRE controller: single-bit (0 or 1) transaction and reset with presence detect. Note that the full-byte operations just repeats the single-bit operation eight times. The relevant points in time are shown as *absolute* time points (in multiples of the time base T_{base}) with the falling edge of the bus as reference points.

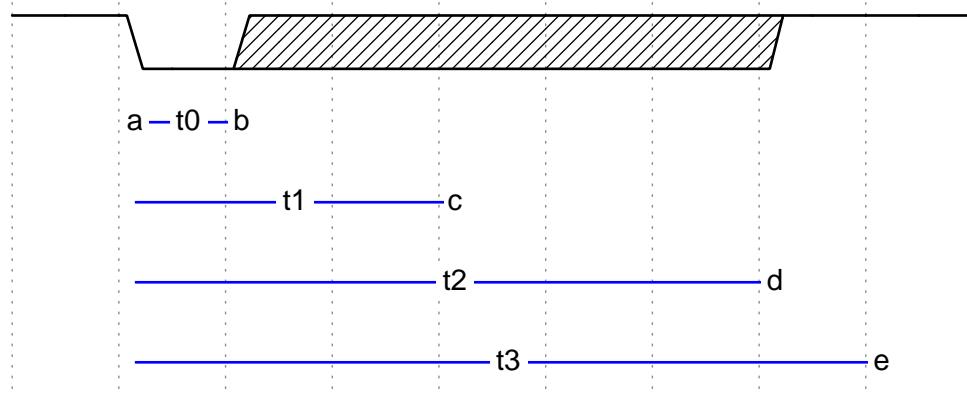


Figure 11. Single-bit data transmission (not to scale)

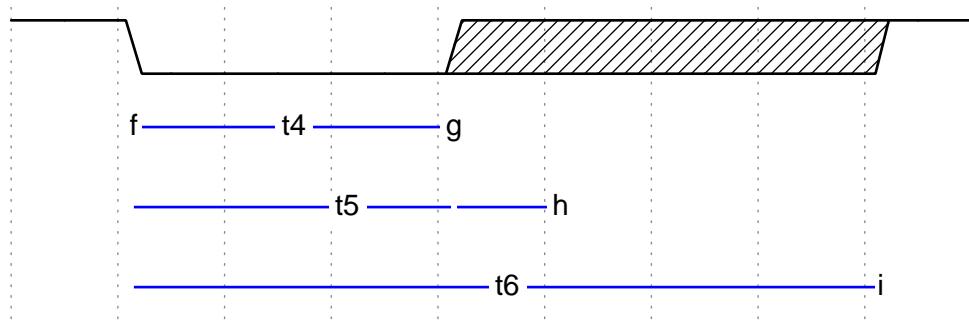


Figure 12. Reset pulse and presence detect (not to scale)

Table 27. Data Transmission Timing

Symbol	Description	Multiples of T_{base}	Time when $T_{base} = 10\mu s$
Single-bit data transmission			
t_0 (a → b)	Time until end of active low-phase when writing a '1' or when reading	1	10μs
t_1 (a → c)	Time until controller samples bus state (read operation)	2	20μs
t_2 (a → d)	Time until end of bit time slot (when writing a '0' or when reading)	7	70μs
t_3 (a → e)	Time until end of inter-slot pause (= total duration of one bit)	9	90μs
Reset pulse and presence detect			
t_4 (f → g)	Time until end of active reset pulse	48	480μs
t_5 (f → h)	Time until controller samples bus presence	55	550μs
t_6 (f → i)	Time until end of presence phase	96	960μs

Default Timing Parameters

The "known-good" default values for base time multiples were chosen for stable and reliable bus operation and not for maximum throughput.

The absolute points in time are hardwired by the VHDL code and cannot be changed during runtime. However, the timing parameter can be customized (if necessary) by editing the ONEWIRE's VHDL source file. The times t0 to t6 correspond to the previous timing diagrams.

Listing 12. Hardwired timing configuration in `neorv32_onewire.vhd`

```
-- timing configuration (absolute time in multiples of the base tick time t_base) --
constant t_write_one_c      : unsigned(6 downto 0) := to_unsigned( 1, 7); -- t0
constant t_read_sample_c    : unsigned(6 downto 0) := to_unsigned( 2, 7); -- t1
constant t_slot_end_c       : unsigned(6 downto 0) := to_unsigned( 7, 7); -- t2
constant t_pause_end_c      : unsigned(6 downto 0) := to_unsigned( 9, 7); -- t3
```

```
constant t_reset_end_c      : unsigned(6 downto 0) := to_unsigned(48, 7); -- t4
constant t_presence_sample_c : unsigned(6 downto 0) := to_unsigned(55, 7); -- t5
constant t_presence_end_c   : unsigned(6 downto 0) := to_unsigned(96, 7); -- t6
```

Overdrive Mode



The ONEWIRE controller does not support the overdrive mode natively. However, it can be implemented by reducing the base time T_{base} (and by eventually changing the hardwired timing configuration in the VHDL source file).

Interrupt

A single interrupt is provided by the ONEWIRE module to signal "idle" condition to the CPU. Whenever the controller is idle (again) and the data/command FIFO is empty, the interrupt becomes active.

Register Map

Table 28. ONEWIRE register map (`struct NEORV32_ONEWIRE`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffff2000	CTRL	0 ONEWIRE_CTRL_EN	r/w	ONEWIRE enable, reset if cleared
		1 ONEWIRE_CTRL_CLEAR	-w	clear RXT FIFO, auto-clears
		3:2 ONEWIRE_CTRL_PRSC1 : ONEWIRE_CTRL_PRSC0	r/w	2-bit clock prescaler select
		11:4 ONEWIRE_CTRL_CLKDIV7 : ONEWIRE_CTRL_CLKDIV0	r/w	8-bit clock divider value
		14:12 -	r-	<i>reserved</i> , read as zero
		18:15 ONEWIRE_CTRL_FIFO_MSB : ONEWIRE_CTRL_FIFO_LSB	r-	FIFO depth; $\log_2(\text{IO_ONEWIRE_FIFO})$
		27:19 -	r-	<i>reserved</i> , read as zero
		28 ONEWIRE_CTRL_TX_FULL	r-	TX FIFO full
		29 ONEWIRE_CTRL_RX_AVAIL	r-	RX FIFO data available
		30 ONEWIRE_CTRL_SENSE	r-	current state of the bus line
		31 ONEWIRE_CTRL_BUSY	r-	operation in progress when set or TX FIFO not empty

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffff20004	DCMD	7:0 ONEWIRE_DCMD_DATA_MSB : ONEWIRE_DCMD_DATA_LSB	r/w	receive/transmit data
		9:8 ONEWIRE_DCMD_CMD_HI : ONEWIRE_DCMD_CMD_LO	-/w	operation command LSBs
		10 ONEWIRE_DCMD_PRESENCE	r/-	bus presence detected
		31:11 -	r/-	<i>reserved</i> , read as zero

2.8.19. Pulse-Width Modulation Controller (PWM)

Hardware source files:	neorv32_pwm.vhd	
Software driver files:	neorv32_pwm.c neorv32_pwm.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>pwm_o</code>	PWM output channels (32-bit)
Configuration generics:	<code>IO_PWM_NUM</code>	number of PWM channels to implement (0..32)
CPU interrupts:	none	

Key Features

- Up to 32 individual channels with up to 16-bit resolution and programmable polarity
- Fast-PWM or phase-correct operation mode with counter-compare and counter-wrap registers
- Toggle-free 0% and 100% duty cycle output rates
- Global clock prescaler

Overview

The PWM module implements a pulse-width modulation controller with up to 32 independent channels. Period length (and thus, the carrier frequency), duty cycle, polarity and operation mode can be programmed individually for each channel. However, the clock for the PWM counter increment is defined by a single global clock prescaler. PWM operation is based on 16-bit wide period counters that are constrained by programmable wrapping values.

The total number of implemented channels is defined by the `IO_PWM_NUM` generic. The PWM output signal `pwm_o` has a static size of 32 bit. Channel 0 corresponds to bit 0, channel 1 to bit 1 and so on. If less than 32 channels are configured, only the LSB-aligned channel bits are connected while the remaining ones are hardwired to zero.

Theory of Operation

The PWM module provides several configuration registers: an `ENABLE` register, a `POLARITY` registers, a `CLKPRSC` register, and a `MODE` register. Each of these register provides one bit (LSB-aligned) for each available PWM channel. Additional, up to `IO_PWM_NUM` channel register `CHANNEL` are provided.

The `ENABLE` register is used to enable individual PWM channels. By using bit-masks several channels can be enabled at once so they operate in perfect lockstep. The PWM counter of a disabled channel is halted and reset to zero. The according PWM output (`pwm_o(i)`) is also reset to 0. The polarity of each channel's output can be inverted by setting the corresponding bit in the `POLARITY` register.

The `CLKPRSC` register configures the *global* clock prescaler that is used for the PWM counter increment of all channels. Eight pre-defined prescaler values are available:

Table 29. PWM prescaler configuration

CLKPRSC	0b000	0b001	0b010	0b011	0b100	0b101	0b110	0b111
Resulting <code>clock_prescaler</code>	2	4	8	64	128	1024	2048	4096

The **MODE** register defines the PWM operation mode of a channel. If bit i is 0 channel i operates in **fast-PWM** mode. If bit i is 1 channel i operates in **phase-correct PWM** mode.

Duty cycle and period length of each channel are defined by the according **CHANNEL**. This register is split in two half-words: the lowest half-word defines the counter-compare value (**CMP**) while the upper half-word defines the counter-wrap value (**TOP**). Regardless of the configured PWM operation mode, the duty cycle of channel i is defined by the following formula:

$$Duty\ Cycle[i] = 100\% * (\text{CMP}[i] / (\text{TOP}[i] + 1))$$

The NEORV32 PWM module supports toggle-free 0% and 100% duty cycle output rates. 0% duty cycle is achieved by setting **CMP = 0**. 100% duty cycle is achieved by setting **CMP = TOP + 1**.

CHANNEL Half-Word Access



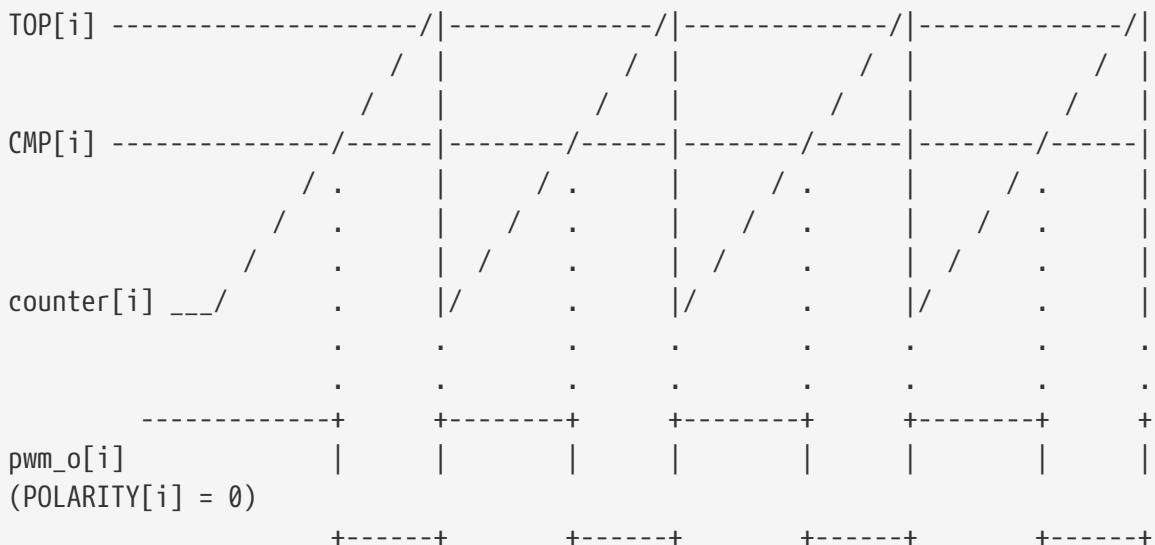
In contrast to many other NEORV32 peripherals, the PWM **CHANNEL[i]** registers can also be accessed using half-word load/store operation. This allows to update the **TOP** and **CMP** sub-registers with a single instruction.

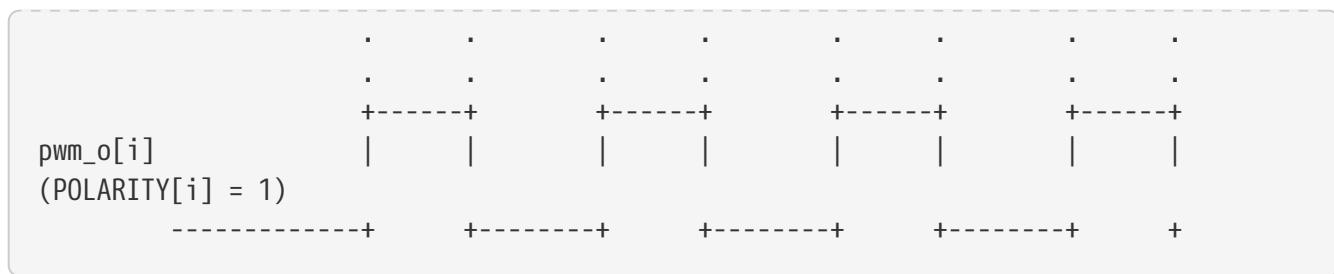
Fast-PWM Mode

When in fast-PWM mode (`MODE[i] = 0`) the channel's PWM counter generates a sawtooth-like waveform. The counter automatically wraps to all-zero when reaching the programmed `TOP[i]` value. Whenever the counter value is less equal to or greater than the `CMP[i]` value the PWM output is set to the programmed polarity (`POLARITY[i]`) value; otherwise it is set to the inverse polarity.

Listing 13. Fast-PWM Operation

MODE[i] = 0





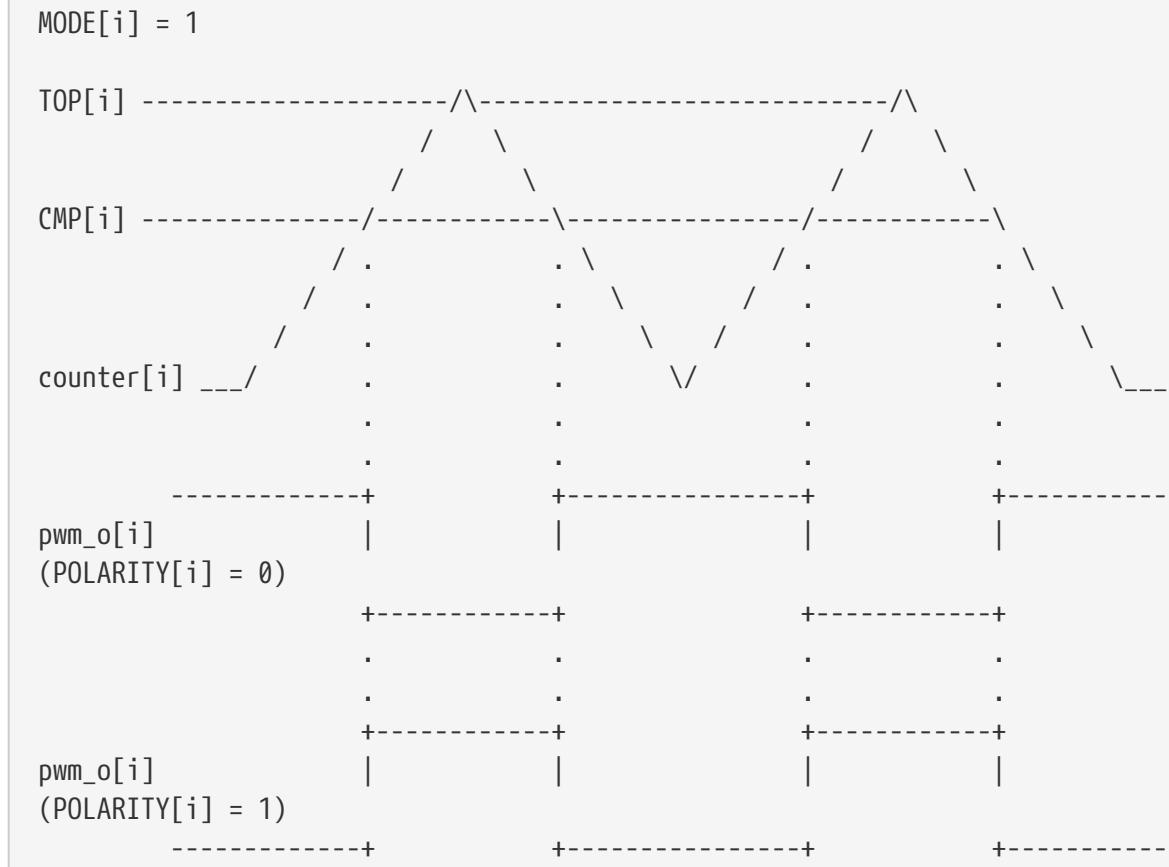
Based on the global clock prescaler and the `TOP` wrap value, the carrier frequency of a channel is defined by the following formula (f_{main} is the main clock frequency of the processor; `clock_prescaler` is the global pre-scaling factor according to the selected `CLKPRSC` value):

$$f_{PWM[i]} = f_{main} / (\text{clock_prescaler} * (\text{TOP}[i] + 1))$$

Phase-Correct PWM Mode

When in phase-correct PWM mode (`MODE[i] = 1`) the channel's PWM counter generates a triangle-like waveform. Phase-correct mode centers the pulse on the same point regardless of the duty cycle. After enabling, the PWM counters counts upward until it reaches the `TOP` value. After that the counter counts downward until it reaches 0 again. Whenever the counter value is equal to or greater than the `CMP[i]` value the PWM output is set to the programmed polarity (`POLARITY[i]`) value; otherwise it is set to the inverse polarity.

Listing 14. Phase-Correct PWM Operation



Based on the global clock prescaler and the **TOP** wrap value, the carrier frequency of a channel is defined by the following formula (f_{main} is the main clock frequency of the processor; **clock_prescaler** is the global pre-scaling factor according to the selected **CLKPRSC** value):

$$f_{PWM[i]} = f_{main} / (\text{clock_prescaler} * 2 * \text{TOP}[i])$$

PWM Dead Time



The phase-correct PWM mode can also be used for controlling modules that require a certain dead time (e.g. H-bridges). This requires two PWM channels, both of which operate in phase-correct mode, but the second channel is configured with an inverse polarity. Both channels need to have the same **TOP** value and must be enabled at once to have a perfect phase and frequency match. The dead time is define by the difference of the channel's **CMP** values:

$$T_{\text{dead}} = 0.5 * T * (\text{CMP}[i] - \text{CMP}[i+1])$$

Register Map

Table 30. PWM register map (struct NEORV32_PWM)

Address	Name [C]	Bit(s)	R/W	Function
0xffff0000	ENABLE	31:0	r/w	Channel enable, one bit per channel
0xffff0004	POLARITY	31:0	r/w	Channel polarity, one bit per channel
0xffff0008	CLKPRSC	2:0	r/w	Global clock prescaler select
0xffff000c	MODE	31:0	r/w	Channel operation mode (0 = fast-PWM, 1 = phase-correct PWM), one bit per channel
0xffff0080	CHANNEL[0].TOPC MP	31:0	r/w	Channel 0 full word access alias (TOP and CMP together at 32-bit value)
0xffff0080	CHANNEL[0].CMP	15:0	r/w	Channel 0 top.wrap value register
0xffff0082	CHANNEL[0].TOP	15:0	r/w	Channel 0 compare value register
...	r/w	...
0xffff00fc	CHANNEL[31].TOP CMP	31:0	r/w	Channel 31 full word access alias (TOP and CMP together at 32-bit value)
0xffff00fc	CHANNEL[31].CMP	15:0	r/w	Channel 31 top.wrap value register
0xffff00fe	CHANNEL[31].TOP	15:0	r/w	Channel 31 compare value register

2.8.20. True Random-Number Generator (TRNG)

Hardware source files:	neorv32_trng.vhd	
Software driver files:	neorv32_trng.c neorv32_trng.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	none	
Configuration generics:	IO_TRNG_EN IO_TRNG_FIFO	implement TRNG when <code>true</code> data FIFO depth, min 1, has to be a power of two
CPU interrupts:	fast IRQ channel 15	data available interrupt (see Processor Interrupts)

Key Features

- True random number generator based on phase noise
- Completely technology-independent architecture
- Optional random data FIFO
- Interrupt based on FIFO status

Overview

The NEORV32 true random number generator provides *physically* true random numbers. It is based on free-running ring-oscillators that generate **phase noise** when being sampled by a constant clock. This phase noise is used as physical entropy source. The TRNG features a platform independent architecture without FPGA-specific primitives, macros or attributes so it can be synthesized for *any* FPGA.

In-Depth Documentation



For more information about the neoTRNG architecture and an analysis of its random quality check out the neoTRNG repository: <https://github.com/stnolting/neoTRNG>

Inferring Latches



The synthesis tool might emit warnings regarding **inferred latches** or **combinatorial loops**. However, this is not design flaw as this is exactly what we want.

Theory of Operation

The TRNG provides two memory mapped interface register. One control register (**CTRL**) for configuration and status check and one data register (**DATA**) for obtaining the random data. The TRNG is enabled by setting the control register's **TRNG_CTRL_EN**. As soon as the **TRNG_CTRL_AVAIL** bit is set a new random data byte is available and can be obtained from the lowest 8 bits of the **DATA**

register. If this bit is cleared, there is no valid data available and the reading **DATA** will return all-zero.

An internal entropy FIFO can be configured using the **IO_TRNG_FIFO** generic. This FIFO automatically samples new random data from the TRNG to provide some kind of *random data pool* for applications which require a larger number of random data in a short time. The random data FIFO can be cleared at any time either by disabling the TRNG or by setting the **TRNG_CTRL_FIFO_CLR** flag. The FIFO depth can be retrieved by software via the **TRNG_CTRL_FIFO_*** bits.

Simulation



When simulating the processor the TRNG is automatically set to "simulation mode". In this mode the physical entropy sources (the ring oscillators) are replaced by a simple **pseudo RNG** based on a LFSR providing only **deterministic pseudo-random** data. The **TRNG_CTRL_SIM_MODE** flag of the control register is set if simulation mode is active.

Interrupt

The TRNG provides a single interrupt request signal that gets triggered when the TRNG is enabled and the data FIFO is completely full (indicating that at least **IO_TRNG_FIFO** bytes of random data are available).

Register Map

Table 31. TRNG register map (`struct NEORV32_TRNG`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffffa0000	CTRL	0 TRNG_CTRL_EN	r/w	TRNG enable
		1 TRNG_CTRL_FIFO_CLR	-/w	flush random data FIFO when set; auto-clears
		5:2 TRNG_CTRL_FIFO_MSB : TRNG_CTRL_FIFO_LSB	r/-	FIFO depth, log2(IO_TRNG_FIFO)
		6 TRNG_CTRL_SIM_MODE	r/-	simulation mode (PRNG!)
		7 TRNG_CTRL_AVAIL	r/-	random data available when set
0xffffa0004	DATA	7:0 TRNG_DATA_MSB : TRNG_DATA_LSB	r/-	random data byte
		31:8 reserved	r/-	reserved, read as zero

2.8.21. Custom Functions Subsystem (CFS)

Hardware source files:	neorv32_cfs.vhd	
Software driver files:	neorv32_cfs.c neorv32_cfs.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>cfs_in_i</code> <code>cfs_out_o</code>	custom input conduit, 256-bit custom output conduit, 256-bit
Configuration generics:	<code>IO_CFS_EN</code>	implement CFS when <code>true</code>
CPU interrupts:	fast IRQ channel 1	CFS interrupt (see Processor Interrupts)

Key Features

- Tightly-coupled co-processor unit for custom hardware accelerators
- Can operate independently of the CPU
- CPU access via memory-mapped registers (16kB)
- Custom 256-bit input and output conduits for custom IO

Overview

The custom functions subsystem is meant for implementing custom tightly-coupled co-processors or interfaces. It provides up to 16384 32-bit memory-mapped read/write registers ([REG](#), see register map below) that can be accessed by the CPU via normal load/store operations. The actual functionality of these register has to be defined by the hardware designer. Furthermore, the CFS provides two input/output conduits to implement custom on-chip or off-chip interfaces.

Just like any other externally-connected IP, logic implemented within the custom functions subsystem can operate *independently* of the CPU providing true parallel processing capabilities. Potential use cases might include dedicated hardware accelerators for en-/decryption (AES), signal processing (FFT) or AI applications (CNNs) as well as custom IO systems like fast memory interfaces (DDR) and mass storage (SDIO), networking (CAN) real-time data transport (I2S) or just replication of existent NEORV32 peripherals.

Custom ISA Instructions



If you like to implement *custom instructions* that are executed right within the CPU's ALU see the [Zxfcuh ISA Extension](#) and the according [Custom Functions Unit \(CFU\)](#).

CFS Template



Take a look at the template CFS VHDL source file ([rtl/core/neorv32_cfs.vhd](#)). The file is highly commented to illustrate all aspects that are relevant for implementing custom CFS-based co-processor designs.

CFS Software Access

The CFS memory-mapped registers can be accessed by software using the provided C-language aliases (see register map table below). Note that all interface registers are defined as 32-bit words of type `uint32_t`.

Listing 15. CFS Software Access Example

```
// C-code CFS usage example
NEORV32_CFS->REG[0] = (uint32_t)some_data_array(i); // write to CFS register 0
int temp = (int)NEORV32_CFS->REG[20]; // read from CFS register 20
```

CFS Custom IOs

The CFS provides two unidirectional input and output conduits `cfs_in_i` and `cfs_out_o`. Both signals are 512 bit wide and are directly propagated to the processor's top entity. These conduits can be used to implement application-specific interfaces like memory or peripheral connections. The actual use case of these signals has to be defined by the hardware designer.

CFS Interrupt

The CFS provides a single high-level-triggered interrupt request signal mapped to the CPU's fast interrupt channel 1.

Register Map

Table 32. CFS register map (`struct NEORV32_CFS`)

Address	Name [C]	Bit(s)	R/W	Function
<code>0xfffeb0000</code>	<code>REG[0]</code>	<code>31:0</code>	(r)/(w)	custom CFS register 0
<code>0xfffeb004</code>	<code>REG[1]</code>	<code>31:0</code>	(r)/(w)	custom CFS register 1
...	...	<code>31:0</code>	(r)/(w)	...
<code>0xfffebffff8</code>	<code>REG[16382]</code>	<code>31:0</code>	(r)/(w)	custom CFS register 16382
<code>0xfffebffffc</code>	<code>REG[16383]</code>	<code>31:0</code>	(r)/(w)	custom CFS register 16383

2.8.22. Smart LED Interface (NEOLED)

Hardware source files:	neorv32_neoled.vhd	
Software driver files:	neorv32_neoled.c neorv32_neoled.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	<code>neoled_o</code>	1-bit serial data output
Configuration generics:	<code>IO_NEOLED_EN</code> <code>IO_NEOLED_TX_FIFO</code>	implement NEOLED controller when <code>true</code> transmission FIFO depth, has to be a power of 2, min 1
CPU interrupts:	fast IRQ channel 9	configurable NEOLED data FIFO interrupt (see Processor Interrupts)

Key Features

- "NeoPixel™"-compatible LED controller
- Hardware-based protocol handling
- Fine-grained programmable bus timings
- Optional data/command FIFO
- Interrupt based on FIFO status

Overview

The NEOLED module provides a dedicated interface for "smart LEDs" like WS2812, WS2811 and other compatible ones. These LEDs provide a single-wire interface that uses an asynchronous serial protocol for transmitting data. The NEOLED module handles the tight timing constraints of the interface entirely in hardware. A configurable data and command buffer (FIFO) allows to utilize (e.g. DMA-backed) block transfer operation without CPU intervention.

LEDs are connected to the `neoled_o` output and can be arbitrarily chained. The module also supports 24-bit and 32-bit LED data. Hence, mixed setups with RGB LEDs (24-bit color) and RGBW LEDs (32-bit color including a dedicated white LED chip) are fully supported.

The NEOLED module provides four accessible interface registers: the control register `CTRL` and the write-only `DATA24`, `DATA32` and `STROBE` registers. The precise interface timing is configured via the control registers. The data and strobe registers are used for sending 24-bit/32-bit data and reset/strobe commands.

Protocol

The interface of the smart LEDs is based on duty-cycle modulation of a specific carrier signal. Data is transmitted in a serial manner starting with the MSB. The intensity for each R/G/B (/W) LED chip is defined via an 8-bit value. The actual data bits are transferred by modifying the duty cycle of the base carriers. A RESET command is sent by pulling the data line LOW for at least 50µs.

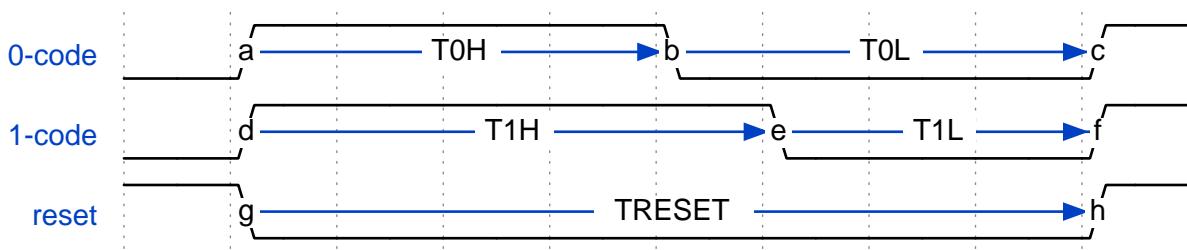


Figure 13. WS2812 bit-level timing (timing does not scale)

Table 33. Example: WS2812 Protocol Timings

T _{total} (T _{carrier})	1.25μs +/- 300ns	period for a single bit
T _{0H}	0.4μs +/- 150ns	high-time for sending a 1
T _{0L}	0.8μs +/- 150ns	low-time for sending a 1
T _{1H}	0.85μs +/- 150ns	high-time for sending a 0
T _{1L}	0.45μs +/- 150 ns	low-time for sending a 0
RESET	Above 50μs	low-time for sending a RESET command

Timing Configuration

The basic carrier frequency (800kHz for the WS2812 LEDs) is configured via a 3-bit clock prescaler (`NEOLED_CTRL_PRSC*`, see table below) that scales the main processor clock f_{main} and a 5-bit cycle multiplier `NEOLED_CTRL_T_TOT_*`.

Table 34. NEOLED Prescaler Configuration

NEOLED_CTRL_PRSC[2:0]	0b000	0b001	0b010	0b011	0b100	0b101	0b110	0b111
Resulting <code>clock_prescaler</code>	2	4	8	64	128	1024	2048	4096

The duty-cycles (or more precisely: the high- and low-times for sending either a '1' bit or a '0' bit) are defined via the 5-bit `NEOLED_CTRL_T_1H_*` and `NEOLED_CTRL_T_0H_*` values, respectively. These programmable timing constants allow to adapt the interface for a wide variety of smart LED protocols.

Timing Configuration - Example (WS2812)

Generate the base clock f_{TX} for the NEOLED TX engine:

- processor clock $f_{\text{main}} = 100 \text{ MHz}$
- `NEOLED_CTRL_PRSCx = 0b001` = $f_{\text{main}} / 4$

$$f_{\text{TX}} = f_{\text{main}}[\text{Hz}] / \text{clock_prescaler} = 100\text{MHz} / 4 = 25\text{MHz}$$

$$T_{\text{TX}} = 1 / f_{\text{TX}} = 40\text{ns}$$

Generate carrier period (T_{carrier}) and **high-times** (duty cycle) for sending **0** (T_{0H}) and **1** (T_{1H}) bits:

- `NEOLED_CTRL_T_TOT = 0b11110` (= decimal 30)
- `NEOLED_CTRL_T_0H = 0b01010` (= decimal 10)
- `NEOLED_CTRL_T_1H = 0b10100` (= decimal 20)

$$T_{carrier} = T_{TX} * \text{NEOLED_CTRL_T_TOT} = 40\text{ns} * 30 = 1.4\mu\text{s}$$

$$T_{0H} = T_{TX} * \text{NEOLED_CTRL_T_0H} = 40\text{ns} * 10 = 0.4\mu\text{s}$$

$$T_{1H} = T_{TX} * \text{NEOLED_CTRL_T_1H} = 40\text{ns} * 20 = 0.8\mu\text{s}$$



The NEOLED SW driver library ([nerv32_neoled.h](#)) provides a simplified configuration function that configures all timing parameters for driving WS2812 LEDs based on the processor clock frequency.

NEOLED Data/Command (FIFO)

The interface features a configurable TX data and command buffer (a FIFO) to allow more CPU-independent operation. The buffer depth is configured via the `IO_NEoled_TX_FIFO` top generic. The FIFO size configuration can be read by software from the `NEOLED_CTRL_FIFO_x` control register bits. If the module is disabled by clearing `NEOLED_CTRL_EN` the entire data and command buffer is cleared.

LED data and commands are written to the `DATA24`, `DATA32` and `STROBE` registers, which are transparently buffered by the optional FIFO. Writing data to `DATA24` will finally transmit the 24 LSB-aligned bits from the written data word to the LED string. Writing data to `DATA32` will finally transmit all 32-bit bits from the written data word to the LED string. Writing any data to `STROBE`, which is also buffered by the FIFO, will send a STROBE (RESET) command to the LED string.

Software can check the FIFO fill level via the control register's `NEOLED_CTRL_TX_EMPTY` and `NEOLED_CTRL_TX_FULL` flags. The `NEOLED_CTRL_TX_BUSY` flag provides additional information if the serial transmit engine is still busy sending data.

NEOLED Interrupt

The NEOLED modules features a single interrupt that gets triggered when the TX data/command FIFO is empty.

Register Map

Table 35. NEOLED register map (`struct NEORV32_NEoled`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
0xffffd0000	CTRL	0 NEOLED_CTRL_EN	r/w	module enable
		3:1 NEOLED_CTRL_PRSC_MSB : NEOLED_CTRL_PRSC_LSB	r/w	clock prescaler select
		8:4 NEOLED_CTRL_T_TOT_MSB : NEOLED_CTRL_T_TOT_LSB	r/w	pre-scaled clock ticks per total single-bit period (T_{total})
		13:9 NEOLED_CTRL_T_0H_MSB : NEOLED_CTRL_T_0H_LSB	r/w	pre-scaled clock ticks per high-time for sending a zero-bit (T_{0H})
		18:14 NEOLED_CTRL_T_1H_MSB : NEOLED_CTRL_T_1H_LSB	r/w	pre-scaled clock ticks per high-time for sending a one-bit (T_{1H})
		24:19 -	r/-	<i>reserved</i> , read as zero
		28:25 NEOLED_CTRL_FIFO_MSB : NEOLED_CTRL_FIFO_LSB	r/w	log(FIFO size)
		29 NEOLED_CTRL_TX_EMPTY	r/-	TX FIFO is empty
		30 NEOLED_CTRL_TX_FULL	r/-	TX FIFO is full
		31 NEOLED_CTRL_TX_BUSY	r/-	TX serial engine is busy when set
0xffffd0004	DATA24	23:0	-/w	write 24-bit RGB data to transmission buffer
0xffffd0008	DATA32	31:0	-/w	write 32-bit RGB data to transmission buffer
0xffffd000c	STROBE	-	-/w	write strobe/reset command to transmission buffer

2.8.23. General Purpose Timer (GPTMR)

Hardware source files:	neorv32_gptmr.vhd	
Software driver files:	neorv32_gptmr.c neorv32_gptmr.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	none	
Configuration generics:	<code>IO_GPTMR_NUM</code>	number of individual GPTMR timer slices (0..16)
CPU interrupts:	fast IRQ channel 12	timer interrupt (see Processor Interrupts)

Key Features

- Up to 16 individual timer slices with 32-bit-wide timer and timer-match registers
- Per-slice enable, mode and interrupt control/status
- Single-shot or interval operation mode
- Global clock prescaler

Overview

The general purpose timer module provides up to 16 individual timers organized as slices (the number of implemented slices is configured by `IO_GPTMR_NUM`). Each slice can operate as single-shot timer or as continuous/interval timer and provides a 32-bit counter (`SLICE[i].CNT`) and a 32-bit threshold registers (`SLICE[i].THR`). Operation is controlled by two 32-bit control and status registers (`CSR0` and `CSR1`). An interrupt is triggered when a slice counter matches the according threshold value.



CSR Access

`CSR0` and `CSR1` can be accessed in full-word mode (`CSR[0/1].WORD`) and also in half-word mode (e.g. `CSR0.ENABLE`) to update configuration sub-words with a single memory operation.

Theory of Operation

A slice i is enabled when the according bit of `CSR0.ENABLE` is set. Only the lowest `IO_GPTMR_NUM` bits are available. When enabled, the slice's counter register `SLICE[i].CNT` starts incrementing at a prescaled clock frequency. This prescaler is globally configured (= for all slices) by the `CSR1.PRSC` register:

Table 36. GPTMR prescaler configuration

<code>CSR1.PRSC</code>	0b000	0b001	0b010	0b011	0b100	0b101	0b110	0b111
Resulting prescaler	2	4	8	64	128	1024	2048	4096

When `SLICE[i].CNT` matches the threshold value in `SLICE[i].THR` the according interrupt-pending bit in `CSR1.IRQ` becomes set. Disabled slices (`CSR0.ENABLE(i) = 0`) cannot generate any new interrupt requests. However, pending interrupt in `CSR1.IRQ` will remain pending even if the according slice is disabled via. Software can clear pending interrupts by writing zero to the according `CSR1.IRQ` bits. A CPU interrupt request is generated if any slice interrupt i is pending and slice i is enabled (`CSR0.ENABLE(i) = 1` and `CSR1.IRQ(i) = 1`).

The time until a slice interrupt is triggered ($T[i]$) depends on the processor clock (f_{main}) global clock prescaler (selected by `CSR1.PRSC`) and the slice threshold value (`SLICE[i].THR`):

$$T[i] = f_{main} / (\text{prescaler} * \text{SLICE}[i].\text{THR})$$

Each slice can operate as single-shot timer or as continuous interval timer: If `CSR0.MODE(i) = 0` the according slice operates in **single-shot mode**: if the counter matches the programmed threshold (`SLICE[i].CNT == SLICE[i].THR`) the slice interrupt becomes pending and the counter stops at the threshold value. If `CSR0.MODE(i) = 1` the according slice operates in **continuous mode**: if the counter matches the programmed threshold (`SLICE[i].CNT == SLICE[i].THR`) the slice interrupt becomes pending and the counter is reset to zero and starts counting again up to the threshold value.

Register Map

Table 37. GPTMR register map (struct NEORV32_GPTMR)

Address	Name [C]	Bit(s), Name [C]	R/W	Description
0xffff10000	<code>CSR0.WORD</code>	31:0	r/w	Control and status register 0 (full-word access).
0xffff10000	<code>CSR0.ENAB</code> <code>LE</code>	15:0	r/w	Slice enable (one bit per implemented slice starting at LSB).
0xffff10002	<code>CSR0.MODE</code>	15:0	r/w	Slice operation mode: 0 = single shot, 1 = continuous (one bit per implemented slice starting at LSB).
0xffff10004	<code>CSR1.WORD</code>	31:0	r/w	Control and status register 1 (full-word access).
0xffff10004	<code>CSR1.IRQ</code>	15:0	r/w	Slice interrupt pending (one bit per implemented slice starting at LSB); write 0 to bit to clear interrupt.
0xffff10006	<code>CSR1.PRSC</code>	2:0	r/w	Global clock prescaler select.
0xffff10080	<code>SLICE[0].</code> <code>CNT</code>	31:0	r/w	Slice 0 counter register.
0xffff10084	<code>SLICE[0].</code> <code>THR</code>	31:0	r/w	Slice 0 threshold register.
...

Address	Name [C]	Bit(s), Name [C]	R/W	Description
0xffff100F8	SLICE[15] .CNT	31:0	r/w	Slice 15 counter register.
0xffff100FC	SLICE[15] .THR	31:0	r/w	Slice 15 threshold register.

2.8.24. Execution Trace Buffer (TRACER)

Hardware source files:	neorv32_tracer.vhd	
Software driver files:	neorv32_tracer.c neorv32_tracer.h	Online software reference (Doxygen) Online software reference (Doxygen)
Top entity ports:	none	
Configuration generics:	<code>IO_TRACER_EN</code> <code>IO_TRACER_BUFFER</code> <code>IO_TRACER_SIMLOG_EN</code>	implement TRACER module when <code>true</code> trace buffer depth, has to be zero or a power of two, min 1 write full trace log to file when <code>true</code> (simulation-only)
CPU interrupts:	fast IRQ channel 5	tracing stop-address match interrupt (see Processor Interrupts)

Key Features

- Delta-tracing of instruction execution
- Auto-stop tracing (and issue interrupt) when reaching a programmable instruction address
- Dump and inspect delta-trace directly from the firmware or via GDB

Overview

The NEORV32 tracer module allows to record program execution in order to find out how the CPU arrived at a certain code point. It is fed by the CPU's [Execution Trace Port](#). In order to keep the data rate low, so-called *instruction delta tracing* is used:

Instruction delta tracing, also known as branch tracing, works by tracking execution from a known start address by sending information about the deltas taken by the program. Deltas are typically introduced by jump, call, return and branch type instructions, although interrupts and exceptions are also types of deltas. Instruction delta tracing provides an efficient encoding of an instruction sequence by exploiting the deterministic way the processor behaves based on the program it is executing.

— RISC-V Trace Specification, github.com/riscv-non-isa/riscv-trace-spec

This means that only non-linear changes of the program counter are recorded. These *deltas* are stored in a local memory (the trace buffer) and can be read by the debugger (e.g. GDB). The size of this trace buffer is configurable via the `IO_TRACER_BUFFER` generic.

The program being executed can also read the trace buffer itself (stand-alone-mode, not debugger required). In the dual-core configuration, for example, one CPU core can analyze the trace of the

other CPU core.



TRACER Demo Program

A demo program for the tracer can be found in [sw/example/demo_tracer](#).

Theory of Operation

The tracer module provides four memory-mapped registers: the status and control register `CTRL`, a "stop-tracing" address register `STOP_ADDR` and two read-only registers to read trace data (`DELTA_SRC` and `DELTA_DST`).

The trace module is enabled by setting the `TRACER_CTRL_EN` bit in the control register. If this bit is cleared, the module is enabled and the internal trace buffer is cleared. Bit `TRACER_CTRL_HSEL` selects the hart / CPU core that shall be traced. This bit is read-only zero for the single-core configuration.

Tracing is started by writing 1 to the `TRACER_CTRL_START` control register bit. Tracing can be stopped at any time by manually writing 1 to the `TRACER_CTRL_STOP` control register bit. Software can check if tracing is in progress by reading the `TRACER_CTRL_RUN`. Tracing is *automatically stopped* when program execution reaches the address in the `STOP_ADDR` register. Automatic trace stopping can be disabled by writing -1 to this register.

Halt Tracing during Debug-Mode



Whenever the CPU enters debug mode, either by a halt request from the debugger or by a (hard-coded) breakpoint or watchpoint, tracing is paused until the CPU resumes normal operation.

During tracing, the module writes the instruction deltas to the internal buffer. The buffer is implemented as FIFO, so that only the last `IO_TRACER_BUFFER` instruction deltas are kept. At least one instruction delta is available in the trace buffer when `TRACER_CTRL_AVAIL` is set. A single entry ("packet") of the trace buffer data can be read via the two registers `DELTA_SRC` and `DELTA_DST`:

- `DELTA_SRC` provides the 31-bit instruction address of the branch source. Bit 0 is set if this trace packet is the very first after trace start; it is cleared for all further packets.
- `DELTA_DST` provides the 31-bit instruction address of the branch destination. Bit 0 is set if this branch was caused by a trap-entry; it is cleared if this branch was caused by a jump/call/branch instruction.

Trace Buffer Read Access



The trace buffer's FIFO read-pointer automatically increments when reading `DELTA_DST`, so make sure to read `DELTA_SRC` before `DELTA_DST`. The implementation as FIFO also means that the **trace data can only be read once**.

Tracer Interrupt

The tracer module features a single interrupt that gets triggered when the traced program reached the address stored in the module's `STOP_ADDR` register. The pending interrupt needs to be explicitly

acknowledged/cleared by writing 1 to the `TRACER_CTRL_IRQ_CLR` control register bit.

Evaluating Trace Data from GDB

Two simple functions for handling the tracer are implemented as GDB script. The script is located in the tracer example program and can be imported into GDB using the `source` command:

Listing 16. Importing the Tracer GDB Script

```
(gdb) source path/to/neorv32/sw/example/demo_tracer/neorv32_tracer.gdb
```

After sourcing, two additional commands are available that can be executed from the GDB command line:

- `tracer_get`: read-out the data from the trace buffer and print the execution history
- `tracer_start arg`: restart the tracer; the CPU ID has to be provided as argument (`arg`, 0 or 1)

Using the tracer from GDB is briefly illustrated in the following example:

Listing 17. Using the Tracer from GDB (GDB started in neorv32/sw/example/demo_tracer)

```
(gdb) make clean elf ①
...
(gdb) load ②
...
(gdb) c ③
...
(gdb) source neorv32_tracer.gdb ④
(gdb) tracer_get ⑤
[0] SRC: 0x3b0 -> DST: 0x2cc <TRACE_START>
Line 128 of "main.c" starts at address 0x3b0 <main+212> and ends at 0x3b4 <main+216>.
Line 68 of "main.c" starts at address 0x2cc <test_code> and ends at 0x2d0
<test_code+4>.
...
```

- ① Compile the demo program.
- ② Upload the compiled ELF.
- ③ Start executing. The program will stop re-entering debug-mode when reaching main's `return` statement.
- ④ Import the tracer helper functions script.
- ⑤ Get trace log and print.

Tracer Simulation Logging



Simulation-Only

This feature is available in simulation only.

The NEORV32 tracer module can also be used to generate a complete log of all executed instructions. Simulation trace logging is enabled by the `IO_TRACER_SIMLOG_EN` top generic. This also requires the tracer is implemented (`IO_TRACER_EN` = true). However, no specific configuration of the control register `CTRL` is required. During simulation, all traced instructions are written to log files in the simulator's home folder:

- `nerv32.tracer0.log` for CPU 0
- `nerv32.tracer1.log` for CPU 1 (only if **Dual-Core Configuration** is enabled)

The trace log is structured line by line where each line describes an executed instruction. The start of an exemplary trace log might look like this:

Listing 18. Exemplary cut-out from a simulation trace log (here: `nerv32.tracer1.log` showing boot of core 1)

51 929749 0x0000009c 0xffff44737 M lui	x14, 0xffff44000 <TRAP_ENTRY> ①
52 929751 0x000000a0 0x00872103 M lw	x2, 2(x14)
53 929759 0x000000a4 0x00c72603 M c.lw	x12, 12(x14)
54 929767 0x000000a6 0xffff40737 M lui	x14, 0xffff40000
55 929770 0x000000aa 0x00072223 M sw	x14, 4(x0)
56 929778 0x000000ae 0x04a0006f M c.jal	x0, 74
57 929798 0x000000f8 0x80000197 M auipc	x3, 0x80000000
58 929800 0x000000fc 0x70818193 M addi	x3, x8, 1800
59 929824 0x00000100 0x0ff0000f M fence	
60 929831 0x00000104 0x0000100f M fence.i	
61 929853 0x00000108 0x30029073 M csrrw	x0, mstatus, x5

① This line is used for explanation (see below).

Column structure:

1. **51**: Instruction index ("order"); a linear increasing counter that starts at zero and increments with each executed instruction; printed as decimal integer.
2. **890622**: Time stamp; a linear increasing counter that starts at zero and increments with each clock cycle; printed as decimal integer.
3. **0x000000c8**: Instruction address (program counter); printed as hexadecimal 32-bit value (`0x` prefix).
4. **0xffff44737**: 32-bit instruction word; printed as hexadecimal 32-bit value (`0x` prefix); compressed 16-bit instructions are shown in their decompressed 32-bit format.
5. **M**: Current operating mode / privilege level (`M` = machine-mode, `U` = user-mode, `D` = debug-mode); printed as single character
6. **lui**: The decoded instruction mnemonic. For unknown instructions `INVALID` is printed instead. Decompressed 16-bit instructions have a prefixed `c.` (e.g. `c.lw`).
7. **x14, 0xffff44000**: Operands of the decoded instruction.
8. If the CPU encounters a trap (synchronous exception or interrupt) `<TRAP_ENTRY>` is printed at the

end of the line corresponding to the first instruction of the trap handler.

Instruction Execution Time



The execution time of instruction i (number of cycles required for retiring) can be calculated by subtracting the current time stamp i from the next time stamp $i+1$.

Register Map

Table 38. TRACER register map (`struct NEORV32_TRACER`)

Address	Name [C]	Bit(s), Name [C]	R/W	Function
<code>0xffff30000</code>	<code>CTRL</code>	<code>0 TRACER_CTRL_EN</code>	r/w	TRACER enable, reset module when 0
		<code>1 TRACER_CTRL_HSEL</code>	r/w	Hart select for tracing (0 = CPU0, 1 = CPU1)
		<code>2 TRACER_CTRL_START</code>	r/w	Start tracing, flag always reads as zero
		<code>3 TRACER_CTRL_STOP</code>	r/w	Manually stop tracing, flag always reads as zero
		<code>4 TRACER_CTRL_RUN</code>	r/-	Tracing in progress when set
		<code>5 TRACER_CTRL_AVAIL</code>	r/-	Trace data available when set
		<code>6 TRACER_CTRL_IRQ_CLR</code>	r/w	Clear pending interrupt when writing 1, flag always reads as zero
		<code>10:7 TRACER_CTRL_TBM_MSB : TRACER_CTRL_TBM_LSB</code>	r/-	<code>log2(IO_TRACER_BUFFER)</code> : trace buffer depth
		<code>31:11 reserved</code>	r/-	<i>reserved</i> , hardwired to zero
<code>0xffff30004</code>	<code>STOP_AD DR</code>	<code>31:0</code>	r/w	Stop-tracing-address register
<code>0xffff30008</code>	<code>DELTA_S RC</code>	<code>31:1</code>	r/-	Branch source address, set to -1 to disable automatic stopping
		<code>0</code>	r/-	1 = very first instruction delta in current trace; 0 = any further instruction delta
<code>0xffff3000c</code>	<code>DELTA_D ST</code>	<code>31:1</code>	r/-	Branch destination address
		<code>0</code>	r/-	1 = branch due trap entry (interrupt or synchronous exception); 0 = branch due to jump/call/branch instruction

2.8.25. System Configuration Information Memory (SYSINFO)

Hardware source files:	neorv32_sysinfo.vhd	
Software driver files:	neorv32_sysinfo.h	Online software reference (Doxygen)
Top entity ports:	none	
Configuration generics:	all	most of the top's configuration generics
CPU interrupts:	none	

Key Features

- Provides detailed SoC configuration information

Overview

The SYSINFO module allows the application software to determine the setting of most of the [Processor Top Entity - Generics](#) that are related to CPU and processor/SoC configuration. This device is always implemented - regardless of the actual hardware configuration since the NEORV32 software framework requires information from this device for correct operation. However, advanced users that do not want to use the default NEORV32 software framework can choose to disable the entire SYSINFO module. This might also be suitable for setups that use the processor just as wrapper for a CPU-only configuration.

Disabling the SYSINFO Module



Setting the `IO_DISABLE_SYSINFO` top entity generic to `true` will remove the SYSINFO module from the design. This option is suitable for advanced uses that wish to use a CPU-only setup that still contains the bus infrastructure. As a result, large parts of the NEORV32 software framework no longer work (e.g. most IO drivers, the RTE and the bootloader). **Hence, this option is not recommended.**

Register Map

All registers of this module are read-only except for the `CLK` register. Upon reset, the `CLK` register is initialized from the `CLOCK_FREQUENCY` top entity generic. Application software can override this default value in order, for example, to take into account a dynamic frequency scaling of the processor.

Table 39. SYSINFO register map (`struct NEORV32_SYSINFO`)

Address	Name [C]	R/W	Description
<code>0xffffe0000</code>	<code>CLK</code>	r/w	clock frequency in Hz (initialized from top's <code>CLOCK_FREQUENCY</code> generic)
<code>0xffffe0004</code>	<code>MISC[4]</code>	r/-	miscellaneous system configurations (see SYSINFO - Miscellaneous Configuration)

Address	Name [C]	R/W	Description
0xffffe0008	SOC	r/-	specific SoC configuration (see SYSINFO - SoC Configuration)
0xffffe000c	CACHE	r/-	cache configuration information (see SYSINFO - Cache Configuration)

SYSINFO - Miscellaneous Configuration

Table 40. SYSINFO MEM Bits

Byte	Name [C]	Description
7:0	SYSINFO_MISC_IMEM_MBS : SYSINFO_MISC_IMEM_LSB	\log_2 (internal IMEM size in bytes), via top's IMEM_SIZE generic
15:8	SYSINFO_MISC_DMEM_MSB : SYSINFO_MISC_DMEM_LSB	\log_2 (internal DMEM size in bytes), via top's DMEM_SIZE generic
19:16	SYSINFO_MISC_HART_MSB : SYSINFO_MISC_HART_LSB	number of physical CPU cores ("harts")
21:20	SYSINFO_MISC_BOOT_MSB : SYSINFO_MISC_BOOT_LSB	boot mode configuration, via top's BOOT_MODE_SELECT generic (see Boot Configuration))
26:22	SYSINFO_MISC_ITMO_MSB : SYSINFO_MISC_ITMO_LSB	\log_2 (internal bus timeout cycles), see Bus Monitor and Timeout
31:27	SYSINFO_MISC_ETMO_MSB : SYSINFO_MISC_ETMO_LSB	\log_2 (external bus timeout cycles), see Bus Monitor and Timeout

SYSINFO - SoC Configuration

Table 41. SYSINFO SOC Bits

Bit	Name [C]	Description
0	SYSINFO_SOC_BOOTLOADER	set if processor-internal bootloader is implemented (via top's BOOT_MODE_SELECT generic; see Boot Configuration)
1	SYSINFO_SOC_XBUS	set if external bus interface is implemented (via top's XBUS_EN generic)
2	SYSINFO_SOC_IMEM	set if processor-internal DMEM is implemented (via top's IMEM_EN generic)
3	SYSINFO_SOC_DMEM	set if processor-internal IMEM is implemented (via top's DMEM_EN generic)
4	SYSINFO_SOC_OCD	set if on-chip debugger is implemented (via top's OCD_EN generic)
5	SYSINFO_SOC_ICACHE	set if processor-internal instruction cache is implemented (via top's ICACHE_EN generic)

Bit	Name [C]	Description
6	SYSINFO_SOC_DCACHE	set if processor-internal data cache is implemented (via top's DCACHE_EN generic)
7	-	<i>reserved</i> , read as zero
8	-	<i>reserved</i> , read as zero
9	-	<i>reserved</i> , read as zero
10	-	<i>reserved</i> , read as zero
11	SYSINFO_SOC_OCD_AUTH	set if on-chip debugger authentication is implemented (via top's OCD_AUTHENTICATION generic)
12	SYSINFO_SOC_IMEM_ROM	set if processor-internal IMEM is implemented as pre-initialized ROM (via top's BOOT_MODE_SELECT generic; see Boot Configuration)
13	SYSINFO_SOC_IO_TWD	set if TWD is implemented (via top's IO_TWD_EN generic)
14	SYSINFO_SOC_IO_DMA	set if direct memory access controller is implemented (via top's IO_DMA_EN generic)
15	SYSINFO_SOC_IO_GPIO	set if GPIO is implemented (via top's IO_GPIO_EN generic)
16	SYSINFO_SOC_IO_CLINT	set if CLINT is implemented (via top's IO_CLINT_EN generic)
17	SYSINFO_SOC_IO_UART0	set if primary UART0 is implemented (via top's IO_UART0_EN generic)
18	SYSINFO_SOC_IO_SPI	set if SPI is implemented (via top's IO_SPI_EN generic)
19	SYSINFO_SOC_IO_TWI	set if TWI is implemented (via top's IO_TWI_EN generic)
20	SYSINFO_SOC_IO_PWM	set if PWM is implemented (via top's IO_PWM_NUM generic)
21	SYSINFO_SOC_IO_WDT	set if WDT is implemented (via top's IO_WDT_EN generic)
22	SYSINFO_SOC_IO_CFS	set if custom functions subsystem is implemented (via top's IO_CFS_EN generic)
23	SYSINFO_SOC_IO_TRNG	set if TRNG is implemented (via top's IO_TRNG_EN generic)
24	SYSINFO_SOC_IO_SDI	set if SDI is implemented (via top's IO_SDI_EN generic)
25	SYSINFO_SOC_IO_UART1	set if secondary UART1 is implemented (via top's IO_UART1_EN generic)
26	SYSINFO_SOC_IO_NEOLED	set if smart LED interface is implemented (via top's IO_NEOLED_EN generic)

Bit	Name [C]	Description
27	SYSINFO_SOC_IO_TRACER	set if execution tracer is implemented (via top's <code>IO_TRACER_EN</code> generic)
28	SYSINFO_SOC_IO_GPTMR	set if GPTMR is implemented (via top's <code>IO_GPTMR_NUM</code> generic)
29	SYSINFO_SOC_IO_SLINK	set if stream link interface is implemented (via top's <code>IO_SLINK_EN</code> generic)
30	SYSINFO_SOC_IO_ONewire	set if ONEWIRE interface is implemented (via top's <code>IO_ONewire_EN</code> generic)
31	-	<i>reserved</i> , read as zero

SYSINFO - Cache Configuration

The SYSINFO cache register provides information about the configuration of the processor caches:

- Processor-Internal Instruction Cache (iCache)
- Processor-Internal Data Cache (dCache)

Table 42. SYSINFO `CACHE` Bits

Bit	Name [C]	Description
3:0	SYSINFO_CACHE_INST_BLOCK_SIZE_3 : SYSINFO_CACHE_INST_BLOCK_SIZE_0	\log_2 (i-cache block size in bytes), via top's <code>ICACHE_BLOCK_SIZE</code> generic
7:4	SYSINFO_CACHE_INST_NUM_BLOCKS_3 : SYSINFO_CACHE_INST_NUM_BLOCKS_0	\log_2 (i-cache number of cache blocks), via top's <code>ICACHE_NUM_BLOCKS</code> generic
11: 8	SYSINFO_CACHE_DATA_BLOCK_SIZE_3 : SYSINFO_CACHE_DATA_BLOCK_SIZE_0	\log_2 (d-cache block size in bytes), via top's <code>DCACHE_BLOCK_SIZE</code> generic
15: 12	SYSINFO_CACHE_DATA_NUM_BLOCKS_3 : SYSINFO_CACHE_DATA_NUM_BLOCKS_0	\log_2 (d-cache number of cache blocks), via top's <code>DCACHE_NUM_BLOCKS</code> generic
16	SYSINFO_CACHE_INST_BURSTS_EN	i-cache burst transfers enabled, via top's <code>CACHE_BURSTS_EN</code> generic
23: 17	0000000	<i>reserved</i>
24	SYSINFO_CACHE_DATA_BURSTS_EN	d-cache burst transfers enabled, via top's <code>CACHE_BURSTS_EN</code> generic
31: 25	0000000	<i>reserved</i>

Chapter 3. NEORV32 Central Processing Unit (CPU)

The NEORV32 CPU is an area-optimized RISC-V core implementing the `rv32i_zicsr_zifencei` base (privileged) ISA and supporting several additional/optional ISA extensions. The CPU's micro architecture is based on a von-Neumann machine build upon a mixture of multi-cycle and pipelined execution schemes. Optionally, the core can be implemented as SMP [Dual-Core Configuration](#).



RISC-V Specifications

This chapter assumes that the reader is familiar with the official RISC-V *User* and *Privileged Architecture* specifications.

Section Structure

- [CPU Top Entity - Signals](#) and [CPU Top Entity - Generics](#)
- [RISC-V Compatibility](#)
- [Architecture](#) and [Full Virtualization](#)
- [Execution Trace Port](#)
- [CPU Tuning Options](#)
- [Instruction Sets and Extensions](#) and [Custom Functions Unit \(CFU\)](#)
- [Control and Status Registers \(CSRs\)](#)
- [Traps, Exceptions and Interrupts](#)
- [Bus Interface](#)

3.1. CPU Top Entity - Signals

The following table shows all interface signals of the CPU top entity `rtl/core/neorv32_cpu.vhd`. The type of all signals is `std_ulogic` or `std_ulogic_vector`, respectively. The "Dir." column shows the signal direction as seen from the CPU.

Table 43. NEORV32 CPU Signal List

Signal	Width/Type	Dir	Description
Clock and Reset			
<code>clk_i</code>	1	in	Global clock line, all registers triggering on rising edge.
<code>rstn_i</code>	1	in	Global reset, low-active.
Status			
<code>trace_o</code>	<code>trace_port_t</code>	out	Execution Trace Port .
<code>sleep_o</code>	1	out	Set while the CPU is in Sleep Mode .
Interrupts (Traps, Exceptions and Interrupts)			
<code>msi_i</code>	1	in	RISC-V machine software interrupt.
<code>mei_i</code>	1	in	RISC-V machine external interrupt.
<code>mti_i</code>	1	in	RISC-V machine timer interrupt.
<code>firq_i</code>	16	in	Custom fast interrupt request signals.
<code>dbi_i</code>	1	in	Request CPU to halt and enter debug mode (RISC-V On-Chip Debugger (OCD)).
Instruction Bus Interface			
<code>ibus_req_o</code>	<code>bus_req_t</code>	out	Instruction fetch bus request.
<code>ibus_rsp_i</code>	<code>bus_rsp_t</code>	in	Instruction fetch bus response.
Data Bus Interface			
<code>dbus_req_o</code>	<code>bus_req_t</code>	out	Data access (load/store) bus request.
<code>dbus_rsp_i</code>	<code>bus_rsp_t</code>	in	Data access (load/store) bus response.



Bus Interface Protocol

See section [Bus Interface](#) for the instruction fetch and data access interface protocol and the according interface types (`bus_req_t` and `bus_rsp_t`).

3.2. CPU Top Entity - Generics

Most of the CPU configuration generics are a subset of the [Processor Top Entity - Generics](#) and are not listed here. However, the CPU provides some *specific* generics that are used to configure the CPU for the NEORV32 processor setup. These generics are assigned by the processor setup only and are not available for user defined configuration. The specific generics are listed below.



Table Abbreviations

The generic type "suv(x:y)" represents a `std::vector<std::uint32_t>`.

Table 44. NEORV32 CPU-Exclusive Generic List

Name	Type	Description
HART_ID	natural	ID of the core (for <code>mhartid</code> CSR).
NUM_HARTS	natural	Total number of cores in the system.
BOOT_ADDR	suv(31:0)	CPU reset address. See section Address Space .
DEBUG_PARK_ADDR	suv(31:0)	"Park loop" entry address for the On-Chip Debugger (OCD) , has to be 4-byte aligned.
DEBUG_EXC_ADDR	suv(31:0)	"Exception" entry address for the On-Chip Debugger (OCD) , has to be 4-byte aligned.
RISCV_ISA_Sdext	boolean	Implement RISC-V-compatible "debug" CPU operation mode required for the On-Chip Debugger (OCD) .
RISCV_ISA_Sdtrig	boolean	Implement RISC-V-compatible Trigger Module. See section On-Chip Debugger (OCD) .
RISCV_ISA_Smpmp	boolean	Implement RISC-V-compatible physical memory protection (PMP). See section Smpmp ISA Extension .
NUM_HW_TRIGGER	natural	Number of hardware triggers for the Trigger Module



Tuning Option Generics

Additional generics that are related to certain *tuning options* are listed in section [CPU Tuning Options](#).

3.3. RISC-V Compatibility

The NEORV32 CPU passes the tests of the [official RISCOF RISC-V Architecture Test Framework](#). This framework is used to check RISC-V implementations for compatibility to the official RISC-V user/privileged ISA specifications. The NEORV32 port of this test framework is available in a separate repository at GitHub: <https://github.com/stnolting/neorv32-riscof>



Unsupported ISA Extensions

Executing instructions or accessing CSRs from yet unsupported ISA extensions will

raise an illegal instruction exception (see section [Full Virtualization](#)).

Incompatibility Issues and Limitations

time[h] CSRs (Wall Clock Time)



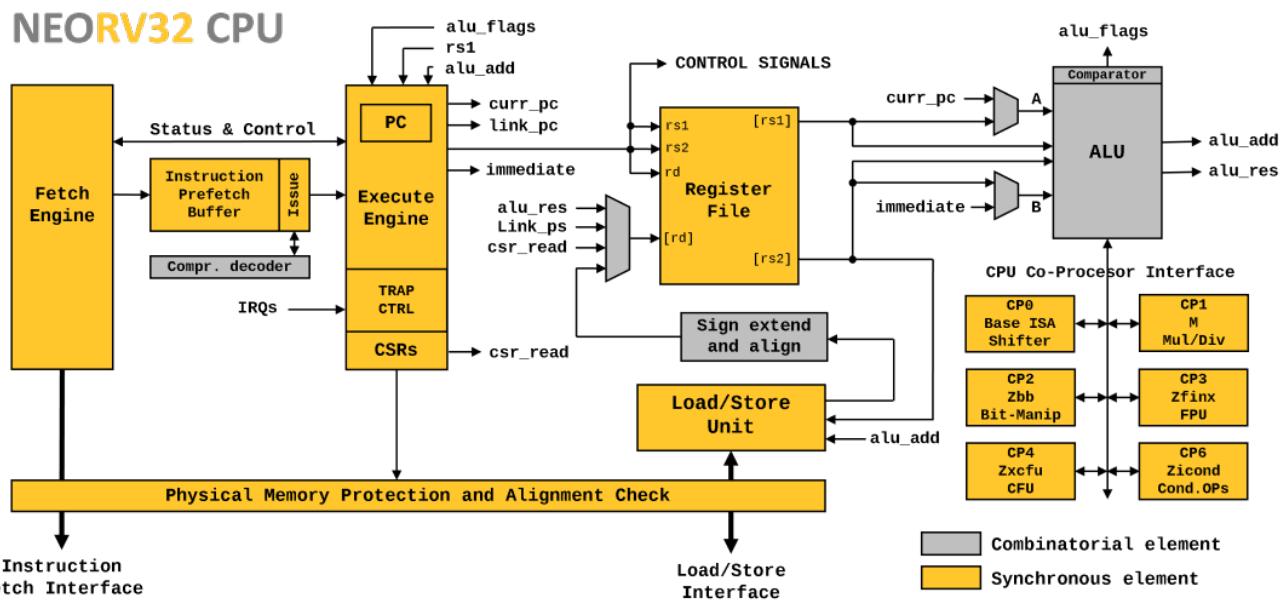
The NEORV32 does not implement the `time[h]` registers. Any access to these registers will trap. It is recommended that the trap handler software provides a means of accessing the machine timer of the [Core-Local Interruptor \(CLINT\)](#).



No Hardware Support for Misaligned Memory Accesses

The CPU does not support resolving unaligned memory access by the hardware (this is not a RISC-V-incompatibility issue but an important thing to know). Any kind of unaligned memory access will raise an exception to allow a *software-based* emulation provided by the application. However, unaligned memory access can be **emulated** using the NEORV32 runtime environment. See section [Application Context Handling](#) for more information.

3.4. Architecture



The CPU implements a pipelined multi-cycle architecture: each instruction is executed as a series of consecutive micro-operations. In order to increase performance, the CPU's front-end (instruction fetch) and back-end (instruction execution) are de-coupled via a FIFO (the instruction prefetch buffer). Thus, the front-end can already fetch new instructions while the back-end is still processing the previously-fetched instructions.

Basically, the CPU's micro architecture is somewhere between a classical pipelined architecture, where each stage requires exactly one processing cycle (if not stalled) and a classical multi-cycle architecture, which executes every single instruction (*including* fetch) in a series of consecutive micro-operations. The combination of these two design paradigms allows an increased instruction execution in contrast to a pure multi-cycle approach (due to overlapping operation of fetch and execute) at a reduced hardware footprint (due to the multi-cycle concept).

As a Von-Neumann machine, the CPU provides independent interfaces for instruction fetch and data access. However, these two bus interfaces are merged into a single processor-internal bus via a prioritizing bus switch (data accesses have higher priority). Hence, *all* memory addresses including peripheral devices are mapped to a single unified 32-bit [Address Space](#).

Linear/In-Order Execution Only



The CPU does not perform any speculative/out-of-order operations at all. Hence, it is not vulnerable to security issues caused by speculative execution (like Spectre or Meltdown).

3.4.1. CPU Register File

The data register file contains the general purpose architecture registers x_0 to x_{31} . For the `rv32e` ISA only the lower 16 registers are implemented. Register zero (x_0/zero) always read as zero and any write access to it has no effect. Up to four individual synchronous read ports allow to fetch up to 4

register operands at once. The write and read accesses are mutually exclusive as they happen in separate cycles. Hence, there is no need to consider things like "read-during-write" behavior.

Memory Tuning Options



The physical implementation of the register file's memory core can be tuned for FPGA or ASIC implementation. See section `CPU_RF_HW_RST_EN` tuning option for more information.

Implementation of the zero Register within FPGA Block RAM



Register `zero` is also mapped to a *physical memory location* within the register file's block RAM. By this, there is no need to add a further multiplexer to "insert" zero if reading from register `zero` reducing logic requirements and shortening the critical path. However, this also requires that the physical storage bits of register `zero` are explicitly initialized (set to zero) by the hardware. This is done transparently by the CPU control requiring no additional processing overhead.

Block RAM Ports



The default register file configuration uses two access ports: a read-only port for reading register `rs2` (second source operand) and a read/write port for reading register `rs1` (first source operand) and for writing processing results to register `rd` (destination register). Hence, a simple dual-port RAM can be used to implement the entire register file. From a functional point of view, read and write accesses to the register file do never occur in the same clock cycle, so no bypass logic is required at all.

3.4.2. CPU Arithmetic Logic Unit

The arithmetic/logic unit (ALU) is used for actual data processing as well as generating memory and branch addresses. All "simple" I ISA Extension computational instructions (like `add` and `or`) are implemented as plain combinatorial logic requiring only a single cycle to complete. More sophisticated instructions like shift operations or multiplications are processed by so-called "ALU co-processors".

The co-processors are implemented as iterative units that require several cycles to complete processing. Besides the base ISA's shift instructions, the co-processors are used to implement all further unprivileged ISA extensions (e.g. M ISA Extension, Zfinx ISA Extension and also custom instructions implemented within the Custom Functions Unit (CFU)). If the CPU logic detects an illegal instruction or an instruction fetch alignment or bus error no co-processor will be triggered to prevent corruption of the co-processor's internal state (like FPU CSRs).

Multi-Cycle Execution Monitor



The CPU control will raise an illegal instruction exception if a multi-cycle functional unit (like the Custom Functions Unit (CFU)) does not complete processing in a bound amount of time (configured via the package's `alu_cp_tmo_c`

constant; default = 512 clock cycles).

3.4.3. CPU Bus Unit

The bus unit takes care of handling data memory accesses via load and store instructions. It handles data adjustment when accessing sub-word data quantities (16-bit or 8-bit) and performs sign-extension for signed load operations. The bus unit also includes the optional [Smpmp ISA Extension](#) that performs permission checks for all data and instruction accesses.

A list of the bus interface signals and a detailed description of the protocol can be found in section [Bus Interface](#). All bus interface signals are driven/buffered by registers; so even a complex SoC interconnection bus network will not effect maximal operation frequency.

Unaligned Accesses



The CPU does not support a hardware-based handling of unaligned memory accesses! Any unaligned access will raise a bus load/store unaligned address exception. The exception handler can be used to *emulate* unaligned memory accesses in software. See the NEORV32 Runtime Environment's [Application Context Handling](#) section for more information.

3.4.4. CPU Control Unit

The CPU control unit is responsible for generating all the control signals for the different CPU modules. The control unit is split into a "front-end" and a "back-end".

Front-End

The front-end is responsible for fetching instructions in chunks of 32-bits. This can be a single aligned 32-bit instruction, two aligned 16-bit instructions or a mixture of those. The instructions including control and exception information are stored to a FIFO queue - the instruction prefetch buffer (IPB). This FIFO has a depth of two entries by default but can be customized via the [ipb_depth_c](#) VHDL package constant.

The FIFO allows the front-end to do "speculative" instruction fetches, as it keeps fetching the next consecutive instruction all the time. This also allows to decouple front-end (instruction fetch) and back-end (instruction execution) so both modules can operate in parallel to increase performance. However, all potential side effects that are caused by this "speculative" instruction fetch are already handled by the CPU front-end ensuring a defined execution stage while preventing security side-channel attacks.

Back-End

Instruction data from the instruction prefetch buffer is decompressed (if the [C](#) ISA extension is enabled) and sent to the CPU back-end for actual execution. Execution is conducted by a state-machine that controls all of the CPU modules. The back-end also includes the [Control and Status Registers \(CSRs\)](#) as well as the trap controller.

3.4.5. Execution Trace Port

The execution trace port provides information to trace the CPU's instruction execution. This can be used for advanced debugging, profiling or verification. Optionally, the trace data can be tracked by the [Execution Trace Buffer \(TRACER\)](#) to perform online branch delta tracing. Furthermore, trace data can also be made available externally by the two `trace_cpu0_o` and `trace_cpu1_o` [top module ports](#).

Trace data is generated by the CPU's trace generator (VHDL file `neorv32_cpu_trace.vhd`), which is enabled if either the [Execution Trace Buffer \(TRACER\)](#) is implemented or if the external trace ports are enabled by the `TRACE_PORT_EN` [top generic](#) (via the CPU core's `CPU_TRACE_EN` tuning option).

The trace port protocol is partly compatible to the [RISC-V Formal Interface \(RVFI\)](#) defined by Yosys (<https://yosyshq.readthedocs.io/projects/riscv-formal/en/latest/rvfi.html>).

The trace ports uses a custom interface type `trace_port_t` (VHDL record) which is defined in the processor's main RTL package file (`neorv32_package.vhd`).

Table 45. Trace Port Signals (`trace_port_t` VHDL type)

Signal	Width	RVFI Equivalent	Description
Instruction Metadata			
<code>valid</code>	1	<code>rvfi_valid</code>	All others signals are valid when set; high for one cycle
Integer Register			
<code>order</code>	32	<code>rvfi_order</code>	Instruction index (linear incrementing counter; resets to zero on overflow)
<code>insn</code>	32	<code>rvfi_insn</code>	Instruction word; compressed instructions are shown in their decompressed 32-bit format (with <code>compr = 1</code>)
<code>trap</code>	1	<code>rvfi_trap</code>	Set if the current instruction causes a synchronous exception
<code>halt</code>	1	<code>rvfi_halt</code>	Set if the current instruction is last instruction before halting
<code>intr</code>	1	<code>rvfi_intr</code>	Set if the current instruction is the first instruction of a trap handler
<code>mode</code>	2	<code>rvfi_mode</code>	Current privilege level; <code>00</code> = user, <code>11</code> = machine
<code>ixl</code>	2	<code>rvfi_ixl</code>	Current <code>XLEN</code> ; hardwired to <code>01</code> = 32-bit
<code>debug</code>	1	-	Set if the current instruction is executed in debug-mode
<code>compr</code>	1	-	Set if the current instruction is a decompressed instruction
<code>rs1_addr</code>	5	<code>rvfi_rs1_addr</code>	register <code>rs1</code> address
<code>rs2_addr</code>	5	<code>rvfi_rs2_addr</code>	register <code>rs2</code> address

Signal	Width	RVFI Equivalent	Description
Program Counter			
<code>rs1_rdata</code>	32	<code>rvfi_rs1_rdata</code>	register <code>rs1</code> read data (<i>before</i> instruction execution)
<code>rs2_rdata</code>	32	<code>rvfi_rs2_rdata</code>	register <code>rs2</code> read data (<i>before</i> instruction execution)
<code>rd_addr</code>	5	<code>rvfi_rd_addr</code>	register <code>rd</code> address; all-zero if no write
<code>rd_rdata</code>	32	<code>rvfi_rd_wdata</code>	register <code>rd</code> write data (<i>after</i> instruction execution)
Control and Status Register			
<code>csr_addr</code>	12	-	CSR access address; all-zero if no CSR access
<code>csr_rdata</code>	32	-	CSR read data (valid if <code>csr_addr != 0</code>)
<code>csr_wdata</code>	32	-	CSR write data
Memory Access			
<code>mem_addr</code>	32	<code>rvfi_mem_addr</code>	memory access address
<code>mem_rmask</code>	4	<code>rvfi_mem_rmask</code>	memory read-mask (byte-wise read-enable; all-zero if no memory read)
<code>mem_wmask</code>	4	<code>rvfi_mem_wmask</code>	memory write-mask (byte-wise write-enable; all-zero if no memory write)
<code>mem_rdata</code>	32	<code>rvfi_mem_rdata</code>	memory read data (valid if <code>mem_rmask != 0</code>)
<code>mem_wdata</code>	32	<code>rvfi_mem_wdata</code>	memory write data (valid if <code>mem_wmask != 0</code>)

3.4.6. CPU Tuning Options

The top module provides several tuning options to optimize the CPU for performance, size and even security. Note that these configuration options have no impact on the actual functionality (e.g. ISA compatibility).



Tuning Options Discovery

Software can check for configured tuning options via specific flags in the `mxcsr` CSR.

CPU_TRACE_EN

Name	Execution trace generation
Type	<code>boolean</code>
Default	<code>false</code> (disabled)
Description	When this option is enabled , the CPU code generates execution trace data that can be consumed by the Execution Trace Buffer (TRACER) for advanced debugging and profiling. When disabled , no trace data is generated.
Note	This option is automatically enabled if the Execution Trace Buffer (TRACER) is implemented or if the <code>TRACE_PORT_EN</code> top generic is <code>true</code> .

CPU_CONSTT_BR_EN

Name	Constant-time branches (data-independent timing)
Type	<code>boolean</code>
Default	<code>false</code> (disabled)
Description	When this option is enabled , all conditional branch instructions have identical execution times for taken and not taken branch conditions. Thus, all branches behave as if they were always taken (including a complete CPU pipeline flush). Enabling this feature makes execution times more predictable and makes timing side-channel attacks more difficult. When disabled , not-taken conditional branches are executed faster without clearing the CPU pipeline. Hence, for maximum performance, this feature should be disabled.

CPU_FAST_MUL_EN

Name	Fast multiplication
Type	<code>boolean</code>
Default	<code>false</code> (disabled)
Description	When enabled the <code>M/Zmmul</code> extension's multiplier is implemented as "plain multiplication" allowing the synthesis tool to infer DSP blocks / multiplication primitives. Multiplication operations only require a few cycles due to the DSP-internal register stages. The execution time is time-independent of the provided operands. When disabled the <code>M/Zmmul</code> extension's multiplier is implemented as bit-serial multiplier that computes one result bit in every cycle. Multiplication operations always require 32 clock cycles.

CPU_FAST_SHIFT_EN

Name	Fast bit-shifting
Type	<code>boolean</code>
Default	<code>false</code> (disabled)
Description	<p>When enabled the ALU's shifter unit is implemented as full-parallel barrel shifter that is capable of shifting a data word by an arbitrary number of positions within a single cycle. Hence, the execution time of any base-ISA shift operation is independent of the provided operands. Furthermore, this feature can help to reduce the risk of timing side-channel attacks. Note that the barrel shifter requires a lot of hardware resources for implementation.</p> <p>When disabled the ALU's shifter unit is implemented as bit-serial shifter that can shift the input data only by one position per cycle. Hence, several cycles might be required to complete any base-ISA shift-related operations. Therefore, the execution time of the serial approach is not time-independent of the provided operands. However, the serial approach requires only a few hardware resources and does not impact the critical path.</p>

CPU_RF_HW_RST_EN

Name	Register file hardware reset
Type	<code>boolean</code>
Default	<code>false</code> (disabled)
Description	<p>When enabled the CPU register file is implemented using single flip flops that provide a full hardware reset. The register file is reset to all-zero after each hardware reset. Note that this options requires a lot of flip flops and LUTs to build the register file. However, timing might be optimized as there is no need to route to far blockRAM resources.</p> <p>When disabled the CPU register file is implemented in a way to allow synthesis to infer memory primitives like blockRAM. Note that these primitives do not provide any kind of hardware reset. Hence, the data content is undefined after reset.</p>

3.4.7. Sleep Mode

The NEORV32 CPU provides a single sleep mode that can be entered to power-down the core reducing dynamic power consumption. Sleep mode is entered by executing the RISC-V `wfi` ("wait for interrupt") instruction.

CPU-external modules like memories, timers and peripheral interfaces are not affected by this. Furthermore, the CPU will continue to buffer/enqueue incoming interrupts. The CPU will leave

sleep mode as soon as any *enabled* interrupt (via `mie`) source becomes *pending* or if a debug session is started.



Sleep during Debugging

When executed in debug-mode or during single-stepping, `wfi` will behave as simple `nop` without entering sleep mode.

3.4.8. Full Virtualization

Just like the RISC-V ISA, the NEORV32 aims to provide *maximum virtualization* capabilities on CPU and SoC level to allow a high standard of **execution safety**. The CPU supports **all** traps specified by the official RISC-V specifications. Thus, the CPU provides defined hardware fall-backs via traps for any expected and unexpected situations (e.g. executing a malformed or not supported instruction or accessing a non-allocated memory address). For any kind of trap the core is always in a defined and fully synchronized state throughout the whole system (i.e. there are no out-of-order operations that might have to be reverted). This allows a defined and predictable execution behavior at any time improving overall execution safety.

3.5. Bus Interface

The NEORV32 CPU provides separated instruction fetch and data access interfaces making it a **Harvard Architecture**: the instruction fetch interface (`i_bus_*` signals) is used for fetching instructions and the data access interface (`d_bus_*` signals) is used to access data via load and store operations. Each of these interfaces can access an address space of up to 2^{32} bytes (4GB).

The bus interface uses two custom interface types: `bus_req_t` is used to propagate the bus access requests downstream from a host to a device. These signals are driven by the request-issuing device (i.e. the CPU core). Vice versa, `bus_rsp_t` is used to return the bus response upstream from a device back to the host and is driven by the accessed device or bus system (i.e. a processor-internal memory or IO device).

The signals of the request bus are split in to two categories: *in-band* signals and *out-of-band* signals. In-band signals always belong to a certain bus transaction and are valid between `stb` (request) being set and the according `err` or `ack` (response) being set. In contrast, the out-of-band signals are not associated with any bus transaction and are always valid when set.

Adding Register Stages



Arbitrary pipeline stages can be added to the request and response buses at any point to relax timing (at the cost of additional latency). However, *all* bus signals (request and response) need to be registered.

Table 46. Bus Interface - Request Bus (`bus_req_t`)

Signal	Width	Description
In-Band Signals		
<code>meta</code>	5	Access meta information. <code>meta(4:3)</code> : core ID from the lowest two bits of the <code>mhartid</code> CSR; the DMA has core ID <code>0b10</code> <code>meta(2)</code> : set when CPU is in debug-mode <code>meta(1)</code> : current privilege mode, <code>0</code> = user-mode, <code>1</code> = machine-mode <code>meta(0)</code> : access source, <code>0</code> = data access, <code>1</code> = instruction fetch
<code>addr</code>	32	Access address using byte-wise addressing. Half-word (16-bit) and word (32-bit) addresses are aligned accordingly (e.g. LSB = 0 for half-word accesses).
<code>data</code>	32	Write data. Writing individual bytes is controlled via <code>ben</code> .
<code>ben</code>	4	Byte-enable for each byte in <code>data</code> . This signal is used for the write-data as well as for the read-data.
<code>stb</code>	1	Request trigger ("strobe"). This signal is high for exactly one cycle starting a bus request. All other <i>in-band</i> signals are valid only if <code>stb</code> is set.
<code>rw</code>	1	Access direction (<code>0</code> = read, <code>1</code> = write).
<code>amo</code>	1	Set if current access is an atomic memory operation.
<code>amoop</code>	4	Actual type of atomic memory operation. Irrelevant if <code>amo</code> is not set.

Signal	Width	Description
burst	1	Set during a burst transaction (always set together with <code>lock</code>).
lock	1	Set if locked access; allow exclusive access to the bus system.
Out-Of-Band Signals		
fence	1	Data (load/store; <code>fence</code>) or instruction (instruction-fetch; <code>fence.i</code>) fence request; single-shot; see Memory Coherence .

Table 47. Bus Interface - Response Bus (`bus_rsp_t`)

Signal	Width	Description
ack	1	Transfer acknowledge, single-shot.
err	1	Transfer error when set, valid if <code>ack = 1</code> .
data	32	Read data, valid if <code>ack = 1</code> .

3.5.1. Bus Interface Protocol

Transactions are triggered by the request bus. A new bus request is initiated by setting the *strobe* signal `stb` high for exactly one cycle. All remaining signals of the bus are set together with `stb` and will remain unchanged until the transaction is completed.

The transaction is completed when the accessed device returns an acknowledgement via the `ack` signal, which is high for exactly one cycle. Together with `ack` the accessed device returns an optional error status (`err`) and read data (`rdata`) in case of a read-access. `err` and `data` are evaluated only when `ack` is high.

The figure below shows three exemplary single-access bus transactions:

1. A read access to address `A_addr` returning `rdata` after several cycles (slow response; `ACK` arrives after several cycles).
2. A write access to address `B_addr` writing `wdata` (fastest response; `ACK` arrives right in the next cycle).
3. A failing read access to address `C_addr` (slow response; `ERR` arrives after several cycles).

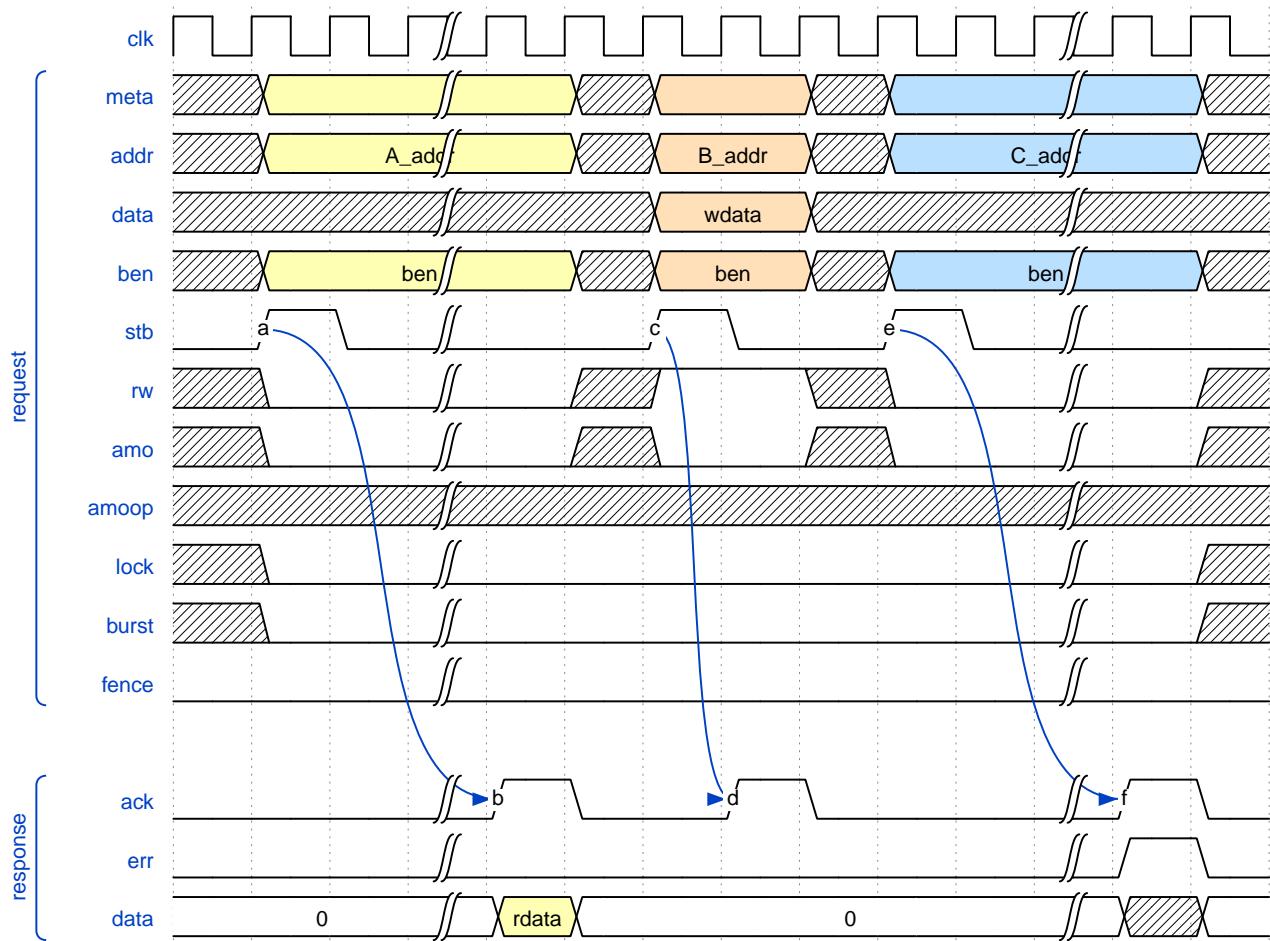


Figure 14. Three Exemplary Single-Access Transactions

3.5.2. Locked Bus Accesses and Bursts

Bus hosts like the CPU or the caches can request exclusive access to the downstream bus system using the "bus lock" mechanism. When the bus is locked the locking host has exclusive access and can issue an arbitrary number of transfers that cannot be interleaved by any other host. This feature is used to implement atomic read-modify-write operations and for burst transfers.

The following figure shows a simplified read-followed-by-write transaction implemented as uninterrupted exclusive bus access and a 4-word incrementing-address read burst.

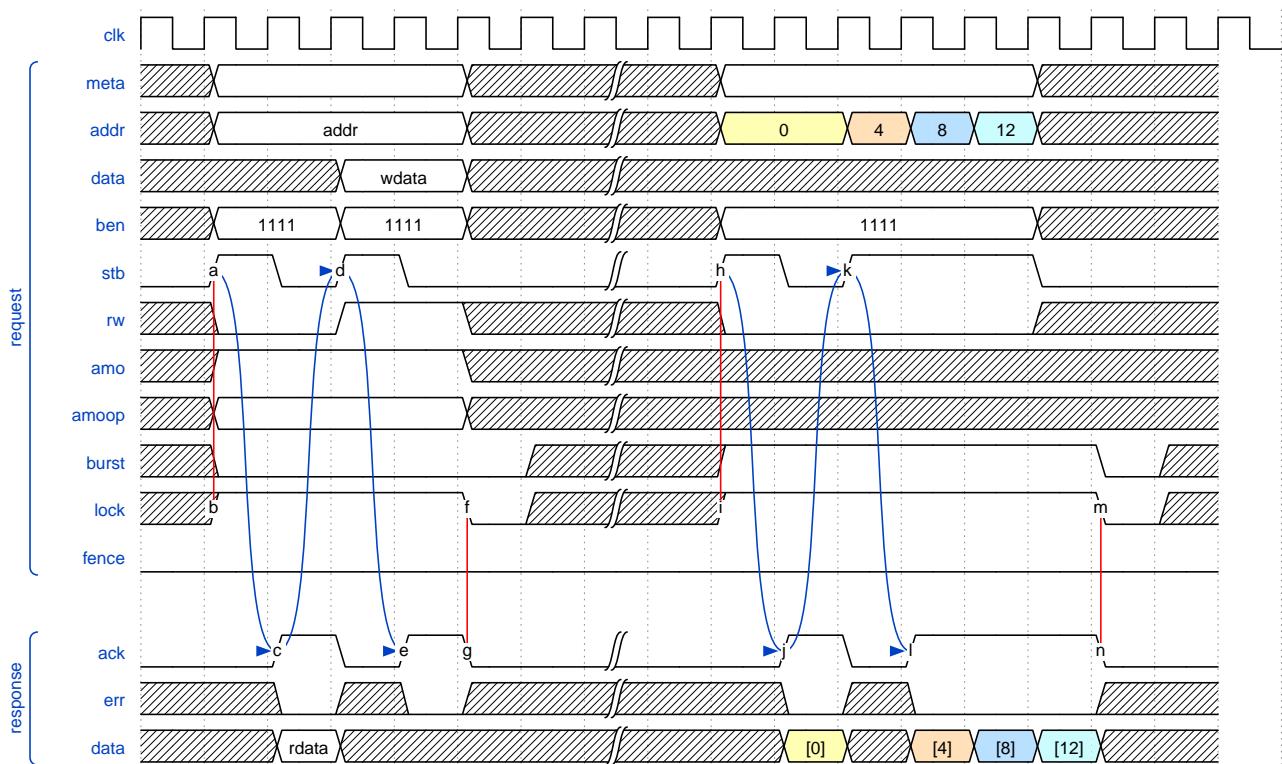


Figure 15. Atomic Access Transfer (left) and 4-Word Incrementing-Address Burst Read Transfer (right)

Atomic Access Transfer

1. The lock request is sent by setting **lock** together with the first **stb** ("a"). **lock** stays high until the end of the entire read-write transfer ("f").
2. The bus is locked when the host received the first **ack** ("c").
3. Now the host can send further requests using **stb** ("d").
4. When the exclusive access is completed the host clears **lock** after receiving the last **ack** of the transfer ("f").
5. **lock** has to be low for at least one cycle after the last **ack** ("f"). This will release the bus lock.

Burst Read Transfer

1. The locked burst request is sent by setting **lock** and **burst** together with the first **stb** ("h"). The address applied to **addr** defines the base/start address of the burst. **lock** and **burst** stay high until the end of the burst ("m").
2. The bus is locked when the host received the first **ack** ("j"). Now the host has exclusive bus access and can request the remaining burst strobes ("k").
3. The burst is completed when all required **stb** pulses are sent and when all corresponding **ack** have been received ("n"). **burst** is cleared. **lock** is cleared and has to be low for at least one cycle to release the bus lock.

Burst Generators



Only the caches (**i-cache / d-cache**) can generate burst transfers (and only if explicitly enabled via **CACHE_BURSTS_EN**).

*Fast Burst Response*

The accessed device/memory can return all data words of the burst already after receiving the first strobe of the locked transfer ("j").

3.6. Instruction Sets and Extensions

The NEORV32 CPU provides several optional RISC-V-compliant and custom/user-defined ISA extensions. The extensions can be enabled/configured via the according [Processor Top Entity - Generics](#). This chapter gives a brief overview of all available ISA extensions.

Table 48. NEORV32 Instruction Set Extensions

Name	Description	Enabled by Generic
A	Atomic memory instructions	<i>Implicitly</i> enabled
B	Bit manipulation instructions	<i>Implicitly</i> enabled
C	Compressed (16-bit) instructions	RISCV_ISA_C
E	Embedded CPU extension (reduced register file size)	RISCV_ISA_E
I	Integer base ISA	Enabled if RISCV_ISA_E is not enabled
M	Integer multiplication and division instructions	RISCV_ISA_M
U	Less-privileged <i>user</i> mode extension	RISCV_ISA_U
X	Platform-specific / NEORV32-specific extension	Always enabled
Zaamo	Atomic read-modify-write memory operations	RISCV_ISA_Zaamo
Zalrsc	Atomic reservation-set memory operations	RISCV_ISA_Zalrsc
Zca (just an alias)	Compressed instruction without floating-point operations	<i>Implicitly</i> enabled via RISCV_ISA_C
Zba	Shifted-add bit manipulation instructions	RISCV_ISA_Zba
Zbb	Basic bit manipulation instructions	RISCV_ISA_Zbb
Zbkb	Scalar cryptography bit manipulation instructions	RISCV_ISA_Zbkb
Zbkc	Scalar cryptography carry-less multiplication instructions	RISCV_ISA_Zbkc
Zbkx	Scalar cryptography crossbar permutation instructions	RISCV_ISA_Zbkx
Zbs	Single-bit bit manipulation instructions	RISCV_ISA_Zbs
Zcb	Additional code size reduction instructions (build upon C)	RISCV_ISA_Zcb
Zfinx	Floating-point instructions using integer registers	RISCV_ISA_Zfinx
Zifencei	Instruction stream synchronization instruction	Always enabled
Zibi	Branches with immediate-comparison	RISCV_ISA_Zibi

Name	Description	Enabled by Generic
Zicntr	Base counters extension	RISCV_ISA_Zicntr
Zicond	Integer conditional operations	RISCV_ISA_Zicond
Zicsr	Control and status register access instructions	Always enabled
Zihpm	Hardware performance monitors extension	RISCV_ISA_Zihpm
Zimop	May-be-operations	RISCV_ISA_Zimop
Zkn	Scalar cryptography NIST algorithm suite	<i>Implicitly</i> enabled
Zknd	Scalar cryptography NIST AES decryption instructions	RISCV_ISA_Zknd
Zkne	Scalar cryptography NIST AES encryption instructions	RISCV_ISA_Zkne
Zknh	Scalar cryptography NIST hash function instructions	RISCV_ISA_Zknh
Zkt	Data independent execution time (of cryptography operations)	<i>Implicitly</i> enabled
Zks	Scalar cryptography ShangMi algorithm suite	<i>Implicitly</i> enabled
Zksed	Scalar cryptography ShangMi block cypher instructions	RISCV_ISA_Zksed
Zksh	Scalar cryptography ShangMi hash instructions	RISCV_ISA_Zksh
Zmmul	Integer multiplication-only instructions	RISCV_ISA_Zmmul
Zcfu	Custom / user-defined instructions	RISCV_ISA_Zcfu
Smpmp	Physical memory protection (PMP) extension	RISCV_ISA_Smpmp
Sdext	External debug support extension (on-chip debugger)	OCD_EN
Sdtrig	Debug trigger module extension (hardware breakpoint)	OCD_NUM_HW_TRIGGER

RISC-V ISA Specification

For more information regarding the RISC-V ISA extensions please refer to the "RISC-V Instruction Set Manual - Volume I: Unprivileged ISA" and "The RISC-V Instruction Set Manual Volume II: Privileged Architecture". A copy of these documents can be found in the projects [docs/references](#) folder.

Discovering ISA Extensions

Software can discover available ISA extensions via the `misa` and `mxisa` CSRs or by executing an instruction and checking for an illegal instruction exception (i.e. [Full Virtualization](#)).

3.6.1. Latency Definitions

This chapter shows the CPI values (cycles per instruction) for each individual instruction/type. For some instructions a placeholder variable from the table below is used. Note that the provided values reflect *optimal bus accesses* (i.e. no congestion/wait states due to dual-core or DMA accesses) and no pipeline wait cycles. To benchmark a certain processor configuration for its setup-specific CPI value please refer to the [sw/example/performance_tests](#) test programs.

Alias	Description
T_shift_latency	1 if <code>CPU_FAST_SHIFT_EN</code> is <code>true</code> ; otherwise <code>T_shift_latency</code> equals the shift offset
T_data_1 latency	1 for accessing the default processor-internal memories; 1 if accessing the data cache results in a hit; 2 for accessing the default processor-internal peripherals; uncached accesses via the Processor-External Bus Interface (XBUS) may have additional wait states
T_inst_1 latency	1 for accessing the default processor-internal memories; 1 if accessing the instruction cache results in a hit; 2 for accessing the default processor-internal peripherals; uncached accesses via the Processor-External Bus Interface (XBUS) may have additional wait states; branches to an unaligned address require additional $2 \times T_{inst_latency}$ cycles
T_mul_latency	1 if <code>CPU_FAST_MUL_EN</code> is <code>true</code> ; 32 otherwise
T_cust_1 latency	latency is defined by the custom Custom Functions Unit (CFU) logic; however, the minimum is 1 and the maximum is 512

3.6.2. A ISA Extension

The **A** ISA extension adds instructions and mechanisms for atomic memory operations that can be used to manage mutually-exclusive access to shared resources. Note that this ISA extension cannot be enabled by a specific generic. Instead, it is enabled if a specific set of sub-extensions is enabled. The **A** extension is shorthand for the following set of other extensions:

- [Zam0 ISA Extension](#) - Atomic read-modify-write instructions.
- [Zalrsc ISA Extension](#) - Atomic reservation-set instructions.

Atomic Variable in C Languages



The C `_Atomic` specifier should be used for atomic variables. See section [Coherence Example](#) for more information.

A ISA Extension Exceptions



Note that load and load-reserved instructions generate load exceptions, whereas store, store-conditional, and AMO instructions generate store exceptions.

3.6.3. B ISA Extension

The **B** ISA extension adds instructions for bit-manipulation operations. This ISA extension cannot be enabled by a specific generic. Instead, it is enabled if a specific set of sub-extensions is enabled. The **B** extension is shorthand for the following set of other extensions:

- [Zba ISA Extension](#) - Address-generation / shifted-add instructions.
- [Zbb ISA Extension](#) - Basic bit manipulation instructions.
- [Zbs ISA Extension](#) - Single-bit operations.

A processor configuration which implements **B** must implement all of the above extensions.

3.6.4. C ISA Extension

The "compressed" ISA extension provides 16-bit encodings of commonly used instructions to reduce code size.

Table 49. Instructions and Timing

Class	Instructions	Execution cycles
ALU	c.addi4spn c.nop c.add[i] c.li c.addi16sp c.lui c.and[i] c.sub c.xor c.or c.mv	2
ALU shifts	c.srl c.srai c.slli	2 + T_shift_latency
Branch	c.beqz c.bnez	not taken: 3 taken: 5 + T_inst_latency
Jump/call	c.jal[r] c.j c.jr	5 + T_inst_latency
Load/store	c.lw c.sw c.lwsp c.swsp	4 + T_data_latency
System	c.break	7 + T_inst_latency

Note that the NEORV32 **C** ISA extension only includes the **Zca** instructions; i.e. all instructions from **C** excluding the single-precision (**F**) and double-precision (**D**) floating-point instructions.

3.6.5. E ISA Extension

The "embedded" ISA extensions reduces the size of the general purpose register file from 32 entries to 16 entries to shrink hardware size. It provides the same instructions as the base **I** ISA extension. Due to the reduced register file size an alternate toolchain ABI ([ilp32e*](#)) is required.

3.6.6. I ISA Extension

The **I** ISA extensions is the base RISC-V integer ISA that is always enabled.

Table 50. Instructions and Timing

Class	Instructions	Execution cycles
ALU	add[i] slt[i] sltu xor[i] or[i] and[i] sub 2 lui auipc	
ALU shifts	sll[i] srl[i] sra[i]	3 + $T_{shift_latency}$
Branch	beq bne blt bge bltu bgeu	not taken: 3 taken: 5 + $T_{inst_latency}$
Jump/call	jal[r]	5 + $T_{inst_latency}$
Load/store	lb lh lw lbu lhu sb sh sw	4 + $T_{data_latency}$
Data fence	fence	2
System	ecall ebreak mret	7 + $T_{inst_latency}$
System	wfi	3

fence Instruction

Analogous to the `fence.i` instruction ([Zifencei ISA Extension](#)) the `fence` instruction triggers a load/store memory synchronization operation by flushing the CPU's data cache. See section [Memory Coherence](#) for more information. NEORV32 ignores the predecessor and successor fields and always executes a conservative fence on all operations.

**wfi Instruction**

The `wfi` instruction is used to enter CPU [Sleep Mode](#).

3.6.7. M ISA Extension

Hardware-accelerated integer multiplication and division operations are available via the RISC-V **M** ISA extension. This ISA extension is implemented as multi-cycle ALU co-process ([rtl/core/neorv32_cpu_cp_multdiv.vhd](#)).

Table 51. Instructions and Timing

Class	Instructions	Execution cycles
Multiplication	mul mulh mulhsu mulhu	3 + $T_{mul_latency}$
Division	div divu rem remu	3 + 32

Multiplication Tuning Options

The physical implementation of the multiplier can be tuned for certain design goals like area or throughput. See section [CPU Tuning Options](#) for more information.

3.6.8. U ISA Extension

In addition to the highest-privileged machine-mode, the user-mode ISA extensions adds a second **less-privileged** operation mode. Code executed in user-mode has reduced CSR access rights. Furthermore, user-mode accesses to the address space (like peripheral/IO devices) can be constrained via the physical memory protection. Any kind of privilege rights violation will raise an exception to allow [Full Virtualization](#).

3.6.9. X ISA Extension

The NEORV32-specific ISA extensions **X** is always enabled. The most important points of the NEORV32-specific extensions are:

- * The CPU provides 16 *fast interrupt* interrupts (**FIRQ**), which are controlled via custom bits in the **mie** and **mip** CSRs. These extensions are mapped to CSR bits, that are available for custom use according to the RISC-V specs.
- Also, custom trap codes for **mcause** are implemented.
- * All undefined/unimplemented/malformed/illegal instructions do raise an illegal instruction exception (see [Full Virtualization](#)).
- * Additional [NEORV32-Specific CSRs](#).

3.6.10. Zaamo ISA Extension

The **Zaamo** ISA extension is a sub-extension of the RISC-V [A ISA Extension](#) and compromises instructions for atomic read-modify-write operations. It is enabled by the top's [RISCV_ISA_Zaamo](#) generic. Atomic read-modify-write operations are handled by the processor's [Atomic Memory Operations Controller](#).

Table 52. Instructions and Timing

Class	Instructions	Execution cycles
Atomic read-modify-write	<code>amoswap.w amoadd.w amoand.w amoor.w amoxor.w</code> <code>amomax[u].w amomin[u].w</code>	$4 + 2 \times T_{data_latency} + 1$



aq and lr Bits

The instruction word's **aq** and **lr** memory ordering bits are not evaluated by the hardware at all.

3.6.11. Zalrsc ISA Extension

The **Zalrsc** ISA extension is a sub-extension of the RISC-V [A ISA Extension](#) and compromises instructions for reservation-set operations. It is enabled by the top's [RISCV_ISA_Zalrsc](#) generic. Atomic reservation-set operations are handled by the processor's [Atomic Reservation-Set Controller](#).

Table 53. Instructions and Timing

Class	Instructions	Execution cycles
Atomic reservation-set	<code>lr.w sc.w</code>	$4 + T_{data_latency}$

*aq and lr Bits*

The instruction word's `aq` and `lr` memory ordering bits are not evaluated by the hardware at all.

3.6.12. Zcb ISA Extension

This ISA extension is part of the "code size reduction" ISA extension `Zc*` and adds additional compressed instruction for common operation. `Zcb` requires the [C ISA Extension](#) to be enabled. Some instructions may require additional ISA (sub-) extensions.

Table 54. Instructions and Timing

Class	Instructions	Depends on ISA Ext.	Execution cycles
Memory	<code>c.lbu c.lh c.lhu c.sb c.sh</code>	-	<code>4 + T_data_latency</code>
Logic	<code>c.not c.zext.b</code>	-	2
Logic	<code>c.sextr.b c.sextr.h c.sextr.w</code>	<code>B</code> or <code>Zbb</code>	3
Arithmetic	<code>c.mul</code>	<code>M</code> or <code>Zmmul</code>	<code>3 + T_mul_latency</code>

RISC-V GCC ISA String



Technically, `Zbc` required the `Zca` extension which is `C` but excluding the floating-point operations. Therefore, `Zca` and `Zbc` must be contained in the ISA string so that the compiler generates `Zbc` instructions.. Example: `MARCH=rv32imc_zca_zcb_...`

3.6.13. Zifencei ISA Extension

This instruction is the only standard mechanism to ensure that stores visible to a hart will also be visible to its instruction fetches. When executed, the CPU flushes the instruction prefetch buffer and reloads the CPU's instruction cache (if enabled). See section [Memory Coherence](#) for more information.

Table 55. Instructions and Timing

Class	Instructions	Execution cycles
Instruction fence	<code>fence.i</code>	2

3.6.14. Zfinx ISA Extension

The `Zfinx` floating-point extension is an *alternative* of the standard `F` floating-point ISA extension. It also uses the integer register file `x` to store and operate on floating-point data instead of a dedicated floating-point register file. Thus, the `Zfinx` extension requires less hardware resources and features faster context changes. This also implies that there are NO dedicated `f` register file-related load/store or move instructions. The `Zfinx` extension's floating-point unit is controlled via dedicated [Floating-Point CSRs](#). This ISA extension is implemented as multi-cycle ALU co-process (`rtl/core/neorv32_cpu_cp_fpu.vhd`).

Fused / Multiply-Add and Division / Square Root Instructions

Fused multiply-add instructions `f[n]m[add/sub].s` and division `fdiv.s` and square root `fsqrt.s` instructions are not supported yet. Special GCC flags are used to prevent the compiler from emitting these operations (see [Default Compiler Flags](#)).

Subnormal Numbers

Subnormal numbers ("de-normalized" numbers, i.e. exponent = 0) are not supported by the NEORV32 FPU. Subnormal numbers are *flushed to zero* setting them to +/- 0 before being processed by **any** FPU operation. If a computational instruction generates a subnormal result it is also flushed to zero during normalization.

Table 56. Instructions and Timing

Class	Instructions	Execution cycles
Add/sub	<code>fadd.s fsub.s</code>	21
Multiply	<code>fmul.s</code>	13
Compare	<code>fmin.s fmax.s feq.s flt.s fle.s</code>	5
Convert (float → int)	<code>fcvt.w.s fcvt.wu.s</code>	10
Convert (int → float)	<code>fcvt.s.w fcvt.s.wu</code>	42
Sign-injection	<code>fsgnj.s fsgnjn.s fsgnjx.s</code>	4
Classify	<code>fclass.s</code>	4

3.6.15. Zibi ISA Extension

The RISC-V **Zibi** ISA extension adds two new "branch with immediate" instructions. The instruction word's `rs2` bit-field selects one out of 32 pre-defined immediate values which is used for the actual comparison.

*Non-Ratified ISA Extension*

The RISC-V **Zibi** ISA extension has not been ratified yet!

Table 57. Instructions and Timing

Class	Instructions	Execution cycles
Branch with immediate	<code>beqi bnei</code>	not taken: 3 taken: 5 + <code>T_inst_latency</code>

3.6.16. Zicntr ISA Extension

The **Zicntr** ISA extension adds the basic `cycle[h]`, `mcycle[h]`, `instret[h]` and `minstret[h]` counter CSRs. Section [\(Machine\) Counter and Timer CSRs](#) shows a list of all **Zicntr**-related CSRs.

*Time CSRs*

The user-mode `time[h]` CSRs are **not implemented**. Any access will trap allowing the trap handler to retrieve system time from the **Core-Local Interruptor (CLINT)**.

*Constrained Access*

User-level access to the counter CSRs can be constrained by the `mcounteren` CSR.

3.6.17. Zicond ISA Extension

The **Zicond** ISA extension adds integer conditional move primitives that allow to implement branchless control flows. It is enabled by the top's `RISCV_ISA_Zicond` generic. This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/nerv32_cpu_cp_cond.vhd`).

Table 58. Instructions and Timing

Class	Instructions	Execution cycles
Conditional	<code>czero.eqz</code> <code>czero.nez</code>	3

3.6.18. Zicsr ISA Extension

This ISA extensions provides instructions for accessing the **Control and Status Registers (CSRs)** as well as further privileged-architecture extensions. This extension is mandatory and cannot be disabled. Hence, there is no generic for enabling/disabling this ISA extension.

*Side-Effects if Destination is Zero-Register*

If `rd=x0` for the `csrrw[i]` instructions there will be no actual read access to the according CSR. However, access privileges are still enforced so these instruction variants *do* cause side-effects (the RISC-V spec. state that these combinations "shall" not cause any side-effects).

Table 59. Instructions and Timing

Class	Instructions	Execution cycles
System	<code>csrrw[i]</code> <code>csrrs[i]</code> <code>csrrc[i]</code>	3

3.6.19. Zihpm ISA Extension

In additions to the base counters the NEORV32 CPU provides up to 13 hardware performance monitors (HPM 3..15), which can be used to benchmark applications. Each HPM consists of an N-bit wide counter (split in a high-word 32-bit CSR and a low-word 32-bit CSR), where N is defined via the top's `HPM_CNT_WIDTH` generic and a corresponding event configuration CSR.

The event configuration CSR defines the architectural events that lead to an increment of the associated HPM counter. See section **Hardware Performance Monitors (HPM) CSRs** for a list of all HPM-related CSRs and event configurations.

Machine-Mode HPMs Only

Note that only the machine-mode hardware performance counter CSR are available (`mhpmcnter*[h]`). Accessing any user-mode HPM CSR (`hpmcounter*[h]`) will raise an illegal instruction exception.

*Increment Inhibit*

The event-driven increment of the HPMs can be deactivated individually via the `mcountinhibit` CSR.

3.6.20. `Zimop` ISA Extension

The `Zimop` ISA extension introduces the concept of instructions that may be operations (MOPs). MOPs are initially defined to simply write zero to the destination register, but are designed to be redefined by later extensions to perform some other system/environment-class actions. NEORV32 can implement this ISA extension for software compatibility. Without the `Zimop` extension, any `mop.r*` instruction will raise an illegal instruction exception.

Table 60. Instructions and Timing

Class	Instructions	Execution cycles
System	<code>mop.r.[0..31] mop.rr.[0..7]</code>	3

3.6.21. `Zba` ISA Extension

The `Zba` sub-extension is part of the *RISC-V bit manipulation* ISA specification ([B ISA Extension](#)) and adds shifted-add / address-generation instructions. It is enabled by the top's `RISCV_ISA_Zba` generic. This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/neorv32_cpu_cp_bitmanip.vhd`).

Table 61. Instructions and Timing

Class	Instructions	Execution cycles
Shifted-add	<code>sh1add sh2add sh3add</code>	4

3.6.22. `Zbb` ISA Extension

The `Zbb` sub-extension is part of the *RISC-V bit manipulation* ISA specification ([B ISA Extension](#)) and adds the basic bit manipulation instructions. It is enabled by the top's `RISCV_ISA_Zbb` generic. This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/neorv32_cpu_cp_bitmanip.vhd`).

Table 62. Instructions and Timing

Class	Instructions	Execution cycles
Logic with negate	<code>andn orn xnor</code>	4

Class	Instructions	Execution cycles
Bit count	<code>clz</code> <code>ctz</code> <code>cpop</code>	$4 + T_{shift_latency}$
Integer maximum/minimum	<code>min[u]</code> <code>max[u]</code>	4
Sign/zero extension	<code>sext.b</code> <code>sext.h</code> <code>zext</code>	4
Bitwise rotation	<code>rol</code> <code>ror[i]</code>	$4 + T_{shift_latency}$
OR-combine	<code>orc.b</code>	4
Byte-reverse	<code>rev8</code>	4

shifter Tuning Options



The physical implementation of the bit-shifter can be tuned for certain design goals like area or throughput. See section [CPU Tuning Options](#) for more information.

3.6.23. Zbs ISA Extension

The **Zbs** sub-extension is part of the *RISC-V bit manipulation* ISA specification ([B ISA Extension](#)) and adds single-bit operations. It is enabled by the top's `RISCV_ISA_Zbs` generic. This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/neorv32_cpu_cp_bitmanip.vhd`).

Table 63. Instructions and Timing

Class	Instructions	Execution cycles
Single-bit	<code>sbsr[i]</code> <code>sbclr[i]</code> <code>sbinv[i]</code> <code>sbext[i]</code>	4

3.6.24. Zbkb ISA Extension

The **Zbkb** sub-extension is part of the *RISC-V scalar cryptography* ISA specification and extends the *RISC-V bit manipulation* ISA extension with additional instructions. It is enabled by the top's `RISCV_ISA_Zbkb` generic. Note that enabling this extension will also enable the **Zbb** basic bit-manipulation ISA extension (which is extended by **Zknb**). This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/neorv32_cpu_cp_bitmanip.vhd`).

Table 64. Instructions and Timing (in addition to Zbb)

Class	Instructions	Execution cycles
Packing	<code>pack</code> <code>packh</code>	4
Interleaving	<code>zip</code> <code>unzip</code>	4
Byte-wise bit reversal	<code>brev8</code>	4

3.6.25. Zbkc ISA Extension

The **Zbkc** sub-extension is part of the *RISC-V scalar cryptography* ISA extension and adds carry-less multiplication instruction. ISA extension with additional instructions. It is enabled by the top's `RISCV_ISA_Zbkc` generic. This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/neorv32_cpu_cp_bitmanip.vhd`).

Table 65. Instructions and Timing

Class	Instructions	Execution cycles
Carry-less multiply	<code>clmul</code> <code>clmulh</code>	4 + 32

3.6.26. Zbkx ISA Extension

The **Zbkx** sub-extension is part of the *RISC-V scalar cryptography* ISA specification and adds crossbar permutation instructions. It is enabled by the top's `RISCV_ISA_Zbkx` generic. This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/neorv32_cpu_cp_crypto.vhd`).

Table 66. Instructions and Timing

Class	Instructions	Execution cycles
Crossbar permutation	<code>xperm8</code> <code>xperm4</code>	4

3.6.27. Zkn ISA Extension

The **Zkn** ISA extension is part of the *RISC-V scalar cryptography* ISA specification and defines the "NIST algorithm suite". This ISA extension cannot be enabled by a specific generic. Instead, it is enabled if a specific set of cryptography-related sub-extensions is enabled.

The **Zkn** extension is shorthand for the following set of other extensions:

- **Zbkb ISA Extension** - Bit manipulation instructions for cryptography.
- **Zbkc ISA Extension** - Carry-less multiply instructions.
- **Zbkx ISA Extension** - Cross-bar permutation instructions.
- **Zkne ISA Extension** - AES encryption instructions.
- **Zknd ISA Extension** - AES decryption instructions.
- **Zknh ISA Extension** - SHA2 hash function instructions.

A processor configuration which implements **Zkn** must implement all of the above extensions.

3.6.28. Zknd ISA Extension

The **Zknd** sub-extension is part of the *RISC-V scalar cryptography* ISA specification and adds NIST AES decryption instructions. It is enabled by the top's `RISCV_ISA_Zknd` generic. This ISA extension is

implemented as multi-cycle ALU co-processor ([rtl/core/neorv32_cpu_cp_crypto.vhd](#)).

Table 67. Instructions and Timing

Class	Instructions	Execution cycles
AES decryption	<code>aes32dsi aes32dsni</code>	6

3.6.29. Zkne ISA Extension

The `Zkne` sub-extension is part of the *RISC-V scalar cryptography* ISA specification and adds NIST AES encryption instructions. It is enabled by the top's `RISCV_ISA_Zkne` generic. This ISA extension is implemented as multi-cycle ALU co-processor ([rtl/core/neorv32_cpu_cp_crypto.vhd](#)).

Table 68. Instructions and Timing

Class	Instructions	Execution cycles
AES decryption	<code>aes32esi aes32esmi</code>	6

3.6.30. Zknh ISA Extension

The `Zknh` sub-extension is part of the *RISC-V scalar cryptography* ISA specification and adds NIST hash function instructions. It is enabled by the top's `RISCV_ISA_Zknh` generic. This ISA extension is implemented as multi-cycle ALU co-processor ([rtl/core/neorv32_cpu_cp_crypto.vhd](#)).

Table 69. Instructions and Timing

Class	Instructions	Execution cycles
sha256	<code>sha256sig0 sha256sig1 sha256sum0 sha256sum1</code>	4
sha512	<code>sha512sig0h sha512sig0l sha512sig1h sha512sig1l sha512sum0r sha512sum1r</code>	4

3.6.31. Zks ISA Extension

The `Zks` ISA extension is part of the *RISC-V scalar cryptography* ISA specification and defines the "ShangMi algorithm suite". This ISA extension cannot be enabled by a specific generic. Instead, it is enabled if a specific set of cryptography-related sub-extensions is enabled.

The `Zks` extension is shorthand for the following set of other extensions:

- [Zbkb ISA Extension](#) - Bit manipulation instructions for cryptography.
- [Zbkc ISA Extension](#) - Carry-less multiply instructions.
- [Zbkx ISA Extension](#) - Cross-bar permutation instructions.
- [Zksed ISA Extension](#) - SM4 block cipher instructions.
- [Zksh ISA Extension](#) - SM3 hash function instructions.

A processor configuration which implements `Zks` must implement all of the above extensions.

3.6.32. Zksed ISA Extension

The `Zksed` sub-extension is part of the *RISC-V scalar cryptography* ISA specification and adds ShangMi block cypher and key schedule instructions. It is enabled by the top's `RISCV_ISA_Zksed` generic. This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/neorv32_cpu_cp_crypto.vhd`).

Table 70. Instructions and Timing

Class	Instructions	Execution cycles
Block cyphers	<code>sm4ed</code>	6
Key schedule	<code>sm4ks</code>	6

3.6.33. Zksh ISA Extension

The `Zksh` sub-extension is part of the *RISC-V scalar cryptography* ISA specification and adds ShangMi hash function instructions. It is enabled by the top's `RISCV_ISA_Zksh` generic. This ISA extension is implemented as multi-cycle ALU co-processor (`rtl/core/neorv32_cpu_cp_crypto.vhd`).

Table 71. Instructions and Timing

Class	Instructions	Execution cycles
Hash	<code>sm3p0 sm3p1</code>	6

3.6.34. Zkt ISA Extension

The `Zkt` sub-extension is part of the *RISC-V scalar cryptography* ISA specification and guarantees data independent execution times of cryptography and cryptography-related instructions. The ISA extension cannot be enabled by a specific generic. Instead, it is enabled implicitly by certain CPU configurations.

The RISC-V `Zkt` specifications provides a list of instructions that are included within this specification. However, not all instructions are required to be implemented. Rather, every one of these instructions that the core does implement must adhere to the requirements of `Zkt`.

Table 72. Zkt instruction listing

Parent extension	Instructions	Data independent execution time?
RVI	<code>lui auipc add[i] slt[i][u] xor[i] or[i]</code> <code>and[i] sub</code>	yes
	<code>sll[i] srl[i] sra[i]</code>	yes if <code>CPU_FAST_SHIFT_EN</code> enabled
RVM	<code>mul[h] mulh[s]u</code>	yes
RVC	<code>c.nop c.addic.lui c.andic.c.sub c.xor</code> <code>c.and c.mv c.add</code>	yes
	<code>c.srlic.c.srai c.slli</code>	yes if <code>CPU_FAST_SHIFT_EN</code> enabled

Parent extension	Instructions	Data independent execution time?
Zcb	c.mul, c.not, c.zext.b RVK	yes aes32ds[m]i aes32es[m]i sha256sig* sha512sig* sha512sum* sm3p0 sm3p1 sm4ed sm4ks
yes	RVB	xperm4 xperm8 andn orn xnor pack[h] brev8 rev8
yes		ror[i] rol

Data-Independent Timing of Branches



To further reduce the possibility of timing side-channel attacks, conditional branches can be executed with constant time (i.e. data-independent execution times). See the [CPU_CONSTT_BR_EN CPU Tuning Options](#) for more information.

3.6.35. Zmmul - ISA Extension

This is a sub-extension of the [M ISA Extension](#) ISA extension. It implements only the multiplication operations of the [M](#) extensions and is intended for size-constrained setups that require hardware-based integer multiplications but not hardware-based divisions, which will be computed entirely in software. Note that the [Zmmul - ISA Extension](#) and [M ISA Extension](#) are mutually exclusive.

3.6.36. Zxcfus ISA Extension

The [Zxcfus](#) presents a NEORV32-specific ISA extension. It adds the [Custom Functions Unit \(CFU\)](#) to the CPU core, which allows to add custom RISC-V instructions to the processor core. For detailed information regarding the CFU, its hardware and the according software interface see section [Custom Functions Unit \(CFU\)](#).

Software can utilize the custom instructions by using *intrinsics*, which are basically inline assembly functions that behave like regular C functions but that evaluate to a single custom instruction word (no calling overhead at all).

CFU Execution Time



The actual CFU execution time depends on the logic being implemented. The CPU architecture requires a minimal execution time of 3 cycles (purely combinatorial CFU operation) and automatically terminates execution after 512 cycles if the CFU does not complete operation within this time window.

Table 73. Instructions and Timing

Class	Instructions	Execution cycles
Custom instructions	Instruction words with <code>custom-0</code> or <code>custom-1</code> opcode	$3 + T_{cust_latency}$

3.6.37. Smpmp ISA Extension

The NEORV32 physical memory protection (PMP) provides an elementary memory protection mechanism that can be used to configure read/write(execute permission of arbitrary memory regions. In general, the PMP can **grant permissions to user mode**, which by default has none, and can **revoke permissions from M-mode**, which by default has full permissions. The NEORV32 PMP is fully compatible to the RISC-V Privileged Architecture Specifications and is configured via several CSRs ([Machine Physical Memory Protection CSRs](#)). Several [Processor Top Entity - Generics](#) are provided to adjust the CPU's PMP capabilities according to the application requirements (pre-synthesis):

1. **PMP_NUM_REGIONS** defines the number of implemented PMP regions (0..16); setting this generic to zero will result in absolutely no PMP logic being implemented
2. **PMP_MIN_GRANULARITY** defines the minimal granularity of each region (has to be a power of 2, minimal granularity = 4 bytes); note that a smaller granularity will lead to wider comparators and thus, to higher area footprint and longer critical path
3. **PMP_TOR_MODE_EN** controls the implementation of the top-of-region (TOR) mode (default = true); disabling this mode will reduce area footprint
4. **PMP_NAP_MODE_EN** controls the implementation of the naturally-aligned-power-of-two (NA4 and NAPOT) modes (default = true); disabling this mode will reduce area footprint and critical path length

PMP Permissions when in Debug Mode



When in debug-mode all PMP rules are bypassed/ignored granting the debugger maximum access permissions.

PMP Time-Multiplex



Instructions are executed in a multi-cycle manner. Hence, data access (load/store) and instruction fetch cannot occur at the same time. Therefore, the PMP hardware uses only a single set of comparators for memory access permissions checks that are switched in an iterative, time-multiplex style reducing hardware footprint by approx. 50% while maintaining full security features and RISC-V compatibility.

PMP Memory Accesses



Load/store accesses for which there are insufficient access permission do not trigger any memory/bus accesses at all. In contrast, instruction accesses for which there are insufficient access permission nevertheless lead to a memory/bus access (causing potential side effects on the memory side=). However, the fetched instruction will be discarded and the corresponding exception will still be triggered precisely.

3.6.38. Sdext ISA Extension

This ISA extension enables the RISC-V-compatible "external debug support" by implementing the

CPU "debug mode", which is required for the on-chip debugger. See section [On-Chip Debugger \(OCD\) / CPU Debug Mode](#) for more information.

Table 74. Instructions and Timing

Class	Instructions	Execution cycles
System	<code>dret</code>	5

3.6.39. **Sdtrig** ISA Extension

This ISA extension implements the RISC-V-compatible "trigger module" which provides support for hardware breakpoints for the on-chip-debugger. See section [On-Chip Debugger \(OCD\) / Trigger Module](#) for more information.

3.7. Custom Functions Unit (CFU)

The Custom Functions Unit (CFU) is the central part of the NEORV32-specific [Zxfcuh ISA Extension](#) and represents the actual hardware module that can be used to implement **custom RISC-V instructions**. These are intended for operations that are inefficient in terms of performance, latency, energy consumption or program memory requirements when implemented entirely in software.

CFU Complexity



The CFU is not intended for complex and **CPU-independent** functional units that implement complete accelerators (like full block-based AES encryption). These kind of accelerators should be implemented as memory-mapped co-processor via the [Custom Functions Subsystem \(CFS\)](#) to allow CPU-independent operation. A comparative survey of all NEORV32-specific hardware extension/customization options is provided in the user guide section [Adding Custom Hardware Modules](#).



Default CFU Hardware/Software Example

The default CFU module ([rtl/core/neorv32_cpu_cp_cfu.vhd](#)) implements the *Extended Tiny Encryption Algorithm (XTEA)* as application example. The according example program is located in [sw/example/demo_cfu](#).

3.7.1. CFU Instruction Formats

The custom instructions executed by the CFU utilize a specific opcode space in the **rv32** 32-bit instruction encoding space that has been explicitly reserved for user-defined extensions by the RISC-V specifications ("Guaranteed Non-Standard Encoding Space"). The NEORV32 CFU uses the **custom-0** and **custom-1** opcodes to identify the instructions implemented by the CFU and to differentiate between the predefined instruction formats. The NEORV32 CFU utilizes these two opcodes to support user-defined **R-type** instructions and **I-type** instructions. Both instruction formats are compliant to the RISC-V specification.

- **custom-0 (0001011)**, used for [CFU R-Type Instructions](#)
- **custom-1 (0101011)**, used for [CFU I-Type Instructions](#)

CFU R-Type Instructions

The R-type CFU instructions operate on two source registers **rs1** and **rs2** and return the processing result to the destination register **rd**. The actual operation can be defined by using the **funct7** and **funct3** bit fields. These immediates can also be used to pass additional data to the CFU like offsets, look-up-tables addresses or shift-amounts. However, the actual functionality is entirely user-defined. Note that all immediate values are always compile-time-static.

Example operation: **rd** \leftarrow **rs1** **xnor** **rs2** (bit-wise logical XNOR)

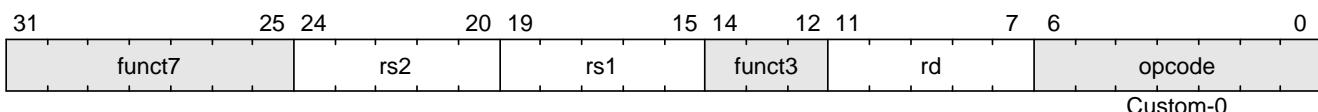


Figure 16. CFU R-type instruction format

- **funct7**: 7-bit immediate (immediate data or function select)
- **rs2**: address of second source register (providing 32-bit source data)
- **rs1**: address of first source register (providing 32-bit source data)
- **funct3**: 3-bit immediate (immediate data or function select)
- **rd**: address of destination register (32-bit processing result)
- **opcode**: **0001011** (RISC-V **custom-0** opcode)

CFU I-Type Instructions

The I-type CFU instructions operate on one source registers **rs1** and a 12-bit immediate value **imm12** and return the processing result to the destination register **rd**. The actual operation can be defined by using the **funct3** bit field. Alternatively, this immediate can also be used to pass additional data to the CFU like offsets, look-up-tables addresses or shift-amounts. However, the actual functionality is entirely user-defined. Note that all immediate values are always compile-time-static.

Example operation: **rd** \leftarrow **rs1** * **imm12** (multiply with immediate)

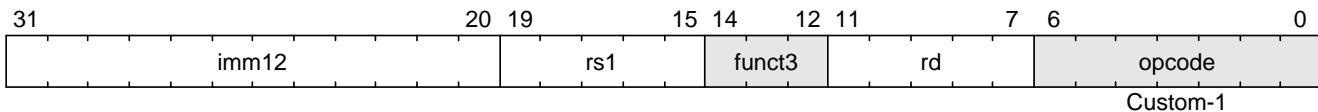


Figure 17. CFU I-type instruction format

- **imm12**: 12-bit immediate
- **rs1**: address of first source register (providing 32-bit source data)
- **funct3**: 3-bit immediate (immediate data or function select)
- **rd**: address of destination register (32-bit processing result)
- **opcode**: **0101011** (RISC-V **custom-1** opcode)

3.7.2. Using Custom Instructions in Software

The custom instructions provided by the CFU can be used in plain C code by using **intrinsics**. Intrinsics behave like "normal" C functions but under the hood they are a set of macros that hide the complexity of inline assembly, which is used to construct the custom 32-bit instruction words. Using intrinsics removes the need to modify the compiler, built-in libraries or the assembler when using custom instructions. Each intrinsic will be compiled into a single 32-bit instruction word without any overhead providing maximum code efficiency. The NEORV32 software framework provides two pre-defined prototypes for custom instructions:

Listing 19. CFU instruction prototypes (defined in [sw/lib/include/neorv32_cfu.h](#))

```
uint32_t neorv32_cfu_r_instr(func7, funct3, rs1, rs2); // R-type instruction
uint32_t neorv32_cfu_i_instr(funct3, imm12, rs1); // I-type instruction
```

The intrinsic functions always return a 32-bit value of type `uint32_t` (the processing result) which can be discarded if not needed. The `funct3`, `funct7` and `imm12` bit-fields are used to pass compile-time-static literals to the CFU. The `rs1` and `rs2` arguments pass actual runtime data to the CFU via register addresses. These register arguments can be populated with variables or literals; the compiler will add the required code to move the data into a register before passing it to the CFU. The following examples shows how to pass arguments:

Listing 20. CFU instruction usage examples

```
uint32_t res = neorv32_cfu_r_instr(0b0000000, 0b101, 8, 123);
uint32_t tmp = neorv32_cfu_i_instr(0b011, 0x47a, res);
neorv32_cfu_r_instr(0b0100100, 0b001, res, tmp);
```

3.7.3. Custom Instructions Hardware

The CFU hardware module ([rtl/core/neorv32_cpu_cp_cfu.vhd](#)) is tightly integrated into to CPU and coupled as ALU co-processor. A simple interface is used to communicate operands, operation definition and the processing result.

Table 75. CFU Interface

Signal	Width	Instruction Type	Description
Request			
<code>start</code>	1	all	Trigger CFU operation; high for exactly one cycle
<code>type</code>	1	all	CFU instruction type (0 = R-type, 1 = I-type); defined by the opcode
<code>funct3</code>	3	all	3-bit function select from the according custom instruction word
<code>funct7</code>	7	R-type	7-bit function select from the according custom instruction word (**)
<code>imm12</code>	12	I-type	12-bit immediate from the custom according instruction word
<code>rs1</code>	32	all	Register operand 1 (addresses by the <code>rs1</code> instruction word bit field)
<code>rs2</code>	32	R-type	Register operand 2 (addresses by the <code>rs2</code> instruction word bit field)
Response			
<code>result</code>	1	all	Processing result that is written to <code>rd</code>

Signal	Width	Instruction Type	Description
valid	32	all	High for exactly one cycle when the CFU operation is done and result is valid

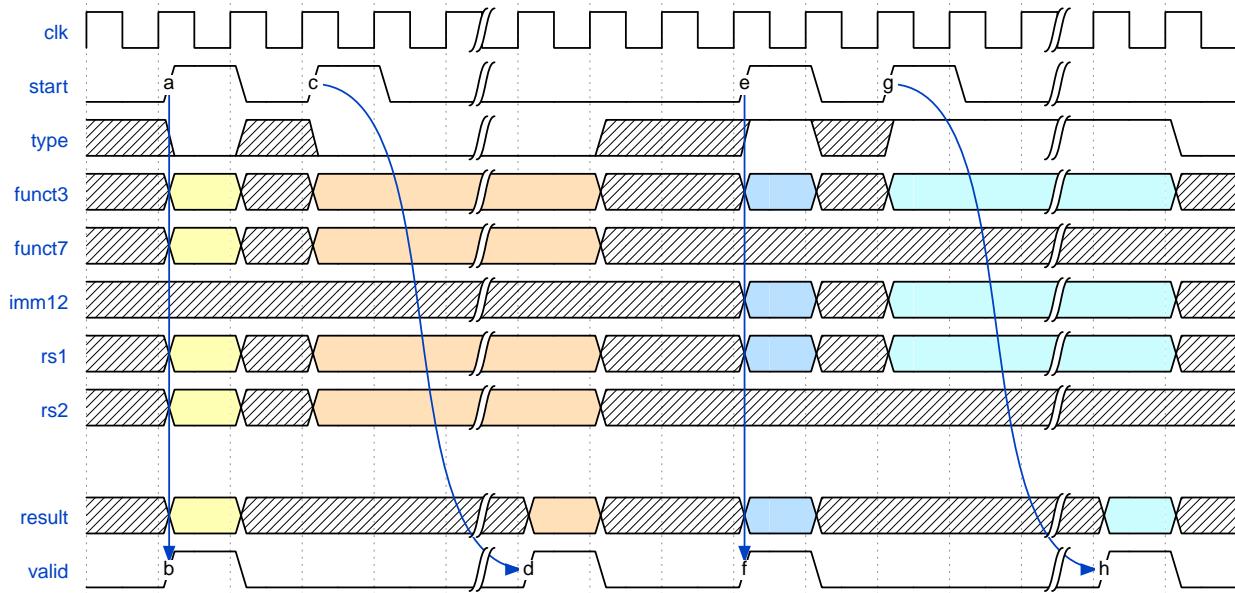


Figure 18. CFU Interface Timing Examples

CFU operations can be entirely combinatorial (e.g. for a bit-reversal operation) so the result is available at the end of the current clock cycle. However, operations can also take several clock cycles to complete (like multiplications) and may also include internal states and memories. However, the CFU has to complete computation within a **bound time window** (default = 512 clock cycles). Otherwise, an illegal instruction exception is raised. See section [CPU Arithmetic Logic Unit](#) for more information.

CFU Exception



The CFU can intentionally raise an illegal instruction exception by not asserting the **valid** signal at all which will cause an execution timeout. For example this can be used to signal invalid configurations/operations to the runtime environment. See the documentation in the CFU's VHDL source file for more information.

3.8. Control and Status Registers (CSRs)

The following table shows a summary of all available NEORV32 CSRs. The address field defines the CSR address for the CSR access instructions. The "Name [ASM]" column provides the CSR name aliases that can be used in (inline) assembly. The "Name [C]" column lists the name aliases that are defined by the NEORV32 core library. These can be used in plain C code. The "Access" column shows the minimal required privilege mode required for accessing the according CSR (**M** = machine-mode, **U** = user-mode, **D** = debug-mode) and the read/write capabilities (**RW** = read-write, **RO** = read-only)

Unused, Reserved, Unimplemented and Disabled CSRs



All CSRs and CSR bits that are not listed in the table below are *unimplemented* and are *hardwired to zero*. Additionally, CSRs that are unavailable ("disabled") because the according ISA extension is not enabled are also considered *unimplemented* and are also hardwired to zero. Any access to such a CSR will raise an illegal instruction exception. All writable CSRs provide **WARL** behavior (write all values; read only legal values). Application software should always read back a CSR after writing to check if the targeted bits can actually be modified.



Read-Only RW CSRs

Some CSRs are listed with "read-write" capabilities but will ignore any value written to them. In the following table these CSRs are highlighted with "█".

Table 76. NEORV32 Control and Status Registers (CSRs)

Address	Name [ASM]	Name [C]	Access	Description
Floating-Point CSRs				
0x001	fflags	CSR_FFLAGS	URW	Floating-point accrued exceptions
0x002	frm	CSR_FRM	URW	Floating-point dynamic rounding mode
0x003	fcsr	CSR_FCSR	URW	Floating-point control and status
Machine Trap Setup CSRs				
0x300	mstatus	CSR_MSTATUS	MRW	Machine status register - low word
0x301	misa	CSR_MISA	MRW	Machine CPU ISA and extensions
0x304	mie	CSR_MIE	MRW	Machine interrupt enable register
0x305	mtvec	CSR_MTVEC	MRW	Machine trap-handler base address for ALL traps
0x306	mcounteren	CSR_MCOUNTERN	MRW	Machine counter-enable register
0x310	mstatush	CSR_MSTATUSH	MRW	Machine status register - high word
Machine Configuration CSRs				
0x30a	menvcfg	CSR_MENVCFG	MRW	Machine environment configuration register - low word

Address	Name [ASM]	Name [C]	Access	Description
0x31a	<code>menvcfgh</code>	<code>CSR_MEVCFGH</code>	MRW	Machine environment configuration register - high word
Machine Counter Setup CSRs				
0x320	<code>mcowntinhibit</code>	<code>CSR_MCOUNTINHIBIT</code>	MRW	Machine counter-inhibit register
Machine Trap Handling CSRs				
0x340	<code>mscratch</code>	<code>CSR_MSCRATCH</code>	MRW	Machine scratch register
0x341	<code>mepc</code>	<code>CSR_MEPC</code>	MRW	Machine exception program counter
0x342	<code>mcause</code>	<code>CSR_MCAUSE</code>	MRW	Machine trap cause
0x343	<code>mtval</code>	<code>CSR_MTVAL</code>	MRW	Machine trap value
0x344	<code>mip</code>	<code>CSR_MIP</code>	MRW	Machine interrupt pending register
0x34a	<code>mtinst</code>	<code>CSR_MTINST</code>	MRW	Machine trap instruction
Machine Physical Memory Protection CSRs				
0x3a0 .. 0x303	<code>pmpcfg0 .. pmpcfg3</code>	<code>CSR_PMPCFG0 .. CSR_PMPCFG3</code>	MRW	Physical memory protection configuration registers
0x3b0 .. 0x3bf	<code>pmpaddr0 .. pmpaddr15</code>	<code>CSR_PMPADDR0 .. CSR_PMPADDR15</code>	MRW	Physical memory protection address registers
Trigger Module CSRs				
0x7a0	<code>tselect</code>	<code>CSR_TSELECT</code>	MRW	Trigger select register
0x7a1	<code>tdata1</code>	<code>CSR_TDATA1</code>	MRW	Trigger data register 1
0x7a2	<code>tdata2</code>	<code>CSR_TDATA2</code>	MRW	Trigger data register 2
0x7a4	<code>tinfo</code>	<code>CSR_TINFO</code>	MRW	Trigger information register
CPU Debug Mode CSRs				
0x7b0	<code>dcsr</code>	-	DRW	Debug control and status register
0x7b1	<code>dpc</code>	-	DRW	Debug program counter
0x7b2	<code>dscratch0</code>	-	DRW	Debug scratch register 0
(Machine) Counter and Timer CSRs				
0xb00	<code>mcycle</code>	<code>CSR_MCYCLE</code>	MRW	Machine cycle counter low word
0xb02	<code>minstret</code>	<code>CSR_MINSTRET</code>	MRW	Machine instruction-retired counter low word
0xb80	<code>mcycleh</code>	<code>CSR_MCYCLEH</code>	MRW	Machine cycle counter high word
0xb82	<code>minstreth</code>	<code>CSR_MINSTRETH</code>	MRW	Machine instruction-retired counter high word
0xc00	<code>cycle</code>	<code>CSR_CYCLE</code>	URO	Cycle counter low word

Address	Name [ASM]	Name [C]	Access	Description
0xc02	instret	CSR_INSTRET	URO	Instruction-retired counter low word
0xc80	cycleh	CSR_CYCLEH	URO	Cycle counter high word
0xc82	instreth	CSR_INSTRETH	URO	Instruction-retired counter high word
Hardware Performance Monitors (HPM) CSRs				
0x323 .. 0x32f	mhpmevent3 .. mhpmevent15	CSR_MHPMEVENT3 .. CSR_MHPMEVENT15	MRW	Machine performance-monitoring event select for counter 3..15
0xb03 .. 0xb0f	mhpmcOUNTER3 .. mhpmcOUNTER15	CSR_MHPMCOUNTER3 .. CSR_MHPMCOUNTER15	MRW	Machine performance-monitoring counter 3..15 low word
0xb83 .. 0xb8f	mhpmcOUNTER3h .. mhpmcOUNTER15h	CSR_MHPMCOUNTER3H .. CSR_MHPMCOUNTER15H	MRW	Machine performance-monitoring counter 3..15 high word
Machine Information CSRs				
0xf11	mvendorid	CSR_MVENDORID	MRO	Machine vendor ID
0xf12	marchid	CSR_MARCHID	MRO	Machine architecture ID
0xf13	mimpid	CSR_MIMPID	MRO	Machine implementation ID / version
0xf14	mhartid	CSR_MHARTID	MRO	Machine hardware thread ID
0xf15	mconfigptr	CSR_MCONFIGPTR	MRO	Machine configuration pointer register
NEORV32-Specific CSRs				
0xbc0	mxcsr	CSR_MXCSR	MRW	Machine status and control register
0xfc0	mxisa	CSR_MXISA	MRO	Extended machine CPU ISA and extensions

3.8.1. Floating-Point CSRs

fflags

Name	Floating-point accrued exceptions
Address	<code>0x001</code>
Reset value	<code>0x00000000</code>
ISA	<code>Zicsr & Zfinx</code>
Description	FPU status flags. on

Table 77. `fflags` CSR bits

Bit	R/W	Function
0	r/w	NX : inexact
1	r/w	UF : underflow
2	r/w	OF : overflow
3	r/w	DZ : division by zero
4	r/w	NV : invalid operation

frm

Name	Floating-point dynamic rounding mode
Address	<code>0x002</code>
Reset value	<code>0x00000000</code>
ISA	<code>Zicsr & Zfinx</code>
Description	The <code>frm</code> CSR is used to configure the rounding mode of the FPU. on

Table 78. `frm` CSR bits

Bit	R/W	Function
2:0	r/w	Rounding mode

fcsr

Name	Floating-point control and status register
------	--

Address **0x003**

Reset **0x00000000**

value

ISA **Zicsr & Zfinx**

Description The **fcsr** provides combined access to the **fflags** and **frm** flags.
on

Table 79. **fcsr** CSR bits

Bit	R/W	Function
4:0	r/w	Accrued exception flags (fflags)
7:5	r/w	Rounding mode (frm)

3.8.2. Machine Trap Setup CSRs

`mstatus`

Name	Machine status register - low word
Address	<code>0x300</code>
Reset value	<code>0x00000000</code>
ISA	<code>Zicsr</code>
Descripti on	The <code>mstatus</code> CSR is used to configure general machine environment parameters.

Table 80. `mstatus` CSR bits

Bit	Name [C]	R/W	Function
3	<code>CSR_MSTATUS_MIE</code>	r/w	MIE : Machine-mode interrupt enable flag
7	<code>CSR_MSTATUS_MPIE</code>	r/w	MPIE : Previous machine-mode interrupt enable flag state
12:11	<code>CSR_MSTATUS_MPP_H</code> : <code>CSR_MSTATUS_MPP_L</code>	r/w	MPP : Previous machine privilege mode, <code>11</code> = machine-mode "M", <code>00</code> = user-mode "U"; other values will fall-back to machine-mode
17	<code>CSR_MSTATUS_MPRV</code>	r/w	MPRV : Effective privilege mode for load/stores; use MPP as effective privilege mode when set; hardwired to zero if user-mode not implemented
21	<code>CSR_MSTATUS_TW</code>	r/w	TW : Trap on execution of <code>wfi</code> instruction in user mode when set; hardwired to zero if user-mode not implemented

`misa`

Name	ISA and extensions
Address	<code>0x301</code>
Reset value	DEFINED , according to enabled ISA extensions
ISA	<code>Zicsr</code>
Descripti on	The <code>misa</code> CSR provides information regarding the availability of basic RISC-V ISA extensions.

The NEORV32 `misa` CSR is read-only. Hence, active CPU extensions are entirely defined by pre-synthesis configurations and cannot be switched on/off during runtime. For compatibility reasons any write access to this CSR is simply ignored and will *not* cause an illegal instruction exception.

Table 81. `misa` CSR bits

Bit	Name [C]	R/W	Function
0	<code>CSR_MISA_A_EXT</code>	r/-	A: CPU extension (atomic memory access) available, set when A ISA Extension enabled
1	<code>CSR_MISA_B_EXT</code>	r/-	B: CPU extension (bit-manipulation) available, set when B ISA Extension enabled
2	<code>CSR_MISA_C_EXT</code>	r/-	C: CPU extension (compressed instruction) available, set when C ISA Extension enabled
4	<code>CSR_MISA_E_EXT</code>	r/-	E: CPU extension (embedded) available, set when E ISA Extension enabled
8	<code>CSR_MISA_I_EXT</code>	r/-	I: CPU base ISA, cleared when E ISA Extension enabled
12	<code>CSR_MISA_M_EXT</code>	r/-	M: CPU extension (mul/div) available, set when M ISA Extension enabled
20	<code>CSR_MISA_U_EXT</code>	r/-	U: CPU extension (user mode) available, set when U ISA Extension enabled
23	<code>CSR_MISA_X_EXT</code>	r/-	X: bit is always set to indicate non-standard / NEORV32-specific extensions
31:30	<code>CSR_MISA_MXL_HI_EXT</code>	r/-	MXL : 32-bit architecture indicator (always 01)
:	<code>CSR_MISA_MXL_LO_EXT</code>		



Machine-mode software can discover available **Z*** sub-extensions (like **Zicsr** or **Zfinx**) by checking the NEORV32-specific **mxisa** CSR.

mie

Name	Machine interrupt-enable register
Address	0x304
Reset value	0x00000000
ISA	Zicsr
Description	The mie CSR is used to enable/disable individual interrupt sources.

Table 82. `mie` CSR bits

Bit	Name [C]	R/W	Function
3	<code>CSR_MIE_MSIE</code>	r/w	MSIE : Machine <i>software</i> interrupt enable (from Core-Local Interruptor (CLINT))

Bit	Name [C]	R/W	Function
7	CSR_MIE_MTIE	r/w	MTIE: Machine <i>timer</i> interrupt enable (from Core-Local Interruptor (CLINT))
11	CSR_MIE_MEIE	r/w	MEIE: Machine <i>external</i> interrupt enable
31:16	CSR_MIE_FIRQ15E : CSR_MIE_FIRQ0E	r/w	Fast interrupt channel 15..0 enable

mtvec

Name	Machine trap-handler base address
Address	0x305
Reset value	0x00000000
ISA	Zicsr
Description	The <code>mtvec</code> CSR holds the trap vector configuration.
	on

Table 83. `mtvec` CSR bits

Bit	R/W	Function
1:0	r/w	MODE: mode configuration, 00 = DIRECT, 01 = VECTORED; other encodings are reserved.
31:2	r/w	BASE: in DIRECT mode = 4-byte-aligned base address of trap base handler, all traps jump to pc = BASE; in VECTORED mode = 128-byte-aligned base address of trap vector table, interrupts cause a jump to pc = BASE + 4 * mcause and exceptions a jump to pc = BASE.

Interrupt Latency



The vectored `mtvec` mode is useful for reducing the time between interrupt request (IRQ) and servicing it (ISR). As software does not need to determine the interrupt cause the reduction in latency can be 5 to 10 times and as low as 26 cycles.

mcounteren

Name	Machine counter enable
Address	0x306
Reset value	0x00000000

ISA [Zicsr](#) & U

Description The `mcounteren` CSR is used to constrain user-mode access to the CPU's counter CSRs.
on

Table 84. `mcounteren` CSR bits

Bit	Name [C]	R/W	Function
0	<code>CSR_MCOUNTREN_CY</code>	r/w	CY: User-mode is allowed to read <code>cycle[h]</code> CSRs when set
1	-	r/-	TM: not implemented, hardwired to zero
2	<code>CSR_MCOUNTREN_IR</code>	r/w	IR: User-mode is allowed to read <code>instret[h]</code> CSRs when set
31:3	-	r/-	HPM: hardwired to zero; hardware performance monitors are accessible by machine-mode software only

`mstatush`

Name Machine status register - high word

Address [0x310](#)

Reset [0x00000000](#)
value

ISA [Zicsr](#)

Description The features of this CSR are not implemented yet. The register is read-only and always
on returns zero.

3.8.3. Machine Trap Handling CSRs

`mscratch`

Name	Scratch register for machine trap handlers
Address	<code>0x340</code>
Reset value	<code>0x00000000</code>
ISA	<code>Zicsr</code>
Descripti on	The <code>mscratch</code> is a general-purpose machine-mode scratch register.

`mepc`

Name	Machine exception program counter
Address	<code>0x341</code>
Reset value	<code>0x00000000</code>
ISA	<code>Zicsr</code>
Descripti on	The <code>mepc</code> CSR provides the instruction address where execution has stopped/failed when an interrupt is triggered / an exception is raised. See section Traps, Exceptions and Interrupts for a list of all legal values. The <code>mret</code> instruction will return to the address stored in <code>mepc</code> by automatically moving <code>mepc</code> to the program counter.



`mepc[0]` is hardwired to zero. If IALIGN = 32 (i.e. [C ISA Extension](#) is disabled) then `mepc[1]` is also hardwired to zero.

`mcause`

Name	Machine trap cause
Address	<code>0x342</code>
Reset value	<code>0x00000000</code>
ISA	<code>Zicsr</code>
Descripti on	The <code>mcause</code> CSRs shows the exact cause of a trap. See section Traps, Exceptions and Interrupts for a list of all legal values.



Read-Only

Note that the NEORV32 `mcause` CSR is updated by the hardware only and cannot be written from software. However, any write-access will be ignored and will not cause any exception to maintain RISC-V compatibility.

Table 85. `mcause` CSR bits

Bit	R/W	Function
4:0	r/w	Exception code: see NEORV32 Trap Listing
31	r/w	Interrupt: 1 if the trap is caused by an interrupt (0 if the trap is caused by an exception)

`mtval`

Name	Machine trap value
Address	0x343
Reset value	0x00000000
ISA	Zicsr
Description	The <code>mtval</code> CSR provides additional information why a trap was entered. See section on Traps, Exceptions and Interrupts for more information.

Read-Only



Note that the NEORV32 `mtval` CSR is updated by the hardware only and cannot be written from software. However, any write-access will be ignored and will not cause any exception to maintain RISC-V compatibility.

`mip`

Name	Machine interrupt pending
Address	0x344
Reset value	0x00000000
ISA	Zicsr
Description	The <code>mip</code> CSR shows currently <i>pending</i> machine-mode interrupt requests. Any write access to this register is ignored.

Table 86. `mip` CSR bits

Bit	Name [C]	R/W	Function
3	CSR_MIP_MSIP	r/-	MSIP : Machine <i>software</i> interrupt pending, triggered by <code>msi_i</code> top port (see CPU Top Entity - Signals); cleared by source-specific mechanism
7	CSR_MIP_MTIP	r/-	MTIP : Machine <i>timer</i> interrupt pending, triggered by <code>mei_i</code> top port (see CPU Top Entity - Signals) or by the processor-internal Core-Local Interruptor (CLINT) ; cleared by source-specific mechanism
11	CSR_MIP_MEIP	r/-	MEIP : Machine <i>external</i> interrupt pending, triggered by <code>mti_i</code> top port (see CPU Top Entity - Signals) or by the processor-internal Core-Local Interruptor (CLINT) ; cleared by source-specific mechanism
31:16	CSR_MIP_FIRQ15P : CSR_MIP_FIRQ0P	r/-	FIRQxP : Fast interrupt channel 15..0 pending, see NEORV32-Specific Fast Interrupt Requests ; cleared by source-specific mechanism

*FIRQ Channel Mapping*

See section [NEORV32-Specific Fast Interrupt Requests](#) for the mapping of the FIRQ channels and the according interrupt-triggering processor module.

mtinst

Name	Machine trap instruction
Address	0x34a
Reset value	0x00000000
ISA	Zicsr
Description	The <code>mtinst</code> CSR provides additional information why a trap was entered. See section Traps, Exceptions and Interrupts for more information.

Read-Only

Note that the NEORV32 `mtinst` CSR is updated by the hardware only and cannot be written from software. However, any write-access will be ignored and will not cause any exception to maintain RISC-V compatibility.

Instruction Transformation

The RISC-V priv. spec. suggests that the instruction word written to `mtinst` by the hardware should be "transformed". However, the NEORV32 `mtinst` CSR uses a simplified transformation scheme: if the trap-causing instruction is a standard 32-bit instruction, `mtinst` contains the exact instruction word that caused the trap. If

the trap-causing instruction is a compressed instruction, `mtinst` contains the de-compressed 32-bit equivalent with bit 1 being cleared while all remaining bits represent the pre-decoded 32-bit instruction equivalent.

3.8.4. Machine Configuration CSRs

menvcfg

Name	Machine environment configuration register - low word
Address	0x30a
Reset value	0x00000000
ISA	Zicsr & U
Description	Currently, the features of this CSR are not supported. Hence, the entire register is hardwired to all-zero.

menvcfg

Name	Machine environment configuration register - high word
Address	0x31a
Reset value	0x00000000
ISA	Zicsr & U
Description	Currently, the features of this CSR are not supported. Hence, the entire register is hardwired to all-zero.

3.8.5. Machine Physical Memory Protection CSRs

The physical memory protection system is configured via the `PMP_NUM_REGIONS` and `PMP_MIN_GRANULARITY` top entity generics. `PMP_NUM_REGIONS` defines the total number of implemented regions. Note that the maximum number of regions is constrained to 16. If trying to access a PMP-related CSR beyond `PMP_NUM_REGIONS` (but below PMP configuration/address register 16) **no illegal instruction exception** is triggered. The according CSRs are read-only (writes are ignored) and always return zero. See section [Smpmp ISA Extension](#) for more information.

pmpcfg

Name	Physical memory protection region configuration registers
Address	<code>0x3a0</code> (<code>pmpcfg0</code>) <code>0x3a1</code> (<code>pmpcfg1</code>) <code>0x3a2</code> (<code>pmpcfg2</code>) <code>0x3a3</code> (<code>pmpcfg3</code>)
Reset value	all <code>0x00000000</code>
ISA	<code>Zicsr</code> & <code>Smpmp</code>
Description	Configuration of physical memory protection regions. Each region provides an individual 8-bit array in these CSRs. Note that, depending on <code>PMP_NUM_REGIONS</code> , a maximum of the lowest 4 configuration registers are physically implemented (<code>pmpcfg0</code> .. <code>pmpcfg3</code>). All remaining register are hardwired to zero.

Table 87. `pmpcfg*` CSR Bits

Bit	Name [C]	R/W	Function
0	<code>PMPCFG_R</code>	r/w	R: Read permission
1	<code>PMPCFG_W</code>	r/w	W: Write permission
2	<code>PMPCFG_X</code>	r/w	X: Execute permission
4:3	<code>PMPCFG_A_MS</code>	r/w	A: Mode configuration (<code>00</code> = OFF, <code>01</code> = TOR, <code>10</code> = NA4, <code>11</code> = NAPOT) B : <code>PMPCFG_A_LS</code> B
7	<code>PMPCFG_L</code>	r/w	L: Lock bit, prevents further write accesses, also enforces access rights in machine-mode, can only be cleared by CPU reset

Implemented Modes



In order to reduce the CPU size certain PMP modes (**A** bits) can be excluded from synthesis. Use the `PMP_TOR_MODE_EN` and `PMP_NAP_MODE_EN` Processor Top Entity - Generics to control implementation of the according modes.

pmpaddr

Name	Physical memory protection region address registers
Address	0x3b0 (pmpaddr1) 0x3b1 (pmpaddr2) 0x3b2 (pmpaddr3) 0x3b3 (pmpaddr4) 0x3b4 (pmpaddr5) 0x3b5 (pmpaddr6) 0x3b6 (pmpaddr6) 0x3b7 (pmpaddr7) 0x3b8 (pmpaddr8) 0x3b9 (pmpaddr9) 0x3ba (pmpaddr10) 0x3bb (pmpaddr11) 0x3bc (pmpaddr12) 0x3bd (pmpaddr13) 0x3be (pmpaddr14) 0x3bf (pmpaddr15)
Reset value	all 0x00000000
ISA	Zicsr & Smpmp
Description	Region address/boundaries configuration. Note that, depending on PMP_NUM_REGIONS, a maximum of the lowest 16 address registers are physically implemented (pmpaddr0 .. pmpaddr15). All remaining register are hardwired to zero.

Table 88. pmpaddr* CSR Bits

Bit	R/W	Description
31:30	r-w	address bits 33 downto 32', hardwired to zero
29:0	r/w	address bits 31 downto 2

3.8.6. (Machine) Counter and Timer CSRs

`time[h]` CSRs (Wall Clock Time)



The NEORV32 does not implement the user-mode `time[h]` registers. Any access to these registers will trap. It is recommended that the trap handler software provides a means of accessing the machine timer off the [Core-Local Interruptor \(CLINT\)](#).



Instruction Retired Counter Increment

The `[m]instret[h]` counter always increments when a instruction enters the pipeline's execute stage no matter if this instruction is actually going to retire or if it causes an exception.

`cycle[h]`

Name	Cycle counter
Address	<code>0xc00 (cycle)</code> <code>0xc80 (cycleh)</code>
Reset value	<code>0x00000000_00000000</code>
ISA	<code>Zicsr & Zicntr</code>
Description	The <code>cycle[h]</code> CSRs are user-mode shadow copies of the according <code>mcycle[h]</code> CSRs. The user-mode counter are read-only. Any write access will raise an illegal instruction exception.

`instret[h]`

Name	Instructions-retired counter
Address	<code>0xc02 (instret)</code> <code>0xc82 (instreth)</code>
Reset value	<code>0x00000000_00000000</code>
ISA	<code>Zicsr & Zicntr</code>
Description	The <code>instret[h]</code> CSRs are user-mode shadow copies of the according <code>minstret[h]</code> CSRs. The user-mode counter are read-only. Any write access will raise an illegal instruction exception.

mcycle[h]

Name	Machine cycle counter
Address	<code>0xb00</code> (<code>mcycle</code>) <code>0xb80</code> (<code>mcycleh</code>)
Reset value	<code>0x00000000_00000000</code>
ISA	<code>Zicsr</code> & <code>Zicntr</code>
Description	If not halted via the <code>mcountinhibit</code> CSR the <code>cycle[h]</code> CSRs will increment with every active CPU clock cycle (CPU not in Sleep Mode). These registers are read/write only for machine-mode software.

minstret[h]

Name	Machine instructions-retired counter
Address	<code>0xb02</code> (<code>minstret</code>) <code>0xb82</code> (<code>minstreth</code>)
Reset value	<code>0x00000000_00000000</code>
ISA	<code>Zicsr</code> & <code>Zicntr</code>
Description	If not halted via the <code>mcountinhibit</code> CSR the <code>minstret[h]</code> CSRs will increment with every retired instruction. These registers are read/write only for machine-mode software

Instruction Retiring

Note that **all** executed instruction do increment the `[m]instret[h]` counters even if they do not retire (e.g. if the instruction causes an exception).

3.8.7. Hardware Performance Monitors (HPM) CSRs

Machine-Mode HPMs Only



Note that only the machine-mode hardware performance counter CSR are available (`mhpcounter*[h]`). Accessing any user-mode HPM CSR (`hpmcounter*[h]`) will raise an illegal instruction exception.

The actual number of implemented hardware performance monitors is configured via the `HPM_NUM_CNTS` top entity generic. Note that always all 13 HPM counter and configuration registers (`mhpcounter*[h]`) are implemented, but only the actually configured ones are implemented as "real" physical registers - the remaining ones will be hardwired to zero. If trying to access an HPM-related CSR beyond `HPM_NUM_CNTS` **no illegal instruction exception is triggered**. These CSRs are read-only, writes are ignored and reads always return zero.

The total counter width of the HPMs can be configured before synthesis via the `HPM_CNT_WIDTH` generic (0..64-bit). If `HPM_NUM_CNTS` is less than 64, all remaining MSB-aligned bits are hardwired to zero.

`mhpmevent`

Name	Machine hardware performance monitor event select
Address	0x233 (<code>mhpmevent3</code>) 0x234 (<code>mhpmevent4</code>) 0x235 (<code>mhpmevent5</code>) 0x236 (<code>mhpmevent6</code>) 0x237 (<code>mhpmevent7</code>) 0x238 (<code>mhpmevent8</code>) 0x239 (<code>mhpmevent9</code>) 0x23a (<code>mhpmevent10</code>) 0x23b (<code>mhpmevent11</code>) 0x23c (<code>mhpmevent12</code>) 0x23d (<code>mhpmevent13</code>) 0x23e (<code>mhpmevent14</code>) 0x23f (<code>mhpmevent15</code>)
Reset value	all 0x00000000
ISA	<code>Zicsr</code> & <code>Zihpm</code>

Descripti on The value in these CSRs define the micro-architectural events that cause an increment of the according `mhpmevent*[h]` counter(s). All available events are listed in the table below. If more than one event is enabled, the according counter will increment if *any* of the enabled events is observed (logical OR). Note that the counter will only increment by 1 step per clock cycle even if more than one trigger event is observed.

Table 89. `mhpmevent*` CSR Bits (Micro-Architectural Counter Events)

Bit	Name [C]	R/W	Event Description
RISC-V-compatible			
0	<code>HPMCNT_EVENT_CY</code>	r/w	active clock cycle (CPU not in Sleep Mode)
1	<code>HPMCNT_EVENT_TM</code>	r/-	<i>not implemented</i> , hardwired to zero
2	<code>HPMCNT_EVENT_IR</code>	r/w	any executed instruction (16-bit/compressed or 32-bit/uncompressed)
NEORV32-specific			
3	<code>HPMCNT_EVENT_COMPR</code>	r/w	any executed 16-bit/compressed (C ISA Extension) instruction
4	<code>HPMCNT_EVENT_WAIT_IS</code>	r/w	instruction dispatch wait cycle (wait for instruction prefetch-buffer refill (CPU Control Unit IPB); caused by a fence instruction, a control flow transfer or a instruction fetch bus wait cycle)
5	<code>HPMCNT_EVENT_WAIT_ALU</code>	r/w	any delay/wait cycle caused by a <i>multi-cycle</i> CPU Arithmetic Logic Unit operation
6	<code>HPMCNT_EVENT_BRANCH</code>	r/w	any executed branch instruction (unconditional, conditional-taken or conditional-not-taken)
7	<code>HPMCNT_EVENT_BRANCH_ED</code>	r/w	any control transfer operation (unconditional jump, taken conditional branch or trap entry/exit)
8	<code>HPMCNT_EVENT_LOAD</code>	r/w	any executed load operation (read-modify-write AMOs are counted as one load and one store operation)
9	<code>HPMCNT_EVENT_STORE</code>	r/w	any executed store operation (read-modify-write AMOs are counted as one load and one store operation)
10	<code>HPMCNT_EVENT_WAIT_LSU</code>	r/w	any memory/bus/cache/etc. delay/wait cycle while executing any load or store operation (caused by a data bus wait cycle))
11	<code>HPMCNT_EVENT_TRAP</code>	r/w	starting processing of any trap (Traps, Exceptions and Interrupts)

Instruction Retiring ("Retired == Executed")



The CPU HPM/counter logic treats all executed instruction as "retired" even if they raise an exception, cause an interrupt, trigger a privilege mode change or were not meant to retire (i.e. claimed by the RISC-V spec.).

Atomic Memory Access

The read-modify-write instructions of the [Zaamo ISA Extension](#) operate as simple load for the CPU hardware. Hence, they will only trigger `HPMCNT_EVENT_LOAD` and only once.

mhpmcouter[h]

Name	Machine hardware performance monitor (HPM) counter
Address	<code>0xb03, 0xb83</code> (<code>mhpmcouter3, mhpmcouter3h</code>) <code>0xb04, 0xb84</code> (<code>mhpmcouter4, mhpmcouter4h</code>) <code>0xb05, 0xb85</code> (<code>mhpmcouter5, mhpmcouter5h</code>) <code>0xb06, 0xb86</code> (<code>mhpmcouter6, mhpmcouter6h</code>) <code>0xb07, 0xb87</code> (<code>mhpmcouter7, mhpmcouter7h</code>) <code>0xb08, 0xb88</code> (<code>mhpmcouter8, mhpmcouter8h</code>) <code>0xb09, 0xb89</code> (<code>mhpmcouter9, mhpmcouter9h</code>) <code>0xb0a, 0xb8a</code> (<code>mhpmcouter10, mhpmcouter10h</code>) <code>0xb0b, 0xb8b</code> (<code>mhpmcouter11, mhpmcouter11h</code>) <code>0xb0c, 0xb8c</code> (<code>mhpmcouter12, mhpmcouter12h</code>) <code>0xb0d, 0xb8d</code> (<code>mhpmcouter13, mhpmcouter13h</code>) <code>0xb0e, 0xb8e</code> (<code>mhpmcouter14, mhpmcouter14h</code>) <code>0xb0f, 0xb8f</code> (<code>mhpmcouter15, mhpmcouter15h</code>)
Reset value	all <code>0x00000000_00000000</code>
ISA	<code>Zicsr & Zihpm</code>
Description	If not halted via the <code>mcountinhibit</code> CSR the HPM counter CSRs increment whenever the configured events from the according <code>mhpmevent</code> CSR occur. The HPM counter registers are read/write for machine mode software and are not accessible for lower-privileged software.

3.8.8. Machine Counter Setup CSRs

mcountinhibit

Name	Machine counter-inhibit register
Address	0x320
Reset value	0x00000000
ISA	Zicsr
Description	Set bit to halt the according counter CSR. on

Table 90. mcountinhibit CSR Bits

Bit	Name [C]	R/W	Description
0	CSR_MCOUNTINHIBIT_I R	r/w	IR: Set to 1 to halt [m]instret[h]; hardwired to zero if Zicntr ISA extension is disabled
1	-	r/-	TM: Hardwired to zero as time[h] CSRs are not implemented
2	CSR_MCOUNTINHIBIT_C Y	r/w	CY: Set to 1 to halt [m]cycle[h]; hardwired to zero if Zicntr ISA extension is disabled
15:3	CSR_MCOUNTINHIBIT_H PM3 : CSR_MCOUNTINHIBIT_H PM15	r/w	HPMx: Set to 1 to halt [m]hpmcount*[h]; hardwired to zero if Zihpm ISA extension is disabled

3.8.9. Machine Information CSRs

mvendorid

Name	Machine vendor ID
Address	<code>0xf11</code>
Reset value	<code>0x00000000</code>
ISA	<code>Zicsr</code>
Description	Read-only and hardwired to zero.
on	

marchid

Name	Machine architecture ID
Address	<code>0xf12</code>
Reset value	<code>0x00000013</code>
ISA	<code>Zicsr</code>
Description	The marchid CSR is read-only and provides the NEORV32 official RISC-V open-source architecture ID (decimal: 19, 32-bit hexadecimal: 0x00000013).
on	

mimpid

Name	Machine implementation ID
Address	<code>0xf13</code>
Reset value	DEFINED
ISA	<code>Zicsr</code>
Description	The mimpid CSR is read-only and provides the version of the NEORV32 as BCD-coded number (example: <code>mimpid = 0x01020312</code> → 01.02.03.12 → version 1.2.3.12).
on	

mhartid

Name	Machine hardware thread ID
Address	<code>0xf14</code>

Reset **DEFINED**
value

ISA **Zicsr**

Description The **mhartid** CSR is read-only and provides the core's hart ID. In a multi-core system, the hard ID of each core is unique and starts at 0 and is incremented continuously.

mconfigptr

Name Machine configuration pointer register

Address **0xf15**

Reset **0x00000000**
value

ISA **Zicsr**

Description The features of this CSR are not implemented yet. The register is read-only and always returns zero.

3.8.10. NEORV32-Specific CSRs

RISC-V-Compliant Mapping



All NEORV32-specific CSRs are mapped to addresses that are explicitly reserved for custom/implementation-specific use.

`mxcsr`

Name	Machine status and control register
Address	<code>0xbc0</code>
Reset value	DEFINED
ISA	<code>Zicsr & X</code>
Description	The <code>mxcsr</code> CSR is a NEORV32-specific read/write CSR that provides additional machine configurations and status information.

Table 91. `mxcsr` CSR Bits

Bit	Name [C]	R/W	Description
25:0	-	r/-	reserved, read as zero
26	<code>CSR_MXCSR_TRACE</code>	r/-	CPU Tuning Options: CPU trace generator and Execution Trace Port enabled (implicitly enabled when enabling Execution Trace Buffer (TRACER))
27	<code>CSR_MXCSR_CONSTTBR</code>	r/-	CPU Tuning Options: constant-time branches enabled when set (<code>CPU_CONSTT_BR_EN</code> top generic)
28	<code>CSR_MXCSR_RFHWRST</code>	r/-	CPU Tuning Options: full hardware reset of register file available when set (<code>CPU_RF_HW_RST_EN</code> top generic)
29	<code>CSR_MXCSR_FASTMUL</code>	r/-	CPU Tuning Options: fast multiplication available when set (<code>CPU_FAST_MUL_EN</code> top generic)
30	<code>CSR_MXCSR_FASTSHIFT</code>	r/-	CPU Tuning Options: fast shifts available when set (<code>CPU_FAST_SHIFT_EN</code> top generic)
31	<code>CSR_MXCSR_IS_SIM</code>	r/-	set if CPU is being simulated (not guaranteed)

`mxisa`

Name	Machine extended ISA and extensions register
Address	<code>0xfc0</code>
Reset value	DEFINED

ISA [Zicsr](#) & [X](#)

Description The [mxisa](#) CSR is a NEORV32-specific read-only CSR that helps machine-mode software to discover additional ISA (sub-)extensions and CPU configuration options.

Table 92. mxisa CSR Bits

Bit	Name [C]	R/W	Description
0	CSR_MXISA_ZICCSR	r/-	Zicsr ISA Extension available
1	CSR_MXISA_ZIFENCEI	r/-	Zifencei ISA Extension available
2	CSR_MXISA_ZMMUL	r/-	Zmmul - ISA Extension available
3	CSR_MXISA_ZXCFU	r/-	Zxfc ISA Extension available
4	CSR_MXISA_ZKT	r/-	Zkt ISA Extension available
5	CSR_MXISA_ZFINX	r/-	Zfinx ISA Extension available
6	CSR_MXISA_ZICOND	r/-	Zicond ISA Extension available
7	CSR_MXISA_ZICNTR	r/-	Zicntr ISA Extension available
8	CSR_MXISA_SMPMP	r/-	Smpmp ISA Extension available
9	CSR_MXISA_ZIHPM	r/-	Zihpm ISA Extension available
10	CSR_MXISA_SDEXT	r/-	Sdext ISA Extension available
11	CSR_MXISA_SDTRIG	r/-	Sdtrig ISA Extension available
12	CSR_MXISA_ZBKX	r/-	Zbkx ISA Extension available
13	CSR_MXISA_ZKND	r/-	Zknd ISA Extension available
14	CSR_MXISA_ZKNE	r/-	Zkne ISA Extension available
15	CSR_MXISA_ZKNH	r/-	Zknh ISA Extension available
16	CSR_MXISA_ZBKB	r/-	Zbkb ISA Extension available
17	CSR_MXISA_ZBKC	r/-	Zbkc ISA Extension available
18	CSR_MXISA_ZKN	r/-	Zkn ISA Extension available
19	CSR_MXISA_ZKSH	r/-	Zksh ISA Extension available
20	CSR_MXISA_ZKSED	r/-	Zksed ISA Extension available
21	CSR_MXISA_ZKS	r/-	Zks ISA Extension available
22	CSR_MXISA_ZBA	r/-	Zba ISA Extension available
23	CSR_MXISA_ZBB	r/-	Zbb ISA Extension available
24	CSR_MXISA_ZBS	r/-	Zbs ISA Extension available
25	CSR_MXISA_ZAAMO	r/-	Zaamo ISA Extension available
26	CSR_MXISA_ZALRSC	r/-	Zalrsc ISA Extension available
27	CSR_MXISA_ZCB	r/-	Zcb ISA Extension available

Bit	Name [C]	R/W	Description
28	CSR_MXISA_ZCA	r/-	"`C` without floating-point", available when C ISA Extension is available
29	CSR_MXISA_ZIBI	r/-	Zibi ISA Extension available
30	CSR_MXISA_ZIMOP	r/-	Zimop ISA Extension available
31	-	r/-	<i>reserved</i> , read as zero

3.9. Traps, Exceptions and Interrupts

In this document the following terminology is used (derived from the RISC-V trace specification available at <https://github.com/riscv-non-isa/riscv-trace-spec>):

- **exception:** an unusual condition occurring at run time associated (i.e. *synchronous*) with an instruction in a RISC-V hart
- **interrupt:** an external *asynchronous* event that may cause a RISC-V hart to experience an unexpected transfer of control
- **trap:** the transfer of control to a trap handler caused by either an *exception* or an *interrupt*

Whenever an exception or interrupt is triggered, the CPU switches to machine-mode (if not already in machine-mode) and continues operation at the address being stored in the **mtvec** CSR. The cause of the trap can be determined via the **mcause** CSR. A list of all implemented **mcause** values and the according description can be found below in section [NEORV32 Trap Listing](#). The address that reflects the current program counter when a trap was taken is stored to **mepc** CSR. Additional information regarding the cause of the trap can be retrieved from the **mtval** and **mtinst** CSRs.

The traps are prioritized. If several *exceptions* occur at once only the one with highest priority is triggered while all remaining exceptions are ignored and discarded. If several *interrupts* trigger at once, the one with highest priority is serviced first while the remaining ones stay *pending*. After completing the interrupt handler the interrupt with the second highest priority will get serviced and so on until no further interrupts are pending.

Interrupts when in User-Mode



If the core is currently operating in less privileged user-mode, interrupts are globally enabled even if **mstatus.mie** is cleared.

Interrupt Signal Requirements - Standard RISC-V Interrupts



All interrupt request signals are **high-active**. Once triggered, a interrupt request line should stay high until it is explicitly acknowledged by a source-specific mechanism (for example by writing to a specific memory-mapped register).

Instruction Atomicity and Forward-Progress



All instructions execute as atomic operations - interrupts can only trigger *between* consecutive instructions. Additionally, if there is a permanent interrupt request, exactly one instruction from the interrupted program will be executed before another interrupt handler can start. This allows program progress even if there are permanent interrupt requests.

3.9.1. Memory Access Exceptions

If a load operation causes any exception, the instruction's destination register is **not written** at all. Furthermore, exceptions caused by a misaligned memory address a physical memory protection

fault do not trigger a memory access request at all.

For 32-bit-only instructions (= no `C` extension) the misaligned instruction exception is raised if bit 1 of the fetch address is set (i.e. not on a 32-bit boundary). If the `C` extension is implemented there will **never** be a misaligned instruction exception at all.

3.9.2. Nested Interrupts

The CPU supports nested interrupt handling in software. It automatically disables interrupts upon entering any trap handler by clearing `mstatus.MIE`. Otherwise, further interrupts during the critical part of the handler (i.e. before saving the context: `mcause`, `mepc`, stack frame, etc.) would cause a context loss. However, software can explicitly enable interrupts again by manually setting `mstatus.MIE` from within the handler to allow another interrupt to trigger.

3.9.3. Custom Fast Interrupt Request Lines

As a custom extension, the NEORV32 CPU features 16 fast interrupt request (FIRQ) lines via the `firq_i` CPU top entity signals. These interrupts have custom configuration and status flags in the `mie` and `mip` CSRs and also provide custom trap codes in `mcause`. These FIRQs are reserved for NEORV32 processor-internal usage only.

3.9.4. NEORV32 Trap Listing

The following tables show all traps that are currently supported by the NEORV32 CPU. It also shows the prioritization and the CSR side-effects.



FIRQ Mapping

See section [NEORV32-Specific Fast Interrupt Requests](#) for the mapping of the FIRQ channels to the according hardware modules.

Table Annotations

The "RTE Trap ID" aliases are defined by the NEORV32 core library and can be used in plain C code when interacting with the pre-defined RTE functions. The `mcause`, `mepc`, `mtval`, and `mtinst` columns show the value being written to the according CSRs when a trap is encountered:

- **I-PC** - address of intercepted instruction (instruction at this address has not been executed yet and must be executed after the trap handler to maintain program flow)
- **PC** - address of instruction that caused the trap (instruction has been executed)
- **ADDR** - bad data memory access address that caused the trap
- **INST** - the actual transformed/decompressed instruction word that caused the trap
- **LAST** - the last transformed/decompressed instruction word that was executed before the trap
- **0** - zero

Table 93. NEORV32 Trap Listing

Priority	mcause	RTE Trap ID	Cause	mepc	mtval	mtinst
Exceptions (synchronous to instruction execution)						
highest	0x00000010	TRAP_CODE_DOUBLE_TRAP	double trap (trap inside trap)	I-PC	0	INST
Interrupts (asynchronous to instruction execution)						
	0x80000010	TRAP_CODE_FIRQ_0	fast interrupt request channel 0	I-PC	0	LAST
	0x80000011	TRAP_CODE_FIRQ_1	fast interrupt request channel 1	I-PC	0	LAST
	0x80000012	TRAP_CODE_FIRQ_2	fast interrupt request channel 2	I-PC	0	LAST
	0x80000013	TRAP_CODE_FIRQ_3	fast interrupt request channel 3	I-PC	0	LAST
	0x80000014	TRAP_CODE_FIRQ_4	fast interrupt request channel 4	I-PC	0	LAST
	0x80000015	TRAP_CODE_FIRQ_5	fast interrupt request channel 5	I-PC	0	LAST
	0x80000016	TRAP_CODE_FIRQ_6	fast interrupt request channel 6	I-PC	0	LAST
	0x80000017	TRAP_CODE_FIRQ_7	fast interrupt request channel 7	I-PC	0	LAST
	0x80000018	TRAP_CODE_FIRQ_8	fast interrupt request channel 8	I-PC	0	LAST
	0x80000019	TRAP_CODE_FIRQ_9	fast interrupt request channel 9	I-PC	0	LAST
	0x8000001a	TRAP_CODE_FIRQ_10	fast interrupt request channel 10	I-PC	0	LAST
	0x8000001b	TRAP_CODE_FIRQ_11	fast interrupt request channel 11	I-PC	0	LAST

Priority	mcause	RTE Trap ID	Cause	mepc	mtval	mtins
lowest	0x80000007	TRAP_CODE_MTI	machine timer interrupt (MTI)	I-PC	0	LAST
worst	0x80000003	TRAP_CODE_MSI	machine software interrupt (MSI)	I-PC	0	LAST
0x8000000b	TRAP_CODE_MEI		machine external interrupt (MEI)	I-PC	0	LAST
0x8000001f	TRAP_CODE_FIRQ_15		fast interrupt request channel 15	I-PC	0	LAST
0x8000001e	TRAP_CODE_FIRQ_14		fast interrupt request channel 14	I-PC	0	LAST
0x8000001d	TRAP_CODE_FIRQ_13		fast interrupt request channel 13	I-PC	0	LAST
0x8000001c	TRAP_CODE_FIRQ_12		fast interrupt request channel 12	I-PC	0	LAST

Table 94. NEORV32 Trap Description

Trap ID [C]	Triggered when ...
TRAP_CODE_I_ACCESS	bus timeout, bus access error or PMP rule violation during instruction fetch
TRAP_CODE_I_ILLEGAL	trying to execute an invalid instruction word (malformed or not supported) or on a privilege violation
TRAP_CODE_I_MISALIGNED	fetching a 32-bit instruction word that is not 32-bit-aligned (see note below)
TRAP_CODE_MENV_CALL	executing <code>ecall</code> instruction in machine-mode
TRAP_CODE_UENV_CALL	executing <code>ecall</code> instruction in user-mode
TRAP_CODE_BREAKPOINT	executing <code>ebreak</code> instruction or if Trigger Module fires
TRAP_CODE_S_MISALIGNED	storing data to an address that is not naturally aligned to the data size (half/word)
TRAP_CODE_L_MISALIGNED	loading data from an address that is not naturally aligned to the data size (half/word)
TRAP_CODE_S_ACCESS	bus timeout, bus access error or PMP rule violation during store data operation
TRAP_CODE_L_ACCESS	bus timeout, bus access error or PMP rule violation during load data operation
TRAP_CODE_FIRQ_*	caused by interrupt-condition of processor-internal modules , see NEORV32-Specific Fast Interrupt Requests

Trap ID [C]	Triggered when ...
TRAP_CODE_MEI	machine external interrupt (via dedicated Processor Top Entity - Signals)
TRAP_CODE_MSI	machine software interrupt (internal Core-Local Interruptor (CLINT) or via dedicated Processor Top Entity - Signals)
TRAP_CODE_MTI	machine timer interrupt (internal Core-Local Interruptor (CLINT) or via dedicated Processor Top Entity - Signals)

*Resumable Exceptions*

Note that not all exceptions are resumable. The "instruction access fault" and the "instruction address misaligned" exceptions are not resumable in most cases.

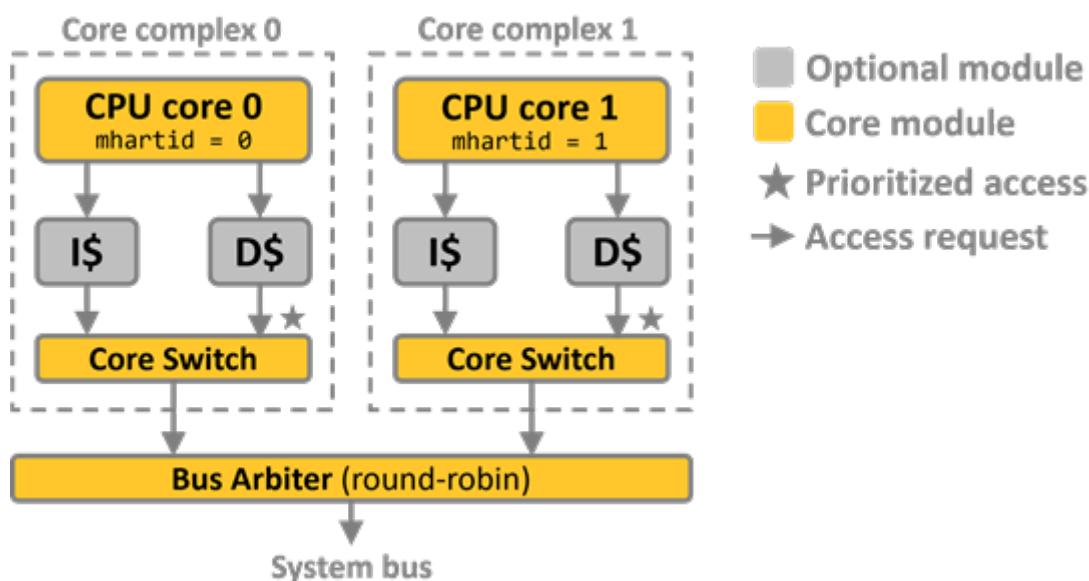
3.10. Dual-Core Configuration

Dual-Core Example Programs



A set of rather simple dual-core example programs can be found in [sw/example/demo_dual_core*](#).

Optionally, the CPU core can be implemented as **symmetric multiprocessing (SMP) dual-core** system. This dual-core configuration is enabled by the [DUAL_CORE_EN top generic](#). When enabled, two *core complexes* are implemented. Each core complex consists of a CPU core and optional instruction (**I\$**) and data (**D\$**) caches. Similar to the single-core [Bus System](#), the instruction and data interfaces are switched into a single bus interface by a prioritizing bus switch. The bus interfaces of both core complexes are further switched into a single system bus using a round-robin arbiter.



Both CPU cores are fully identical and use the same ISA, tuning and cache configurations provided by the according [top generics](#). However, each core can be identified by the according "hart ID" that can be retrieved from the [mhartid](#) CSR. CPU core 0 (the *primary core*) has [mhartid = 0](#) while core 1 (the *secondary core*) has [mhartid = 1](#).

The following table summarizes the most important aspects when using the dual-core configuration.

CPU configuration	Both cores use the same cache, CPU and ISA configuration provided by the according top generics.
Debugging	A special SMP openOCD script (sw/openocd/openocd_neorv32.dual_core.cfg) is required to debug both cores at once. SMP-debugging is fully supported by the RISC-V gdb port.
Clock and reset	Both cores use the same global processor clock and reset.
Address space	Both cores have full access to the same physical Address Space .

Interrupts	All Processor Interrupts are routed to both cores. Hence, each core has access to all NEORV32-Specific Fast Interrupt Requests (FIRQs). Additionally, the RISC-V machine-level <i>external interrupt</i> (via the top <code>mext_irq_i</code> port) is also send to both cores. In contrast, the RISC-V machine level <i>software</i> and <i>timer</i> interrupts are core-exclusive (provided by the Core-Local Interruptor (CLINT)).
RTE	The NEORV32 Runtime Environment can be used for both cores. However, the RTE needs to be explicitly initialized on each core (executing <code>neorv32_rte_setup()</code>). Note that the installed trap handlers apply to both cores. The installed user-defined trap handlers can check the according core's ID via the <code>mhartid</code> CSR to perform core-specific trap handling.
Memory	Each core has its own stack. The top of stack of core 0 is defined by the Linker Script while the top of stack of core 1 has to be explicitly defined by core 0 (see Dual-Core Boot). Both cores share the same heap, <code>.data</code> and <code>.bss</code> sections. Hence, only core 0 setups the <code>.data</code> and <code>.bss</code> sections at boot-up.
Constructors and destructors	Constructors and destructors are executed by core 0 only (see section C Standard Library).
Cache layout	If enabled, each CPU core has its own data and/or instruction cache.
Cache coherency	Be aware that there is no cache snooping available. If any CPU1 cache is enabled care must be taken to prevent access to outdated data - either by using cache synchronization (<code>fence</code> / <code>fence.i</code> instructions) or by using atomic memory accesses. See Memory Coherence for more information.
Bootloader	Only core 0 will boot and execute the bootloader while core 1 is held in standby.
Booting	See section Dual-Core Boot .

3.10.1. Dual-Core Boot

After reset, both cores start booting. However, core 1 will - regardless of the **Boot Configuration** - always enter **Sleep Mode** right inside the default **Start-Up Code (crt0)**. The primary core (core 0) will continue booting, executing either the **Bootloader** or the pre-installed image from the internal instruction memory (depending on the boot configuration).

To boot-up core 1, the primary core has to use a special library function provided by the NEORV32 software framework:

Listing 21. CPU Core 1 Launch Function Prototype (note that this function can only be executed on core 0)

```
int neorv32_smp_launch(int (*entry_point)(void), uint8_t* stack_memory, size_t
stack_size_bytes);
```

When executed, core 0 uses the two 32-bit `MTIMECMP` registers of the **Core-Local Interruptor (CLINT)** to store the *launch configuration* at a defined address location. This launch configuration consists of the stack configuration (via `stack_memory` and `stack_size_bytes`) and the actual entry point for core

- After these registers have been populated, core 1 will trigger core 1's software interrupt (also via the CLINT) to wake it from sleep mode. After that, core 1 will fetch the launch configuration and will start execution at the configured entry point.

Listing 22. CPU Core 1 Main Function

```
int core1_main(void) { // return 'int', no arguments
    return 0; // return to crt0 and go to sleep mode
}
```

Core 1 Stack Memory



The memory for the stack of core 1 (`stack_memory`) can be either statically allocated (i.e. a global volatile memory array; placed in the `.data` or `.bss` section of core 0) or dynamically allocated (using `malloc`; placed on the heap of core 0). In any case the memory should be aligned to a 16-byte boundary.

After that, the primary core triggers the *machine software interrupt* of core 1 using the **Core-Local Interruptor (CLINT)**. Core 1 wakes up from sleep mode, consumes the configuration structure and finally starts executing at the provided entry point. When `neorv32_smp_launch()` returns (with no error code) the secondary core is online and running.

Chapter 4. Software Framework

The NEORV32 project comes with a complete software ecosystem called the "software framework" which is based on the C-language RISC-V GCC port and consists of the following parts:

- Compiler Toolchain
- Core Libraries
- System View Description File (SVD)
- Application Makefile
- Default Compiler Flags
- Linker Script
- C Standard Library
- Start-Up Code (crt0)
- Executable Image Formats
- NEORV32 Runtime Environment
- Bootloader

Software Documentation



All core libraries and example programs are documented *in-code* using **Doxygen**. The API documentation is automatically built and deployed to GitHub pages: <https://stnolting.github.io/neorv32/sw/files.html>.

Example Programs



A collection of annotated example programs illustrating how to use certain NEORV32 functions and peripheral/IO modules can be found in [sw/example](#).

4.1. Compiler Toolchain

The toolchain for this project is based on the free and open RISC-V GCC-port. You can find the compiler sources and build instructions in the official RISC-V GNU toolchain GitHub repository: <https://github.com/riscv-collab/riscv-gnu-toolchain>.

Toolchain Installation



More information regarding the toolchain (building from scratch or downloading prebuilt ones) can be found in the user guide section [Software Toolchain Setup](#).

4.2. Core Libraries

The NEORV32 project provides a set of pre-defined C libraries that allow an easy integration of the processor/CPU features (also called "HAL" - *hardware abstraction layer*). All driver and runtime-related files are located in `sw/lib`. These library files are automatically included and linked by adding the following include statement:

```
#include <neorv32.h> // NEORV32 HAL, core and runtime libraries
```

The NEORV32 HAL consists of the following files.

Table 95. NEORV32 Hardware Abstraction Layer File List

C source file	C header file	Description
-	neorv32.h	Main NEORV32 library file
neorv32_aux.c	neorv32_aux.h	General auxiliary/helper function
neorv32_cfs.c	neorv32_cfs.h	Custom Functions Subsystem (CFS) HAL
neorv32_clint.c	neorv32_clint.h	Core-Local Interruptor (CLINT) HAL
neorv32_cpu.c	neorv32_cpu.h	NEORV32 Central Processing Unit (CPU) HAL
	neorv32_csr.h	Control and Status Registers (CSRs) definitions
neorv32_cfu.c	neorv32_cfu.h	Custom Functions Unit (CFU) HAL
neorv32_dma.c	neorv32_dma.h	Direct Memory Access Controller (DMA) HAL
neorv32_gpio.c	neorv32_gpio.h	General Purpose Input and Output Port (GPIO) HAL
neorv32_gptmr.c	neorv32_gptmr.h	General Purpose Timer (GPTMR) HAL
-	neorv32_intrinsics.h	Macros for intrinsics and custom instructions
-	neorv32_legacy.h	Legacy / backwards-compatibility wrappers (do not use for new designs)
neorv32_neoled.c	neorv32_neoled.h	Smart LED Interface (NEOLED) HAL
neorv32_onewire.c	neorv32_onewire.h	One-Wire Serial Interface Controller (ONEWIRE) HAL
neorv32_pwm.c	neorv32_pwm.h	Pulse-Width Modulation Controller (PWM) HAL
neorv32_rte.c	neorv32_rte.h	NEORV32 Runtime Environment
neorv32_sdi.c	neorv32_sdi.h	Serial Data Interface Controller (SDI) HAL
neorv32_slink.c	neorv32_slink.h	Stream Link Interface (SLINK) HAL
neorv32_smp.c	neorv32_smp.h	HAL for the SMP Dual-Core Configuration
neorv32_spi.c	neorv32_spi.h	Serial Peripheral Interface Controller (SPI) HAL

C source file	C header file	Description
	neorv32_sysinfo.h	System Configuration Information Memory (SYSINFO) HAL
neorv32_tracer.c	neorv32_tracer.h	Execution Trace Buffer (TRACER) HAL
neorv32_trng.c	neorv32_trng.h	True Random-Number Generator (TRNG) HAL
neorv32_twd.c	neorv32_twd.h	Two-Wire Serial Device Controller (TWD) HAL
neorv32_twi.c	neorv32_twi.h	Two-Wire Serial Interface Controller (TWI) HAL
neorv32_uart.c	neorv32_uart.h	Primary Universal Asynchronous Receiver and Transmitter (UART0) and UART1 HAL
neorv32_wdt.c	neorv32_wdt.h	Watchdog Timer (WDT) HAL
neorv32_newlib.c	-	Platform-specific system calls for <i>newlib</i>

*Doxygen API Documentation*

The Doxygen-based documentation of all core libraries is available online at <https://stnolting.github.io/neorv32/sw/files.html>.

4.3. System View Description File (SVD)

A CMSIS-SVD-compatible **System View Description (SVD)** file including all peripherals is available in [sw/svd](#).

4.4. Application Makefile

Application compilation is based on a centralized GNU makefile ([sw/common/common.mk](#)). Each software project (for example the ones in [sw/example](#) folder) should provide a local makefile that just includes the central makefile:

```
# Set path to NEORV32 root directory
NEORV32_HOME ?= ../../..
# Include the main NEORV32 makefile
include $(NEORV32_HOME)/sw/common/common.mk
```

Thus, the functionality of the central makefile (including all targets) becomes available for the project. A project-local makefile should be used to define all setup-relevant configuration options instead of changing the central makefile to keep the code base clean. Setting variables in the project-local makefile will override the default configuration. Most example projects already provide a makefile that list all relevant configuration options.

The following example shows the configuration of a local Makefile:

```
# Override the default CPU ISA
MARCH = rv32imc_zicsr_zifencei

# Override the default RISC-V GCC prefix
RISCV_PREFIX ?= riscv-none-elf-

# Override default optimization goal
EFFORT = -Os

# Add extended debug symbols for Eclipse
USER_FLAGS += -ggdb -gdwarf-3

# Additional sources
APP_SRC += $(wildcard ./*.c)
APP_INC += -I .

# Adjust processor IMEM size and base address
USER_FLAGS += -Wl,--defsym,__neorv32_rom_size=16k
USER_FLAGS += -Wl,--defsym,__neorv32_rom_base=0x00000000

# Adjust processor DMEM size and base address
USER_FLAGS += -Wl,--defsym,__neorv32_ram_size=8k
USER_FLAGS += -Wl,--defsym,__neorv32_ram_base=0x80000000

# Adjust maximum heap size
USER_FLAGS += -Wl,--defsym,__neorv32_heap_size=2k

# Set path to NEORV32 root directory
```

```
NEORV32_HOME ?= ../../..  
  
# Include the main NEORV32 makefile  
include $(NEORV32_HOME)/sw/common/common.mk
```

Setup of a New Project



When creating a new project, copy an existing project folder or at least the makefile to the new project folder. It is recommended to create new projects also in `sw/example` to keep the file dependencies. However, these dependencies can be manually configured via makefile variables if the new project is located somewhere else. For more complex projects, it may be useful to use explicit `source` and `include` folders. See `sw/example/coremark` for an example.

4.4.1. Makefile Targets

Invoking a project-local makefile (executing `make` or `make help`) will show the help menu that lists all available targets as well as all variable including their *current* setting.

```
neorv32/sw/example/hello_world$ make  
NEORV32 Software Makefile  
Find more information at https://github.com/stnolting/neorv32  
Use make V=1 or set BUILD_VERBOSE to increase build verbosity
```

Targets:

```
help      show this text  
check     check toolchain and list supported ISA extensions  
info      show project/makefile configuration  
gdb       start GNU debugging session  
asm       compile and generate <main.asm> assembly listing file for manual  
debugging  
elf       compile and generate <main.elf> ELF file  
exe       compile and generate <neorv32_exe.bin> executable image file for  
bootloader upload (includes a HEADER!)  
bin       compile and generate <neorv32_raw_exe.bin> executable memory image  
hex       compile and generate <neorv32_raw_exe.hex> executable memory image  
coe       compile and generate <neorv32_raw_exe.coe> executable memory image  
mem       compile and generate <neorv32_raw_exe.mem> executable memory image  
mif       compile and generate <neorv32_raw_exe.mif> executable memory image  
image     compile and generate VHDL IMEM application boot image  
<neorv32_application_image.vhd> in local folder  
install   compile, generate and install VHDL IMEM application boot image  
<neorv32_application_image.vhd>  
sim       in-console simulation using default/simple testbench and GHDL  
hdl_lists regenerate HDL file-lists (*.f) in NEORV32_HOME/rtl  
all      exe + install + hex + bin + asm  
elf_info  show ELF layout info
```

```

elf_sections show ELF sections
clean      clean up project home folder
clean_all   clean up project home folder and image generator
bl_image    compile and generate VHDL BOOTROM bootloader boot image
<neorv32_bootloader_image.vhd> in local folder
bootloader  compile, generate and install VHDL BOOTROM bootloader boot image
<neorv32_bootloader_image.vhd>

```

Variables:

```

BUILD_VERBOSE Set to increase build verbosity: 0
USER_FLAGS    Custom toolchain flags [append only]: "-ggdb -gdwarf-3 -Wl,
--defsym,__neorv32_rom_size=16k -Wl,--defsym,__neorv32_ram_size=8k"
USER_LIBS     Custom libraries [append only]: ""
EFFORT        Optimization level: "-Os"
MARCH         Machine architecture: "rv32i_zicsr_zifencei"
MABI          Machine binary interface: "ilp32"
APP_INC       C include folder(s) [append only]: "-I ."
APP_SRC       C source folder(s) [append only]: "./main.c"
APP_OBJ       Object file(s) [append only]
ASM_INC       ASM include folder(s) [append only]: "-I ."
RISCV_PREFIX  Toolchain prefix: "riscv-none-elf-"
NEORV32_HOME  NEORV32 home folder: "../../../"
GDB_ARGS      GDB (connection) arguments: "-ex target extended-remote
localhost:3333"
GHDL_RUN_FLAGS GHDL simulation run arguments: ""

```

Build Artifacts



All *intermediate* build artifacts (like object files and binaries) will be places into a (new) project-local folder named `build`. The *resulting* build artifacts (like executable, the main ELF and all memory initialization/image files) will be placed in the root project folder.



Increase Verbosity

Use `make V=1` or set `BUILD_VERBOSE` in your environment to increase build verbosity.

4.4.2. Default Compiler Flags

The central makefile uses specific compiler flags to tune the code to the NEORV32 hardware. Hence, these flags should not be altered. However, experienced users can modify them to further tune compilation.

Table 96. Compiler Options (CC_OPTS)

<code>-Wall</code>	Enable all compiler warnings.
--------------------	-------------------------------

<code>-ffunction-sections</code>	Put functions in independent sections. This allows a code optimization as dead code can be easily removed.
<code>-fdata-sections</code>	Put data segment in independent sections. This allows a code optimization as unused data can be easily removed.
<code>-nostartfiles</code>	Do not use the default start code. Instead, the NEORV32-specific start-up code (sw/common/crt0.S) is used (pulled-in by the linker script).
<code>-mno-fdiv</code>	Use built-in software functions for floating-point divisions and square roots (since the according instructions are not supported yet).
<code>-mstrict-align</code>	Unaligned memory accesses cannot be resolved by the hardware and require emulation.
<code>-mbranch-cost=10</code>	Branching costs a lot of cycles.
<code>-Wl,--gc-sections</code>	Make the linker perform dead code elimination.
<code>-ffp-contract=off</code>	Disable floating-point expression contraction (fused multiply-add/sub; not supported by the NEORV32 FPU).
<code>-g</code>	Add (simple) debug information.

Table 97. Linker Libraries (LD_LIBS)

<code>-lm</code>	Include/link with math.h .
<code>-lc</code>	Search for the standard C library when linking.
<code>-lgcc</code>	Make sure we have no unresolved references to internal GCC library subroutines.

Advanced Debug Symbols

By default, only "simple" symbols are added to the ELF (`-g`). Extended debug flags (e.g. for Eclipse) can be added using the `USER_FLAGS` variable (e.g. `USER_FLAGS += -ggdb -gdwarf-3`). Note that other debug flags may be required depending of the GCC/GDB version

4.5. Linker Script

The NEORV32-specific linker script (`sw/common/neorv32.ld`) is used to link the compiled sources according to the processor's [Address Space](#)). For the final executable, only two memory segments are required:

Table 98. Linker script - Memory Segments

Memory section	Description
<code>rom</code>	Instruction memory address space (processor-internal Instruction Memory (IMEM) and/or external memory)
<code>ram</code>	Data memory address space (processor-internal Data Memory (DMEM) and/or external memory)

These two sections are configured by several variables defined in the linker script and exposed to the build framework (aka the makefile). Those variable allow to customized the RAM/ROM sizes and base addresses. Additionally, a certain amount of the RAM can be reserved for the software-managed heap (see [RAM Layout](#)).

Table 99. Linker script - Configuration

Memory section	Description	Default
<code>__neorv32_rom_size</code>	"ROM" size (instruction memory / IMEM)	16kB
<code>__neorv32_ram_size</code>	"RAM" size (data memory / DMEM)	8kB
<code>__neorv32_rom_base</code>	"ROM" base address (instruction memory / IMEM)	0x00000000
<code>__neorv32_ram_base</code>	"RAM" base address (data memory / DMEM)	0x80000000
<code>__neorv32_heap_size</code>	Maximum heap size; part of the "RAM"	0kB

Each variable provides a default value (e.g. "16K" for the instruction memory /ROM /IMEM size). These defaults can be overridden by setup-specific values to take the user-defined processor configuration into account (e.g. a different IMEM size). The `USER_FLAGS` variable provided by the [Application Makefile](#) can also be used to customize the memory configuration. For example, the following line can be added to a project-specific local makefile to adjust the memory sizes:

Listing 23. Overriding Default Memory Sizes (configuring 64kB IMEM and 32kB DMEM)

```
USER_FLAGS += "-Wl,--defsym,__neorv32_rom_size=64k -Wl,
--defsym,__neorv32_ram_size=32k"
```

Memory Configuration Constraints

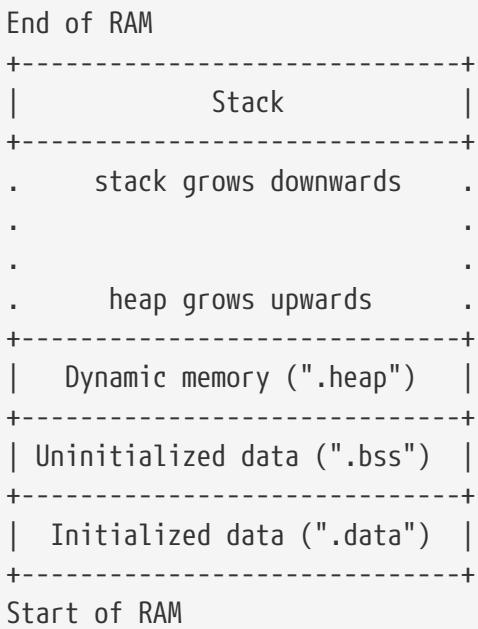


Memory sizes have to be a multiple of 4 bytes. Memory base addresses have to be 32-bit-aligned.

4.5.1. RAM Layout

The default NEORV32 linker script uses the defined RAM (**DMEM**) size to map several sections. Note that depending on the application some sections might have zero size.

Listing 24. Default RAM Layout



- Initialized data (.data):** The data section is placed right at the beginning of the RAM. For example, this contains *explicitly initialized* global variables which are initialized by the [Start-Up Code \(crt0\)](#).
- Non-initialized data (.bss):** This dynamic data section contains data like global variables *without explicit initialization*. This memory section is cleared by the [Start-Up Code \(crt0\)](#).
- Heap (.heap):** The heap is used for dynamic memory that is managed by functions like `malloc()` and `free()`. The heap grows upwards. This section is not initialized at all.
- Stack:** The stack starts at the end of the RAM at the last 16-byte aligned address. According to the RISC-V ABI / calling convention the stack is 128-bit-aligned before procedure entry. The stack grows downwards.

Heap Size



The maximum size of the heap is defined by the `__neorv32_heap_size` variable. This variable has to be **explicitly defined** in order to define a heap size (and to use dynamic memory allocation at all).

Heap-Stack Collision



Take care when using dynamic memory to avoid collision of the heap and stack memory areas. There is no compile-time protection mechanism available as the actual heap and stack size are defined by *runtime* data. The [physical memory protection](#) extension can be used to implement a guarding mechanism.

4.5.2. ROM Layout

The default NEORV32 linker script uses the defined ROM (IMEM) size to map several sections. Note that depending on the application some sections might have zero size.

Listing 25. Default ROM Layout

```
End of ROM
+-----+
| Initial data for ".data" |
+-----+
| Read-only data (" .rodata") |
+-----+
| Program code (" .text") |
+-----+
Start of ROM
```

1. **Program code (`.text`)**: This section contains the actual instructions of the program.
2. **Read-only data (`.rodata`)**: Read-only data, for example for constants and strings.
3. **Heap (`.heap`)**: Initialization data that is copied at runtime to the RAM's `data` section by the **Start-Up Code (crt0)**.

Binary Executable Image

The main **Application Makefile** extracts the `.text`, `.rodata`, and `.data` sections from the final ELF and concatenates them in order to build the final binary executable image (`main.bin`). Once placed in memory, this image can be executed right away. Furthermore, this image is used by the image generator auxiliary tool to generate different **Executable Image Formats**.



4.6. C Standard Library

The default software framework relies on **newlib** as default C standard library. Newlib provides hooks for common "system calls" (like file handling and standard input/output) that are used by other C libraries like **stdio**. These hooks are available in [sw/lib/source/newlib.c](#) and were adapted for the NEORV32 processor.

Standard Input/Output Streams / Consoles



The **UART0** is used to implement all the standard input, output and error consoles (**STDIN** = file number 0, **STDOUT** = file number 1, **STDERR** = file number 2). All other input/output streams (other file number than 0,1,2) are redirected to **UART1**.

Constructors and Destructors



Constructors and destructors for plain C code or for C++ applications are supported by the software framework. See [sw/example/hello_cpp](#) for a minimal example. Note that constructor and destructors are only executed by core 0 (primary core) in the SMP [Dual-Core Configuration](#).



Newlib Test/Demo Program

A simple test and demo program that uses some of newlib's system functions (like **malloc/free** and **read/write**) is available in [sw/example/demo_newlib](#).



Executing System Functions on a Host Computer

The NEORV32 on-chip debugger supports **Semihosting** which allow to use the input/output facilities on a host computer.

4.7. Start-Up Code (`crt0`)

The CPU and also the processor require a minimal start-up and initialization code to bring the hardware into an operational state. Furthermore, the C runtime requires an initialization before compiled code can be executed. This setup is done by the start-up code ([sw/common/crt0.S](#)) which is automatically linked with *every* application program and gets mapped before the actual application code so it gets executed right after boot.

The `crt0.S` start-up performs the following operations:

1. Setup the stack pointer and the global pointer according to the [RAM Layout](#) provided by the [Linker Script](#) symbols.
2. Initialize `mstatus` CSR disabling machine-level interrupts.
3. Install an endless loop as trap handler to `mtvec` CSR: the core will halt if any trap occurs.
4. Clear `mie` CSR disabling all interrupt sources.
5. Initialize all integer register `x1 - x31` (only `x1 - x15` if the `E` CPU extension is enabled).
6. If the executing CPU core is not core 0, an SMP-specific code is executed and the CPU is halted in sleep mode. See section [Dual-Core Boot](#) for more information.
7. Setup `.data` section to configure initialized variables.
8. Clear the `.bss` section.
9. Call constructors (if there are any).
10. Call the application's `main()` function (with no arguments; `argc = argv = 0`).
11. If `main()` returns:
 - All interrupt sources are disabled by clearing `mie` CSR.
 - The return value of `main()` is copied to the `mscratch` CSR to allow inspection by the debugger.
 - The core will halt if any trap occurs.
 - Call all *destructors* (if there are any).
 - The CPU halts: enter sleep mode executing the `wfi` instruction in an endless loop.

4.8. Executable Image Formats

The binary executable image (see [ROM Layout](#)) is further processed by the NEORV32 image generator ([sw/image_gen](#)) to generate a final executable file. The image generator can generate several types of executable file formats selected by a flag when calling the generator. **Note that all these options are managed by the makefile (see [Makefile Targets](#)).**

<code>app_bin</code>	Generates an executable binary file (including a bootloader header) for upload via the bootloader.
<code>app_vhd</code>	Generates an executable VHDL memory initialization image for the processor-internal IMEM.
<code>bld_vhd</code>	Generates an executable VHDL memory initialization image for the processor-internal BOOT ROM.
<code>raw_hex</code>	Generates a raw 8x ASCII hex-char file for custom purpose.
<code>raw_bin</code>	Generates a raw binary file for custom purpose.
<code>raw_coe</code>	Generates a raw COE file for FPGA memory initialization.
<code>raw_mem</code>	Generates a raw MEM file for FPGA memory initialization.
<code>raw_mif</code>	Generates a raw MIF file for FPGA memory initialization.

Image Generator Compilation



The sources of the image generator are automatically compiled when invoking the makefile (requiring a *native* GCC installation).

Executable Header



For the `app_bin` option the image generator adds a small header to the executable. This header is required by the [Bootloader](#) to identify and manage the executable. The header consists of three 32-bit words located right at the beginning of the file. The first word of the executable is the signature word and is always `0xB007C0DE`. Based on this word the bootloader can identify a valid image file. The next word represents the size in bytes of the actual program image in bytes. A simple complement checksum of the actual program image is given by the third word. This provides a simple protection against data transmission or storage errors. **Note that this executable format cannot be used for *direct* execution.**

VHDL Memory Initialization Files



The generated VHDL images are plain VHDL packages providing program data as constant array. Note that this array is always extended to a power of 2. Unused elements are set to zero.

Optionally, the NEORV32 image generator can also be used stand-alone:

```
neorv32/sw/image_gen$ ./image_gen  
NEORV32 executable image generator
```

Usage: image_gen [options]
Example: image_gen -i main.bin -o main_exe.bin -t app_bin

Options:

- h Show this help text and exit
- i file_name Input binary file name; mandatory
- o file_name Output file name; mandatory
- t type Type of image to generate; default is 'app_bin'

Image type:

- app_bin Application executable for bootloader upload (binary file with header)
- app_vhd Application memory image (IMEM VHDL package file)
- bld_vhd Bootloader memory image (BOOTROM VHDL package file)
- raw_hex ASCII hex file (raw executable)
- raw_bin Binary file (raw executable)
- raw_coe COE file (raw executable)
- raw_mem MEM file (raw executable)
- raw_mif MIF file (raw executable)

4.9. Bootloader

The NEORV32 bootloader provides an optional built-in firmware that allows to upload new application firmware at any time without the need to re-synthesize the FPGA's bitstream. The bootloader is automatically executed after reset when enabled via the [Boot Configuration BOOT_MODE_SELECT](#) generic. The bootloader provides the following key features:

- interactive user console via UART
- upload executable via UART
- load executable from SPI flash; store executable to SPI flash
- load executable from TWI flash; store executable to TWI flash (*option disabled by default*)
- load executable from SD card (*option disabled by default*)
- automatic boot sequence: automatically boot from TWI/SPI flash / SD card
- highly configurable to adapt to application requirements

Pre-Built Bootloader Image



This section refers to the **default** NEORV32 bootloader. The [Bootloader ROM \(BOOTROM\)](#) image file already contains the pre-compiled bootloader ([sw/bootloader/bootloader.c](#)) in its default configuration.

Minimal Hardware Requirements



The default bootloader image was compiled for a minimal [rv32e_zicsr_zifencei](#) ISA configuration and requires a RAM (DMEM) size of at least 256 bytes. These constraints ensure that the bootloader can be executed on any CPU/processor configuration. It is recommended to enable at least [UART0](#), the [CLINT](#) and the [GPIO](#) controller. See [Customizing the Internal Bootloader](#) for more information.

4.9.1. Bootloader Console

The default bootloader provides a serial console via [UART0](#) for user interaction using the following terminal settings:

- 19200 Baud, 8 data bits, no parity bit, 1 stop bit ([19200-8-N-1](#))
- line breaks: carriage return + newline ([\r\n](#))

Terminal Program



Any terminal program that can connect to a serial port should work. However, make sure the program can transfer data in *raw* byte mode without any protocol overhead (e.g. XMODEM). Note that some terminal programs struggle with transmitting files larger than 4kB (see <https://github.com/stnolting/neorv32/pull/215>). **For Windows I can recommend TeraTerm (GUI) and SimplySerial (Window terminal). For Linux I can recommend picocom.**

By default the bootloader uses the LSB of the top entity's GPIO output port (`gpio_o(0)`) for an high-active status LED. All other output pins are set to low. After reset, the status LED will start blinking at 2Hz and the [Auto Boot Sequence](#) is started. This auto-boot sequence can be skipped within 10s by pressing any key.

Listing 26. Default Bootloader Console

```
NEORV32 Bootloader
build: Sep 3 2025 ①

Auto-boot in 10s. Press any key to abort. ②
Aborted. ③

Type 'h' for help.
CMD:> ④
```

- ① Bootloader version (compile date).
- ② Start of auto-boot sequence.
- ③ Auto-boot sequence aborted due to user console input.
- ④ Command prompt.

To see a list of all available commands press `h` to see the help menu:

Listing 27. Bootloader Help Menu

```
CMD:> h
Available CMDs:
h: Help ①
i: System info ②
r: Restart ③
u: Upload via UART ④
t: TWI flash - load ⑤
w: TWI flash - program ⑥
l: SPI flash - load ⑦
s: SPI flash - program ⑧
c: SD card - load ⑨
e: Start executable ⑩
x: Exit ⑪
CMD:>
```

- ① Show "Available CMDs" help text again.
- ② Show hardware configuration information.
- ③ Restart bootloader and auto-boot sequence.
- ④ Upload new executable (`neorv32_exe.bin`) via UART.
- ⑤ Load executable from TWI flash (*this option is disabled by default*).

- ⑥ Store executable to TWI flash (*this option is disabled by default*).
- ⑦ Load executable from SPI flash.
- ⑧ Program previously-uploaded executable to SPI flash.
- ⑨ Load executable from SD card (*this option is disabled by default*).
- ⑩ Start the (up)loaded executable.
- ⑪ Raise a breakpoint exception: If a debugger is connected this will transfer control to the debugger. If no debugger is connected this will print an exception error ([Bootloader Error Codes](#)) shutting down the processor.

Unavailable Commands



Note that not all options are enabled in the default bootloader configuration in order to keep the bootloader ROM size below 4kB. However, these options can be enabled in `config.h` (see [Customizing the Internal Bootloader](#)).

Available commands can be executed by typing the according letter. For example, when executing the `i` command the general system configuration is printed:

Listing 28. System Configuration Information

```
CMD:> i
HWV: 0x01120101 ①
CLK: 0x05f5e100 ②
MISA: 0x40901107 ③
XISA: 0x0fc06fd3 ④
SOC: 0x38efc87b ⑤
MISC: 0x0a010d00 ⑥
CMD:>
```

- ① Processor hardware version in BCD format (`mimpid` CSR).
- ② Processor clock speed in Hz (`CLK` register of [System Configuration Information Memory \(SYSINFO\)](#)).
- ③ RISC-V CPU extensions (`misa` CSR).
- ④ NEORV32-specific CPU extensions (`mxisa` CSR).
- ⑤ Processor configuration (`SOC` register of [System Configuration Information Memory \(SYSINFO\)](#)).
- ⑥ Miscellaneous memory and SoC configuration (`MISC` register of [System Configuration Information Memory \(SYSINFO\)](#)).

4.9.2. Auto Boot Sequence

After reset, the bootloader waits 10 seconds for a UART console input before it starts the automatic boot sequence. Depending on the configuration ([Customizing the Internal Bootloader](#)) the bootloader will try to fetch a valid executable from different sources:

1. Try to load an executable from TWI flash (default device ID is `0xA0`).
2. Try to load an executable from SPI flash (default SPI chip select line is `spi_csn_o(0)`).
3. Try to load file `boot.bin` from SD flash (default SPI chip select line is `spi_csn_o(1)`).

If a valid boot image is loaded it will be immediately started. If no valid executable can be fetched the interactive bootloader console is started.

4.9.3. Uploading an Executable

1. Connect the primary UART (UART0) interface of the processor to a serial port of your host computer.
2. Start a serial terminal program.
3. Open a connection to the the serial port your UART is connected to.
4. Press the NEORV32 reset button to restart the bootloader. The status LED starts blinking and the bootloader intro screen appears in the console. Press any key to abort the automatic boot sequence and to start the actual bootloader user interface console.
5. Execute the "Upload" command by typing `u`. Now the bootloader is waiting for a binary executable to be send: `Awaiting neorv32_exe.bin...`
6. Use the "send file" option of your terminal program to send a valid NEORV32 executable (`neorv32_exe.bin`). Make sure the terminal sends the executable in raw binary mode.
7. If everything went fine, `Awaiting neorv32_exe.bin... OK` is printed in the terminal.
8. The executable is now in the instruction memory of the processor. To execute the program right now run the "start executable" command by typing `e`.

4.9.4. Programming an SPI (/TWI) Flash

This guide shows how to write an executable to the SPI flash via the bootloader so it can be automatically fetched and executed after processor reset. If the TWI flash option is enabled, an according command for programming the TWI flash is available.

1. Reset the NEORV32 processor and wait until the bootloader start screen appears.
2. Abort the auto boot sequence and start the user console by pressing any key.
3. Press `u` to upload the executable that shall be programmed to the flash.
4. Send the binary via the terminal program. When the upload is completed and "OK" appears, press `s` to trigger the SPI flash programming:

```
CMD:> u
Awaiting neorv32_exe.bin... OK
CMD:> s
Write 0x00001614 bytes to SPI flash @0x00400000 (y/n)?
```

5. The bootloader shows the size of the executable and the base address of the SPI flash where the executable will be stored. A prompt appears: type **y** to start the programming or type **n** to abort.

```
CMD:> u
Awaiting neorv32_exe.bin... OK
CMD:> s
Write 0x00001614 bytes to SPI flash @0x00400000 (y/n)?
Flashing... OK
CMD:>
```

6. Note that flash programming can take some time (depending on the TWI/SPI clock configuration in `config.h`). If "OK" appears, the programming process was successful and the flash can be used for the auto-boot sequence.

4.9.5. Booting from SD Card

The SD card is accessed in SPI mode. The card has to be formatted using the **FAT32** file system with a sector size of at least 512 bytes. The executable has to be placed in the card's root directory. By default the bootloader will fetch the `boot.bin` file, which is just a renamed copy of the default `neorv32_exe.bin` file that is generated by the application compilation flow. The name of the SD card's boot file can be changed to a custom file name (see [Customizing the Internal Bootloader](#)), but it has to use the 8.3 DOS format (max 8 character for the name plus dot plus 3 characters suffix).

SD-card and FAT32 support is provided by the great **Petit FatFs** library by Elm-Chan: https://elm-chan.org/fsw/ff/00index_p.html

4.9.6. Customizing the Internal Bootloader

The NEORV32 bootloader provides several options to configure it for a custom setup. It is configured via set of C-language `defines` in `sw/bootloader/config.h`. All defines provide default value that can be edited or overridden by Makefile directives.

Table 100. Bootloader Configuration Parameters

Parameter	Default	Legal values	Description
Memory layout			
<code>EXE_BASE_ADDR</code>	<code>0x00000000</code>	any 4-byte-aligned address	Memory base address for the executable; also the boot address for the application.
Serial console - requires Primary Universal Asynchronous Receiver and Transmitter (UART0)			
<code>UART_EN</code>	<code>1</code>	<code>0, 1</code>	Set to <code>1</code> to enable the serial console.
<code>UART_BAUD</code>	<code>19200</code>	any	Baud rate of UART0.
Status LED - requires General Purpose Input and Output Port (GPIO)			
<code>STATUS_LED_EN</code>	<code>1</code>	<code>0, 1</code>	Enable bootloader status led ("heart beat") at <code>GPIO</code> output port pin <code>STATUS_LED_PIN</code> when <code>1</code> .

Parameter	Default	Legal values	Description
STATUS_LED_PIN	0	0..31	GPIO output pin used for the high-active status LED.
Auto-boot - requires Core-Local Interruptor (CLINT)			
AUTO_BOOT_EN	1	0,1	Auto-boot enabled when 1.
AUTO_BOOT_TIMEOUT	10	any	Timeout in seconds after which the auto-boot sequence starts (if there is no UART input by the user).
TWI flash - requires Two-Wire Serial Interface Controller (TWI)			
TWI_FLASH_EN	1	0,1	Set to 1 to enable the SPI flash (boot) options.
TWI_FLASH_PROG_EN	1	0,1	Set to 1 to enable programming the TWI flash from the bootloader menu.
TWI_FLASH_CLK_PRSC	CLK_PRSC_C_1024	CLK_PRSC_2 CLK_PRSC_4 CLK_PRSC_8 CLK_PRSC_64 CLK_PRSC_128 CLK_PRSC_1024 CLK_PRSC_2048 CLK_PRSC_4096	TWI clock prescaler.
TWI_FLASH_CLK_DIV	0	0..15	TWI clock divider value.
TWI_FLASH_ID	0xA0	any	8-bit TWI device address (with R/W bit cleared).
TWI_FLASH_BASE_ADDR	0x00000000	32-bit	Defines the TWI flash base address for the executable.
TWI_FLASH_ADDR_BYTES	2	1,2,3,4	Number of TWI flash address bytes.
SPI flash - requires Serial Peripheral Interface Controller (SPI)			
SPI_FLASH_EN	1	0,1	Set to 1 to enable the TWI flash (boot) options.
SPI_FLASH_PROG_EN	1	0,1	Set to 1 to enable programming the SPI flash from the bootloader menu.
SPI_FLASH_CS	0	0..7	SPI chip select line (port <code>spi_csn_o</code>) for selecting the SPI flash.
SPI_FLASH_CLK_PRSC	CLK_PRSC_C_64	CLK_PRSC_2 CLK_PRSC_4 CLK_PRSC_8 CLK_PRSC_64 CLK_PRSC_128 CLK_PRSC_1024 CLK_PRSC_2048 CLK_PRSC_4096	SPI clock prescaler.

Parameter	Default	Legal values	Description
SPI_FLASH_CLK_DIVIDER	0	0..15	SPI clock divider value.
SPI_FLASH_BASE_ADDRESS_DDR	0x00400000	32-bit	Defines the SPI flash base address for the executable.
SPI_FLASH_ADDRESS_SIZE	3	1,2,3,4	SPI flash address size in number of bytes.
SPI_FLASH_SECTOR_SIZE	16*1024	any	Number of SPI flash address bytes.
SPI SD card - requires Serial Peripheral Interface Controller (SPI)			
SPI_SDCARD_ENABLE	0	0,1	Set to 1 to enable booting from SD card.
SPI_SDCARD_CHIP_SELECT	1	0..7	SPI chip select line (port <code>spi_csn_o</code>) for selecting the SD card.
SPI_SDCARD_CLOCK_PRESCALER	CLK_PRSC_1_C_64	CLK_PRSC_2 CLK_PRSC_4 CLK_PRSC_8 CLK_PRSC_64 CLK_PRSC_128 CLK_PRSC_1024 CLK_PRSC_2048 CLK_PRSC_4096	SPI clock prescaler.
SPI_SDCARD_CLOCK_DIVIDER	0	0..15	SPI clock divider value.
SPI_SDCARD_BOOT_IMAGE	"boot.b_in"	8.3 DOS format	File name of the boot image. Has to be located in the root directory.
Branding - for text printed via serial console			
THEME_INTRO	"NEORV32 Bootloader"	any string	Intro text that is shown in the bootloader console.
THEME_EXECUTABLE	"nerv32_exe.b_in"	any string	Name of executable that is shown in the console menu.

4.9.7. Bootloader Error Codes

ERROR_DEVICE	A device-accessing function returned an error code. Make sure that the device is properly connected and that all required processor modules/interface are actually enabled (by the according Processor Top Entity - Generics).
ERROR_SIGNATURE	The signature that indicates a valid NEORV32 executable of the accessed executable is incorrect. This can be caused by a temporary transmission error or by an invalid or corrupted executable.

ERROR_CHECKSUM	The checksum of the loaded executable is incorrect. This can be caused by a temporary transmission error or by an invalid or corrupted executable.
ERROR_EXCEPTION	An unexpected exception has occurred. This can be caused by an invalid bootloader configuration (non-available processor modules, memory layout, ...). For debugging purpose the error message will also display the content of the mcause , mepc , mtinst and mtval CSRs. Example: ERROR_EXCEPTION 0x00000003 0xffe00cb4 0x00100073 0x00000000

4.10. NEORV32 Runtime Environment

The NEORV32 software framework provides a minimal runtime environment ("RTE") that takes care of a stable and *safe* execution environment by providing a unified interface for handling of *all* traps (exceptions and interrupts). Once initialized, the RTE provides **Default RTE Trap Handlers** that catch all possible traps. These default handlers just output a message via UART when a certain trap has been triggered. The default handlers can be overridden by the application code to install application-specific handler functions for each trap.

Using the RTE is **optional but highly recommended** for bare-metal / non-OS applications. The RTE provides a simple and comfortable way of delegating traps to application-specific handlers while making sure that all traps (even though they are not explicitly used by the application) are handled correctly. Performance-optimized applications or embedded operating systems may not use the RTE at all in order to increase response time.

4.10.1. RTE Operation

The RTE manages the trap-related CSRs of the CPU's privileged architecture (see [Machine Trap Handling CSRs](#)). It initializes the `mtvec` CSR in DIRECT mode, which provides the base entry point for *all* traps. The address stored to this register defines the address of the **first-level trap handler**, which is provided by the NEORV32 RTE. Whenever an exception or interrupt is triggered this first-level trap handler is executed.

The first-level handler performs a complete context save, analyzes the source of the trap and calls the according **second-level trap handler**, which takes care of the actual exception/interrupt handling. The RTE manages an internal look-up table to track the addresses of the according second-level trap handlers.

After the initial RTE setup, each entry in the RTE's trap handler look-up table is initialized with a **Default RTE Trap Handlers**. These default handler do not execute any trap-related operations - they just output a debugging message via the primary UART (UART0) (if enabled) to inform the user that a trap has occurred that is not (yet) handled by a proper application-specific trap handler. After sending this message, the RTE tries to resume normal execution by moving on to the next linear instruction.

Dual-Core Configuration



The RTE's internal trap handler look-up table is used globally for **both** cores. If a core-specific handling is required, the according user-defined trap handler need to retrieve the core's ID from `mhartid` and branch accordingly.

4.10.2. Using the RTE

The NEORV32 runtime environment is part of the default NEORV32 software framework. The links to the according software references are listed below.

neorv32_ [Online software reference \(Doxygen\)](#)

rte.c

neorv32_ [Online software reference \(Doxygen\)](#)

rte.h

The RTE has to be explicitly enabled by calling the according setup function. It is recommended to do this right at the beginning of the application's `main` function. For the SMP [Dual-Core Configuration](#) the RTE setup functions has to be called on each core that wants to use the RTE.

Listing 29. RTE Setup Right at the Beginning of "main"

```
int main() {
    neorv32_rte_setup(); // setup NEORV32 runtime environment
    ...
}
```

After setup, all traps will trigger execution of the RTE's [Default RTE Trap Handlers](#) at first. In order to use application-specific trap handlers the default debug handlers can be overridden by installing user-defined ones:

Listing 30. Installing an Application-Specific Trap Handler (Function Prototype)

```
int neorv32_rte_handler_install(uint8_t id, void (*handler)(void));
```

The first argument `id` defines the "trap ID" (for example a certain interrupt request) that shall be handled by the user-defined handler. These IDs are defined in `sw/lib/include/neorv32_rte.h`. However, more convenient device-specific aliases are also defined in `sw/lib/include/neorv32.h`. The second argument `handler` is the actual function that implements the user-defined trap handler. The custom handler functions must have a specific type without any arguments and with no return value:

Listing 31. Custom Trap Handler (Function Prototype)

```
void custom_trap_handler_xyz(void) {
    // handle trap...
}
```

Custom Trap Handler Attributes



Do NOT use the `interrupt` attribute for the application trap handler functions! This would place an `mret` instruction at the end of the handler making it impossible to return to the first-level trap handler of the RTE core.



`mscratch` CSR

The **mscratch** CSR should not be used inside an application trap handler as this register is used by the RTE to provide the base address of the application's stack frame **Application Context Handling** (i.e. modifying the registers of application code that caused a trap).

The following example shows how to install trap handlers for exemplary traps.

Listing 32. Installing Custom Trap Handlers Examples

```
neorv32_rte_handler_install(RTE_TRAP_MTI, machine_timer_irq_handler); // handler for
machine timer interrupt
neorv32_rte_handler_install(RTE_TRAP_MENV_CALL, environment_call_handler); // handler
for machine environment call exception
neorv32_rte_handler_install(SLINK_RX_RTE_ID, slink_rx_handler); // handler for SLINK
receive interrupt
```

4.10.3. Default RTE Trap Handlers

The default RTE trap handlers are executed when a certain trap is triggered that is not (yet) handled by an application-defined trap handler. The default handler will output a message giving additional debug information via the **Primary Universal Asynchronous Receiver and Transmitter (UART0)** to inform the user and it will also try to resume normal program execution (exemplary RTE outputs are shown below). The specific message right at the beginning of the debug trap handler message corresponds to the trap code obtained from the **mcause** CSR (see **NEORV32 Trap Listing**).

In most cases the RTE can successfully continue operation - for example if it catches an **interrupt** request that is not handled by the actual application program. However, if the RTE catches an unhandled **trap** like a bus access fault exception, continuing execution will most likely fail making the CPU crash.

Listing 33. RTE Default Trap Handler UART0 Output Examples

```
<NEORV32-RTE-PANIC> [cpu0|M] Illegal instruction MEPC=0x000002d6 MTINST=0x000000FF
MTVAL=0x00000000 </NEORV32-RTE-PANIC> ①
<NEORV32-RTE-PANIC> [cpu0|U] Illegal instruction MEPC=0x00000302 MTINST=0x00000000
MTVAL=0x00000000 </NEORV32-RTE-PANIC> ②
<NEORV32-RTE-PANIC> [cpu0|U] Load address misaligned MEPC=0x00000440 MTINST=0x01052603
MTVAL=0x80000101 </NEORV32-RTE-PANIC> ③
<NEORV32-RTE-PANIC> [cpu1|M] FIRQ channel 0x00000003 MEPC=0x00000820 MTINST=0x00000000
MTVAL=0x00000000 </NEORV32-RTE-PANIC> ④
<NEORV32-RTE-PANIC> [cpu1|M] Instruction access fault MEPC=0x90000000
MTINST=0x42078b63 MTVAL=0x00000000 FATAL! HALTING CPU </NEORV32-RTE-PANIC>\n ⑤
```

- ① Illegal 32-bit instruction **MTINST=0x000000FF** at address **MEPC=0x000002d6** while CPU 0 was in machine-mode (**[cpu0|M]**).
- ② Illegal 16-bit instruction **MTINST=0x00000000** at address **MEPC=0x00000302** while CPU 0 was in user-mode (**[cpu0|U]**).

- ③ Misaligned load access at address `MEPC=0x00000440` caused by instruction `MTINST=0x01052603` (trying to load a full 32-bit word from address `MTVAL=0x80000101`) while CPU 0 was in user-mode (`[cpu0|U]`).
- ④ Fast interrupt request from channel 3 before executing instruction at address `MEPC=0x00000820` while CPU 1 was in machine-mode (`[cpu0|M]`).
- ⑤ Instruction bus access fault at address `MEPC=0x90000000` while executing instruction `MTINST=0x42078b63` while CPU 1 was in machine-mode (`[cpu0|M]`). This is a fatal/non-resumable exception halting the CPU.

4.10.4. Application Context Handling

Upon trap entry the RTE backups the entire application context (i.e. all `x` general purpose registers) to the stack. The context is restored automatically after trap completion. The base address of the according stack frame is copied to the `mscratch` CSR. By having this information available, the RTE provides dedicated functions for accessing and altering the application context:

Listing 34. RTE Context Access Functions

```
// Prototypes
uint32_t neorv32_rte_context_get(int x); // read register
void    neorv32_rte_context_put(int x, uint32_t data); // write data to register

// Examples
uint32_t tmp = neorv32_rte_context_get(9); // read register 'x9'
neorv32_rte_context_put(28, tmp); // write 'tmp' to register 'x28'
```

The `x` argument is used to specify one of the RISC-V general purpose register `x0` to `x31`. Note that registers `x16` to `x31` are not available if the RISC-V [E ISA Extension](#) is enabled. For the SMP [Dual-Core Configuration](#) the provided context functions will access the stack frame of the interrupted application code that was running on the specific CPU core that caused the trap entry.

The context access functions can be used by application-specific trap handlers to *emulate* unsupported CPU / SoC features like unimplemented IO modules, unsupported instructions and even unaligned memory accesses.

Demo Program: Emulate Unaligned Memory Access



A demo program, which showcases how to emulate unaligned memory accesses using the NEORV32 runtime environment can be found in [sw/example/demo_emulate_unaligned](#).

Chapter 5. On-Chip Debugger (OCD)

The NEORV32 Processor features an *on-chip debugger* (OCD) compatible to the **Minimal RISC-V Debug Specification** implementing the **execution-based debugging** scheme. A copy of the specification is available in [docs/references](#). The on-chip debugger is implemented if the `OCD_EN` processor top generic is set to `true`. Optionally, up to 16 hardware triggers can be implemented ([Sdtrig ISA Extension](#)) to support hardware-assisted break- and watchpoints. Furthermore, the OCD supports an optional [Debug Authentication](#) module to constrain OCD access to authorized parties.

Hands-On Tutorial



A simple example on how to use NEORV32 on-chip debugger in combination with OpenOCD and the GNU debugger is shown in section [Debugging using the On-Chip Debugger](#) of the User Guide.

Section Structure

- [openOCD](#)
- [Semihosting](#)
- [Debug Transport Module \(DTM\)](#)
- [Debug Module \(DM\)](#)
- [Debug Authentication](#)
- [CPU Debug Mode](#)
- [Trigger Module](#)

Key Features

- standard 4-wire JTAG access port
- debugging of up to 4 CPU cores ("harts")
- full control of the CPU: halting, single-stepping and resuming
- indirect access to all core registers and the entire processor address space (via program buffer)
- execution of arbitrary programs via the program buffer
- compatible with upstream OpenOCD and GDB
- optional trigger module for up to 16 hardware break- and watchpoints
- optional authentication for advanced security

Configuration Options

Table 101. NEORV32 OCD Configuration Generics

Name	Type	Default	Description
OCD_EN	boolean	false	Implement the on-chip debugger and the CPU debug mode..
OCD_NUM_HW_TRIGGER_S	natural	0	Number of implemented HW triggers (Trigger Module / Sdtrig ISA Extension) for hardware break-/watchpoints (0..16).
OCD_AUTHENTICATION	boolean	false	Implement optional Debug Authentication module.
OCD_JEDEC_ID	suv(10:0)	"0000000000 00"	JEDEC ID; continuation codes plus vendor ID (passed to the JTAG Debug Transport Module (DTM)).

Overview

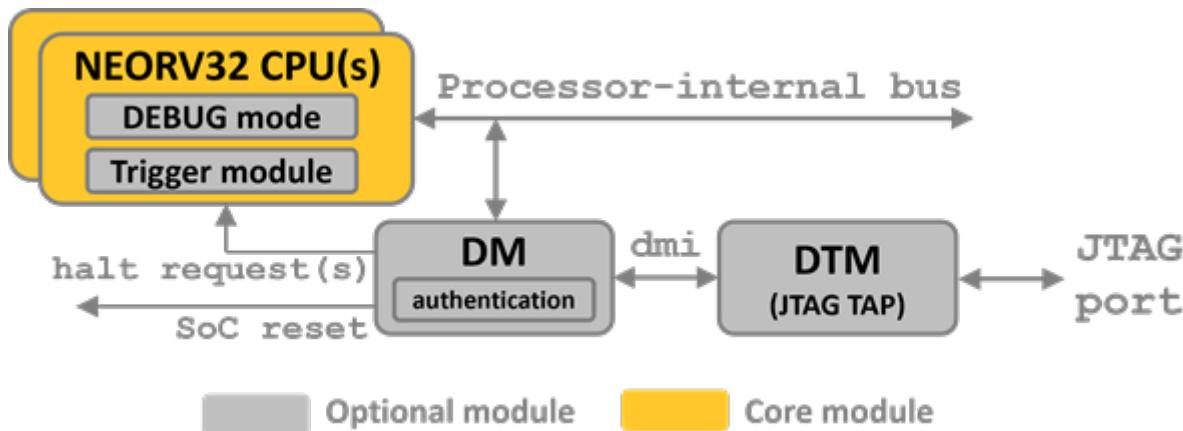


Figure 19. NEORV32 on-chip debugger complex

The NEORV32 on-chip debugger is based on five hardware modules:

1. **Debug Transport Module (DTM):** JTAG access tap to allow an external adapter to interface with the *debug module (DM)*.
2. **Debug Module (DM):** The RISC-V debug module is the main bridge between the external debugger and the processor being debugged. It provides a *data buffer* for data transfer from/to the DM, a *code ROM* containing the "park loop" code, a *program buffer* to allow the debugger to execute small programs defined by the DM and a *status register* that is used to communicate *exception*, *halt*, *resume* and *execute* requests/acknowledges between the debugger and the CPU.
3. **Debug Authentication:** Authenticator module to secure on-chip debugger access. By default this module implements a very simple authentication mechanism as example. Users can modify/replace this default logic to implement arbitrary authentication mechanism.
4. **CPU Debug Mode ISA extension:** This ISA extension provides the "debug execution mode" as another CPU operation mode that is used to execute the park loop code from the DM. This mode also provides additional CSRs and instructions.
5. **CPU Trigger Module:** This module provides up to 16 hardware triggers that can be used as hardware-assisted break- or watchpoints.

Theory of Operation

When debugging the system using the OCD, the external debugger (e.g. GDB) issues a halt request to the CPU to make it enter so-called *debug mode*. In this mode the application-defined architectural state of the system/CPU is "frozen" so the debugger can monitor it without interfering with the actual application. However, the OCD can also modify the entire architectural state at any time. While in debug mode, the debugger has full control over the entire CPU core.

After halting, the CPU executes the "park loop" code from the code ROM of the debug module (DM). This park loop implements an endless loop that is used to poll a memory-mapped [Status Register](#) of the DM. The flags in this register are used to communicate requests from the DM and to acknowledge their processing by the CPU: trigger execution of the program buffer or resume the halted application. Furthermore, the CPU uses this register to signal that the CPU has halted after a halt request or to signal that an exception has been raised while being in debug mode.

5.1. openOCD

By default, the free and open-source openOCD (<https://openocd.org/>) is used to establish a debugger connection. However, other interfaces can be used (like Segger or Lauterbach tools) if they provide RISC-V support.

openOCD is configured by a set of configuration scripts which are located in `neorv32/sw/openocd`. Two top scripts are provided: one for the single-core processor configuration (`neorv32.cfg`) and another one for the SMP dual-core configuration (`neorv32.dual-core.cfg`). Both scripts have the same format and include several helper scripts.

Listing 35. NEORV32 openOCD single-core CFG file (`neorv32.cfg`)

```
# configuration
set PATH [file dirname [file normalize [info script]]] ①
set CORENAME neorv32 ②
set NUMCORES 1 ③

# configure JTAG interface, setup target, authenticate and start
source [file join $PATH lib/interface.cfg] ④
source [file join $PATH lib/target.cfg] ⑤
source [file join $PATH lib/authenticate.cfg] ⑥
source [file join $PATH lib/start.cfg] ⑦
```

- ① Get the absolute path of the script.
- ② Set the core name that will show up in openOCD and GDB. This can be changed to any custom name.
- ③ Number of CPU cores (1 for the single-core setup, 2 for the dual-core setup).
- ④ Interface adapter (JTAG) configuration script. Replace by custom adapter setup.
- ⑤ NEORV32-specific target configuration and initialization.
- ⑥ Optional authentication process; the [Default Authentication Mechanism](#) is implemented as example; Replace this when using a custom authentication mechanism.
- ⑦ Reset and halt the target.

5.2. Semihosting

Semihosting is a mechanism developed by ARM that enables code running on the NEORV32 to communicate and use the input/output facilities on a host computer that is running a debugger. Examples of these facilities include keyboard input, screen output, and disk I/O. For example, you can use this mechanism to enable functions in the C library, such as `printf()` and `scanf()`, to use the screen and keyboard of the host instead of having a screen and keyboard on the target system.
(source: [https://developer.arm.com/documentation/dui0471/i/semitesting/what-is-semitesting?
lang=en](https://developer.arm.com/documentation/dui0471/i/semitesting/what-is-semitesting?lang=en))

For example, with semihosting you can

- print to the host's `stdout` console
- read from the host's `stdin` console
- execute arbitrary command in the host's shell
- retrieve the host's system time
- read and write files on the host system
- ...



Semihosting Example Program

A simple semihosting example program can be found in [sw/example/demo_semitesting](#).



RISC-V Semihosting Specification

A copy of the implemented RISC-V semihosting specification is available in the documentation references folder: [docs/references/riscv-semitesting.pdf](#)

The RISC-V semihosting builds upon the ARM standard. Hence, semihosting has to be **explicitly enabled** on the host side (in `GDB`) using the ARM command set:

```
(gdb) monitor arm semihosting enable  
semihosting is enabled
```

File accesses need to be explicitly enabled. Additionally, the base folder for accessing those file should be defined:

```
(gdb) monitor arm semihosting_fileio enable  
(gdb) monitor arm semihosting_basedir path/to/neorv32/sw/example/demo_semitesting
```

The NEORV32 software framework provides a build-in library for semihosting primitive ([sw/lib/include/neorv32_semitesting.h](#)). Additionally, accesses to the standard IO streams (`stdin` and `stdout`) can be automatically mapped to the host's console. Functions such as `printf` and `puts`

can then print right to the host's `stdout` console. Vice versa, functions like `scanf` and `fgets` will read from the host's `stdin`. To enable this automatic mapping, the `define STDIO_SEMIHOSTING` needs to be defined and the application firmware needs to be recompiled.

Listing 36. Forward `stdio.h` Calls to Host Computer

```
USER_FLAGS += -DSTDIO_SEMIHOSTING
```

When `STDIO_SEMIHOSTING` is defined *all* file accesses provided by `stdio.h` (via Newlib) will be redirected to the host computer via semihosting services.



Redirecting all UART data via Semihosting

You can redirect **all** physical UART data via semihosting by defining `UART_SEMIHOSTING` (e.g. "`USER_FLAGS+=-DSTDIO_SEMIHOSTING -DUART_SEMIHOSTING`").



Semihosting Services Without a Host

If any semihosting request is issued without a host being connected, an environment breakpoint exception is raised.

Further references:

- A great overview: <https://interrupt.memfault.com/blog/arm-semihosting>
- Implementing semihosting on RISC-V: <https://embeddedinn.com/articles/tutorial/understanding-riscv-semihosting/>
- Description of the service calls by ARM: <https://developer.arm.com/documentation/dui0203/j/semihosting/semihosting-operations?lang=en>

5.3. Debug Transport Module (DTM)

The debug transport module "DTM" (VHDL module: [rtl/core/neorv32_debug_dtm.vhd](#)) provides a bridge between a standard 4-wire JTAG test access port ("tap") and the internal debug module interface.

Table 102. JTAG Top Level Signals of the DTM

Name	Width	Direction	Description
jtag_tck_i	1	in	serial clock
jtag_tdi_i	1	in	serial data input
jtag_tdo_o	1	out	serial data output
jtag_tms_i	1	in	mode select

Maximum JTAG Clock



All JTAG signals are synchronized to the processor's clock domain. Hence, no additional clock domain is required for the DTM. However, this constraints the maximal JTAG clock frequency (`jtag_tck_i`) to be less than or equal to **1/5** of the processor clock frequency (`clk_i`).



JTAG TAP Reset

The NEORV32 JTAG TAP does not provide a dedicated reset signal ("TRST"). However, JTAG-level resets can be triggered using TMS signaling.



Maintaining the JTAG Chain

If the on-chip debugger is disabled the JTAG serial input `jtag_tdi_i` is directly connected to the JTAG serial output `jtag_tdo_o` to maintain the JTAG chain.

The DTM implement a single 5-bit *instruction register* **IR** and several *data registers* **DR** with different sizes. The individual data registers are accessed by writing the according address to the instruction register. The following table shows all available data registers and their addresses:

Table 103. JTAG TAP registers

Address (via IR)	Name	Size (bits)	Description
00001	IDCODE	32	identification code (see below)
10000	DTMCS	32	debug transport module control and status register (see below)
10001	DMI	41	debug module interface (see below)
others	BYPASS	1	default JTAG bypass register

Table 104. IDCODE - DTM Identification Code Register

Bit(s)	Name	R/W	Description
31:28	<code>version</code>	r/-	version ID, hardwired to zero
27:12	<code>partid</code>	r/-	part ID, hardwired to zero
11:1	<code>manid</code>	r/-	JEDEC manufacturer ID, assigned via the JEDEC_ID generic
0	-	r/-	hardwired to 1

Table 105. [DTMCS](#) - DTM Control and Status Register

Bit(s)	Name	R/W	Description
31:21	-	r/-	<i>reserved</i> , hardwired to zero
20:18	<code>errinfo</code>	r/-	not implemented; hardwired to zero
17	<code>dmihardreset</code>	r/w	setting this bit will reset the debug module interface; this bit auto-clears
16	<code>dmireset</code>	r/w	setting this bit will clear the sticky error state; this bit auto-clears
15	-	r/-	<i>reserved</i> , hardwired to zero
14:12	<code>idle</code>	r/-	recommended idle states (= 0, no idle states required)
11:10	<code>dmistat</code>	r/-	read-only alias of DMI.op ; hardwired to 00 (NOP)
9:4	<code>abits</code>	r/-	number of address bits in DMI register (= 6)
3:0	<code>version</code>	r/-	0001 = DTM is compatible to RISC-V debug spec. versions v0.13 and v1.0

Table 106. [DMI](#) - DTM Debug Module Interface Register

Bit(s)	Name	R/W	Description
40:34	<code>addr</code>	r/w	7-bit address, see DM Registers
33:2	<code>data</code>	r/w	32-bit data to write/read to/from the addresses DM register
1:0	<code>op</code>	r/w	2-bit operation (00 = NOP; 10 = write; 01 = read)

5.4. Debug Module (DM)

The debug module "DM" (VHDL module: [rtl/core/neorv32_debug_dm.vhd](#)) acts as a translation interface between abstract operations issued by the debugger application (like GDB) and the platform-specific debugger hardware. It supports the following features:

- Gives the debugger necessary information about the implementation.
- Allows the hart to be halted/resumed/reset and provides the current status.
- Provides abstract read and write access to the halted hart's general purpose registers.
- Provides access to a reset signal that allows debugging from the very first instruction after reset.
- Provides a *program buffer* to force the hart to execute arbitrary instructions.
- Allows memory accesses (to the entire address space) from a hart's point of view.
- Optionally implements an authentication mechanism to secure on-chip debugger access.

The NEORV32 DM follows the "Minimal RISC-V External Debug Specification" to provide full debugging capabilities while keeping resource/area requirements at a minimum. It implements the **execution based debugging scheme** for up to four individual CPU cores ("harts") and provides the following architectural core features:

- program buffer with 2 entries and an implicit `ebreak` instruction at the end
- indirect bus access via the CPU using the program buffer
- abstract commands: "access register" plus auto-execution
- halt-on-reset capability
- optional authentication



DM Spec. Version

The NEORV32 DM complies to the RISC-V DM spec version 1.0.

From the DTM's point of view, the DM implements a set of **DM Registers** that are used to control and monitor the debugging session. From the CPU's point of view, the DM implements several memory-mapped registers that are used for communicating data, instructions, debugging control and status (**DM CPU Access**).

External Reset Output

The entire processor can be reset at any time by the debugger via the `ndmreset` bit of the `dmcontrol` register. This signal is also available as processor top signal (**Processor Top Entity - Signals: `rstn_ocd_o`**) and can be used to reset processor-external modules via the on-chip debugger. This signal is low-active and synchronous to the processor clock. It is available if the on-chip debugger is actually implemented; otherwise it is hardwired to 1. Note that the signal also becomes active (low) when the processor's main reset signal is active (even if the on-chip debugger is deactivated or disabled for synthesis).

5.4.1. DM Registers

The DM is controlled via a set of registers that are accessed via the DTM. The following registers are implemented:

Unimplemented Registers



Write accesses to registers that are not implemented are simply ignored and read accesses to these registers will always return zero. In both cases no error condition is signaled to the DTM.

Table 107. Available DM registers

Address	Name	Description
0x04	data0	Abstract data register 0
0x10	dmcontrol	Debug module control
0x11	dmstatus	Debug module status
0x12	hartinfo	Hart information
0x16	abstracts	Abstract control and status
0x17	command	Abstract command
0x18	abstractauto	Abstract command auto-execution
0x1d	nextdm	Base address of next DM; reads as zero to indicate there is only one DM
0x20	progbuf0	Program buffer 0
0x21	progbuf1	Program buffer 1
0x30	authdata	Data to/from the authentication module
0x38	sbc	System bus access control and status; reads as zero to indicate there is no system bus access
0x40	haltsum0	Hart halt summary

data0

0x04 **Abstract data 0**

data0

Reset value: **0x00000000**

Basic read/write data exchange register to be used with abstract commands (for example to read/write data from/to CPU GPRs).

dmcontrol

0x10 **Debug module control register**

dmcontrol

Reset value: **0x00000000**

Control of the overall debug module and the hart. The following table shows all implemented bits. All remaining bits/bit-fields are configured as "zero" and are read-only. Writing '1' to these bits/fields will be ignored.

Table 108. `dmcontrol` Register Bits

Bit	Name [RISC-V]	R/W	Description
31	<code>haltreq</code>	-/w	set/clear hart halt request
30	<code>resumereq</code>	-/w	request hart to resume
28	<code>ackhavereset</code>	-/w	write 1 to clear *havereset flags
27	-	r/-	reserved, hardwired to zero
26	<code>hasel</code>	r/-	0: only a single hart can be selected at once
25:16	<code>hartsello</code>	r/w	hart select; only the lowest 3 bits are implemented
15:6	<code>hartselhi</code>	r/-	hardwired to zero
5:4	-	r/-	reserved, hardwired to zero
3	<code>setresethaltreq</code>	r/-	0: halt-on-reset not implemented
2	<code>clrresethaltreq</code>	r/-	0: halt-on-reset not implemented
1	<code>ndmreset</code>	r/w	put whole system (except OCD) into reset state when 1
0	<code>dmactive</code>	r/w	DM enable; writing 0-1 will reset the DM

`dmstatus`

0x11 **Debug module status register**

`dmstatus`

Reset value: `0x00400083`

Current status of the overall debug module and the hart. The entire register is read-only.

Table 109. `dmstatus` Register Bits

Bit	Name [RISC-V]	Description
31:25	<code>reserved</code>	reserved; zero
24	<code>ndmresetpendning</code>	DM in reset state
23	<code>stickyunavail</code>	*unavail bits reflect the current state
22	<code>impebreak</code>	1: indicates an implicit <code>ebreak</code> instruction after the last program buffer entry
21:20	<code>reserved</code>	reserved; zero

Bit	Name [RISC-V]	Description
19	allhavereset	1 when the selected hart is in reset state
18	anyhavereset	
17	allresumeack	1 when the selected hart has acknowledged a resume request
16	anyresumeack	
15	allnonexistent	1 when the selected hart is not available
14	anynonexistent	
13	allunavail	1 when the DM is disabled to indicate the selected hart is unavailable
12	anyunavail	
11	allrunning	1 when the selected hart is running
10	anyrunning	
9	allhalted	1 when the selected hart is halted
8	anyhalted	
7	authenticated	set if authentication passed; see Debug Authentication
6	authbusy	set if authentication is busy, see Debug Authentication
5	hasresethaltr eq	0: halt-on-reset is not supported (directly)
4	confstrptrval id	0: no configuration string available
3:0	version	0011: DM compatible to debug spec. version v1.0

hartinfo0x12 **Hart information**[hartinfo](#)Reset value: *see below*

This register gives information about the hart. The entire register is read-only.

Table 110. **hartinfo** Register Bits

Bit	Name [RISC-V]	Description
31:24	reserved	reserved; zero
23:20	nscratch	0001: number of dscratch* CPU registers = 1
19:17	reserved	reserved; zero
16	dataaccess	0: the data registers are shadowed in the hart's address space
15:12	datasize	0001: number of 32-bit words in the address space dedicated to shadowing the data registers (1 register)

Bit	Name [RISC-V]	Description
11:0	dataaddr	= <code>dm_data_base_c(11:0)</code> , signed base address of <code>data</code> words (see address map in DM CPU Access)

abstracts

0x16 Abstract control and status

abstracts

Reset value: `0x02000801`

Command execution info and status.

Table 111. `abstracts` Register Bits

Bit	Name [RISC-V]	R/W	Description
31:29	<code>reserved</code>	r/-	reserved; zero
28:24	<code>progbuftype</code>	r/-	<code>0010</code> : size of the program buffer (<code>progbuf</code>) = 2 entries
23:11	<code>reserved</code>	r/-	reserved; zero
12	<code>busy</code>	r/-	set when a command is being executed
11	<code>relaxedpriv</code>	r/-	<code>1</code> : PMP rules are ignored when in debug mode
10:8	<code>cmderr</code>	r/w	error during command execution (see below); has to be cleared by writing <code>111</code>
7:4	<code>reserved</code>	r/-	reserved; zero
3:0	<code>datacount</code>	r/-	<code>0001</code> : number of implemented <code>data</code> registers for abstract commands = 1

Error codes in `cmderr` (highest priority first):

- `000` - no error
- `100` - command cannot be executed since hart is not in expected state
- `011` - exception during command execution
- `010` - unsupported command
- `001` - invalid DM register read/write while command is/was executing

command

0x17 Abstract command

command

Reset value: `0x00000000`

Writing this register will trigger the execution of an abstract command. New command can only be executed if `cmderr` is zero. The entire register is write-only (reads will return zero).



The NEORV32 DM only supports **Access Register** abstract commands. These commands can only access the hart's GPRs x0 - x15/31 (abstract command register

index `0x1000 - 0x101f`).

Table 112. command Register Bits

Bit	Name [RISC-V]	R/W	Description / required value
31:24	cmdtype	-/w	00000000: indicates "access register" command
23	reserved	-/w	reserved, has to be 0 when writing
22:20	aarsize	-/w	010: indicates 32-bit accesses
21	aarpostincrement	-/w	0: post-increment is not supported
18	postexec	-/w	set if the program buffer is executed <i>after</i> the command
17	transfer	-/w	set if the operation in write is conducted
16	write	-/w	1: copy data0 to [regno], 0: copy [regno] to data0
15:0	regno	-/w	GPR-access only; has to be 0x1000 - 0x101f

abstractauto

0x18 Abstract command auto-execution abstractauto
Reset value: **0x00000000**
Register to configure if a read/write access to a DM register re-triggers execution of the last abstract command.

Table 113. abstractauto Register Bits

Bit	Name [RISC-V]	R/W	Description
17	autoexecprogbu f[1]	r/w	when set reading/writing from/to <code>proobuf1</code> will execute <code>command</code> again
16	autoexecprogbu f[0]	r/w	when set reading/writing from/to <code>proobuf0</code> will execute <code>command</code> again
0	autoexecdata[0]	r/w	when set reading/writing from/to <code>data0</code> will execute <code>command</code> again

protobuf

0x20	Program buffer 0	probuf0
0x21	Program buffer 1	probuf1

Reset value: 0x00000013 ("NOP")

Program buffer (two entries) for the DM.

authdata

0x30 Authentication data authdata

Reset value: *user-defined*

This register serves as a 32-bit serial port to/from the authentication module. See [Debug Authentication](#).

haltsum0

0x30 **Halt summary 0**

[haltsum0](#)

Reset value: [0x00000000](#)

Each bit corresponds to a hart being halted. Only the lowest four bits are implemented.

5.4.2. DM CPU Access

From the CPU's perspective the DM acts like another memory-mapped peripheral. It occupies 512 bytes of the CPU's address space starting at address [base_io_dm_c](#) ([0xfffff000](#)). This address space is divided into four sections with 64 bytes each to provide access to the *park loop code ROM*, the *program buffer*, the *data buffer* and the *status register*. The program buffer, the data buffer and the status register do not fully occupy the 64-byte-wide sections and are mirrored several times across the according section.

Table 114. DM CPU Access - Address Map

Base address	Physical size	Description
0xfffffffff00	64 bytes	ROM for the "park loop" code (Code ROM)
0xfffffffff40	16 bytes	Program buffer (progbuf)
0xfffffffff80	4 bytes	Data buffer (data0)
0xfffffffffc0	4 bytes	Status Register

DM Register Access



All memory-mapped registers of the DM can only be accessed by the CPU when in debug mode. Hence, the DM registers are not accessible for normal CPU operations. Any CPU access outside of debug mode will raise a bus access fault exception.

Code ROM

The code ROM contain the minimal OCD firmware that implements the debuggers part loop.

Park Loop Code Sources ("OCD Firmware")



The assembly sources of the park loop code are available in [sw/ocd-firmware/park_loop.S](#).

The park loop code provides two entry points where code execution can start. These are used to enter the park loop either when an explicit debug-entry/halt request has been issued (for example a

halt request) or when an exception has occurred while executing code in debug mode (from the program buffer).

Table 115. Park Loop Entry Points

Address	Description
<code>dm_exc_entry_c</code> (<code>base_io_dm_c + 0</code>)	Exception entry address
<code>dm_park_entry_c</code> (<code>base_io_dm_c + 4</code>)	Normal entry address (halt request)

When the CPU enters (via an explicit halt request from the debugger) or re-enters debug mode (for example via an `ebreak` in the DM's program buffer), it jumps to the **normal entry point** that is configured via the `CPU_DEBUG_PARK_ADDR` CPU generic. By default, this address is set to `dm_park_entry_c`, which is defined in the main package file. If an exception is encountered during debug mode, the CPU jumps to the address of the **exception entry point** configured via the `CPU_DEBUG_EXC_ADDR` CPU generic. By default, this address is set to `dm_exc_entry_c`, which is also defined in the main package file.

Status Register

The status register provides a direct communication channel between a CPU executing the park loop and the debugger-controlled DM. This register is used to communicate **requests** which are issued by the DM, and the according **acknowledges** which are issued by a CPU core.

From the CPU side the register is accessed using byte-wide load and store operations. Each of the four bytes is associated to a specific request when read or to an according acknowledge when written.

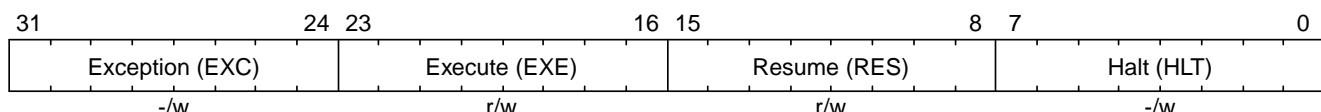


Figure 20. DM Status Register Layout

If the park-loop code of a CPU reads a non-zero value from `RES` then it is requested to resume operation. If it reads a non-zero value from `EXE` then it is requested to execute the program buffer. Accordingly, the park-loop code acknowledges these requests by writing *any* value (including zero) to the according byte. The `HAT` and `EXC` bytes are used to acknowledge the external halt request and to indicate that executing the program buffer has encountered an exception, respectively.

Table 116. DM Status Register (byte-wise access)

Name	Bits	CPU Read Access (Request)	CPU Write Access (Acknowledge)
<code>EXC</code>	31:24	-	CPU has encountered an exception
<code>EXE</code>	23:16	DM requests program buffer execution	CPU has started to execute program buffer
<code>RES</code>	15:8	DM requests resuming	CPU has started to resume
<code>HLT</code>	7:0	-	CPU has halted

In the SMP [Dual-Core Configuration](#) all cores use the same request/acknowledge communication mechanism by reading and writing the DM status register. The CPU-specific requests and acknowledges can still be clearly assigned, as the CPU [Bus Interface](#) implicitly transmits the ID of the accessing CPU core. Therefore, CPU core i only sees the requests intended for it and can only set the acknowledges originating from itself when accessing the DM status register.

5.5. Debug Authentication

Optionally, the on-chip debugger DM can be equipped with an *authenticator module* to secure debugger access. This authentication is enabled by the `OCD_AUTHENTICATION` top generic. When disabled, the debugger is always authorized and has unlimited access. When enabled, the debugger is required to authenticate in order to gain access.

The authenticator module is implemented as individual RTL module (`rtl/core/neorv32_debug_auth.vhd`). By default, it implements a very simple authentication mechanism. Note that this default mechanism is not secure in any way - it is intended as example logic to illustrate the interface and authentication process. Users can modify the default logic or replace the entire module to implement a more sophisticated custom authentication mechanism.

The authentication interface is compliant to the RISC-V debug spec and is based on a single CSR and two additional status bits:

- **`authdata`** CSR: this 32-bit register is used to read/write data from/to the authentication module. It is hardwired to all-zero if authentication is not implemented.
- **`dmstatus`** CSR:
 - The `authenticated` bit (read-only) is set if authentication was successful. The debugger can access the processor only if this bit is set. It is automatically hardwired to `1` (always authenticated) if the authentication module is not implemented.
 - The `authbusy` bit (read-only) indicates if the authentication module is busy. When set, no data should be written/read to/from `authdata`. This bit is automatically hardwired to `0` (never busy) if the authentication module is not implemented.

openOCD provides dedicated commands to exchange data with the authenticator module:

Listing 37. openOCD RISC-V Authentication Commands

```
riscv authdata_read      // read 32-bit from authdata CSR
riscv authdata_write value // write 32-bit value to authdata CSR
```

Based on these two primitives arbitrary complex authentication mechanism can be implemented.

5.5.1. Default Authentication Mechanism



The default authentication mechanism is not secure at all. Replace it by a custom design.

The default authenticator hardware implements a very simple authentication mechanism: a single read/write bit is implemented that directly corresponds to the `authenticated` bit in `dmstatus`. This bit can be read/written as bit zero (LSB) of the `authdata` CSR. Writing `1` to this register will result in a successful authentication.

The default openOCD configuration script provides a helper script for authentication. This script

also provides several helper functions for interaction with the RISC-V debug mechanism. Additionally, the default authentication mechanism is implemented there (as example):

Listing 38. Default authentication process ([sw/openocd/authentication.cfg](#))

```
# read challenge
set CHALLENGE [authenticator_read]
# compute response (default authenticator module)
set RESPONSE [expr {$CHALLENGE | 1}]
# send response
authenticator_write $RESPONSE
# success?
authenticator_check
```

5.6. CPU Debug Mode

The NEORV32 CPU Debug Mode is compatible to the **Minimal RISC-V Debug Specification 1.0 Sdext** (external debug) ISA extension. When enabled via the CPU's **Sdext ISA Extension** generic and/or the processor's **OCD_EN** it adds a new CPU operation mode ("debug mode"), three additional **CPU Debug Mode CSRs** and one additional instruction (**dret**) to the core.

Debug-mode is entered on any of the following events:

1. The CPU executes an **ebreak** instruction (when in machine-mode and **ebreakm** in **dcsr** is set OR when in user-mode and **ebreaku** in **dcsr** is set).
2. A debug halt request is issued by the DM (via CPU **db_halt_req_i** signal, high-active).
3. The CPU completes executing of a single instruction while being in single-step debugging mode (**step** in **dcsr** is set).
4. A hardware trigger from the **Trigger Module** fires (if **exe** in **tdata1 / mcontrol** is set).



From a hardware point of view these debug-mode-entry conditions are normal traps (synchronous exceptions or asynchronous interrupts) that are handled transparently by the control logic.

Whenever the CPU enters debug-mode it performs the following operations:

- wake-up CPU if it was send to sleep mode by the **wfi** instruction
- switch to debug-mode privilege level
- move the current program counter to **dpc**
- copy the hart's current privilege level to the **prv** flags in **dcsr**
- set **cause** in **dcsr** according to the cause why debug mode is entered
- **no update** of **mtval**, **mcause**, **mtval** and **mstatus[h]** CSRs
- load the address configured via the CPU's (**CPU_DEBUG_PARK_ADDR**) generic to the program counter jumping to the "debugger park loop" code stored in the debug module (DM)

When the CPU is in debug-mode:

- while in debug mode, the CPU executes the parking loop and - if requested by the DM - the program buffer
- effective CPU privilege level is **machine** mode; any active physical memory protection (PMP) configuration is bypassed
- the **wfi** instruction acts as a **nop** (also during single-stepping)
- if an exception occurs while being in debug mode:
 - if the exception was caused by any debug-mode entry action the CPU jumps to the normal entry point (defined by the **CPU_DEBUG_PARK_ADDR** generic) of the park loop again (for example when executing **ebreak** while in debug-mode)

- for all other exception sources the CPU jumps to the exception entry point (defined by the `CPU_DEBUG_EXC_ADDR` generic) to signal an exception to the DM; the CPU restarts the park loop again afterwards
- interrupts are disabled; however, they will remain pending and will get executed after the CPU has left debug mode and is not being single-stepped
- if the DM makes a resume request, the park loop exits and the CPU leaves debug mode (executing `dret`)
- the standard counters (**Machine**) Counter and Timer CSRs `[m]cycle[h]` and `[m]instret[h]` are stopped
- all **Hardware Performance Monitors (HPM)** CSRs are stopped

Debug mode is left either by executing the `dret` instruction or by performing a hardware reset of the CPU. Executing `dret` outside of debug mode will raise an illegal instruction exception.

Whenever the CPU leaves debug mode it performs the following operations:

- set the hart's current privilege level according to the `prv` flags of `dcsr`
- restore the original program counter from `dpc` resuming normal operation

5.6.1. CPU Debug Mode CSRs

Two additional CSRs are required by the "Minimal RISC-V Debug Specification": the debug mode control and status register `dcsr` and the debug program counter `dpc`. An additional general purpose scratch register for debug-mode-only (`dscratch0`) allows faster execution by having a fast-accessible backup register. These CSRs are only accessible if the CPU is in debug mode. If these CSRs are accessed outside of debug mode an illegal instruction exception is raised.

`dcsr`

Name	Debug control and status register
Address	<code>0x7b0</code>
Reset value	<code>0x40000410</code>
ISA	<code>Zicsr & Sdext</code>
Descripti on	This register is used to configure the debug mode environment and provides additional status information.

Table 117. Debug control and status register `dcsr` bits

Bit	Name [RISC-V]	R/W	Description
31:28	<code>xdebugver</code>	r/-	<code>0100</code> : CPU debug mode is compatible to spec. version 1.0
27:16	-	r/-	<code>000000000000</code> : reserved

Bit	Name [RISC-V]	R/W	Description
15	<code>ebreakm</code>	r/w	<code>ebreak</code> instructions in <code>machine</code> mode will <i>enter</i> debug mode when set
14	<code>ebreakh</code>	r/-	<code>0</code> : hypervisor mode not supported
13	<code>ebreaks</code>	r/-	<code>0</code> : supervisor mode not supported
12	<code>ebreaku</code>	r/w	<code>ebreak</code> instructions in <code>user</code> mode will <i>enter</i> debug mode when set
11	<code>stepie</code>	r/-	<code>0</code> : IRQs are disabled during single-stepping
10	<code>stopcount</code>	r/-	<code>1</code> : standard counters and HPMs are stopped when in debug mode
9	<code>stoptime</code>	r/-	<code>0</code> : timers increment as usual
8:6	<code>cause</code>	r/-	cause identifier: why debug mode was entered (see below)
5	-	r/-	<code>0</code> : reserved
4	<code>mprvn</code>	r/-	<code>1</code> : <code>mprv</code> in <code>mstatus</code> is also evaluated when in debug mode
3	<code>nmiip</code>	r/-	<code>0</code> : non-maskable interrupt is pending
2	<code>step</code>	r/w	enable single-stepping when set
1:0	<code>prv</code>	r/w	CPU privilege level before/after debug mode

Cause codes in `dcsr.cause` (highest priority first):

- `010` - triggered by hardware [Trigger Module](#)
- `001` - executed [EBREAK](#) instruction
- `011` - external halt request (from DM)
- `100` - return from single-stepping

dpc

Name Debug program counter

Address `0x7b1`

Reset `0x00000000`

value

ISA `Zicsr & Sdext`

Descripti on The register is used to store the current program counter when debug mode is entered. The `dret` instruction will return to the address stored in `dpc` by automatically moving `dpc` to the program counter.



`dpc[0]` is hardwired to zero. If `IALIGN` = 32 (i.e. [C ISA Extension](#) is disabled) then `dpc[1]` is also hardwired to zero.

dscratch0

Name	Debug scratch register 0
Address	<code>0x7b2</code>
Reset value	<code>0x00000000</code>
ISA	<code>Zicsr & Sdext</code>
Description	The register provides a general purpose debug mode-only scratch register.
on	

5.7. Trigger Module

The RISC-V [Sdtrig](#) ISA extension adds a *trigger module* to the CPU core. The number of hardware triggers is configured via the [OCD_NUM_HW_TRIGGER](#)s top generic. Up to 16 hardware triggers can be implemented. If [OCD_NUM_HW_TRIGGER](#)s is set to a value greater than 0 the [Sdtrig ISA Extension](#) is automatically enabled. The trigger module implements a subset of the features described in the "RISC-V Debug Specification / Trigger Module ([Sdtrig](#))" and complies to version v1.0 of that spec. Note that the NEORV32 trigger module only supports "type 6 - instruction address match" triggers. These triggers can be used as hardware-assisted breakpoints or as hardware-assisted watchpoints.

- A **breakpoint** stops the program whenever a particular point in the program is reached (i.e. after executing the instruction at the specified address): GDB's [hbreak](#) command.
- A **watchpoint** stops the program whenever the value of a variable or expression changes (i.e. after a load and/or store operation has accessed data at the specified address): GDB's [\[a|r\]watch](#) commands.

From a hardware point of view the trigger modules raises an *asynchronous* exception right after the triggering operation has been retired. The [dpc](#) CSR shows the instruction address where normal execution must be resumed to preserve the program flow

Debug-Mode Only



The trigger module can be used by debug-mode only. Write accesses from machine-mode are constrained as [tdata1.dmode](#) is hardwired to one. However, the [Smpmp ISA Extension](#) can be used for machine-mode instruction or data address traps.

5.7.1. Trigger Module CSRs

The [Sdtrig](#) ISA extension adds 4 additional CSRs. These CSRs can also be read from machine-mode, but write-accesses are restricted because [tdata1.dmode](#) is hardwired to one.

[tselect](#)

Name Trigger select register

Address [0x7a0](#)

Reset [0x00000000](#)

value

ISA [Zicsr](#) & [Sdtrig](#)

Description This CSR is used to select which hardware trigger is available via the [tdata1](#), [tdata2](#) and [tinfo](#) CSRs. Only the lowest $\log_2(\text{OCD_NUM_HW_TRIGGERS})+1$ bit are writable. However, software should check if [tinfo](#) is non-zero for the written [tselect](#) value to determine if the selected trigger exists.

tdata1

Name	Trigger data register 1, visible as trigger "type 6 match control" (mcontrol6)		
Address	0x7a1		
Reset value	0x60000048		
ISA	Zicsr & Sdtrig		
Description	This CSR provides access to the configuration bits of the currently selected trigger (via tselect). Note that only "type 6" triggers are supported. The according configuration bits are listed below. Write accesses from machine-mode are ignored (because tdata1.dmode is hardwired to one).		

Table 118. Match control CSR (tdata1) bits of the trigger selected via tselect

Bit	Name [RISC-V]	R/W	Description
31:28	type	r/-	0110 : address match trigger type 6
27	dmode	r/-	'1': ignore write accesses to tdata1 and tdata2 from machine-mode
26	uncertain	r/-	0: trigger satisfies the configured conditions
25	hit1	r/c	see hit0 bit
24	vs	r/-	0: VS-mode not supported
23	vu	r/-	0: VU-mode not supported
22	hit0	r/c	hit1 (MSB) and hit0 (LSB) combine into a single 2-bit field; this field is set to 0b11 when the trigger just fired; timing: the trigger fires immediately after the instruction that triggered has retired, but before any subsequent instructions were executed; the debugger has to write zeros to acknowledge the firing trigger
21	select	r/-	0: only address matching is supported
20:19	reserved	r/-	00 : hardwired to zero
18:16	size	r/-	000 : match accesses of any size
15:12	action	r/-	0001 : enter debug-mode on trigger match
11	chain	r/-	0: chaining is not supported
10:6	match	r/-	0000 : equal-match only
6	m	r/-	1: trigger enabled when in machine-mode
5	uncertainen	r/-	0: feature not supported, hardwired to zero
4	s	r/-	0: S-mode not supported
3	u	r/-	trigger enabled when in user-mode, set if U ISA extension is enabled

Bit	Name [RISC-V]	R/W	Description
2	execute	r/w	set to enable this trigger to fire on instruction address match (i.e. "breakpoint")
1	store	r/w	set to enable this trigger to fire on data store address match (i.e. "watchpoint")
0	load	r/w	set to enable this trigger to fire on data load address match (i.e. "watchpoint")

tdata2

Name	Trigger data register 2
Address	0x7a2
Reset	0x00000000
value	
ISA	Zicsr & Sdtrig
Descripti on	This CSR provides access to the address bits (instruction address for breakpoints; data access address for watchpoints) of the currently selected trigger (via tselect). Write accesses from machine-mode are ignored (because tdata1.dmode is hardwired to one).

tinfo

Name	Trigger information register
Address	0x7a4
Reset	0x01000040
value	
ISA	Zicsr & Sdtrig
Descripti on	The CSR shows additional trigger information of the currently selected trigger (via tselect). This CSR is read-only, all write access is ignored.

Table 119. Trigger info CSR (**tinfo**) bits of the trigger selected via **tselect**

Bit	Name [RISC-V]	R/W	Description
31:24	version	r/-	0x01: compatible to Sdtrig spec. version v1.0
23:15	reserved	r/-	0x00: hardwired to zero
15:0	info	r/-	0x0040: only "type 6" (= bit 6 set) trigger is supported; for invalid value of tselect this bit-field read as 0x0001 ("type 0" = none)

Chapter 6. Legal

About

The NEORV32 RISC-V Processor

<https://github.com/stnolting/neorv32>

Stephan Nolting, M.Sc.

EU European Union

stnolting[ät]gmail[dot]com



DOI

This project provides a *digital object identifier* provided by [zenodo](#):
DOI 10.5281/zenodo.5018888

License

BSD 3-Clause License

Copyright (c) NEORV32 contributors. Copyright (c) 2020 - 2025, Stephan Nolting. All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

1. Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
2. Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.
3. Neither the name of the copyright holder nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT HOLDER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.



SPDX Identifier

SPDX-License-Identifier: BSD-3-Clause

Proprietary Notice

- "GitHub" is a subsidiary of Microsoft Corporation.
- "Vivado" and "Artix" are trademarks of AMD Inc.
- "ARM", "AMBA", "AXI", "AXI4", "AXI4-Lite" and "AXI4-Stream" are trademarks of Arm Holdings plc.
- "ModelSim" is a trademark of Mentor Graphics – A Siemens Business.
- "Quartus [Prime]" and "Cyclone" are trademarks of Intel Corporation.
- "iCE40", "UltraPlus" and "Radian" are trademarks of Lattice Semiconductor Corporation.
- "GateMate" is a trademark of Cologne Chip AG.
- "Windows" is a trademark of Microsoft Corporation.
- "Tera Term" copyright by T. Teranishi.
- "NeoPixel" is a trademark of Adafruit Industries.
- "Segger Embedded Studio" and "J-Link" are trademarks of Segger Microcontroller Systems GmbH.
- Images/figures made with *Microsoft Power Point*.
- Diagrams made with *WaveDrom*.
- Documentation made with [asciidoc](#).

All further/unreferenced projects/products/brands belong to their according copyright holders. No copyright infringement intended.

Disclaimer

This project is released under the BSD 3-Clause license. NO COPYRIGHT INFRINGEMENT INTENDED. Other implied or used projects/sources might have different licensing – see their according documentation for more information.

Limitation of Liability for External Links

This document contains links to the websites of third parties ("external links"). As the content of these websites is not under our control, we cannot assume any liability for such external content. In all cases, the provider of information of the linked websites is liable for the content and accuracy of the information provided. At the point in time when the links were placed, no infringements of the law were recognizable to us. As soon as an infringement of the law becomes known to us, we will immediately remove the link in question.

Acknowledgments

A big shout-out to the community and all the [contributors](#), who helped improving this

project! This project would not be where it is without them. ☺

RISC-V - instruction sets want to be free!

Continuous integration provided by [GitHub Actions](#) and powered by [GHDL](#).